

BigData Spark – Hands-On

Optimisations

SQL Execution Plan, DAG, SparkUI
Performances

Arnaud Nauwynck

Oct 2022

Objectives of Hands-On



Reminders:

Local Install, spark-shell

Spark File IO

MetaStore tables

1/ Execution Plan, SQL Explain / dataset.explain

2/ Spark UI, DAG, Narrow/Wide Transformations

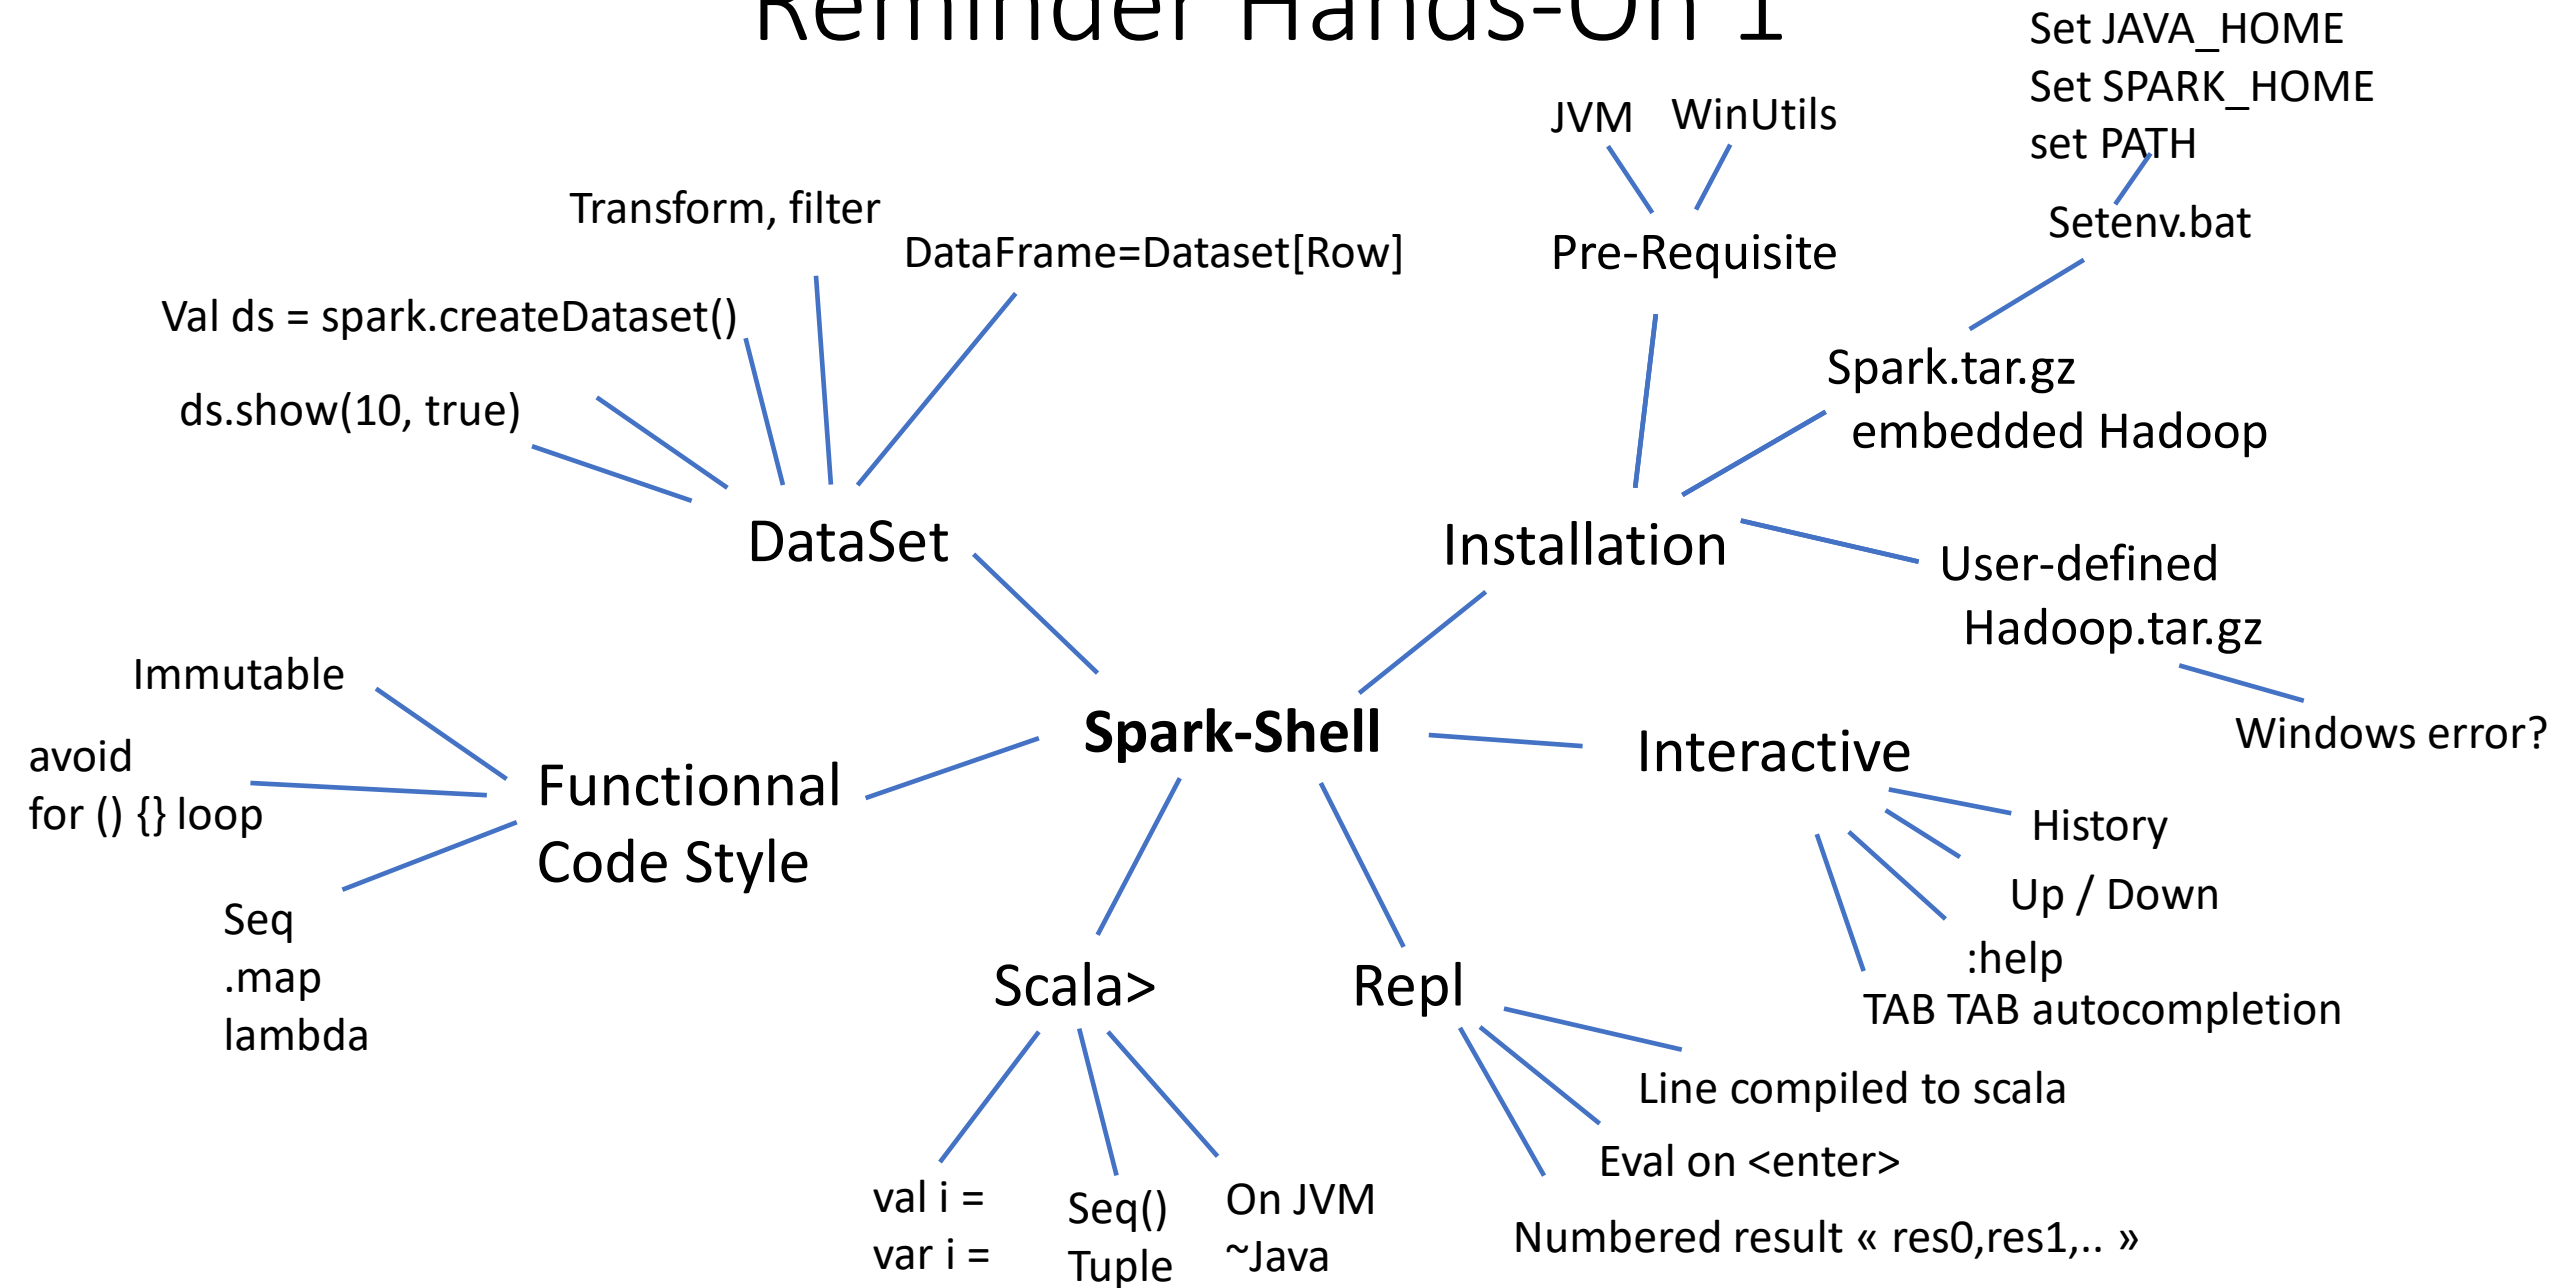
3/ RDD cache / checkpoint

4/ Partition Pruning / Columns Pruning / Predicate-Push-Down

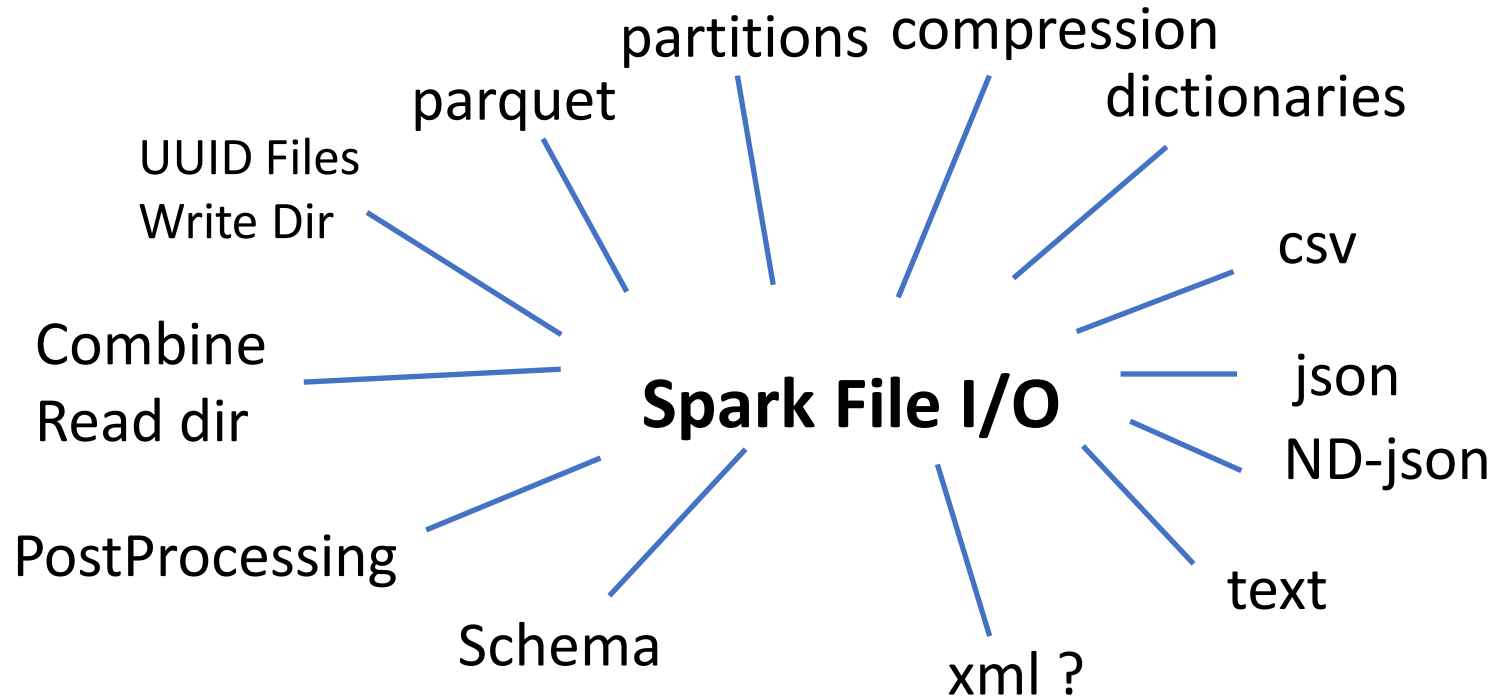
5/ Parquet Optims (Sort, Stats, Dictionary, Bloom, BlockSize)

6/ perfs: SpillToDisk, JOINS, hint

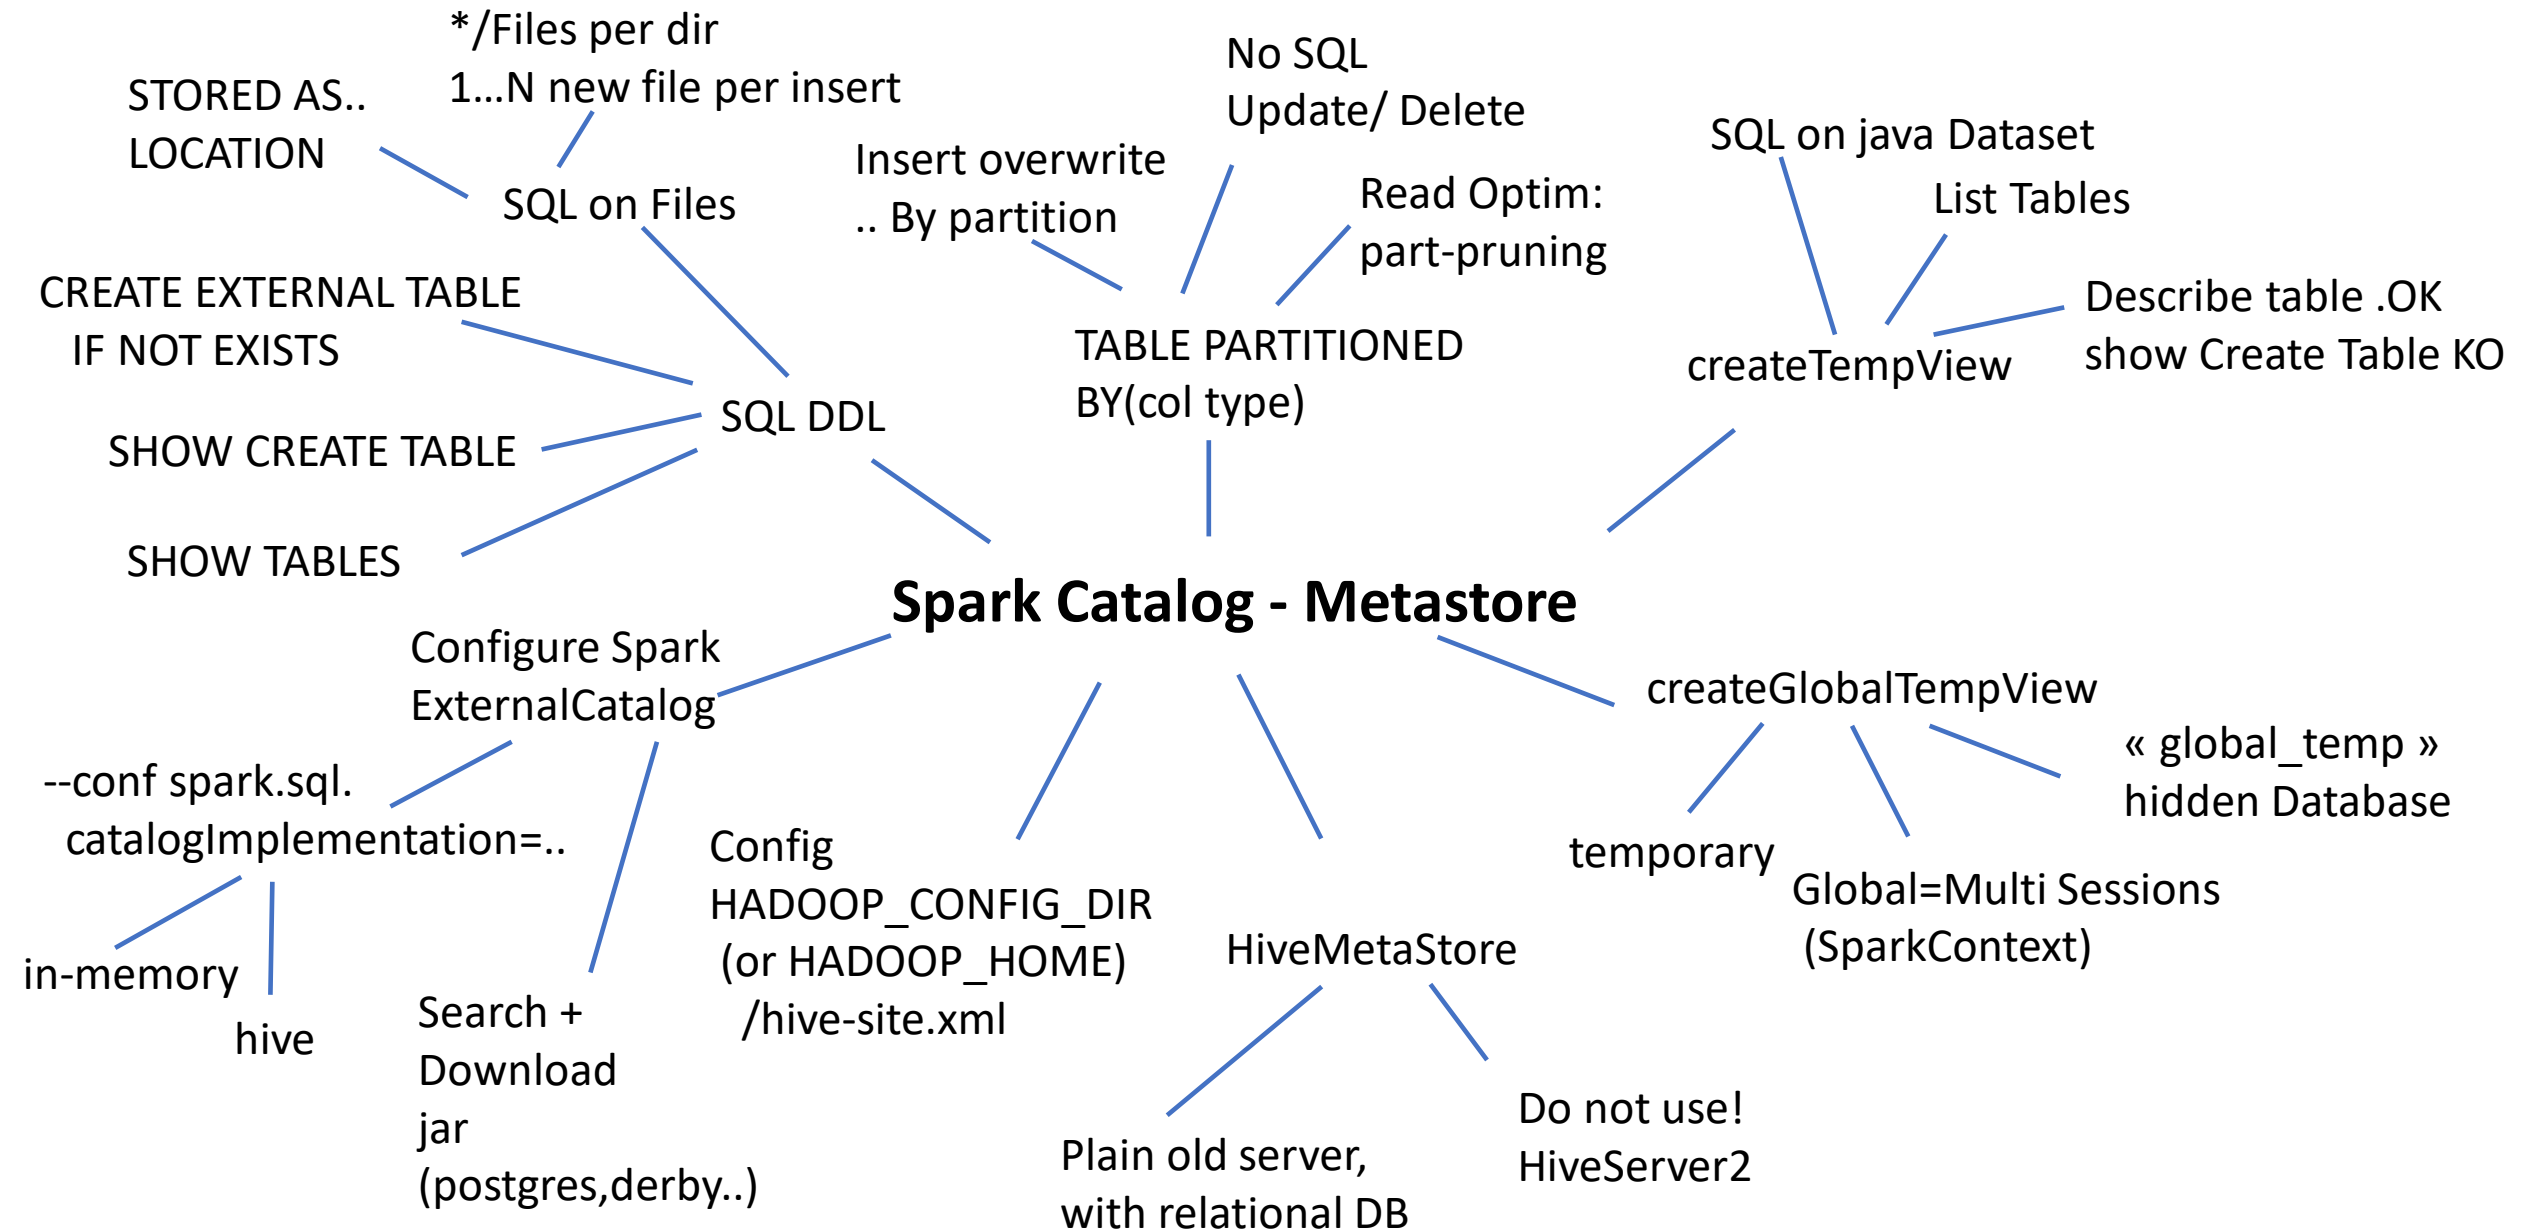
Reminder Hands-On 1



Reminder Hands-On 2



Reminder Hands-On 3



Objectives of Hands-On



1/ Execution Plan, SQL Explain / dataset.explain

2/ Spark UI, DAG, Narrow/Wide Transformations

3/ RDD cache / checkpoint

4/ Partition Pruning / Columns Pruning / Predicate-Push-Down

5/ Parquet Optims (Sort, Stats, Dictionary, Bloom, BlockSize)

6/ perfs: SpillToDisk, JOINS, hint

Exercise 1: Query Partitioned Table

WHERE partitionColumn=..

a/ remind on previous Hands-On

SHOW CREATE TABLE db1.address => .. Un-partitioned table, PARQUET

SHOW CREATE TABLE db1.address_by_dept => .. Partitioned table, PARQUET

b/ execute queries

SELECT * FROM address_by_dept WHERE dept=92

SELECT * FROM address_by_dept WHERE dept in (75,78,91,92)

SELECT * FROM address_by_dept WHERE commune_name = 'Nanterre'

c/ Remind which dir / files should be read by spark... which query is faster

Exercise 2: Execute « EXPLAIN <<SQL>> »

a/ Execute SQL... prefixed by “EXPLAIN” keyword

```
spark.sql("EXPLAIN select count(*) FROM db1.address_by_dept WHERE dept=92")
```

Hint: display nicely, using `.show(false)` OR `.foreach(println(_))`

b/ do you see “PartitionFilters: [..]” ?

c/ read Tree from depth-first :
start from bottom (leaf) line,
when understood then read line above (operator)

Exercise 3 : compare « EXPLAIN » queries

a/ Compare both

```
spark.sql("EXPLAIN select count(*) FROM db1.address_by_dept WHERE dept=92")
```


```
spark.sql("EXPLAIN select count(*) FROM db1.address WHERE dept=92")
```

b/ what changed ?

... Exercise 3: EXPLAIN

[EXTENDED | CODEGEN | COST | FORMATTED]

<https://spark.apache.org/docs/latest/sql-ref-syntax-qry-explain.html>

 Overview Programming Guides ▾ API Docs ▾ Deploying ▾ More ▾

Spark SQL Guide

- Getting Started
- Data Sources
- Performance Tuning
- Distributed SQL Engine
- PySpark Usage Guide for Pandas with Apache Arrow
- Migration Guide
- SQL Reference
 - ANSI Compliance
 - Data Types
 - Datetime Pattern
 - Number Pattern
 - Functions
 - Identifiers
 - Literals
 - Null Semantics
 - SQL Syntax
 - Data Definition Statements
 - Data Manipulation

EXPLAIN

Description

The EXPLAIN statement is used to provide logical/physical plans for an input statement. By default, this clause provides information about a physical plan only.

Syntax

```
EXPLAIN [ EXTENDED | CODEGEN | COST | FORMATTED ] statement
```

Parameters

- **EXTENDED**

Generates parsed logical plan, analyzed logical plan, optimized logical plan and physical plan. Parsed Logical plan is a unresolved plan that extracted from the query. Analyzed logical plans transforms which translates unresolvedAttribute and unresolvedRelation into fully typed objects. The optimized logical plan transforms through a set of optimization rules, resulting in the physical plan.
- **CODEGEN**

Generates code for the statement, if any and a physical plan.
- **COST**

Exercise 4: dataset.explain() API

```
val ds = spark.sql("select count(*) FROM db1.address_by_dept WHERE dept=92")  
ds.explain
```

```
ds.explain(false)
```

```
ds.explain(true)
```

Objectives of Hands-On



1/ Execution Plan, SQL Explain / dataset.explain

2/ Spark UI, DAG, Narrow/Wide Transformations

3/ RDD cache / checkpoint

4/ Partition Pruning / Columns Pruning / Predicate-Push-Down

5/ Parquet Optims (Sort, Stats, Dictionary, Bloom, BlockSize)

6/ perfs: SpillToDisk, JOINS, hint

Exercise 5 : Open Spark-UI

browse to SQL -> last Query -> Detailed Plan

a/ Open Spark-UI at **http://localhost:4040/**

b/ go in last Tab « SQL / DataFrame »

c/ click on line for SQL Query

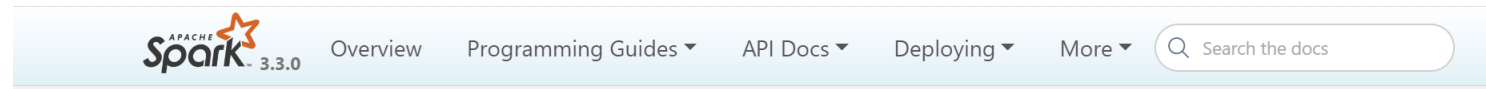
read the Execution Plan as Blue rectangles and Arrows(DAG)

d/ search carefully the PartitionFilters: [..]

it appears twice: as mouse-over tooltip, and in Detailed text

Exercise 5 : Spark UI Documentation...

<https://spark.apache.org/docs/latest/web-ui.html>



Web UI

Apache Spark provides a suite of web user interfaces (UIs) that you can use to monitor the status and resource consumption of your Spark cluster.

Table of Contents

- [Jobs Tab](#)
 - [Jobs detail](#)
- [Stages Tab](#)
 - [Stage detail](#)
- [Storage Tab](#)
- [Environment Tab](#)
- [Executors Tab](#)
- [SQL Tab](#)
 - [SQL metrics](#)
- [Structured Streaming Tab](#)
- [Streaming \(DStreams\) Tab](#)
- [JDBC/ODBC Server Tab](#)

Jobs Tab

The Jobs tab displays a summary page of all jobs in the Spark application and a details page for each job. The summary page shows high-level information, such as the status, duration, and progress of all jobs and the overall event timeline. When you click on a job on the summary page, you see the details page for that job. The details page further shows the event timeline, DAG visualization, and all stages of the job.

Exercise 6 : dataset parallelism (N partitions)
=> show N tasks in Job->Stage->Tasks

a/ browse from SQL logical plan to execution plan: Jobs->Stages

b/ take SparkUI screenshot showing detailed table results
(time, IO statistics) PER partition.

c/ Check table results has N rows,
where $N = ds.toJavaRDD.getNumPartitions$
Is the dataset well balanced (or skewed) ?

d/ take screenshot of (min-average-max) partition results in DAG

Objectives of Hands-On



1/ Execution Plan, SQL Explain / dataset.explain



2/ Spark UI, DAG,
Narrow/Wide Transformations

3/ RDD cache / checkpoint

4/ Partition Pruning / Columns Pruning / Predicate-Push-Down

5/ Parquet Optims (Sort, Stats, Dictionary, Bloom, BlockSize)

6/ perfs: SpillToDisk, JOINS, hint

Reminder: Narrow/Wide Transformations

<https://spark.apache.org/docs/latest/rdd-programming-guide.html#transformations>

Transformations

The following table lists some of the common transformations supported by Spark. Refer to the RDD API doc ([Scala](#), [Java](#), [Python](#), [R](#)) and pair RDD functions doc ([Scala](#), [Java](#)) for details.

Transformation	Meaning
map (<i>func</i>)	Return a new distributed dataset formed by passing each element of the source through a function <i>func</i> .
filter (<i>func</i>)	Return a new dataset formed by selecting those elements of the source on which <i>func</i> returns true.
flatMap (<i>func</i>)	Similar to map, but each input item can be mapped to 0 or more output items (so <i>func</i> should return a Seq rather than a single item).
mapPartitions (<i>func</i>)	Similar to map, but runs separately on each partition (block) of the RDD, so <i>func</i> must be of type <code>Iterator<T> => Iterator<U></code> when running on an RDD of type T.
mapPartitionsWithIndex (<i>func</i>)	Similar to mapPartitions, but also provides <i>func</i> with an integer value representing the index of the partition, so <i>func</i> must be of type <code>(Int, Iterator<T>) => Iterator<U></code> when running on an RDD of type T.
sample (<i>withReplacement</i> , <i>fraction</i> , <i>seed</i>)	Sample a fraction <i>fraction</i> of the data, with or without replacement, using a given random number generator seed.
union (<i>otherDataset</i>)	Return a new dataset that contains the union of the elements in the source dataset and the argument.
intersection (<i>otherDataset</i>)	Return a new RDD that contains the intersection of elements in the source dataset and the argument.
distinct ([<i>numPartitions</i>])	Return a new dataset that contains the distinct elements of the source dataset.
groupByKey ([<i>numPartitions</i>])	When called on a dataset of (K, V) pairs, returns a dataset of (K, Iterable<V>) pairs. Note: If you are grouping in order to perform an aggregation (such as a sum or average) over each key, using <code>reduceByKey</code> or <code>aggregateByKey</code> will yield much better performance. Note: By default, the level of parallelism in the output depends on the number of partitions of the parent RDD. You can pass an optional <code>numPartitions</code> argument to set a different number of tasks.

reduceByKey (<i>func</i> , [<i>numPartitions</i>])	When called on a dataset of (K, V) pairs, returns a dataset of (K, V) pairs where the values for each key are aggregated using the given reduce function <i>func</i> , which must be of type <code>(V,V) => V</code> . Like in <code>groupByKey</code> , the number of reduce tasks is configurable through an optional second argument.
aggregateByKey (<i>zeroValue</i>)(<i>seqOp</i> , <i>combOp</i> , [<i>numPartitions</i>])	When called on a dataset of (K, V) pairs, returns a dataset of (K, U) pairs where the values for each key are aggregated using the given combine functions and a neutral "zero" value. Allows an aggregated value type that is different than the input value type, while avoiding unnecessary allocations. Like in <code>groupByKey</code> , the number of reduce tasks is configurable through an optional second argument.
sortByKey ([<i>ascending</i>], [<i>numPartitions</i>])	When called on a dataset of (K, V) pairs where K implements Ordered, returns a dataset of (K, V) pairs sorted by keys in ascending or descending order, as specified in the boolean <i>ascending</i> argument.
join (<i>otherDataset</i> , [<i>numPartitions</i>])	When called on datasets of type (K, V) and (K, W), returns a dataset of (K, (V, W)) pairs with all pairs of elements for each key. Outer joins are supported through <code>leftOuterJoin</code> , <code>rightOuterJoin</code> , and <code>fullOuterJoin</code> .
cogroup (<i>otherDataset</i> , [<i>numPartitions</i>])	When called on datasets of type (K, V) and (K, W), returns a dataset of (K, (Iterable<V>, Iterable<W>)) tuples. This operation is also called <code>groupWith</code> .
cartesian (<i>otherDataset</i>)	When called on datasets of types T and U, returns a dataset of (T, U) pairs (all pairs of elements).
pipe (<i>command</i> , [<i>envVars</i>])	Pipe each partition of the RDD through a shell command, e.g. a Perl or bash script. RDD elements are written to the process's stdin and lines output to its stdout are returned as an RDD of strings.
coalesce (<i>numPartitions</i>)	Decrease the number of partitions in the RDD to <code>numPartitions</code> . Useful for running operations more efficiently after filtering down a large dataset.
repartition (<i>numPartitions</i>)	Reshuffle the data in the RDD randomly to create either more or fewer partitions and balance it across them. This always shuffles all data over the network.
repartitionAndSortWithinPartitions (<i>partitioner</i>)	Repartition the RDD according to the given partitioner and, within each resulting partition, sort records by their keys. This is more efficient than calling <code>repartition</code> and then sorting within each partition because it can push the sorting down into the shuffle machinery.

Exercise 7: combine several **Narrow** transformations

filter, sample, select, withColumn ...

a/ Execute query like

```
val addressDs = .....
```

```
.filter(a).filter(b) // => check Spark combine as .filter(a && b) !!
```

b/ ..like

```
addressDs
```

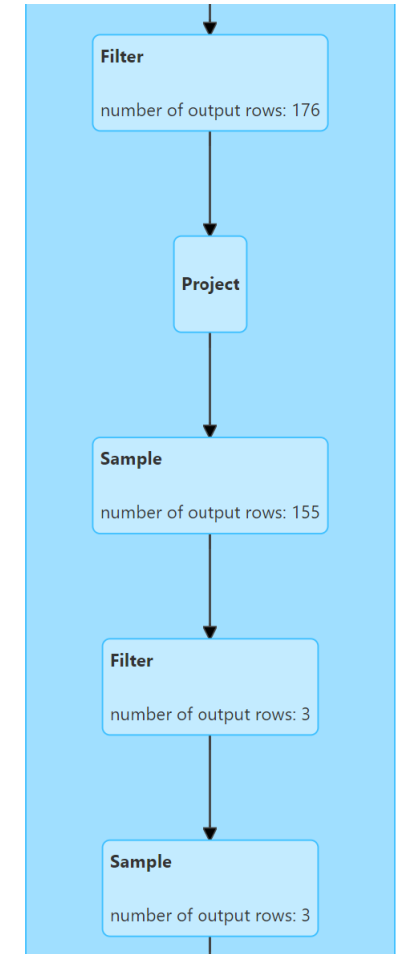
```
.withColumn(..)
```

```
.filter( .. ) .sample(0.9)
```

```
.filter( .. ) .sample (0.9)
```

```
.count
```

b/ Check in Spark UI that there is only 1 « **WholeStageCodegen (1)** » containing several instructions



Exercise 8: WholeStageCodegen java code for(;;) on Dataset<Row>...

a/ Show the corresponding generated Java Code of WholeStageCodegen

use **ds.queryExecution.debug.codegen**

b/ read it ...

find « .. extends org.apache.spark.sql.execution.BufferedRowIterator »

<https://github.com/apache/spark/blob/master/sql/core/src/main/java/org/apache/spark/sql/execution/BufferedRowIterator.java#L97>

c/ find method @Override **protected void processNext()** {

.. It is supposed to map copy all filtered rows x columns to new RDD[Row]

d/ do you find your filter conditions on your columns ?

Exercise 9: cascade several Wide Transformations: repartition, groupByKey, aggregateByKey, aggregate, join, distinct

a/ Execute query like

addressDs

.repartition(3)

.repartition(2, \$"commune_insee")

.repartition(4)

.repartition(2, \$"commune_nom")

.count

b/ Check in Spark UI that there are N stages (=N shuffles)

Exercise 10: Check on 2 consecutive Stages that previous Shuffle Write = next Shuffle Read

▼ Completed Stages (4)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Stage Id ▼	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
9	show at <console>:23 +details	2022/10/06 17:17:44	41 ms	<div>1/1</div>				
5	show at <console>:23 +details	2022/10/06 17:17:43	0,6 s	<div>2/2</div>			78.8 MiB	36.2 MiB
2	show at <console>:23 +details	2022/10/06 17:17:41	2 s	<div>1/1</div>			39.6 MiB	78.8 MiB
0	show at <console>:23 +details	2022/10/06 17:17:39	2 s	<div>8/8</div>	27.2 MiB			39.6 MiB

Exercise 11: Why do you see Greyed boxes... « Skipped Stages » ?

- a/ If you don't see... check you executed a wide transformation (not narrow)
- b/ why all except the last Stage are greyed, and the last is blue?
- c/ what is the difference «with green » point like in `.cache()` ?
(cf next Exercise on `dataset.cache()` and `.checkpoint()`)

Exercise 12: mix cascade of Narrow and Wide

Example

```
spark.sql("select * from address3")  
  .repartition(1).sample(0.9)  
  .repartition(2).sample(0.8).filter("commune_nom like '%a%'")  
  .repartition(3).sample(0.7).filter("commune_nom like '%t%'").sample(0.6)  
  .count
```

a/ Explain how many Shuffles you expect

b/ How many instructions you expect in each

WholeStageCodegen (1), WholeStageCodegen (2), WholeStageCodegen (3)

Objectives of Hands-On



1/ Execution Plan, SQL Explain / dataset.explain



2/ Spark UI, DAG, Narrow/Wide Transformations



3/ RDD cache / checkpoint

4/ Partition Pruning / Columns Pruning / Predicate-Push-Down

5/ Parquet Optims (Sort, Stats, Dictionary, Bloom, BlockSize)

6/ perfs: SpillToDisk, JOINS, hint

Exercise 13: Several Actions on same Dataset

a/ Execute several actions on same Dataset

```
ds.show
```

```
ds.show
```

b/ equivalently... several SQL

```
select * from address
```

```
select * from address
```

b/ Check in SparkUI that ds is recomputed each time
(show screenshots with job ids, stage ids)

c/ Check Spark File IO Statistics, Shuffle, Time elapsed

Exercise 14: Avoiding re-computation

dataset .cache()

a/ Same as Exercise 12... but use before

ds.cache() // or equivalent: .persist()

Execute several actions on cached Dataset

ds.show

ds.show

b/ Check in SparkUI that ds is NOT recomputed

... But Full lineage is still displayed several time in greyed / with **green point**

give screenshots showing job ids / stage ids

c/ check in SparkUI > Storage > RDD

d/ and after ... **ds.unpersist()**

Exercise 15: Execute complex DAG on dataset with lot of duplicates

Example: dataset repeated, with small variations, then union ...

```
val ds1 = spark.sql("select * from address3").repartition(2).filter("commune_nom like '%n%'").repartition(2).sample(0.9);
val ds2 = ds1.union(ds1).limit(1000).repartition(2).sample(0.9);
val ds3 = ds2.union(ds2).limit(1000).repartition(2).sample(0.9).union(ds2);
val ds4 = ds3.union(ds3).limit(1000).repartition(2).sample(0.9).union(ds3);
val ds5 = ds4.union(ds4).limit(1000).repartition(2).sample(0.9).union(ds3);
ds5.count;
```

Open corresponding DAG in Spark UI

Exercise 16: How to simplify DAG display ?checkpoint() !

a/ ensure

```
sc.setCheckpointDir(« c:/data/checkpoint-dir »)
```

b/ Same as Exercise 15... but use checkpoint:

```
ds3=..
```

```
val ds3 = ds3.checkpoint(); // IMPORTANT TO RE-ASSIGN
```

```
ds4=..
```

c/ Check in SparkUI that ds3 is persisted,
... AND lineage no more displayed

d/ there is NO « unCheckpoint() » as there is for unpersist()

10 mn pause

Objectives of Hands-On



1/ Execution Plan, SQL Explain / dataset.explain



2/ Spark UI, DAG, Narrow/Wide Transformations



3/ RDD cache / checkpoint



4/ Partition Pruning / Columns Pruning / Predicate-Push-Down

5/ Parquet Optims (Sort, Stats, Dictionary, Bloom, BlockSize)

6/ perfs: SpillToDisk, JOINS, hint

Exercise 17: Reminder Partition Pruning

There are 3 « Pruning » Optimizations done by Spark

1/ Column Pruning

2/ Partition Pruning

3/ Predicate-Push-Down

Questions: (interrogation on Spark Lesson)

a/ explain in few words what is 1/, 2/, 3/

b/ what are pros/cons of many (small) partitions ?

Exercise 18: Column Pruning

Reminding that Parquet is a « Columnar File Format »

Check by reading only 1 column, that Spark does not read fully Parquet Files
... only Page block of selected column.

This is « Column Pruning »

Example:

```
spark.sql(« select distinct commune_nom from address »).count
```

Check in SparkUI the File Bytes read statistics.

Compare with total file size

Check in SparkUI the IO read statistics per partition tasks

Give screenshots of SparkUI

Exercise 19: Predicate-Push-Down

a/ Execute following Queries

```
select count(*) from db1.address where commune_nom = 'Nanterre'
```

```
select count(*) from db1.address where UPPER(commune_nom) = 'NANTERRE'
```

```
select count(*) from db1.address where commune_nom like '%Nanterre%'
```

b/ Check in SparkUI the execution Plan

Search for « PushedFilters:[..] »

c/ For which query Spark is able to push down « some/all conditions » to parquet library ?

d/ explain how parquet files should be saved for having good read performances

Hint: use `sc.setCallSite(« query comment »)` to find easily your query in Spark UI

Objectives of Hands-On



1/ Execution Plan, SQL Explain / dataset.explain



2/ Spark UI, DAG, Narrow/Wide Transformations



3/ RDD cache / checkpoint



4/ Partition Pruning / Columns Pruning / Predicate-Push-Down



5/ Parquet Optims (Sort, Stats, Dictionary, Bloom, BlockSize)

6/ perfs: SpillToDisk, JOINS, hint

Exercise 20: prepare sorted parquet table for Predicate-Push-Down

We want to have few PARQUET files (1 or several, but not hundred)

Each PARQUET File split in several blocks of 16Mo (default=256Mo)

Each block sorted to have efficient Dictionaries and Min-Max Statistics

```
val allAdressCsvDs = spark.read.options(Map("header" -> "true", "delimiter" -> ";",  
"inferSchema" -> "true")).format("csv").load("C:/data/OpenData-gouv.fr/bal/adresses-  
france.csv")  
.withColumn("dept", regexp_replace(col("commune_insee"), "0*(.*)...", "$1").cast("int"))  
.withColumn("code", col("commune_insee").cast("int"))
```

```
allAdressCsvDs.repartition(1)  
  .sort("dept").sortWithinPartitions("dept", "code", "voie_nom")  
  .write.format("parquet").option("parquet.block.size", 16*1024*1024)  
  .saveAsTable("db1.address_sorted")
```

```
allAdressCsvDs.repartition(10, col("dept"))  
  .sortWithinPartitions("dept", "code", "voie_nom")  
  .write.format("parquet").option("parquet.block.size", 16*1024*1024)  
  .saveAsTable("db1.address_sorted10")
```

Exercise 21: check.. getNumPartitions

Check that

- a/ table db1.address_sorted is saved
on 1 parquet file,
and file is split into 55 partitions
giving a total of numPartition=55
- b/ table db1.address_sorted10 is saved
on 10 parquet files
and files are split (differently..)
giving a total of numPartition=57
- c/ is there a big difference between a/ and b/ ?
Which is « better » ?

Hint: use **dataset.toJavaRDD.getNumPartitions**

Exercise 22: Predicate-Push-Down

Execute several Queries with WHERE clauses « field=value »
and see efficiency of Skipped/Read bytes compared to total file size

select count(*) from db1.address_sorted where ...

a/ where commune_nom='Nanterre'

b/ where code=92050

c/ where commune_nom='_UneCommuneQuiNExistePas'

d/ where code=0

Compare Bytes read for queries... Are they same?

Compare Plan « PushedFilters: [..] » / Compare jobs->stage tasks counts

Hint: use « **sc.setCallSite**(« DisplayName »); » to find queries more easily in Spark UI

Objectives of Hands-On



1/ Execution Plan, SQL Explain / dataset.explain



2/ Spark UI, DAG, Narrow/Wide Transformations



3/ RDD cache / checkpoint



4/ Partition Pruning / Columns Pruning / Predicate-Push-Down



5/ Parquet Optims (Sort, Stats, Dictionary, Bloom, BlockSize)



6/ perfs: SpillToDisk, JOINS, hint

Exercise 23: Detect SpillToDisk problems

When Spark has not enough RAM memory, it uses « swapping » to disk.
In Spark, this is called « SpillToDisk ».

This is a dreadful signal, to ask for more memory (or better algorithm)

a/ execute a task taking too much memory

... example a « `.repartition(2).sortWithinPartitions(« commune_insee »)` »

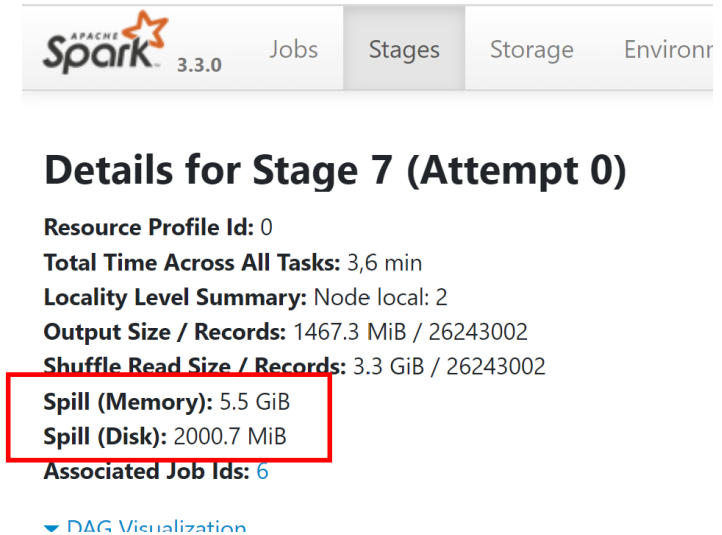
b/ detect the slowness and diagnostic it in SparkUI

c/ retry executing same request with more memory, and compare performances

Exercise 23: SpillToDisk

Example:

```
spark.sql("select * from db1.addr3")  
  .repartition(2).sortWithinPartitions("commune_insee")  
  .write.saveAsTable("db1.addr_repart2")
```



The screenshot shows the Apache Spark UI interface. At the top, there's a navigation bar with tabs for 'Jobs', 'Stages', 'Storage', and 'Environr'. The 'Stages' tab is selected. Below the navigation bar, the title 'Details for Stage 7 (Attempt 0)' is displayed. The main content area shows various metrics for this stage:

- Resource Profile Id: 0
- Total Time Across All Tasks: 3,6 min
- Locality Level Summary: Node local: 2
- Output Size / Records: 1467.3 MiB / 26243002
- Shuffle Read Size / Records: 3.3 GiB / 26243002
- Spill (Memory): 5.5 GiB**
- Spill (Disk): 2000.7 MiB**
- Associated Job Ids: 6

At the bottom, there is a link for 'DAG Visualization'.

Find other (more detailed) infos in SparkUI

Exercise 24: Join

- a/ Extract and save from table address a new table « city »
Containing « name, code, dept, average_address_longitude, average_address_latitude »

- b/ Join table « address » and « city »,
and compute for each address the offset longitude/latitude to the city average center

- c/ study Execution plan
check that Spark is already efficiently « Broadcasting » the small 'city' table
=> no need to « SortMerge » the big 'addr' table

Exercise 25: HINT to force join strategy

compare with previous Exercise 24/, but with SQL hinting

SELECT /*+ MERGE(a) */ ... FROM addr a ...

Cf <https://spark.apache.org/docs/latest/sql-ref-syntax-qry-select-hints.html>

Examples

```
-- Join Hints for broadcast join
SELECT /*+ BROADCAST(t1) */ * FROM t1 INNER JOIN t2 ON t1.key = t2.key;
SELECT /*+ BROADCASTJOIN (t1) */ * FROM t1 left JOIN t2 ON t1.key = t2.key;
SELECT /*+ MAPJOIN(t2) */ * FROM t1 right JOIN t2 ON t1.key = t2.key;

-- Join Hints for shuffle sort merge join
SELECT /*+ SHUFFLE_MERGE(t1) */ * FROM t1 INNER JOIN t2 ON t1.key = t2.key;
SELECT /*+ MERGEJOIN(t2) */ * FROM t1 INNER JOIN t2 ON t1.key = t2.key;
SELECT /*+ MERGE(t1) */ * FROM t1 INNER JOIN t2 ON t1.key = t2.key;

-- Join Hints for shuffle hash join
SELECT /*+ SHUFFLE_HASH(t1) */ * FROM t1 INNER JOIN t2 ON t1.key = t2.key;

-- Join Hints for shuffle-and-replicate nested loop join
SELECT /*+ SHUFFLE_REPLICATE_NL(t1) */ * FROM t1 INNER JOIN t2 ON t1.key = t2.key;
```

Optional Exercise 26: Default Shuffle = 200 ?!

Study number of partitions after a (« merge ») join (not a « broadcastjoin »)

a/ is it 200 ? Why

b/ How do you change default ?

c/ What should you always do before saving to File(s)?
(to avoid 200 small files)

Exercise 27 : MindMap

Draw a MindMap to summarize
what you did and learn from this Hands-On session

Your MindMap should
start with word « Spark – Optimizations & RDD DAGs» in the middle
Then draw star edges to other word chapters and sub-chapters

Objectives of Hands-On



1/ Execution Plan, SQL Explain / dataset.explain



2/ Spark UI, DAG, Narrow/Wide Transformations



3/ RDD cache / checkpoint



4/ Partition Pruning / Columns Pruning / Predicate-Push-Down



5/ Parquet Optims (Sort, Stats, Dictionary, Bloom, BlockSize)



6/ perfs: SpillToDisk, JOINS, hint

Questions ?

Take-Away

What You learned ?

Next Steps

Unfortunately,
NO More Lessons
NO More Hands-On

But still, many more Spark concepts to learn :

- Spark Clustering
- Spark on Kubernetes, in Cloud
- Java binding, UDF, map
- Analytical Queries
- Machine Learning
- Spark Streaming
- ...