Introduction to BigData – Spark – Processing

Concrete Examples Experience at SG

Arnaud Nauwynck

Oct 2022

Objectives

4 x 2 hours of lessons
4 x 2 hours of Hands-On

Objective Lesson 1 : Overview

What is BigData ?
What is Spark ?
What is Data processing ?
… concrete examples from experience in a company

Objective Lesson 2,3,4 : Deep Dive in Spark

Outline

- What is BigData ?
  - Order of Magnitudes for « Big »
  - History
  - Evolution of Softwares, Spark
- Description of a Datalake
  - Content, data feeding
  - How it is organized
  - Who uses it
- Example of Data-Processing
  - RAW to LAKE, Reports
- Change Storage-Compute, Evolution to Cloud

About Me

NOT using Twitter/LinkedIn/FaceBook/

Only arnaud.nauwynck@gmail.com

IT Passionate since 20 years

Working at Société Générale (La Defense)
In BigData team since 4 years

Expert Java + BigData + more

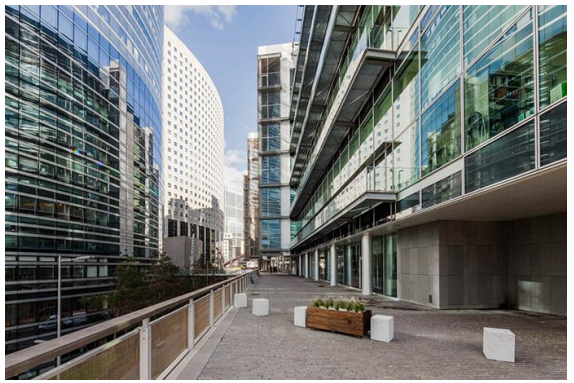

Github repos

https://github.com/Arnaud-Nauwynck

http://arnaud-nauwynck.github.io/

DISCLAIMER

I do not know everything. I am biaised with what I know.
My opinions do not represent any company

About You …
Quick Survey ?

Raise your Hands…

You are developper ?

In Java ? SQL ?

You know Hadoop / Hdfs / Spark ?

You want to be Data Engineer / Data Scientist ?

Outline

- What is BigData ?
  - Order of Magnitudes for « Big »
  - History
- Description of a Datalake
  - Content, data feeding
  - How it is organized
  - Who uses it
- Example of Data-Processing
  - RAW to LAKE, Reports
- Change Storage-Compute, Evolution to Cloud

How Big ? => Then Which Tool

1 .. 1 000

50 000

1 Million (e6)

(**Mega** Bytes)

1 Billion(e9)

(**Tera** Bytes)

>= e9 : BIG Data

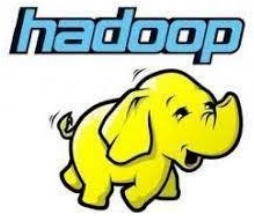

(**Peta** Bytes)

Compute Resources
Vertical Scaling

Historically:

have Bigger needs => buy BIGGER Server

Very Frequent in Bank OLTP Databases

128 cores x 500Go RAM + SAN Bay of Disks

= Millions of Euros

.. BUT price is exponential to performance, and limited

BigData definition

Data does not fit on 1 server … even Huge, because of

3 V = Volume / Velocity / Variety

Bigger Hardware, different Software
Vertical Scaling -> Horyzontal

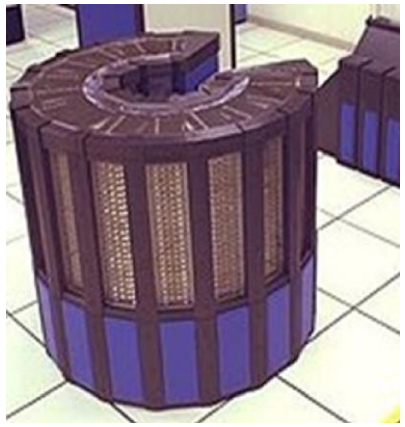

1 OS, 1 process, N Threads

Distributed computing
… Sotfwares need
Network, Fault Tolerant

BigData At Scale
Google Datacenter

Not Only Gafas

Example of Smaller Companies, obviously doing BigData

Criteo (internet advertisment)
Cern (particule accelerators)
Gouvernment
Social Networks
Auchan, Carrefour, .. (E-Commerce)
MeteoFrance (Computing, Satelite Images)
Uber, Taxis, Cars (IOT)
Medical, Pharmaceutic, Genetics
Industries
Banks
..

BigData In Banks ?

Corporate Investment Bank

(for private / owned trading)

Network Bank (for everybody customers)

Example of Société Générale:
At La Defense

Example of Société Générale:
At Val de Fontenay

Data for
- Historical Market datas
- Risk / Var management
- Aggregation of all Trading departments
- Regulatory Reporting
- Business analysis
- ..

Data for
- RGPD… warn personnal data!
- Anti Fraud detection,
- customers-oriented
- Business analysis
- ..

BigData at SG (Investment Bank)

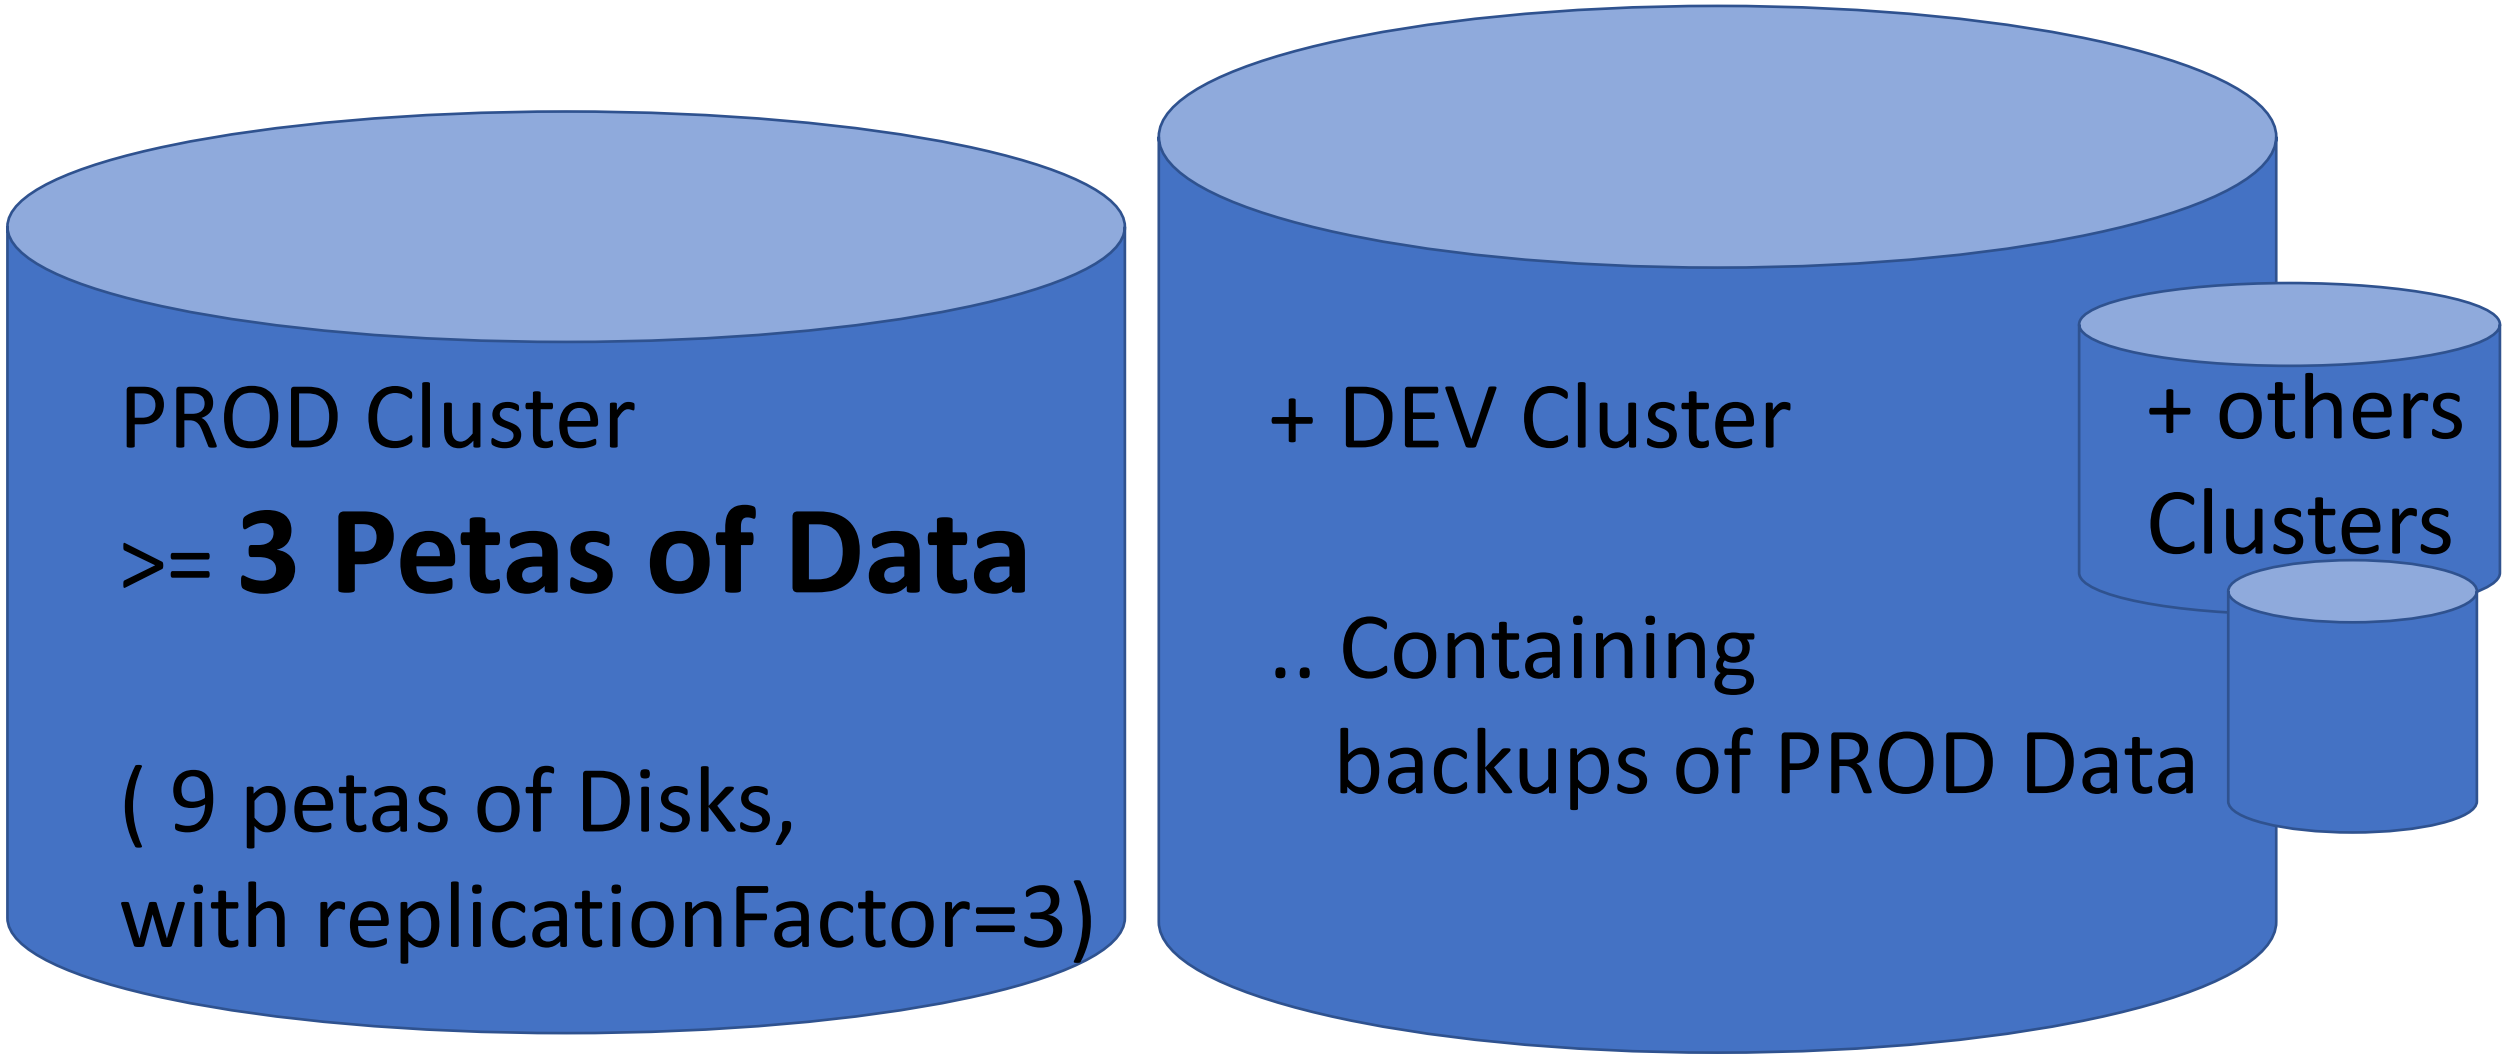

Typical Datacenter Hardware
= N Racks x 10 blades (1U / 2U)

Example of a Server for Datalake

1 server = 42 cores + 256 Go RAM + 8 disks + network 1Gb/s

« Small » Datalake ?

100 servers for PROD
(+ 120 servers for DEV)

= 4200 cores
+ 25600 Go RAM : 25 Tera of RAMS
+ 800 disks : 2 Petas of Data (6 peta of Disks)

… now 3 Petas on Azure

Hadoop Cluster Architecture

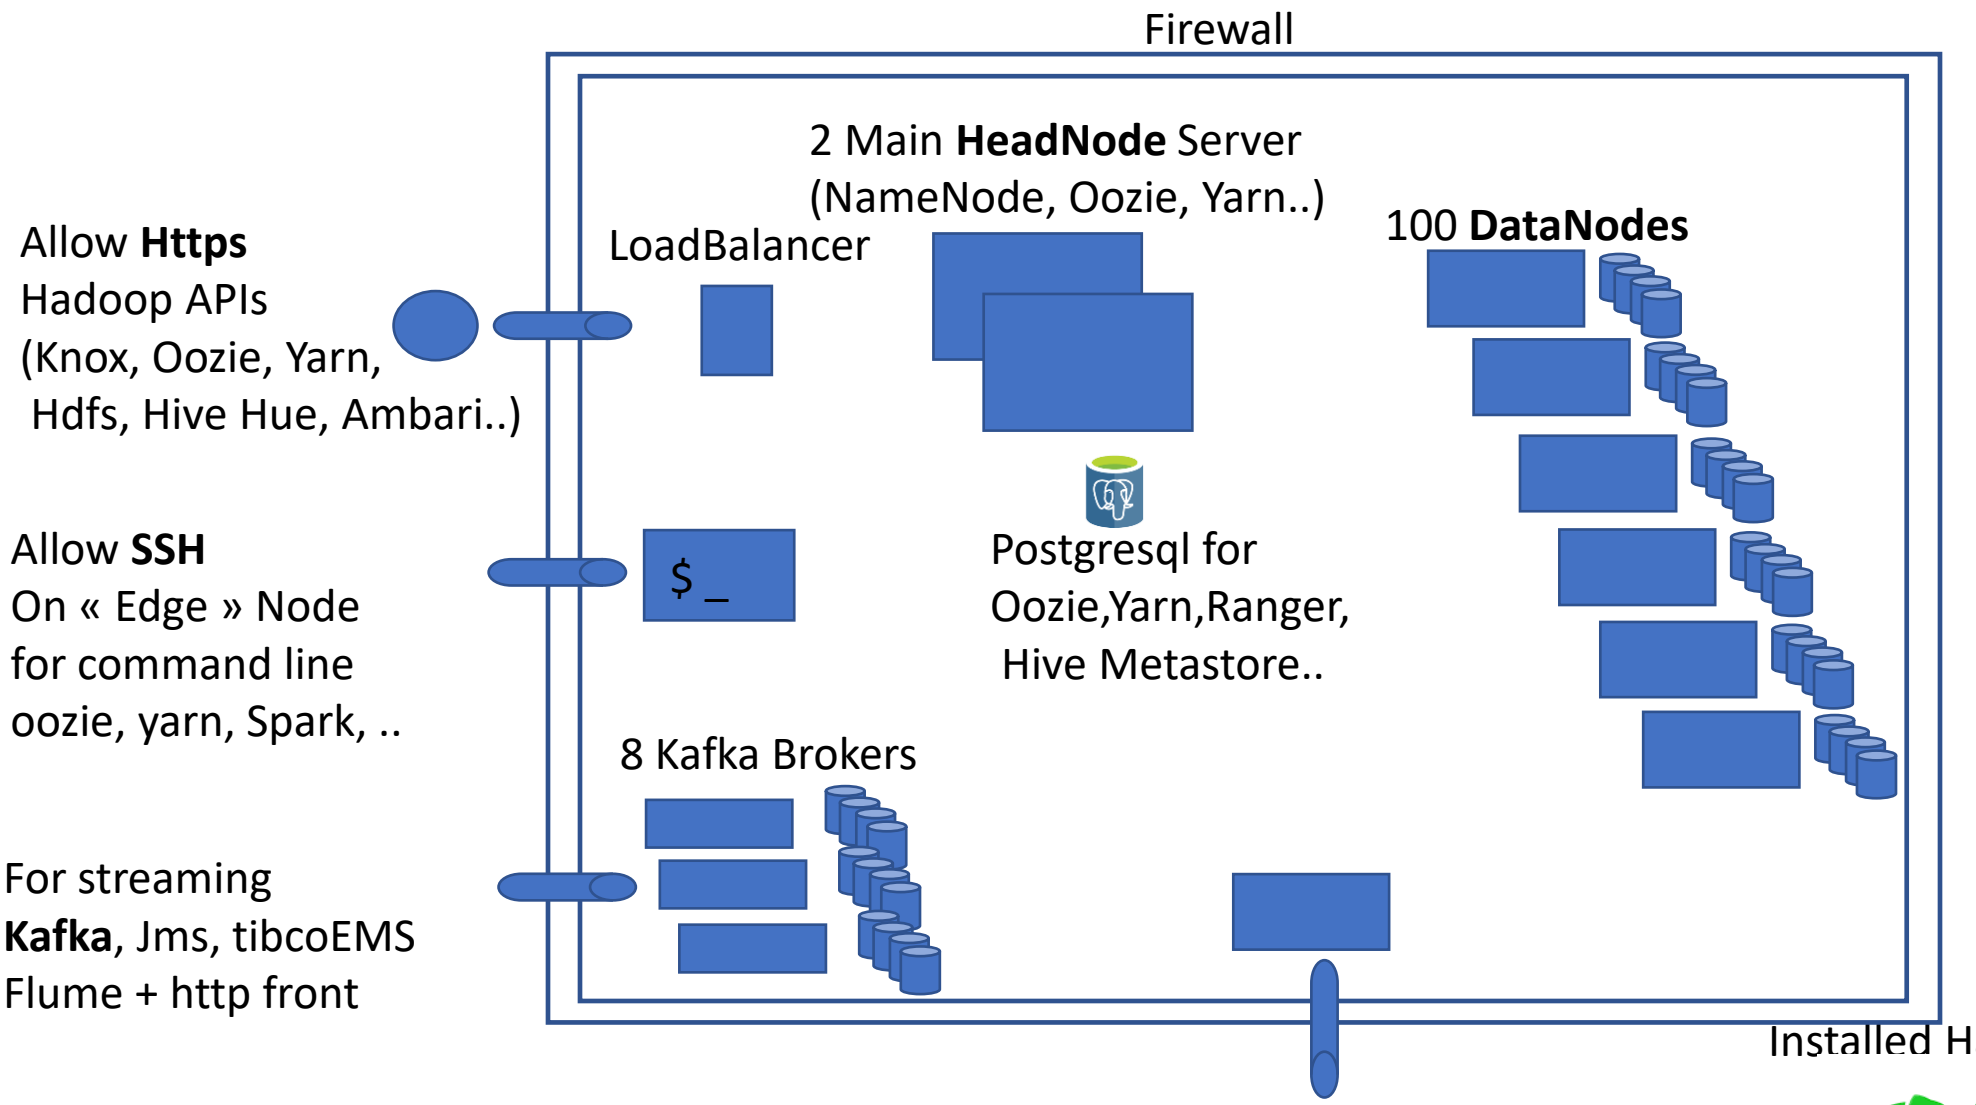

Installed Hadoop Distribution

HORTONWORKS®

How Softwares evolve to handle all this ?

Which ones to use today?

Hadoop Ecosystem « Explosion »

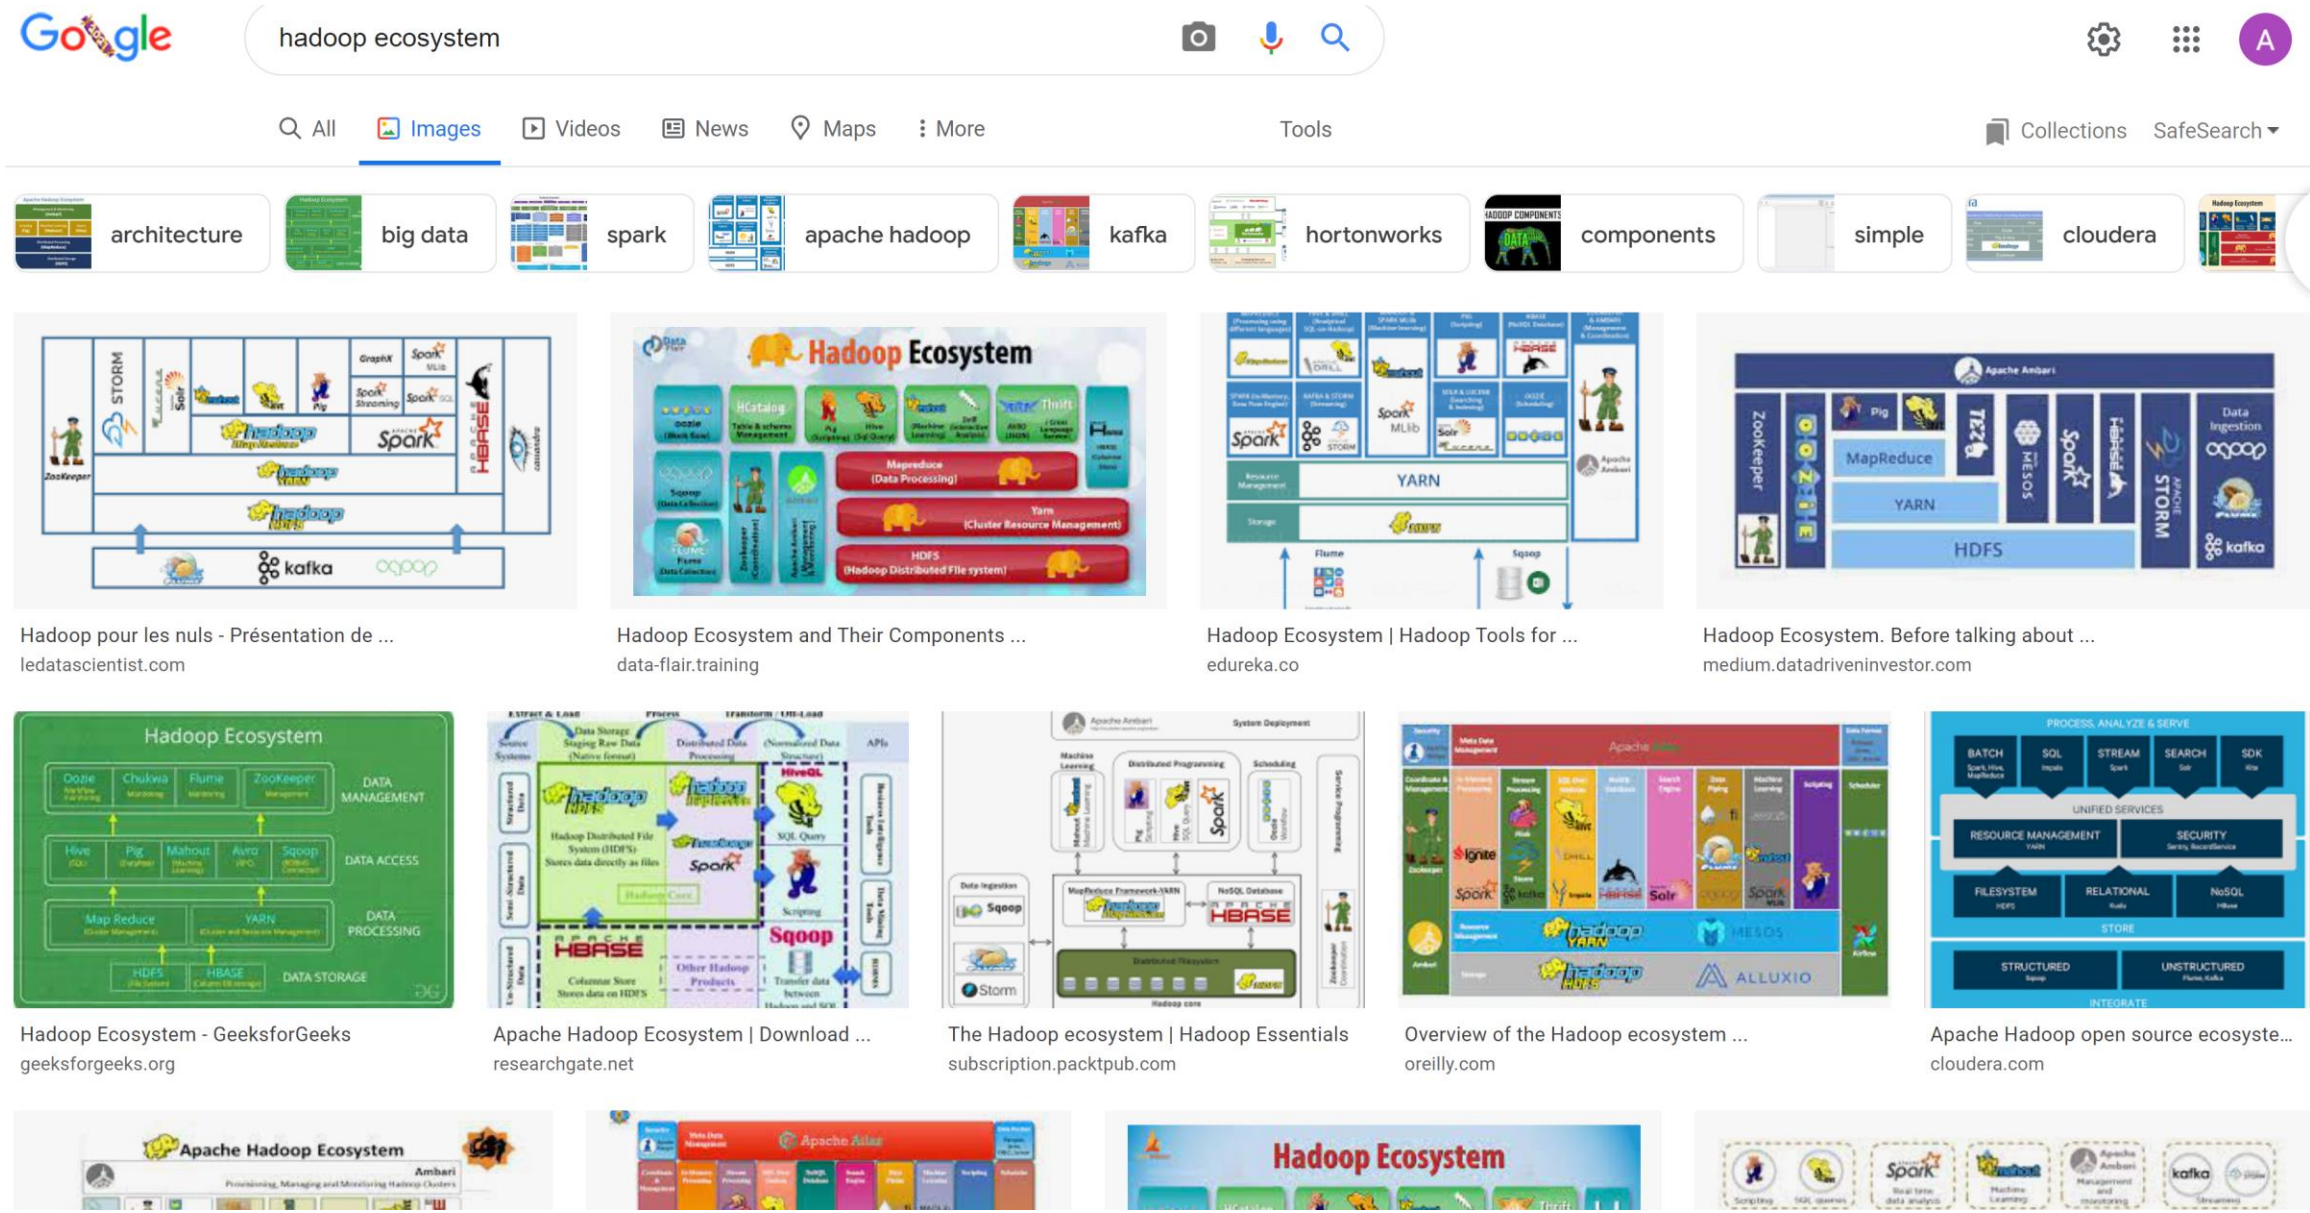

Spark (Recent) History & Ancestors

MapReduce @ Google

2002 @Google

2004 Google Paper published

2014 Google No more used of MapReduce

2006 @Yahoo Hadoop implementation

2008 Apache Open-Source

2012 Yarn (v2)

2021 MapReduce bashing … HDFS & Hadoop also HortonWork bought by Cloudera HDInsight @Azure…very bad choice .. To be abandonned

MPI

1995 Message Passing Interface

2010 Spark paper

2013 Apache top-level

??

2015 Kubernetes

2020 **Spark on K8s**

Datalake History At SG

~8 years ago
2014

~2 years ago (2020)
Decision to migrate to Azure
+ Dismantle On-Prem

Dec 2022
OFF

Init BigData
cluster at SGCIB

2 major PROD projects:
- MiFid reports (all trades)
- Var Risk

> 150 projects

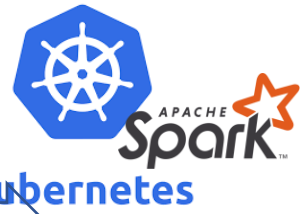

~ 1 year ago (2021)
>= 100 projects migrate to Azure

Internal Re-organizations
Another Cluster..

Point in common
Despite all changes

MapReduce @Yahoo = Hadoop .. 2006

Constraint
=>
Architecture
Choice

Commodity Hardwares (datacenters):
Only HDD + RAM
Data Locality : co-host Storage near Compute
use RAM to cache
avoid network

Think different?

2006

2020

MapReduce -> Spark + other changes

Constraint

=>

Architecture

Choice

Faster Disk (SSD)

Faster Network

Compute != Storage

Cloud

Think different?

2006

2022

Simple => Many Specific Systems => Unified

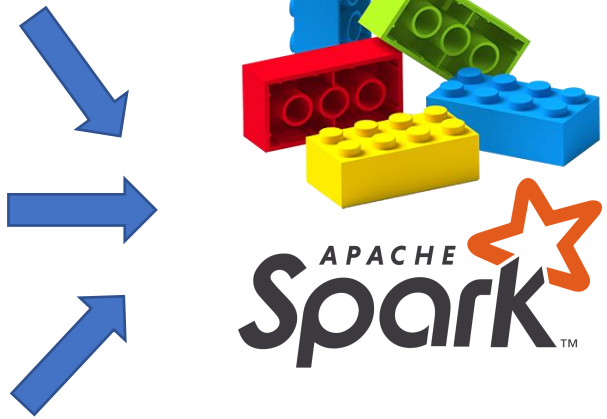

« **Simple** » ecosystem
(verbose inneficient & complex java code)

« **Bazard** » ecosystem
(**Too MANY TOO** SPECIFIC redundant, complexes)

“Unified” ecosystem
Simple
+ extensible modules

At The end, Only 1 will remain (French TV Game: Koh-Lanta)


APACHE
Spark
TM

Spark = « **Unified Engine** »

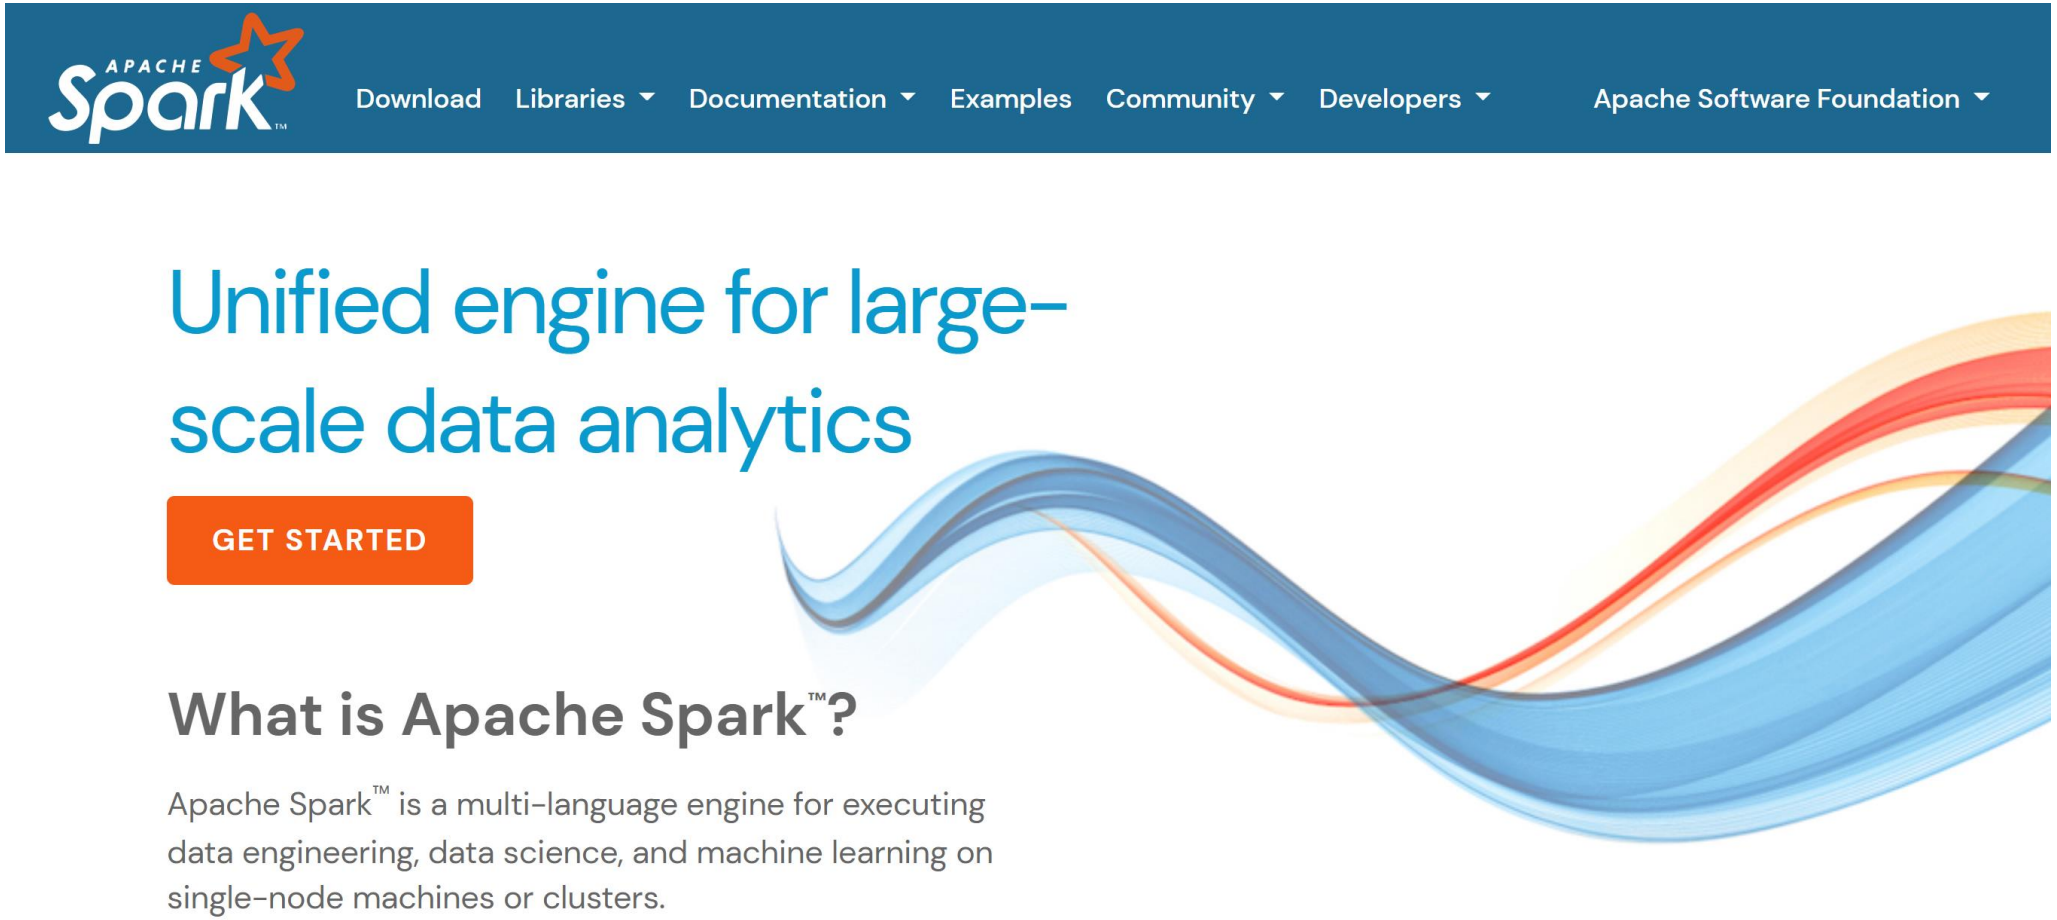

Multi Purposes – Multi Langages

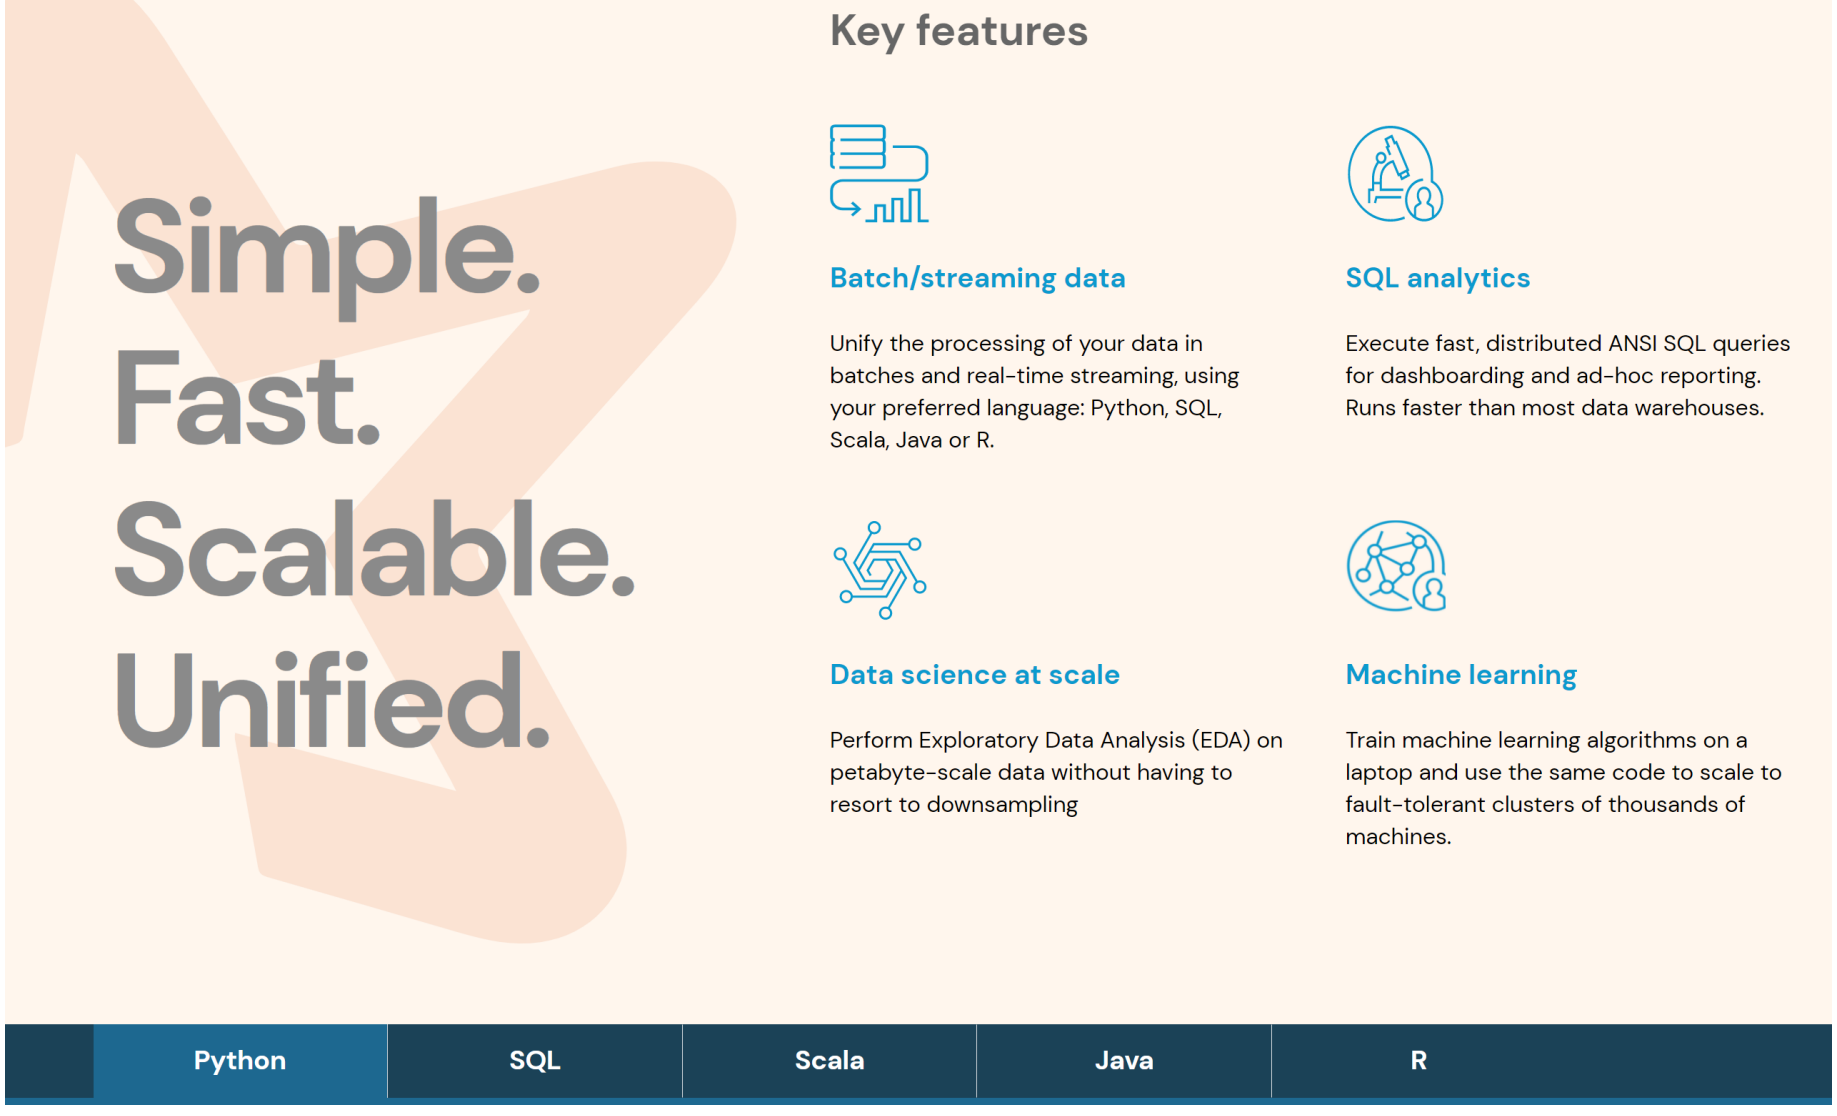

Spark-Core + ...

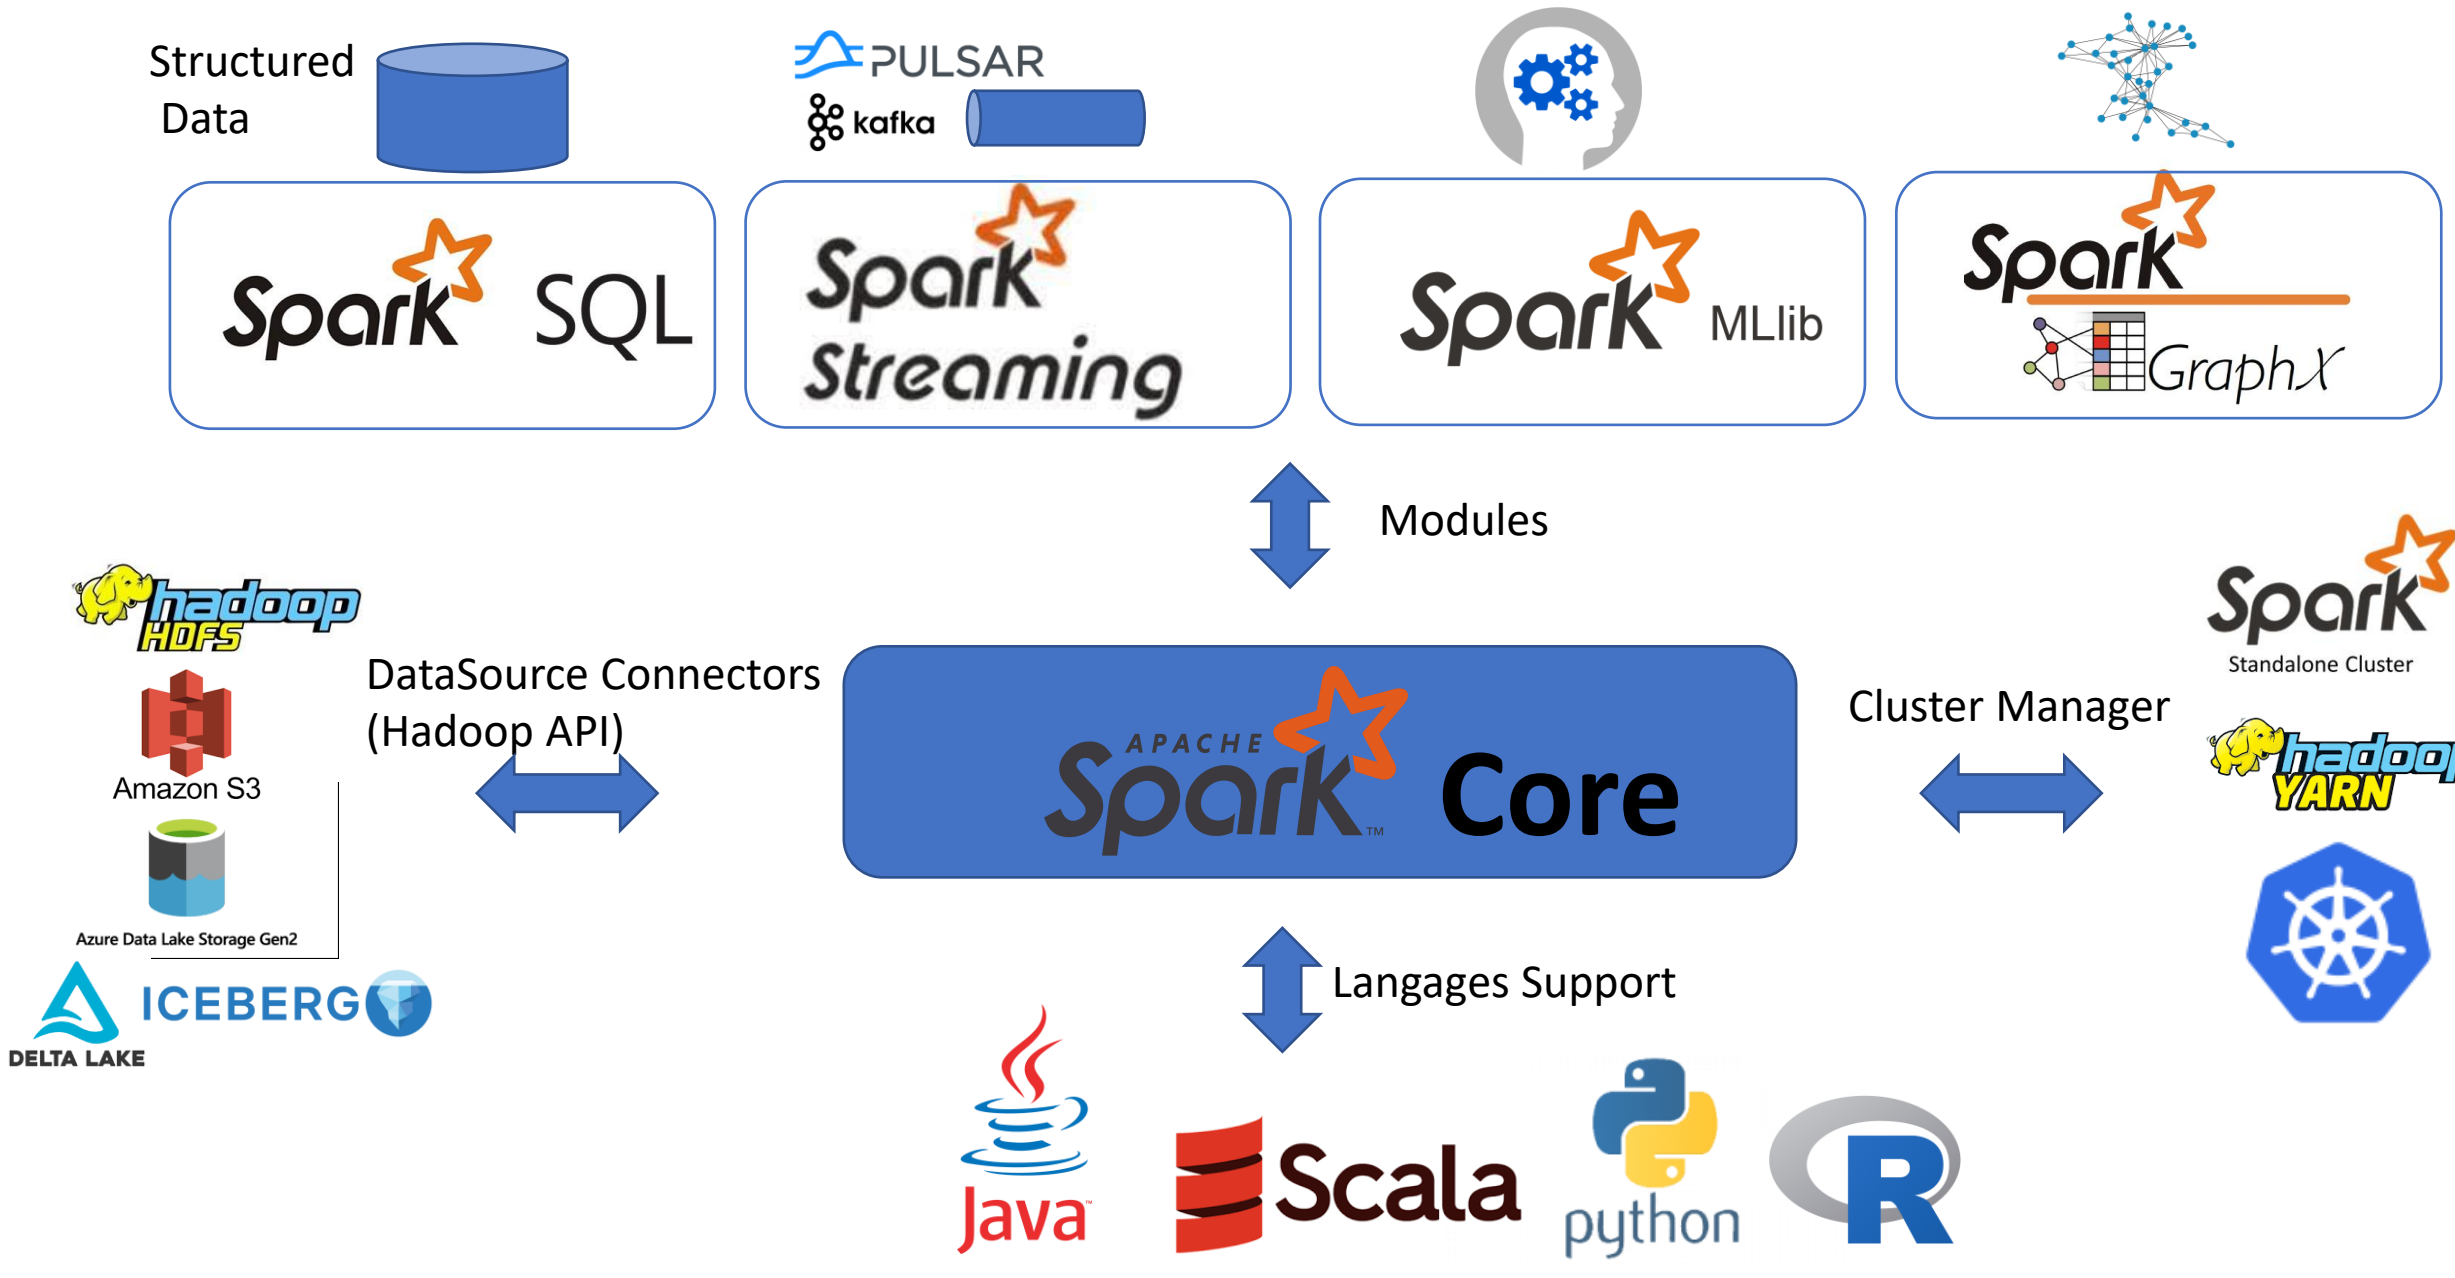
Structured Data
PULSAR
kafka
Spark SQL
Spark Streaming
Spark MLlib
Spark GraphX
Modules
hadoop HDFS
Amazon S3
Azure Data Lake Storage Gen2
ICEBERG
DELTA LAKE
DataSource Connectors (Hadoop API)
APACHE Spark Core
Cluster Manager
Spark Standalone Cluster
hadoop YARN
Langages Support
Java
Scala
python
R

Outline

- What is BigData ?
  - Order of Magnitudes for « Big »
  - History
  - Evolution of Softwares, Spark
- Description of a Datalake
  - Content, data feeding
  - How it is organized
  - Who uses it
- Example of Data-Processing
  - RAW to LAKE, Reports
- Change Storage-Compute, Evolution to Cloud

Not Only Machines

Tower of Babel

מִגְדַּל בָּבֶל

The Tower of Babel by Pieter Bruegel the Elder (1563)

General information	
Type	Tower
Location	Babylon
Height	See § Height

~ 200 different Teams / 10 departments
… for 1000 developpers / data engineers / users

In Paris.. But not Only (Bangalore, Montreal, London, New-York, ..)

On ~6000 tables,
described in 1500 DataSets

Lot of Financial / Functional aspects

Feeding Data Inputs to Datalake: aggregating >= 150 Systems

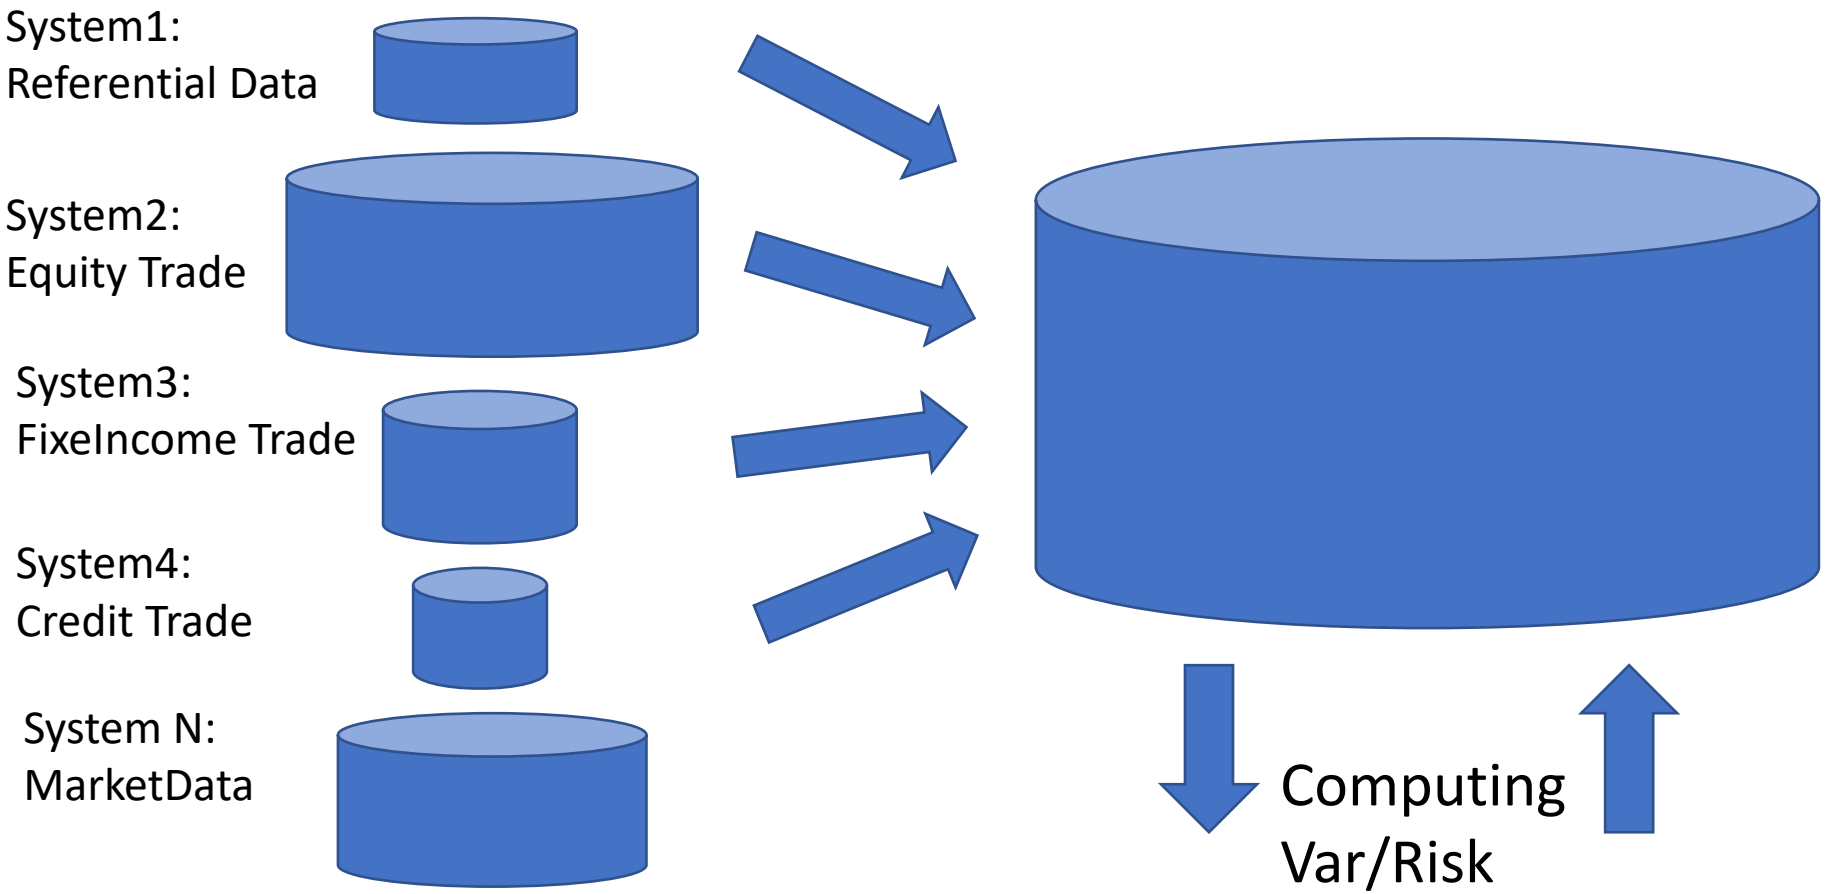

Feeding Data / Streaming

BigData at SG is mostly « uploading daily files »
then launching « daily batches »

Many Trade extractions are Real-Time Streaming
… but aggregated in Files per Hours / per Day

Var Computing during night … Huge streaming of events
from dedicated pricing cluster (GPUs)

Feeding Types & Volumetry

Daily batches: upload Files

Streaming then converted to Files

Teras/day
(millions of files)

Streaming to Key-Values Databases

Billions events/day

Peeks to 4000 events/seconds
on each 100 x servers
During 4hours nightly Var batches

Content of a Datalake

FileSystem
= Directories + Files

For batches

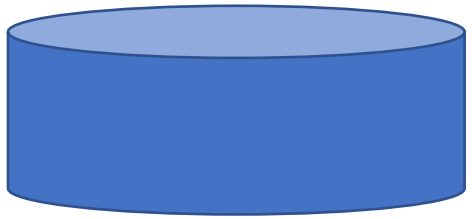

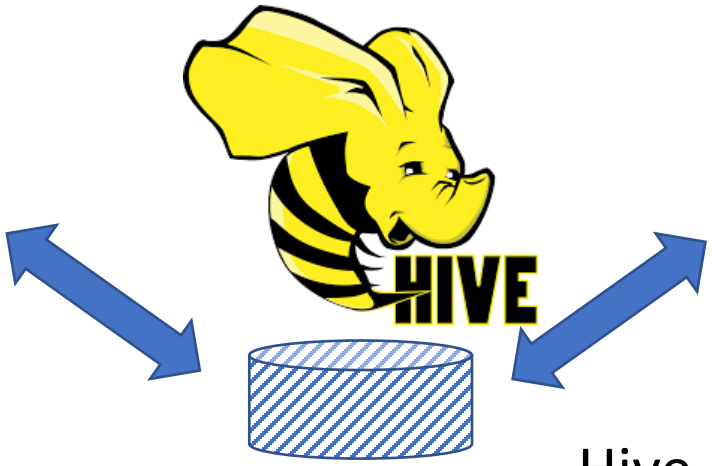

Key-Value database

FAST get/put access by single key

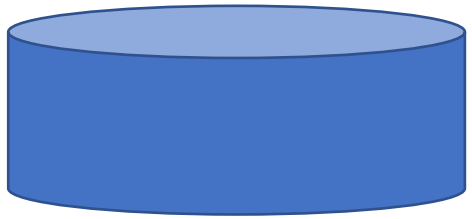

Hive Metastore = SQL View
« Table » abstraction for HDFS / Hbase

Storage = Files and Directories

How would you organize

1 Billions Files

in **50 Millions Directories** ?

How do you ensure data respect directories organization ?

Try Scanning Directories ?

A Very « Simple » Specific Sub-System

« Tick » per seconds
From Market Exchange Places
(Bourse Paris, London Stock Exchange, …)

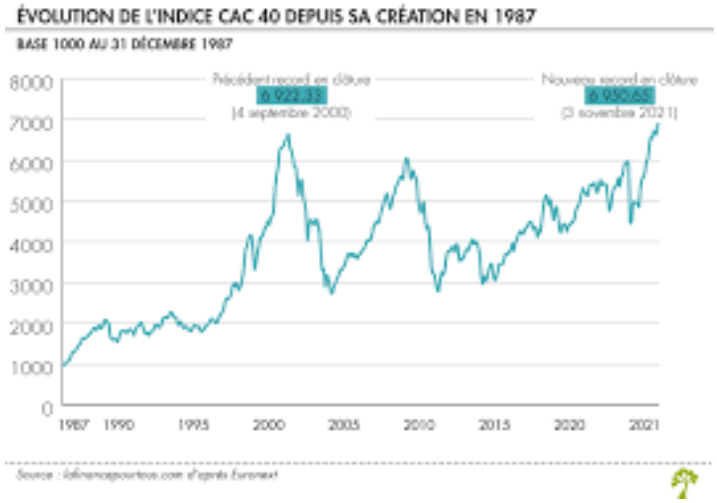

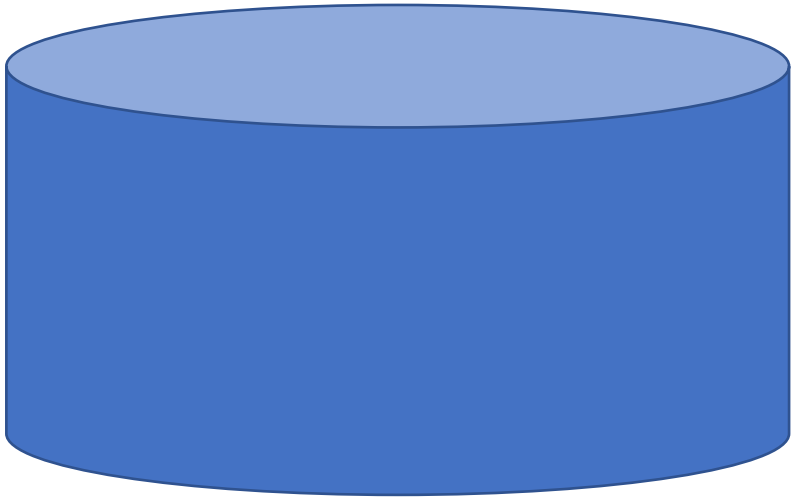

1 binary file per Day/Place/Instrument
… 2M files per day
… 40 Giga per day
… since 22 years … 25 Billions Files / Petas

Volume, Velocity … but not Variety
(NO spark, No Hadoop)

Governance
Data Directories Organization
Read-Write Permissions, Quota, RDPG …

How would you organize Read-Write Permissions on
6000 tables x 1000 jobs ?

Project ask access
AND describe « Datasets »
(data owner)
=> Validation by DataOffice
=> Create domain sub-dirs
+ granted recursive
WRITE permission

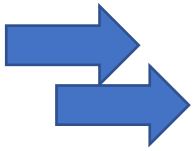

/ <root HDFS>
/env
/ {raw|lake}
/dir =by Team
/subDir =by domain
/sub-subDir=Hive table
/hive partitions
/*.parquet

Project ask for a « Usage » Job
Explain Inputs/Output
=> Validation by DataOffice
=> Grant recursive
READ permission for inputs

Hadoop hates too many Small Files

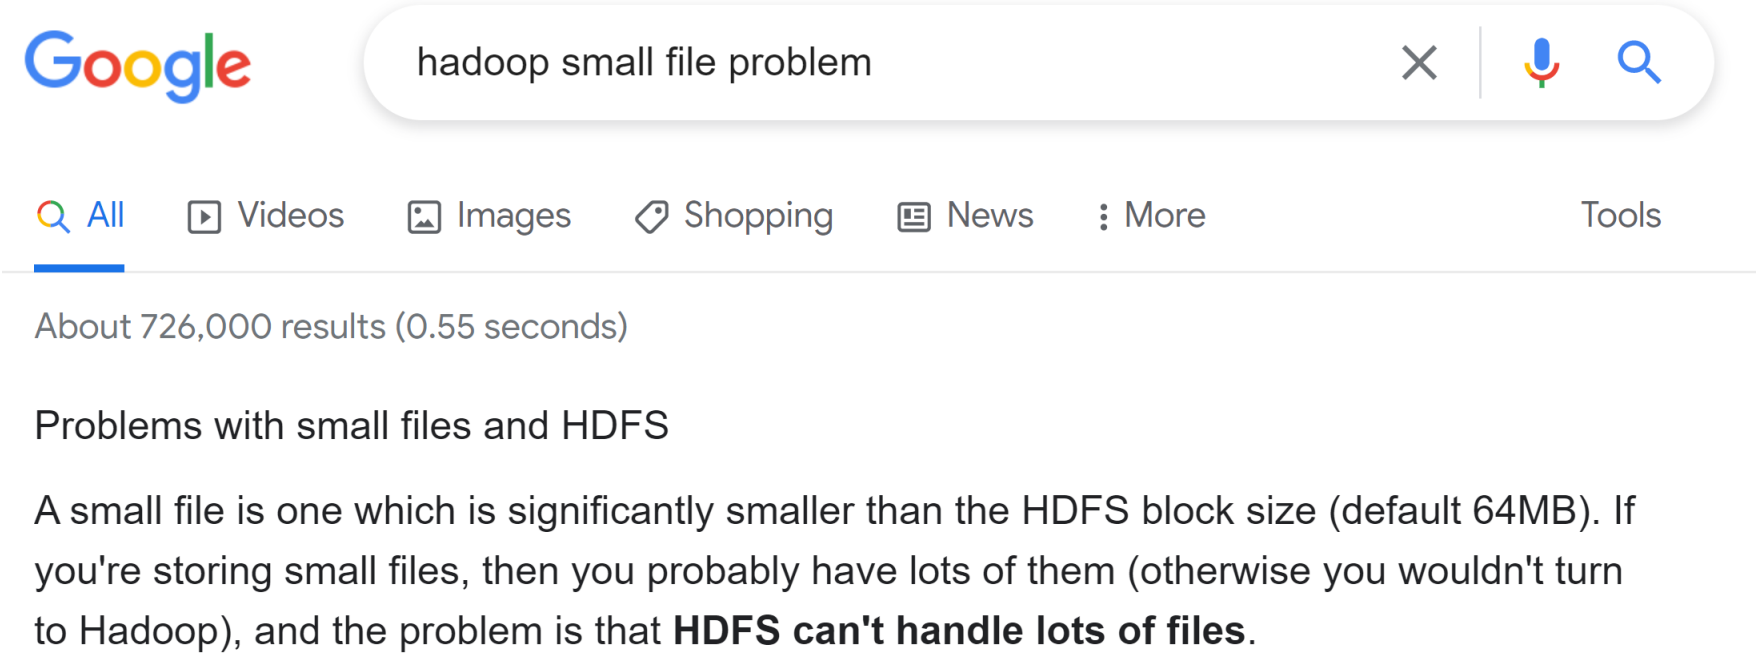

BIG Files ? More than Giga ?
YES … BUT PARQUET Partitioned

Doing BigData = using 

Files (correctly with 

)

End Of Story.

LAKE = Parquet
Starting from RAW: Postgresql or {CSV|Json|Avro..}
STAGING Dirs

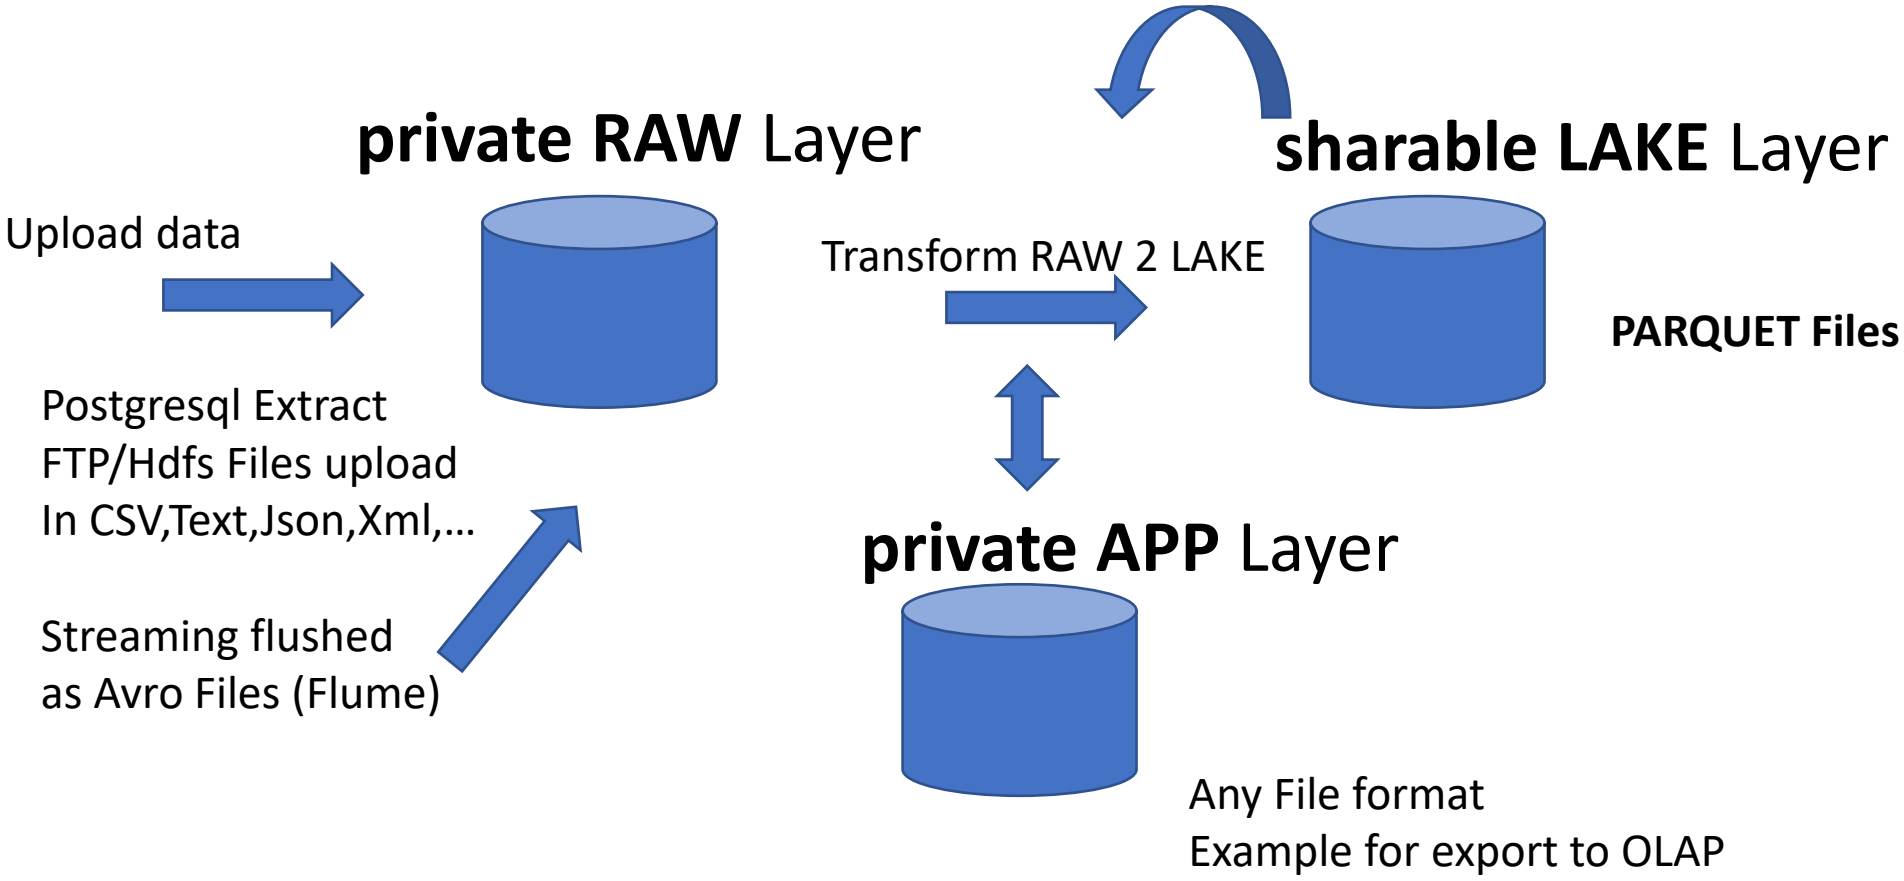

Feeding – Transforming – Consuming Data

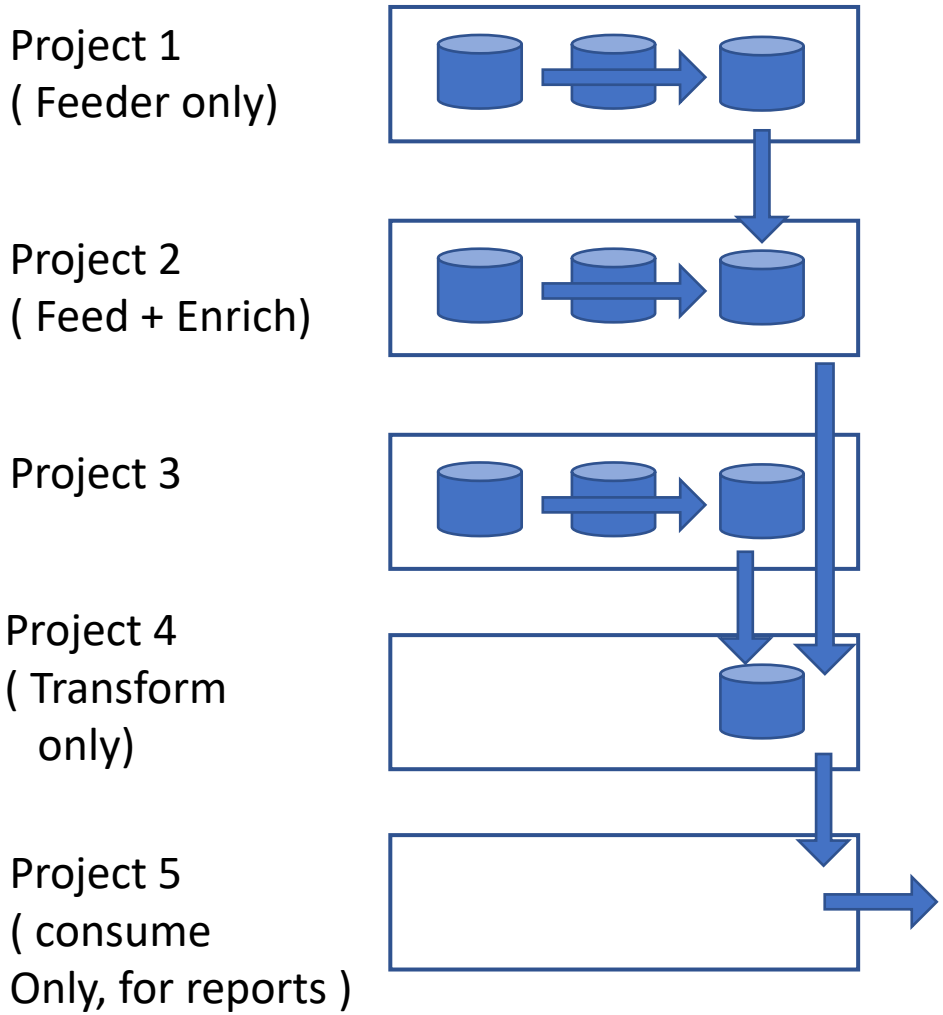

Every project is granted
3 Read-Write Databases: RAW, APP, LAKE

LAKE are sharable
… designed to be efficiently RE-read after (for analytical)

Every project can request Read access on others LAKE
(after validation by Data-Office)

A Dataset is writable by only 1 Project (the owner)

Big Processing

>= 1000 periodic jobs running daily / or every hours

Some processing take >= 15 Hours

Some processing are distributed on more than 200 yarn containers

Some take more than 1000 Go of RAM

Few Interactive Queries Processing screenshot Zeppelin

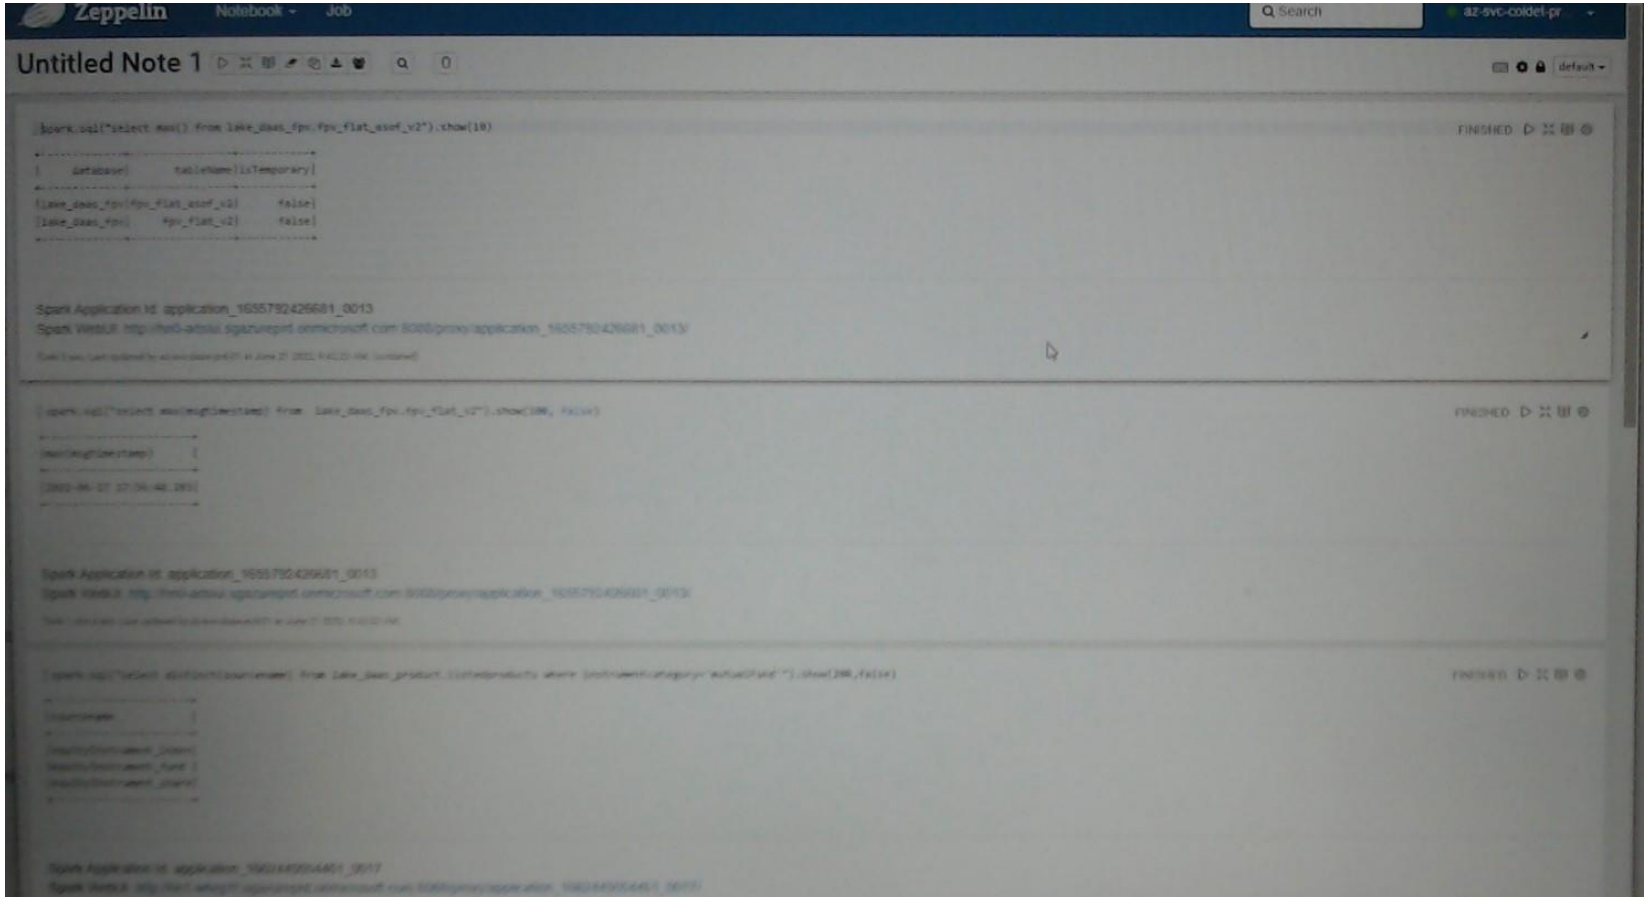

Orchestration / DataLineage
Data and Workflows are linked

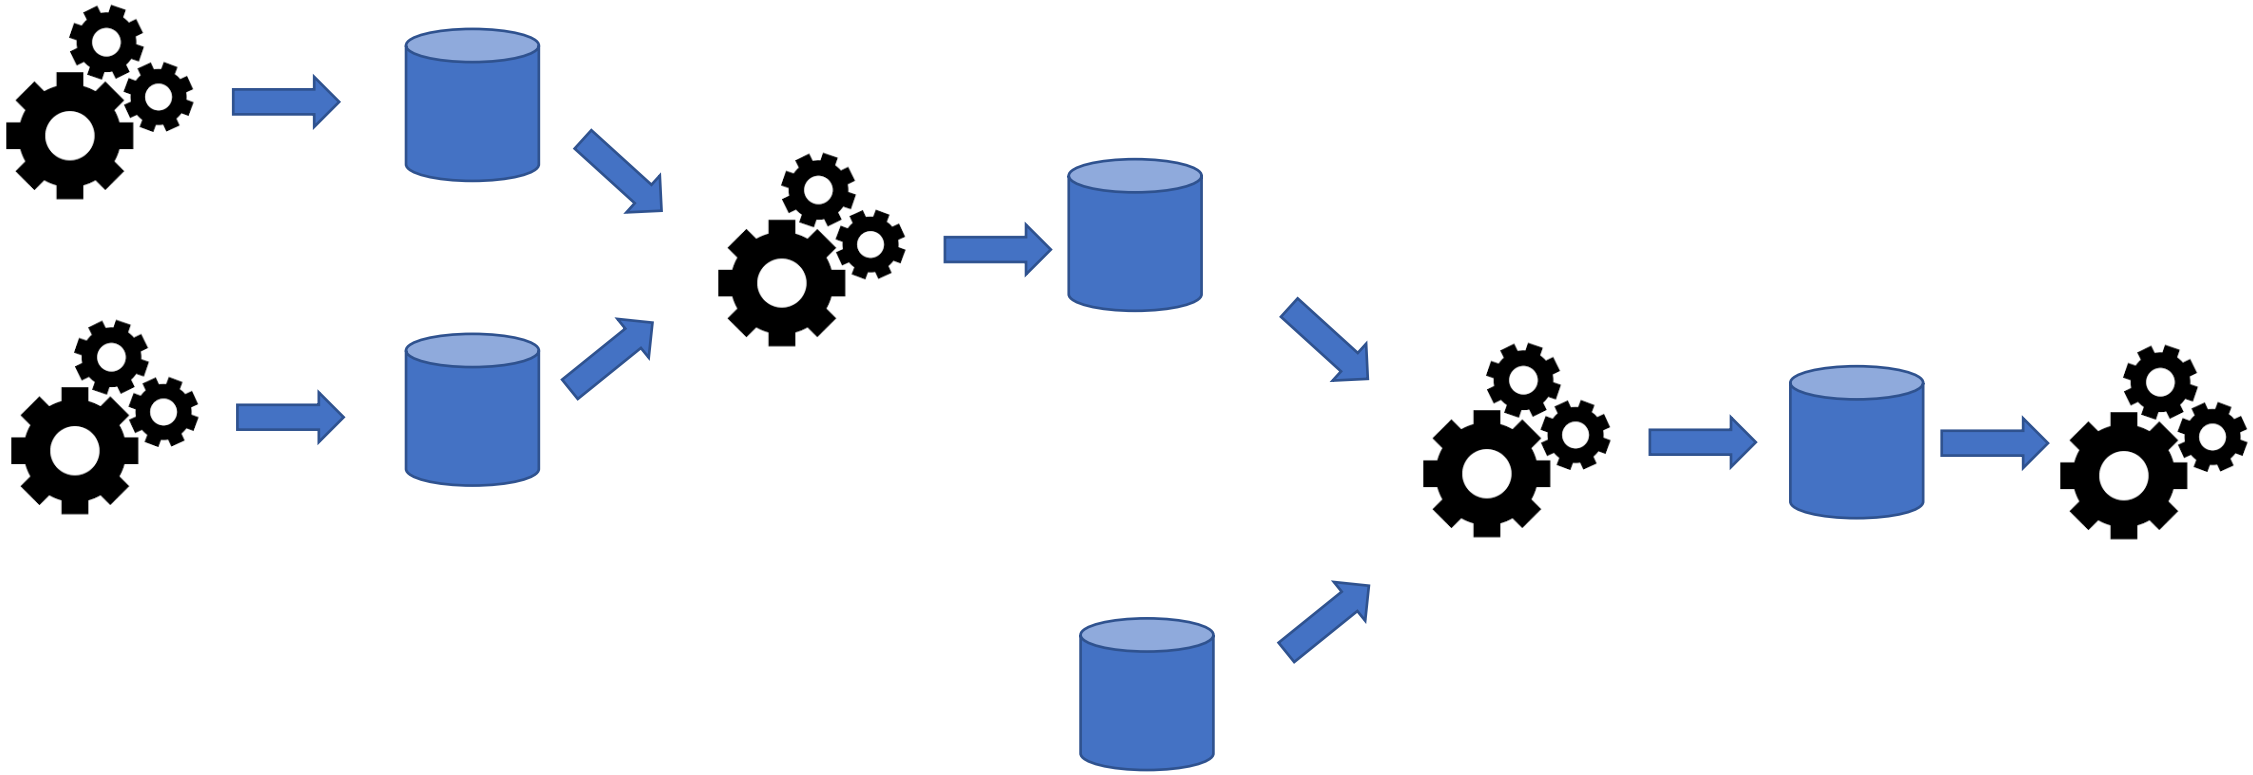

Read-Write Permissions from DataOffice

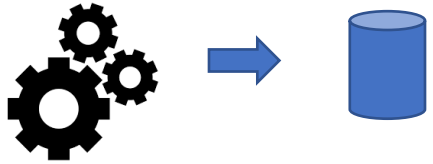

Dataset « Feeding » declaration:

= exclusive Write Permission : Project -> canWrite -> DB

For validation by DataOffice, DataSet must be described :
confidentiality, source, business owner, quality, RDPG legal

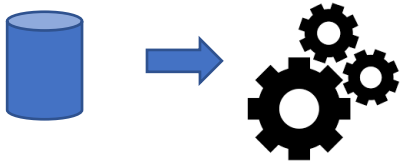

Consumer declaration:

= authorization to read : Project -> canRead -> DB
… for a specific Usage

For validation by DataOffice, use-case must be described,
usage must be « legitimate » , RDI legal

DataLineage
Data depends on other Sources Data

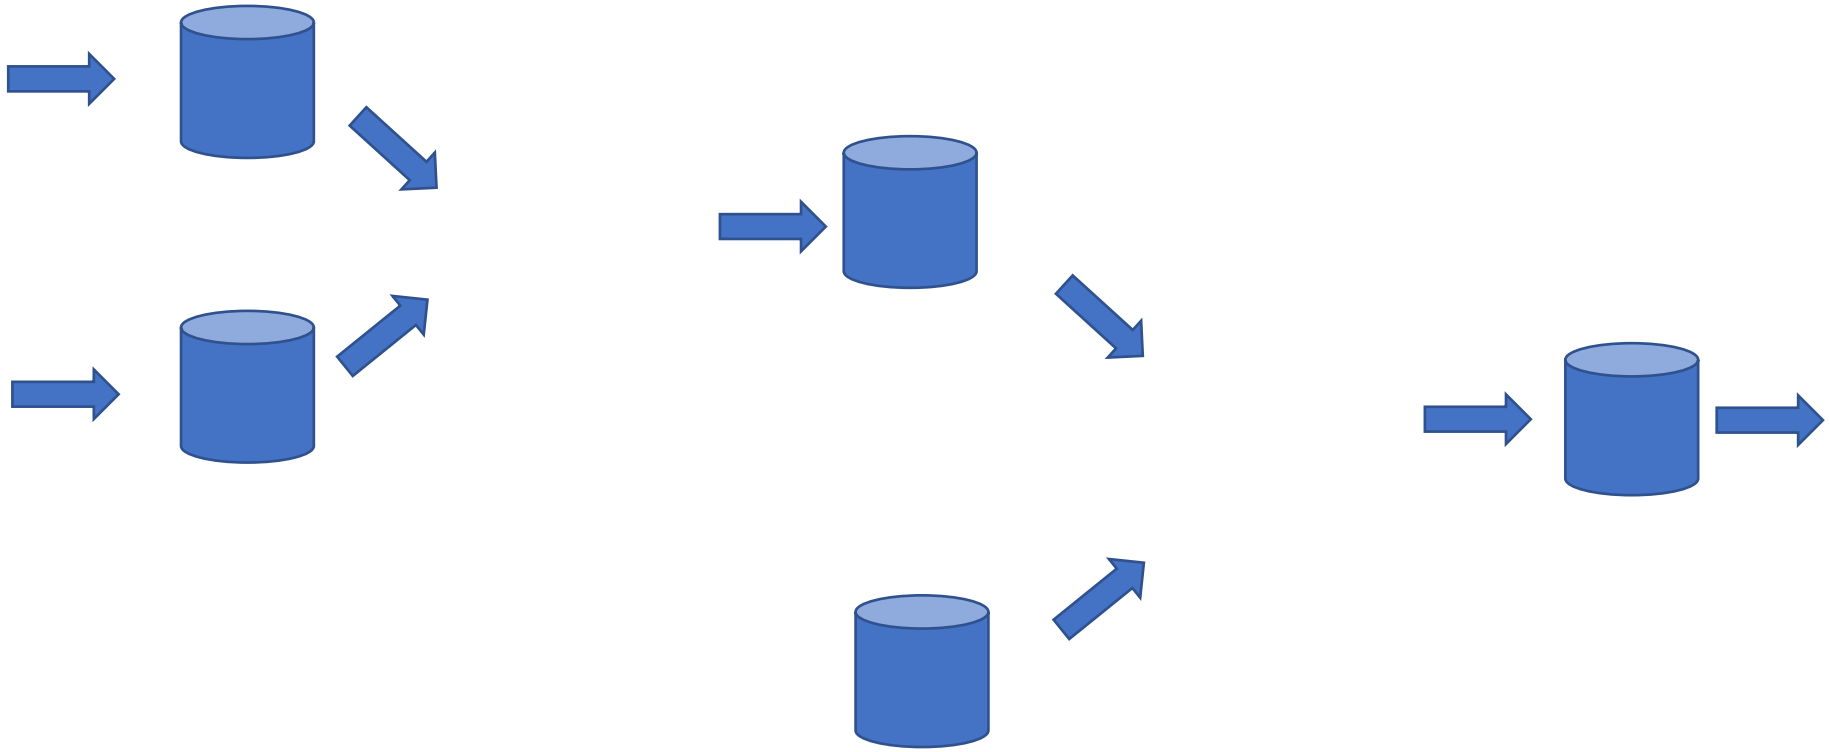

Screenshots (Fuzzy.. On purpose) Read/Write Dependencies Hive Database / Job…

Zooming on 4 out of 250 Hive databases

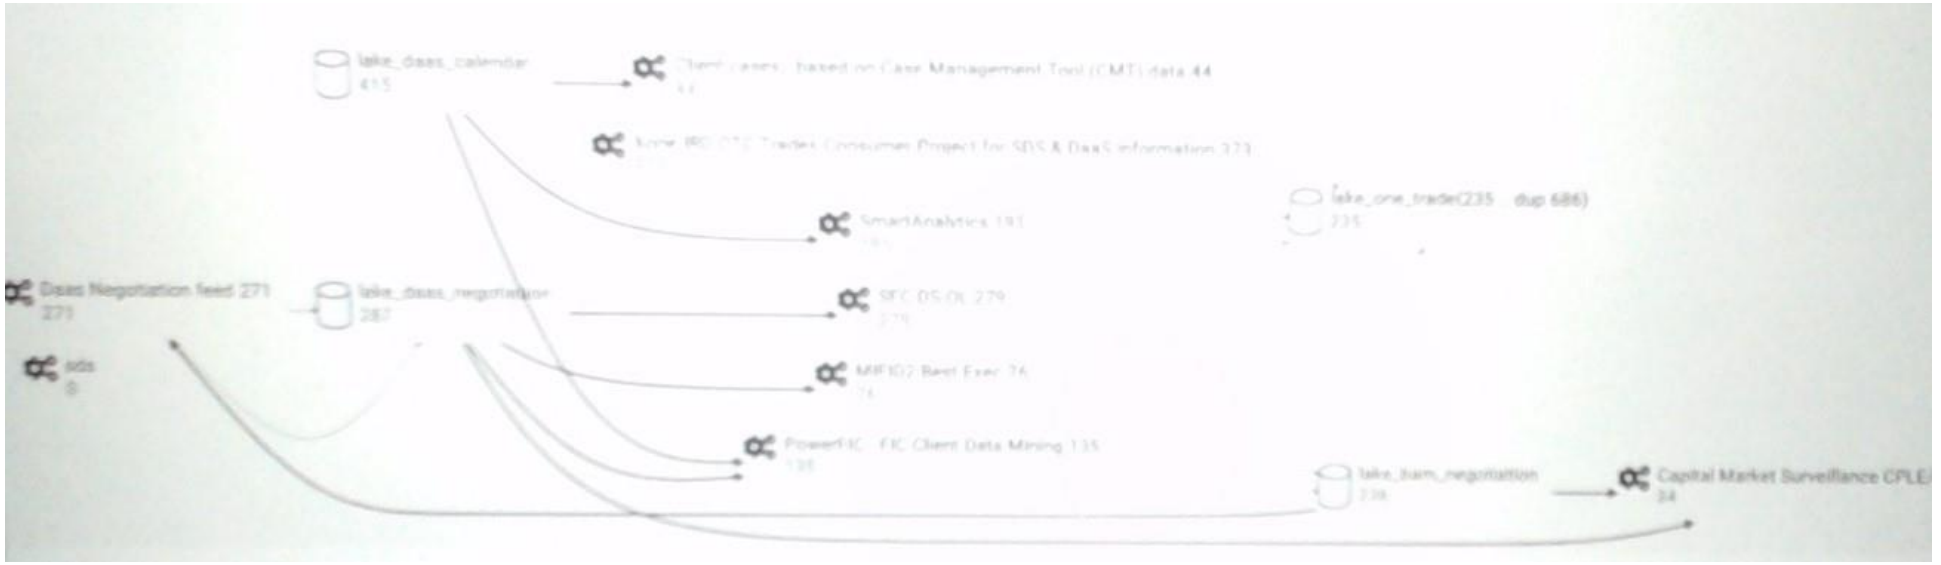

Orchestration
Start process2 after process1 finished

Start condition=
CRON 1 * * * * * *

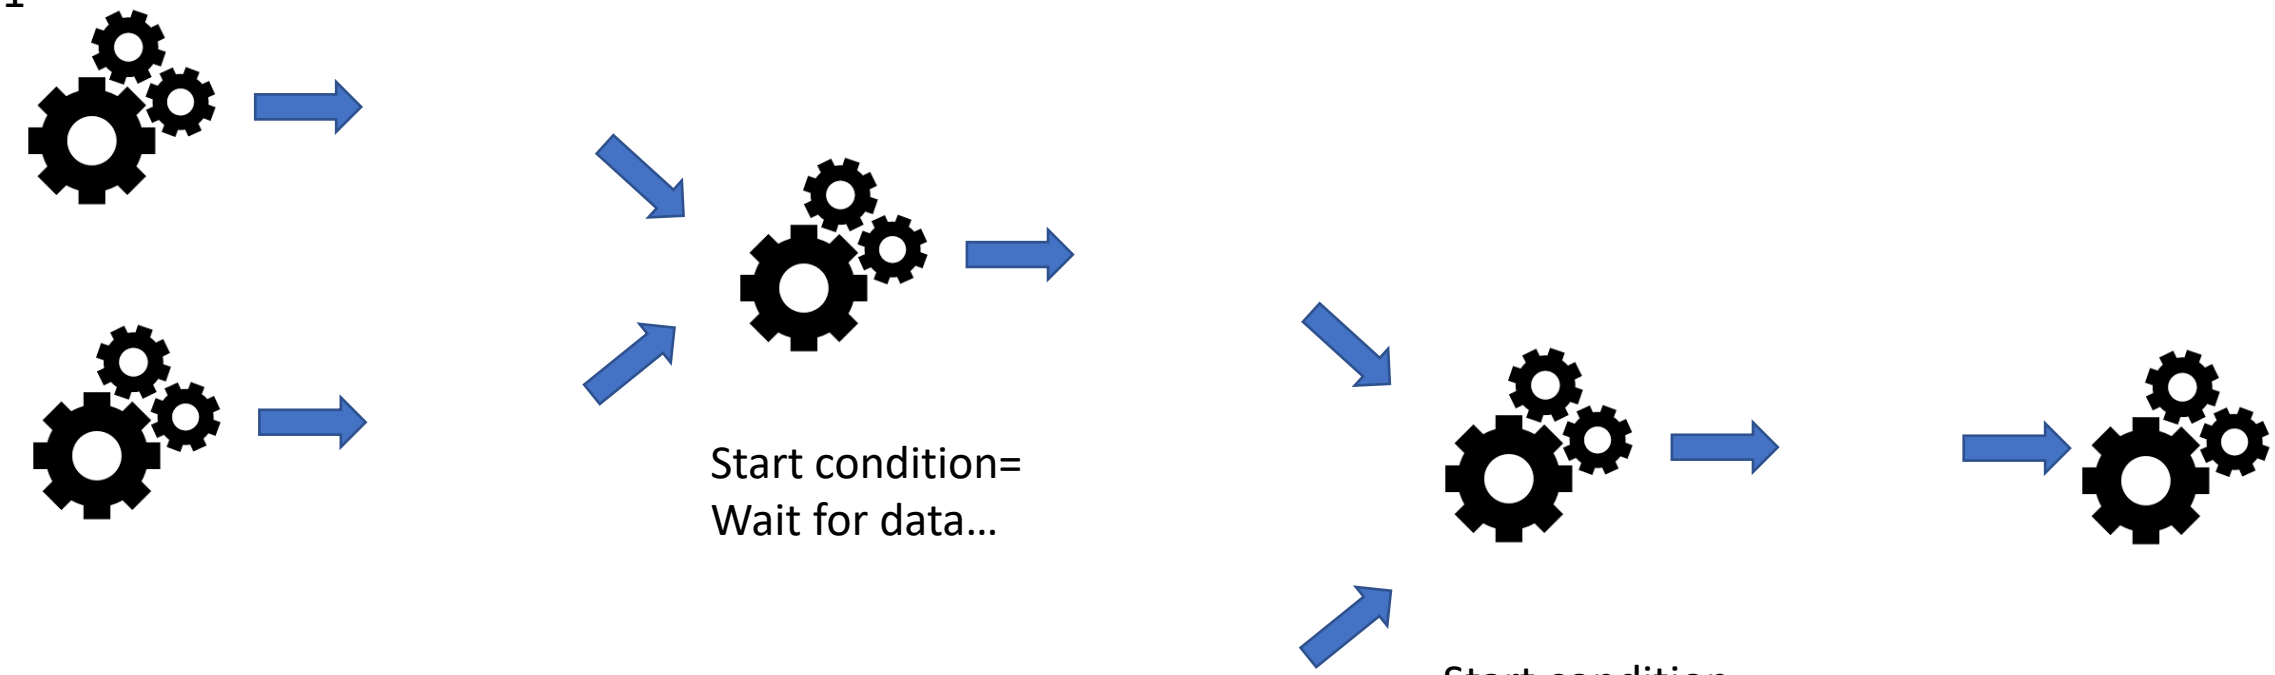

Start condition=
Wait for data…

Start condition=
CRON 10 * * * * * *

Prod Screenshot (Fuzzy.. on purpose) Running Yarn Applications

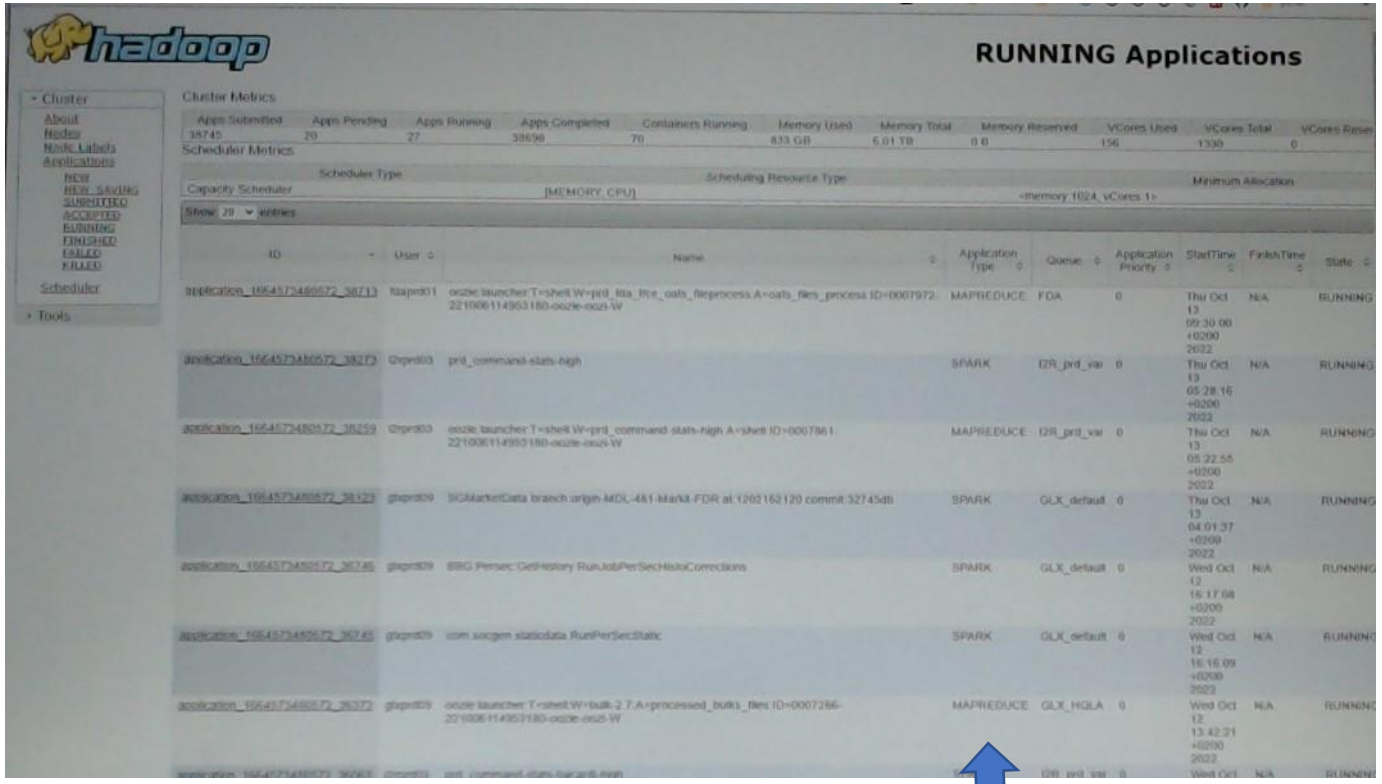

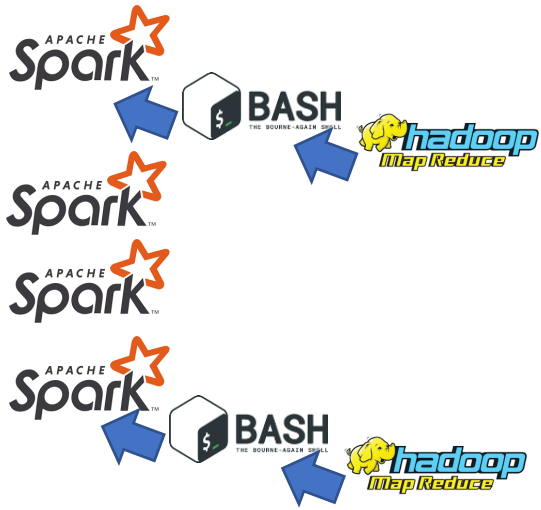

MAPREDUCE ?

… only internally to launch shell commands.. To spark

ALL others = SPARK

Prod Screenshot : Yarn Queues (Scheduler, CPU+Mem usages)

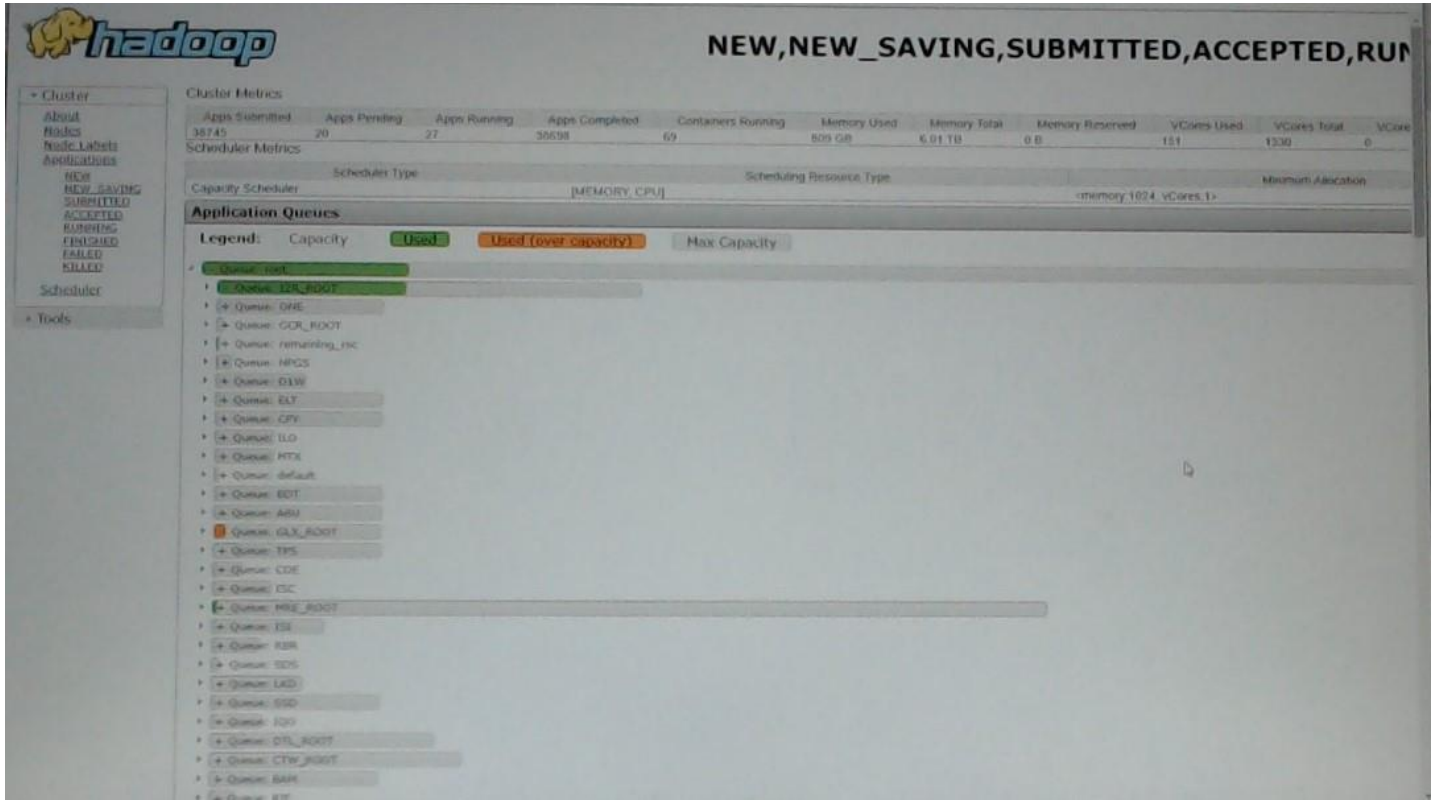

You see scrollbar?...
Hundreds or Yarn Queues
+ dozen of sub-queues

Screenshot
Oozie (Coordinators / Workflows)

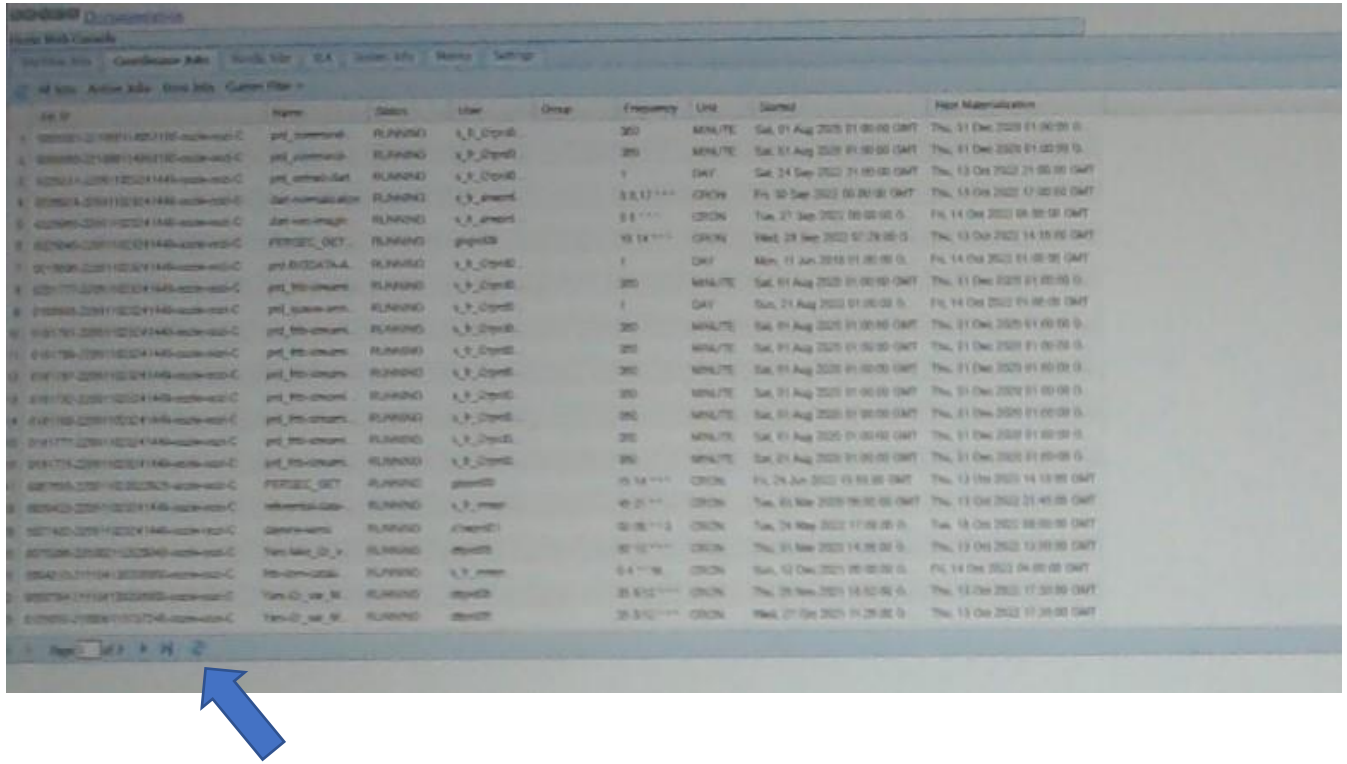

You see pagination?...
Thousands of CRON jobs

Oozie display only last 50 000 workflows
… so only 2 days of history

Orchestration « By CRON: 10 * * * * * * » Example of Possible Failures

batch1 started at 1 *** … but takes today (anormally) 15h instead of 10h

expected

actual

batch2 started at 12 *** …. data not yet present

=> NO « today » data (only yesterday)
Finished FAST, with NO data

batch2 started, during critical section of batch1 … table dropped/recreated after

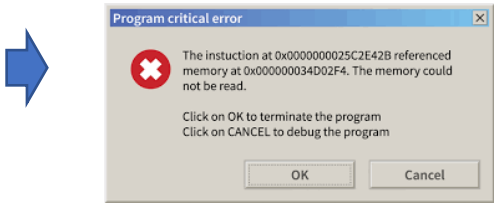

Outline

- What is BigData ?
  - Order of Magnitudes for « Big »
  - History
  - Evolution of Softwares, Spark
- Description of a Datalake
  - Content, data feeding
  - How it is organized
  - Who uses it
- Example of Data-Processing
  - RAW to LAKE, Reports
- Change Storage-Compute, Evolution to Cloud

Typical Example of adding 1 daily partition

```
$ hdfs dfs –mkdir hdfs://raw/team/domain/table/date=2022-10-12

$ hdfs dfs –put localFile hdfs://raw/team/domain/table/date=2022-10-12/

$ spark-sql –master local[1] \
        –e « ALTER TABLE raw_lake_domain.table
              ADD PARTITION ( date='2022-10-12' ) »
```

Typical RAW to LAKE processing with Spark as Java code

read

```
spark.read
    .format(« csv »)
    .option(«schema », « col1 type1, … colN typeN »)
    .load(« hdfs://raw/team/domain/table/date=2022-10-12 »)
```

transform

```
    .map(row ->  transformData(row) )
```

write

```
    .repartition(3, « col1 »)
    .sortWithinPartition(« col1, col2, col3 »)
    .write
    .format(« parquet »)
    .save(« hdfs://lake/team/domain/table/date=2022-10-22 »)
    ;
```

Typical RAW to LAKE processing with Spark as SQL code

write

```
INSERT OVERWRITE
      lake_team_domain.table
```

transform

```
SELECT /* +REPARTITION(col1, 3) */
   col1, col2,
   udf_func1(col3, col4)  as col3,
   udf_func2(col4, col5)  as col4,
   ..
```

read

```
FROM
      raw_team_domain.table
```

transform

```
JOIN
      lake_anotherTeam_domain.anotherTable x ON x.ID=id
```

read

```
WHERE date='2022-10-22' AND ..
```

write

```
SORT BY col1, col2, col3     -- idem sortWithinPartition
```

Example of LAKE Aggregation

```
INSERT OVERWRITE
      lake_team_domain.table
SELECT * FROM (
    SELECT * FROM table1 WHERE ..
  UNION
    SELECT * FROM table2 WHERE ..
  UNION
    SELECT * FROM table3 WHERE ..
  UNION
    SELECT * FROM table4 WHERE ..
)
SORT BY col1, col2, col3     -- idem sortWithinPartition
```

Example of « latest value » cristalisation analytical query « over(partition by) »

```
INSERT OVERWRITE
      lake_team_domain.table
SELECT
  col1,col2,…. colN    -- idem * EXCEPT rank   (cf issue SPARK-33164)
FROM (
  SELECT *,
    RANK() OVER (PARTITION BY id ORDER BY update_time DESC) as rank
  FROM lake_team_domain.event_table
)
WHERE rank=1
SORT BY col1, col2, col3     -- idem sortWithinPartition
```

Example of a « Small » Data Processing resource needed = RAM

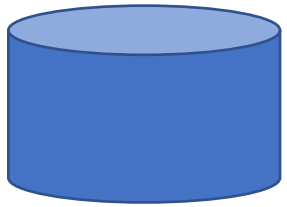

Extract all the Trades on perimeter « X »
Since <= 3 Years

Enrich « JOIN »
With other tables

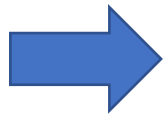

Basic Compute
SQL
Select .. where ..

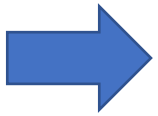

Save Result
Export to OLAP Cube

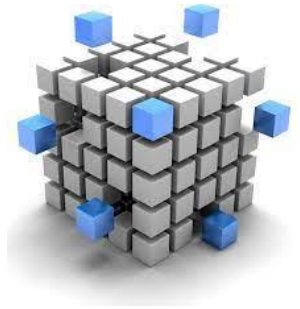

Processing on Azure
using cluster 10 x D14 (100 000 euros/year)
ONLY 600 Giga of RAM data …
NOT Enough => Swapping => take 4 Hours

Example of a Slow « Un-parallelized » Processing resource needed = CPU x Time x ..

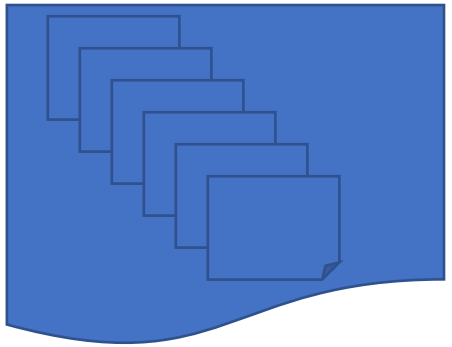

Appending every day
~1 Avro file, containing
~50 Millions of <XML>..</XML>
~70 Mega bytes

<XML>

For each XML, extract
xpath(« a/b/c/d/e ») as column1,
xpath(« a/b/f/g/h ») as column2,
…. as column150,

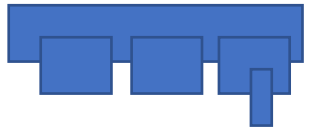

Append to History
(Audit Events)

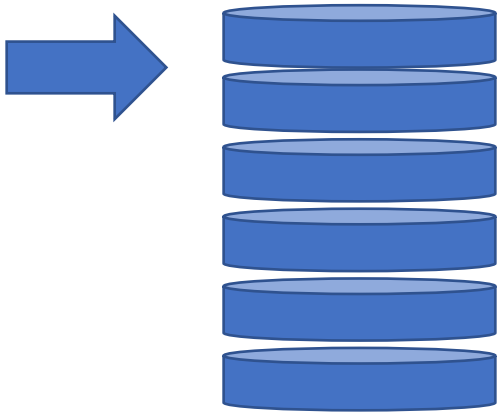

Synthesised Latest View

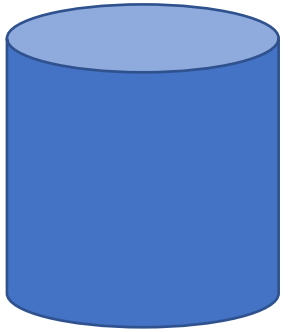

(updates on EquityOption trade description)

A Slow Processing ?

Legacy « On-Prem »

Daily Batch takes >= 15 Hours
But is mostly SINGLE-THREADED !!!

Some Parallelized intermediate parts take
>= 200 yarn containers … 10% of Cluster
… Result = ULTRA FAST !!

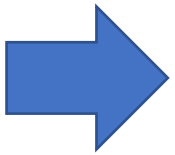

Backported code « as-is » to Azure HDInsight
Used Cluster of 10 x D14 cluster (40 000 euros/month)

Run NEVER Finished
… stopped in « echo count lines before insert $(…) »

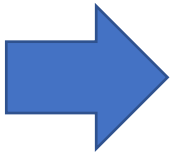

Fully re-implemented in Spark + Java
instead of Hive SQL + UDF xpath
+ Optimized …

Now run in 40 minutes on 2 x D13 cluster

Outline

- What is BigData ?
  - Order of Magnitudes for « Big »
  - History
  - Evolution of Softwares, Spark
- Description of a Datalake
  - Content, data feeding
  - How it is organized
  - Who uses it
- Example of Data-Processing
  - RAW to LAKE, Reports
- Change Storage-Compute, Evolution to Cloud

Appendix:
Cloud

… More Detailed on

« From OnPrem to Azure Cloud »

see Document « -cloud.pdf »

Appendix: Cloud

Fundamental Resources TradeOffs

Network (Bandwith, Latency)

Horyzontal Scale

X

CPU (HOT, Watts)

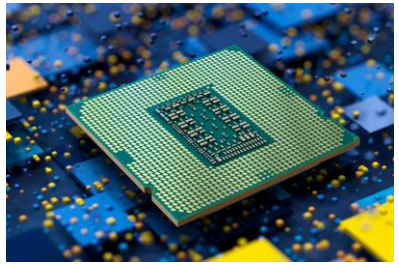

Development price / Run cost

€€€€€€$$$$$

Cold Storage (SLOW)

Storage (~Small)

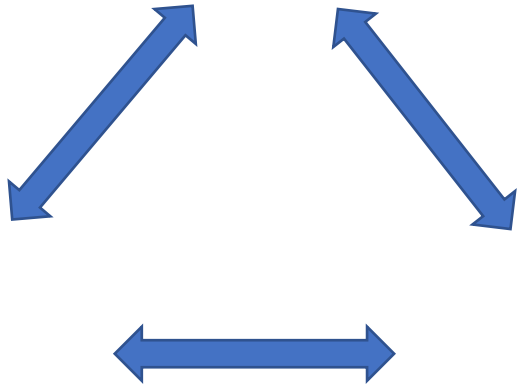

RAM (FAST .. Expensive)

\+ + + + + +

\- - - - - -

Appendix: Cloud

Performance Changes -> Architecture Changes

Historically:
SLOW Disks
Disk Collocated with RAM+CPU

MapReduce « send program » collocated to data

Now:
FAST(ER) SSD Disks + Networks

Can « send» data to distributed programs

Storage and Compute can be SPLIT … OK !

Higher Density for Huge Storage, lower prices

48 * disks in 2U blade

1 Peta SSD … 200 000 €

Cloud Provider : AWS, Azure, Google
offer Storage solutions (S3, Azure Storage, ..)

Azure Storage

Amazon S3

+ Managed Services
+ Serverless
+ Kubernetes

Appendix: Cloud

Migration to Cloud
Azure - AWS - Google

Goals:

- Elasticity of Storage
  (No more fear of HDFS FileSystem Full)
- Elasticity of Compute
  (adapt CPU to workload… Pay only what you use)

- Clear visibility of cost per Projects / internal refacturation
- By-Pass internal IT department

- Easier (?) Self-Service API for Provisionning
- No More Multi-tenant (no risk of 1 project crashing/consuming whole cluster)

Since 2020 SGCIB is moving its Datalake to Azure

Appendix: Cloud

Architecture of Datalake on Clouds

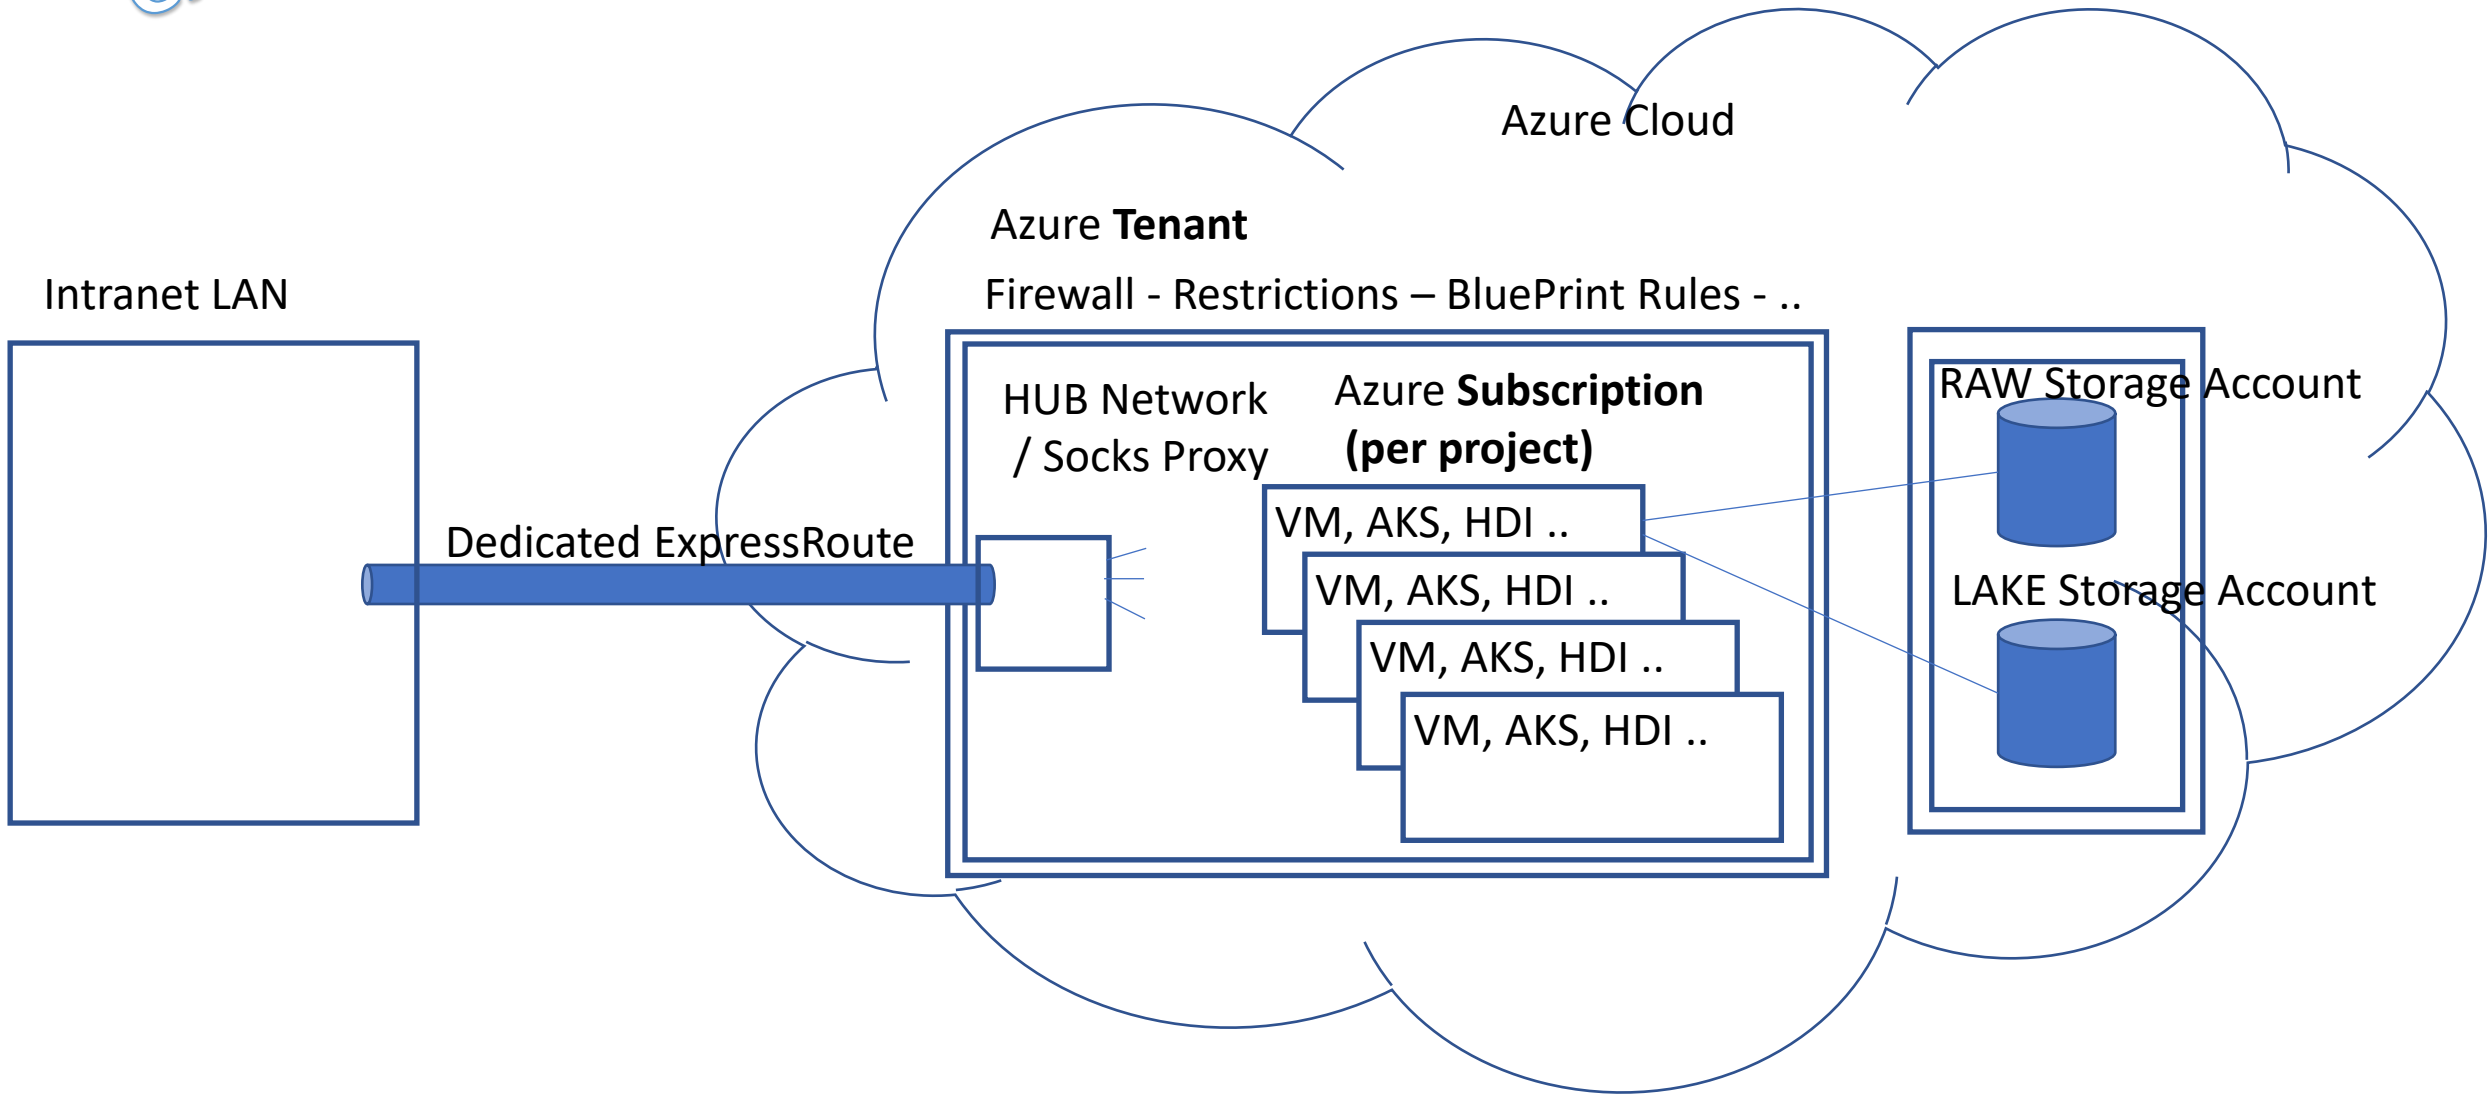

Appendix: Cloud

BigData Engineering from OnPrem to Azure

Development of custom Yarn/Oozie/Ranger tools

1 Huge cluster, used as Multi-tenant (several users)

Used at 100%

Almost Disk-Full
Many report on disk usage / data purges

Clear view of all Running workloads
(1 Ambari screen)

Lot of Network & Security
Development of Provisionning « sudo » tools

NO more multi-tenant system… but LOT of small clusters

Hundred of clusters, each used at 5% !!!
COST COST COST
Necessity to Optimize Performance
(batches too long… 100x slower / too expensive)

No more Disk worry
… data is growing (no more purges?)

NO admin central view
NO Orchestration of workloads

Appendix: Cloud

What About Spark in Azure Migration?

COST COST COST

Necessity to Optimize Performance

(batches too long… 100x slower / too expensive)

=> Many projects have migrated + optimized

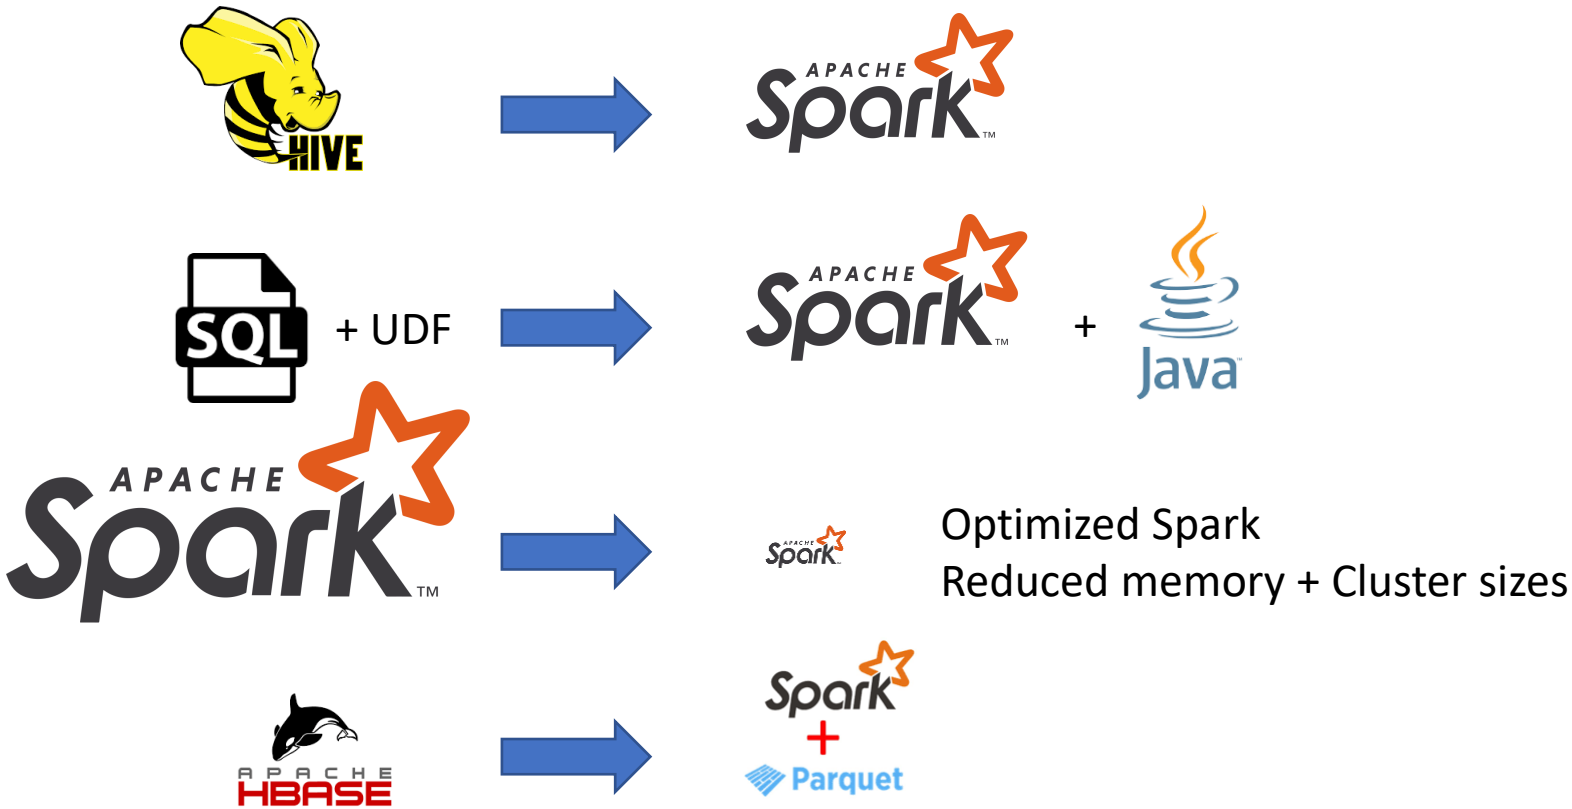

Appendix: Cloud

Managed Spark ?
Goals: Autoscaling / « serverless »

Managed by Azure :

Azure Synapse Analytics

= Apache Spark

Managed via Kubernetes

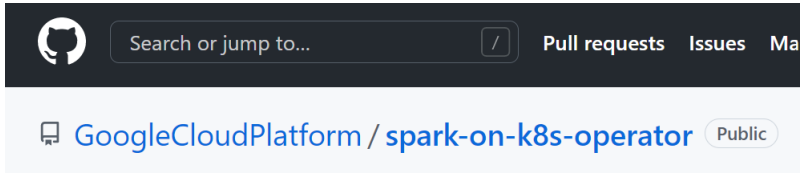

= kubernetes

Apache Spark

databricks

databricks

= Apache Spark

AWS : EMR

GCP : BigQuery

Questions?

Take Away

What is BigData ? Horyzontal Scaling
- compute: cluster with Tera of RAM used by Spark apps
- storage: Petas of Files, in parquet

What is Spark ?
- Simple unified Sql/Java engine
- for distributed compute (Yarn/Kube)
  - distributed storage (HDFS/cloud)

What is Processing ?
- mostly spark batches
- Feeding RAW
- Transforming RAW to LAKE
- Consuming SQL analytics

Hadoop ecosystem is complex, Spark brings simplicity
Ecosystem is evolving (Cloud, Kubernetes)