

Failures and Resiliency Principles of Distributed Computing

course 2024
arnaud-nauwynck@gmail.com

This document:
[https://github.com/Arnaud-Nauwynck/presentations/pres-bigdata/
3-failure-and-resiliency-pinciples-distributed-computing](https://github.com/Arnaud-Nauwynck/presentations/pres-bigdata/3-failure-and-resiliency-pinciples-distributed-computing)

MTBF

M.T.B.F = Mean Time Between Failure

For HDD ~ 500 000 hours \geq 50 years

May looks good at home, to save your data
("Mean" is probabilistic average...
Real errors will occur before)

MTBF « at Scale »

In DataCenter with 10 000 servers x 4 disks

=> 1 failure every 1h 15mn

= 19 failures per day

becomes a recurrent task / job



When Failure(s) Happens ?

=> Program fails ?

=> System fails ?

=> need relaunch manually ?

=> diagnostic hardware/software ?

=> interrupt all, repair ?



Studying « When Failure » 1/5

=> Program fails ?

=> System fails ?

=> need relaunch manually ?


=> diagnostic hardware/software ?

=> interrupt all, repair ?

When Failure..

=> Program fails ?

Individual Component : Obviously FAIL



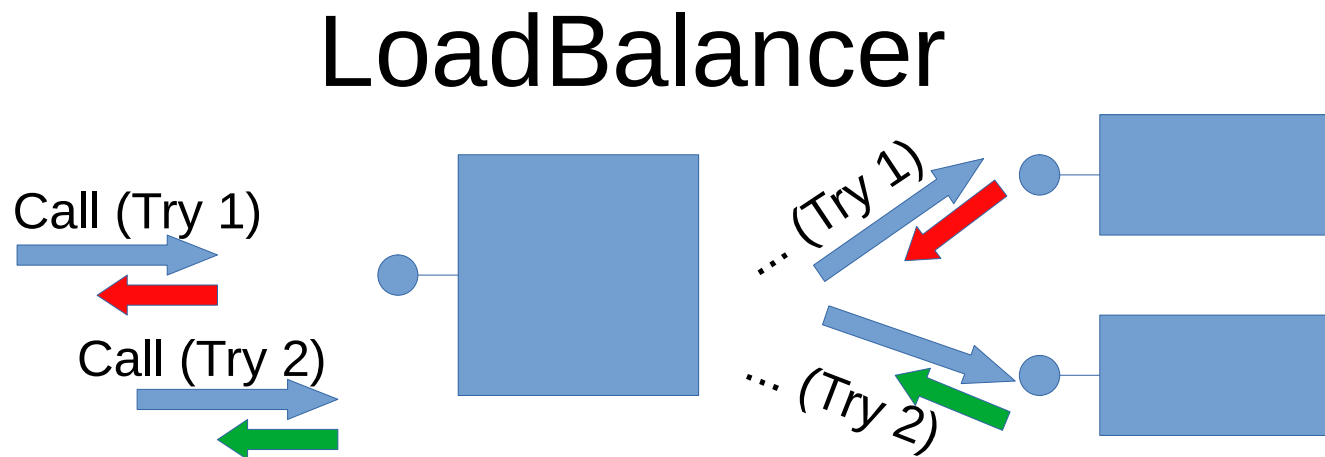
Distributed Architecture : RESILIENT

resist (hopefully) to (some) failure

Failure is not « Exceptional »
.. is a « normal » path
consider it everywhere in code

```
public void doSomething() throws Exception {  
    int maxRetry = 5;  
    for(int retry = 0; retry < maxRetry; retry++) {  
        try {  
            anyTreatmentCanFail();  
            break; // success!  
        } catch(Exception ex) {  
            Log.warn("Failed .. retry " + retry + "/" + maxRetry, ex);  
            sleep(100 * Math.pow(2, maxRetry)); // wait a little  
        }  
        if (retry+1>=maxRetry) throw ex;  
    }  
}
```

Retrying with a LoadBalancer .. might just work




Dispatch to underlying servers
Using Round-Robin


Server Health Check

Temporary evict from RoundRobin Pool

Health Check



Client-Side « LB »



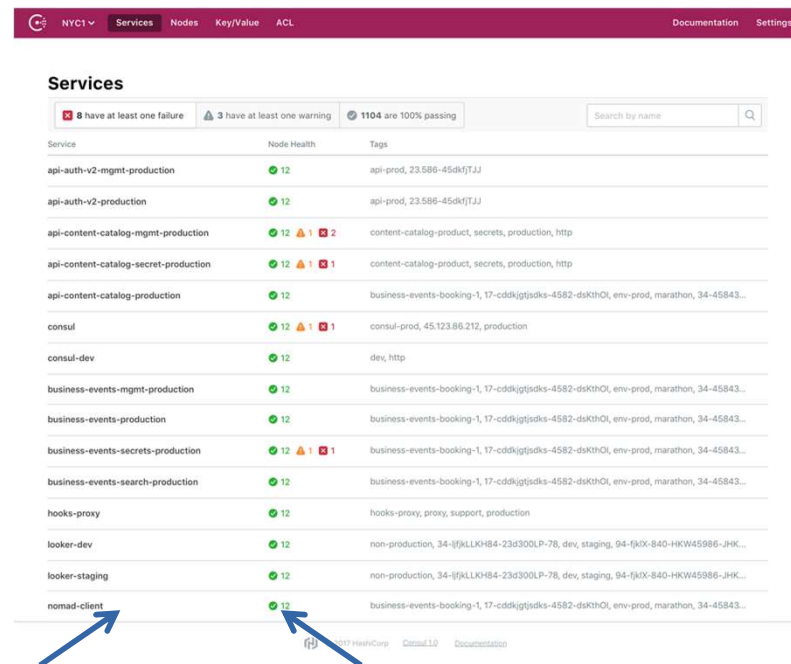
NO « LB »

Need discovery mechanism on client

Service >= Nodes

Examples of Service Discovery

DNS, ServiceMesh,
Zookeeper, Consul.io,
HAProxy, Kubernetes, ..



The screenshot shows the Consul.io web interface. At the top, there's a navigation bar with 'Services', 'Nodes', 'Key/Value', and 'ACL'. Below this, the 'Services' page is displayed. It features a summary bar indicating '8 have at least one failure', '3 have at least one warning', and '1104 are 100% passing'. A search bar is also present. The main content is a table with three columns: 'Service', 'Node Health', and 'Tags'. The table lists various services such as 'api-auth-v2-mgmt-production', 'api-content-catalog-mgmt-production', 'consul', 'business-events-mgmt-production', and 'nomad-client'. Each service row shows a green circle with a checkmark for healthy nodes and a red circle with an exclamation mark for nodes with warnings or failures. For example, 'api-content-catalog-mgmt-production' shows 12 healthy nodes, 1 warning, and 2 failures. 'nomad-client' shows 12 healthy nodes. Two blue arrows point from the text labels below to the 'Node Health' column: one points to the 'nomad-client' row and the other points to the 'api-content-catalog-mgmt-production' row.

Service	Node Health	Tags
api-auth-v2-mgmt-production	12	api-prod, 23.586-45dktTJJ
api-auth-v2-production	12	api-prod, 23.586-45dktTJJ
api-content-catalog-mgmt-production	12 1 2	content-catalog-product, secrets, production, http
api-content-catalog-secret-production	12 1 1	content-catalog-product, secrets, production, http
api-content-catalog-production	12	business-events-booking-1, 17-cddkigtjdsks-4582-dskhOI, env-prod, marathon, 34-45843...
consul	12 1 1	consul-prod, 45.123.86.212, production
consul-dev	12	dev, http
business-events-mgmt-production	12	business-events-booking-1, 17-cddkigtjdsks-4582-dskhOI, env-prod, marathon, 34-45843...
business-events-production	12	business-events-booking-1, 17-cddkigtjdsks-4582-dskhOI, env-prod, marathon, 34-45843...
business-events-secrets-production	12 1 1	business-events-booking-1, 17-cddkigtjdsks-4582-dskhOI, env-prod, marathon, 34-45843...
business-events-search-production	12	business-events-booking-1, 17-cddkigtjdsks-4582-dskhOI, env-prod, marathon, 34-45843...
hooks-proxy	12	hooks-proxy, proxy, support, production
looker-dev	12	non-production, 34-ijfjLLKH84-23d300LP-78, dev, staging, 94-fkIX-840-HKW45986-JHK...
looker-staging	12	non-production, 34-ijfjLLKH84-23d300LP-78, dev, staging, 94-fkIX-840-HKW45986-JHK...
nomad-client	12	business-events-booking-1, 17-cddkigtjdsks-4582-dskhOI, env-prod, marathon, 34-45843...

Consul.io Services

Nodes

Studying « When Failure » 2/5

=> Program fails ?

=> System fails ?

=> need relaunch manually ?

=> diagnostic hardware/software ?

=> interrupt all, repair ?

System = Union of independent
components

Each component can be killed
... The system must survive each failure

Be Confident in « System »

Walk on a « Ant » ...



... « Anthill » not in danger



How to Check System does not « Fail » ?

Very difficult to « prove » system correctness

Easy to test : « kill and see »

Not an exhaustive test... Repeat + Changes

Test Kill with Chaos Engineering

Initiated by Netflix

Goal : Randomly kill Process / VM / Datacenter



SLA : 99.99 Up-time ?



= Daily: 8s

Weekly: 1m 0s

Monthly: 4m 22s

Quarterly: 13m 8s

Yearly: 52m 35s

3 nines : 99.999 = Yearly : 5mn15s

4 nines : 99.9999 = Yearly : 31s

Studying « When Failure » 3/5

=> Program fails ?

=> System fails ?

=> need relaunch manually ?


=> diagnostic hardware/software ?

=> interrupt all, repair ?

When Failure ...

=> need relaunch manually ?

Hard-Coded Launch to specific server.. fail



Auto Distributed: RESILIENT

don't launch manually on a specific server

Let the system select one

Pet vs Cattle




Pet have id=name
You take extra care of it



Cattle ≥ 1000 : anonymous
Id=Number, interchangeable

Launching => Scheduling on Allocated Resource



You submit
Something to run

You don't run yourself

internally, it may
wait/allocate resource
do retries...
Not your concern.

Studying « When Failure » 4/5

=> Program fails ?

=> System fails ?

=> need relaunch manually ?

=> diagnostic hardware/software ?

=> interrupt all, repair ?

When Failure ...

=> diagnostic hardware/software ?

NO ? ... difficult anyway

Bugs may exist

BUT failure is « not » a bug

Studying « When Failure » 5/5

=> Program fails ?

=> System fails ?

=> need relaunch manually ?

=> diagnostic hardware/software ?

=> interrupt all, repair ?

When Failure ...

=> interrupt all, repair ?

NO ...

Immutable-Infrastructure

Do not edit infra once created

Drop and re-create new VMs

When Failure ...

Hardware HOT-PLUG :
unplug old disk, and plug new one
Without interruption



When Failure ...

Idem for Software :
add/remove servers to cluster
(no hard-coded topology/configs)

Duplicate Failable Component for Fewer System Failure

If a component has 0.01 chance to Fail today

Adding another (independent) component...

=> Probability that 2 fail today = $0.01 * 0.01 = 0.0001$

=> Probability that 3 fail today = $0.01^3 = 0.000001 = 1e-6$

« If » your system still works with 1 working component out of 2 ... better

System \geq Component

If No Correlated / Dispatch / No Spof

Electric power : when fail ... all fails (correlated failures)

Need reliable network + dispatching to working components (retry/detect failed ones)

Everything between the components can be a « SPOF »

Network may be a « Single Point Of Failure »

as all forgotten components not redundant

Story of Arienne 501

Duplicated / « Correlated » Errors...




Flight computation performed twice in parallel on 2 isolated hardwares

But using same program...

Both programs throw same «overflow exception » at exact same millis

=> Lessons learned : use 2 hardwares + 2 independent softwares + ..




Single Point Of Failure ?



High Availability Servers...

But Shared State / Database .. SPOF

Examples of SPOF




NO SPOF : Duplicate Everything


But how to duplicate a single Source of Truth « Data » ?

Copy => stale data / replication / distributed lock ?


Stateless, Spof, Sharded (easy) vs Statefull (difficult)




Stateless
(or stale Read-Only
cached data)



Spof



Sharded... 2 spofs
(ex : id modulo 2)
Horizontal scale : OK
But not replicated




Statefull
Concurrent Distributed Data


Distributed Lock ?
Replication ?
Source of truth ?

« HA » High Availability

Active – Standby + Replication



Transparent Delegate to Active



Leader Election System

Curr Leader = node1

Synonym Terms...

Master – Slaves (not politically correct)

Leader – Followers

Active – StandBys

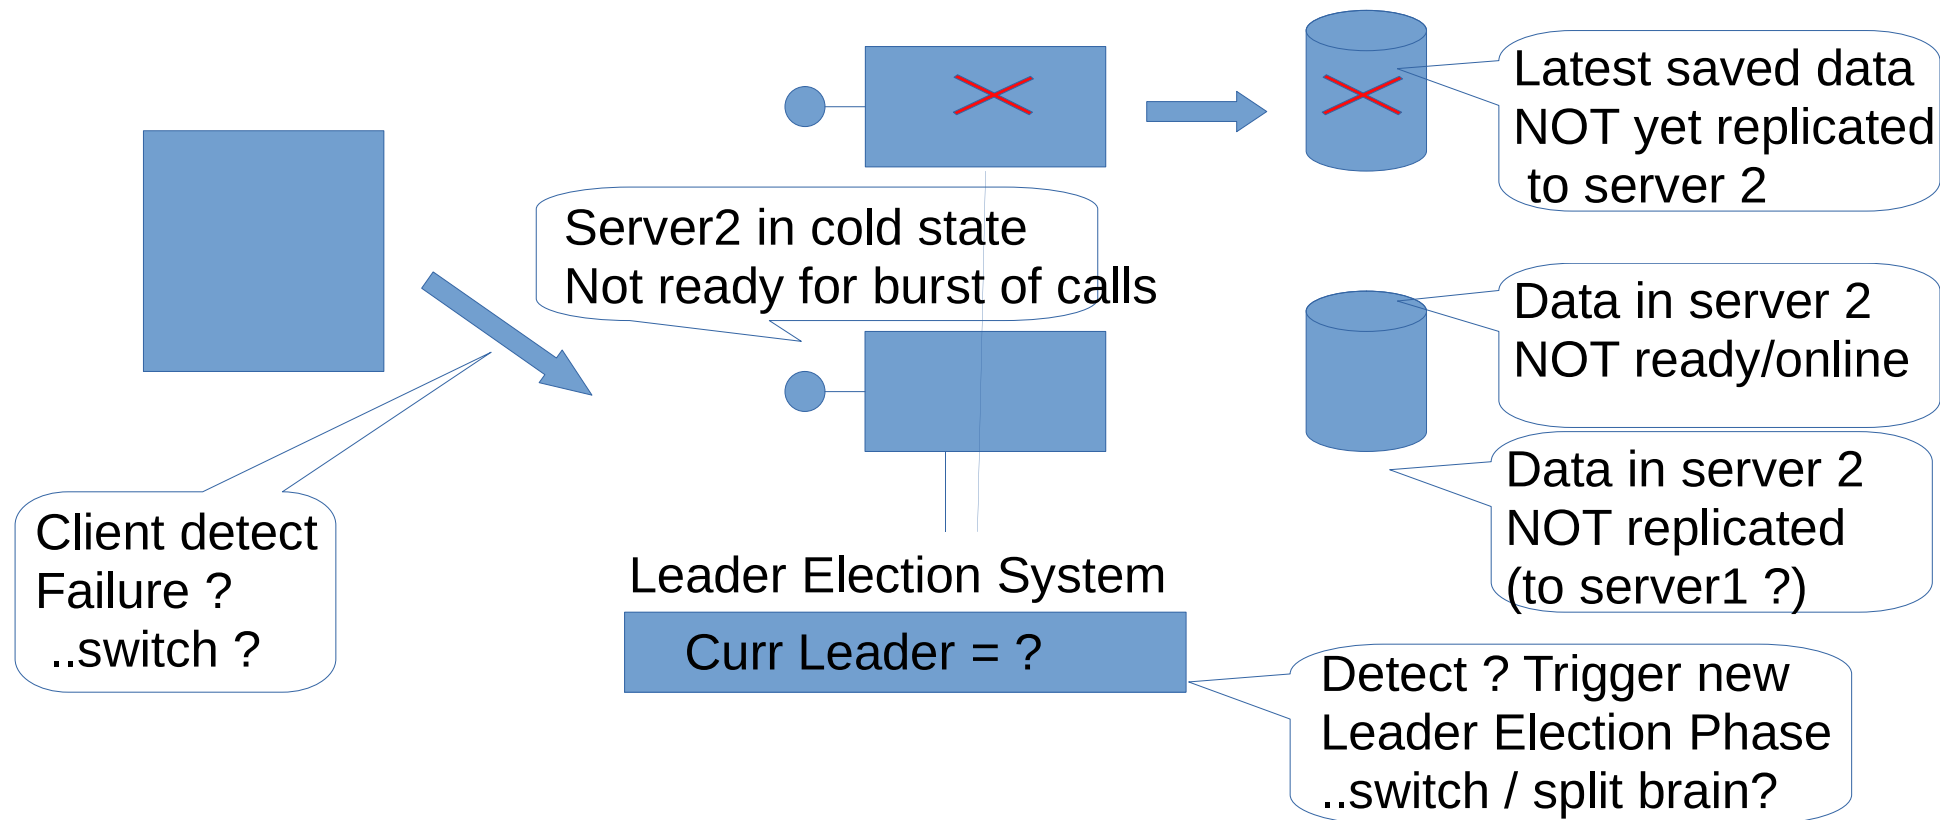
Primary – Secondaries

Main – Replica

MainSite - Disaster Recovery

MainDisk - BackupDisk

Switching from Active to StandBy



Problem with Leaders same as in Political

works only when there is exactly 1 leader

≥ 2 leaders ... conflicts / Split-Brains / Network Partitioning

0 Leader ... nothing works

Imposters pretend to be Leader

Leader staying leader too long after new election

Slow to organize election

Trust the result of election ?

Who can organize if there is no leader?


Can't decide election in 50 % / 50 % equality ... need an odd number : 3, 5, 7..

There must be at least 50 % participation (no Abstention)

Otherwise both candidates pretend to majority (not absolute majority) with 50 %

Leader fail to « Start Lead » ... re-electing too Often ?

Example of Catastrophic scenario :



1/ Server « N » become Leader

2/ It needs huge memory to serve requests

There are too many requests (burst of waiting requests)

3/ « Full GC » occurs, take $\geq 30s$, and « stop-the-world »



30s is the default Timeout for socket connection / read response

... During this time, Leader fails to answer to « Health check »


4/ Leader is considered not healthy => not leader any more

=> 5/ Electing another Leader « N+1 » => back to beginning 1/

Split Brain / Network Partitioning



1 cluster, 1 master
status=OK
... but network partition happens



2 clusters fully isolated each
Each have 1 master
Each have status=OK
... data diverge, will discover conflict on merge

Quorum of 50 %



Quorum of >50 % reaches => elect candidate1



NO Quorum of 50 % reaches
=> re-organize new election
(proba=0.5 that 2 of 3 vote for same)



Majority but NO Quorum
... high risk of Split-Brain
(all others may vote and you are not aware)
=> wait enough voters (re-organize / accept vote)

Waiting Quorum

= system « Not Available »

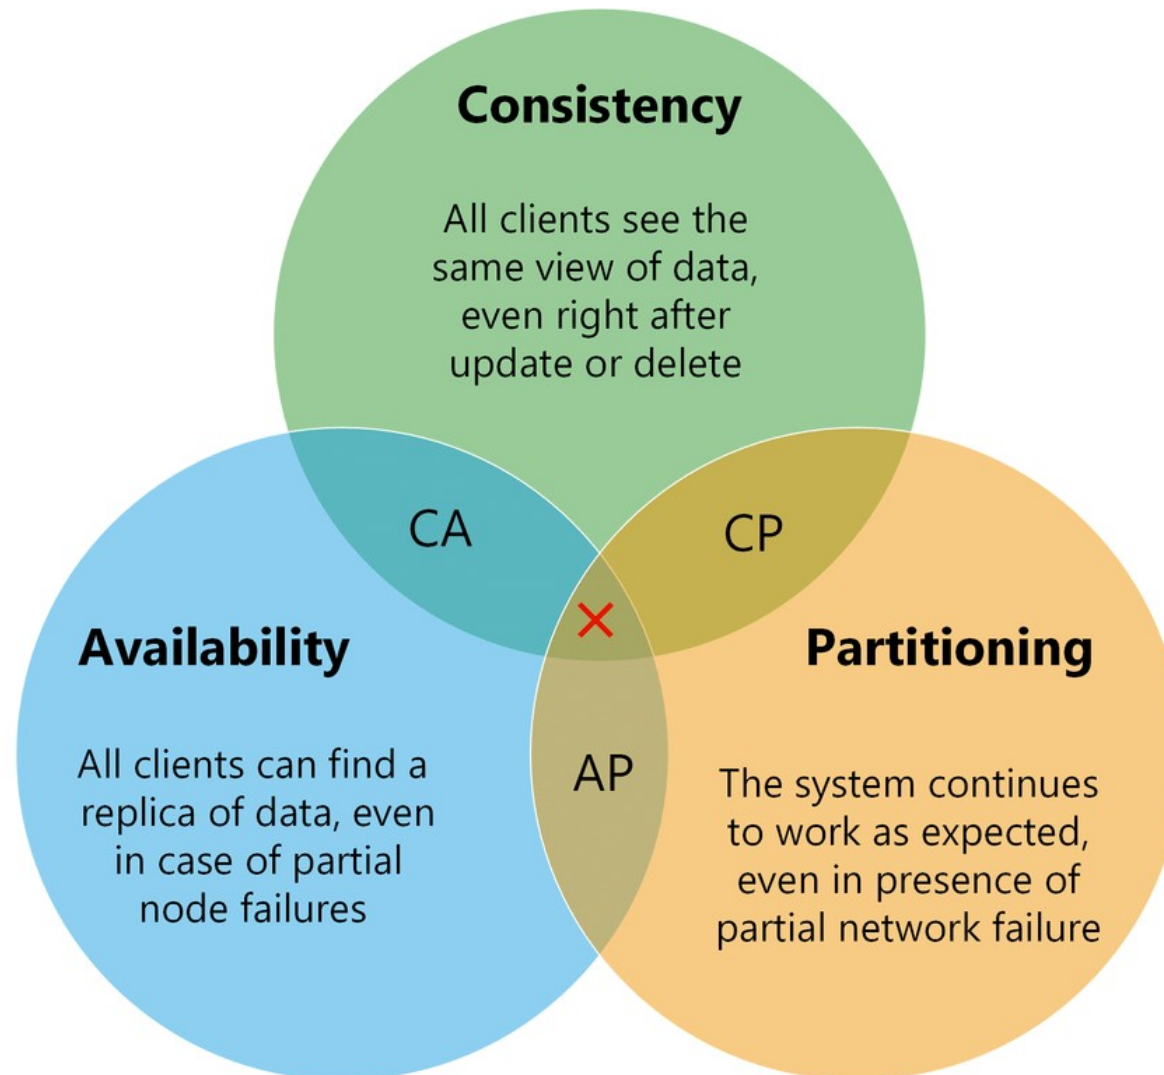


Waiting Quorum

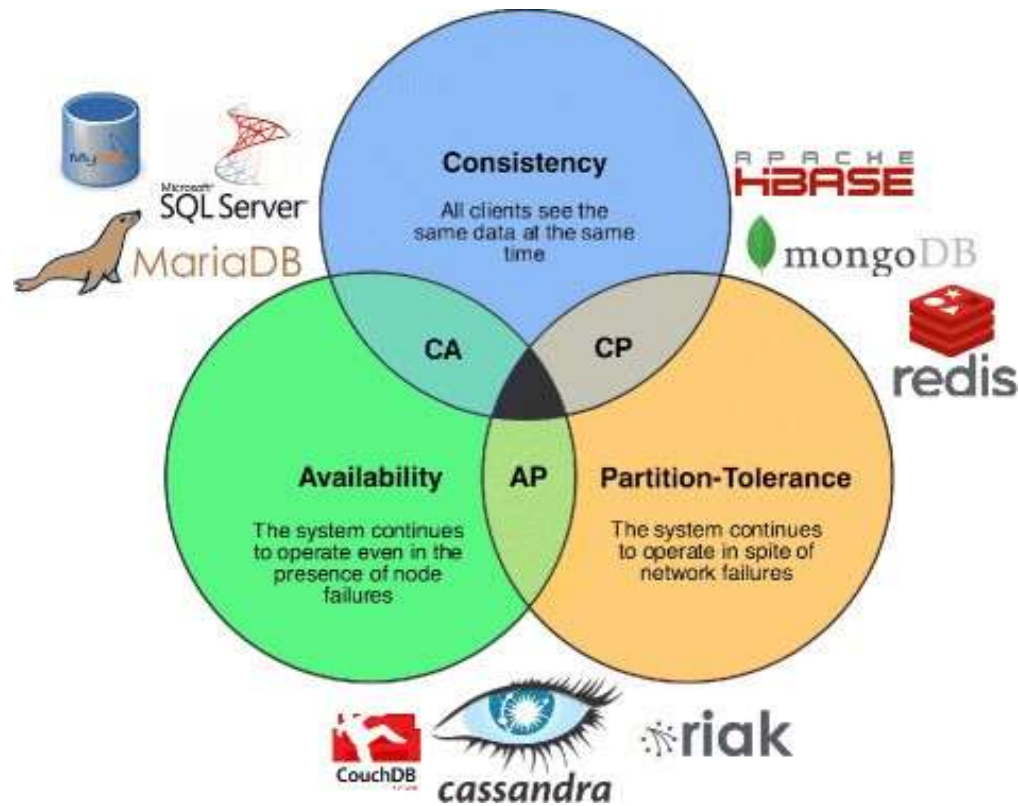
$\leq 50\%$

... means cluster 0 % working

CAP Theorem .. Choose 2 out of 3




Databases choices




« Eventually Consistent »

.. Means « NOT » always consistent ... but ends to be

Distributed Coordination Server(s)



Difficult to synchronize data on N servers
 $N \text{ servers} \Rightarrow \frac{1}{2} * N * (N-1)$ connections .. too many
Slow to organize election





Small subsystem (3,5,7 nodes)
for quorum / election
Storing coordination data
 $N \text{ servers} \Rightarrow N+3$ connections

ZooKeeper

Because Coordinating Distributed Systems is a Zoo

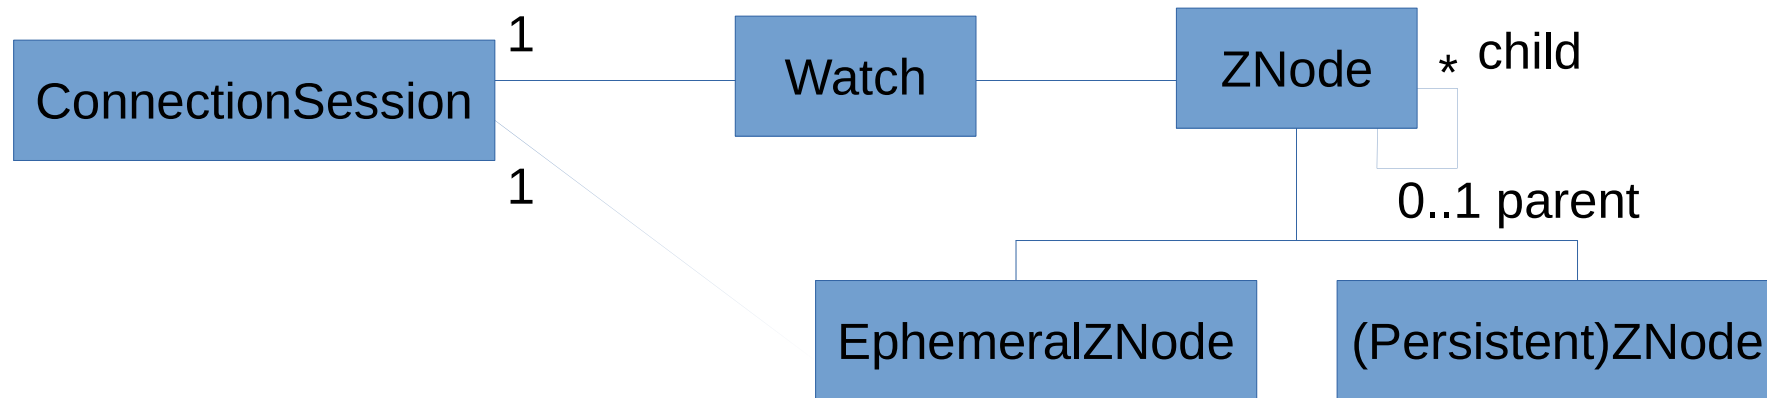


Zookeeper is **CP** (Consistent and Partition Tolerant)
Not **A** (Available)




All writes are serialized to Leader
(exclusive lock)
+ replicated to quorum for committing

ZooKeeper Features



Zookeeper kernel : « key=value » database, with atomic features
Organized as directories hierarchy, called Znode
With Listening capabilities
And Ephemeral Znode for connections

ZooKeeper



Get key..
Set key=value
List key/*
Listen Key

Znode « / »

Znode « /data »


« /prop » = value

Znode « /servers »

EphemeralZNode « /node1 »


EphemeralZNode « /node2 »

EphemeralZNode « /node3 »



Example for service discovery :
Listen « /servers/* »

ZooKeeper : for Master Election + Persist Topology/Metadata Infos



1 Master for All => does not scale



Followers IDLE
in StandBy Mode
(waste of cpu)


Sharding for Horizontal Scaling

Sharding = choose responsible server from data Id
(shared rule for clients and servers)


ServerId = Id modulo 3




Example
id=5 → Hash: 2



When $Id \% 3 = 0$



When $Id \% 3 = 1$



When $Id \% 3 = 2$



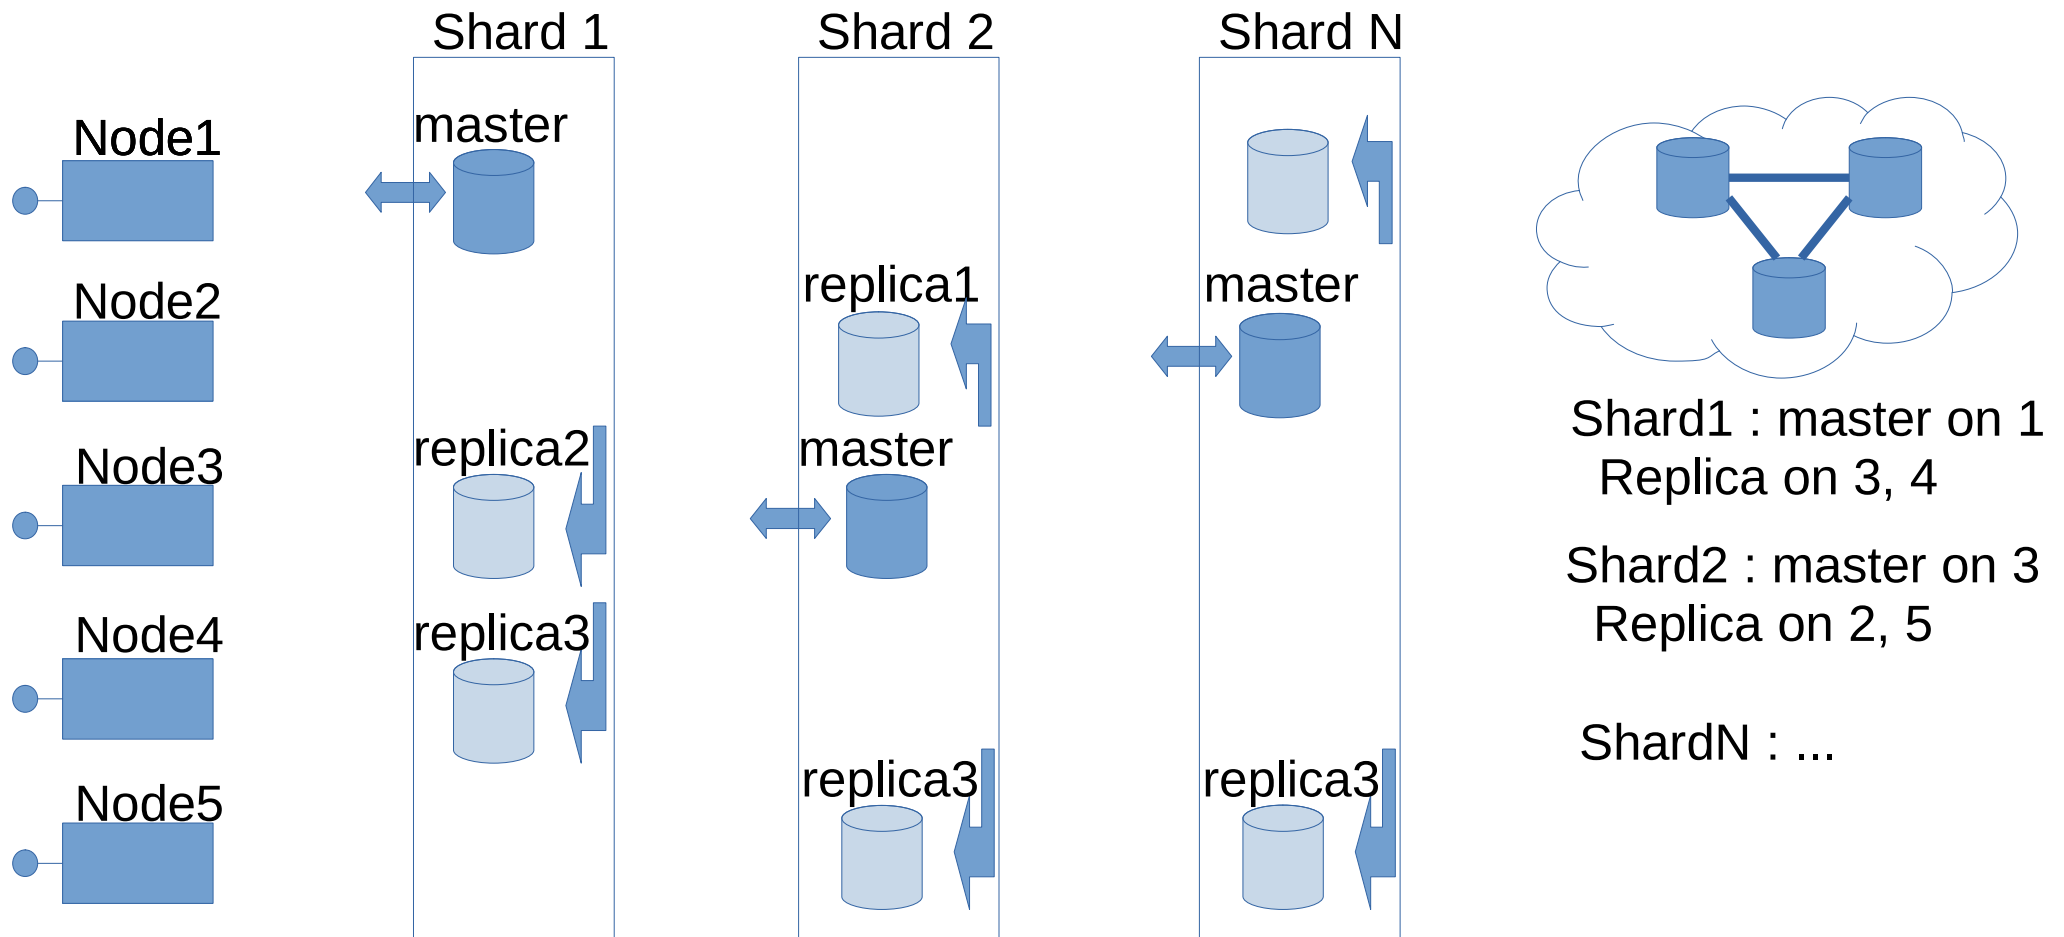
Sharding .. Pros/Cons

Scale very well « linearly » with number of servers

ONLY for request with known « ID »
(not for search / full scan)

Difficult to redimension « N » at runtime... need re-shuffle all data !

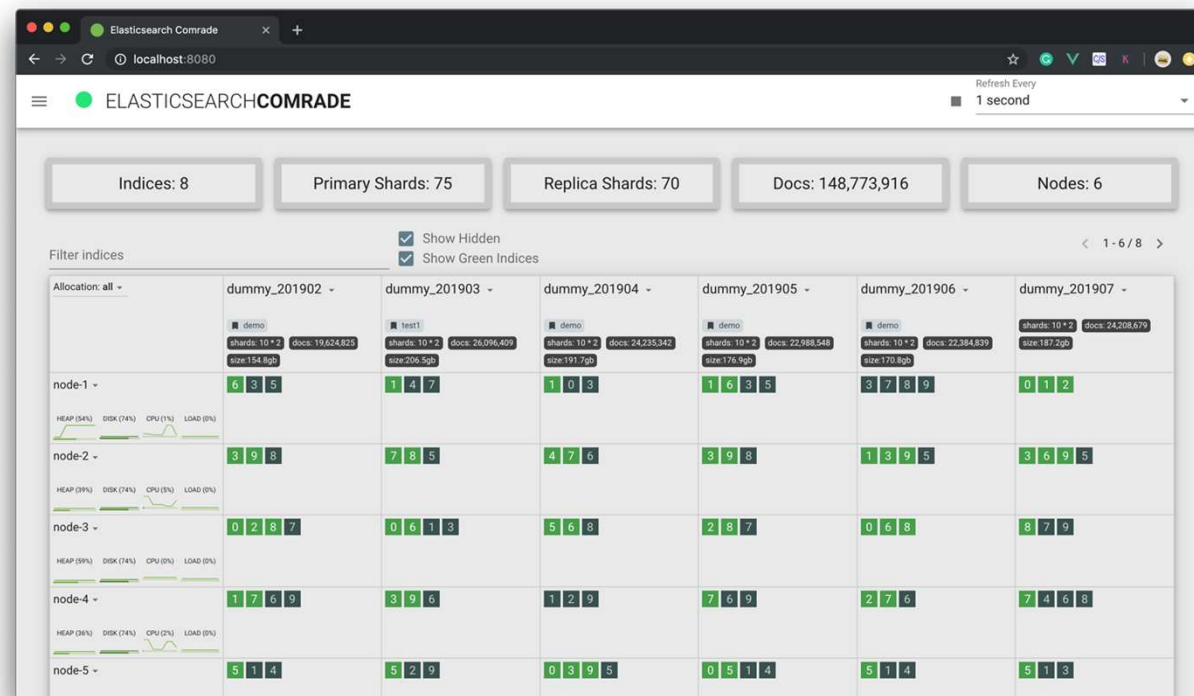
1 Master per Shard



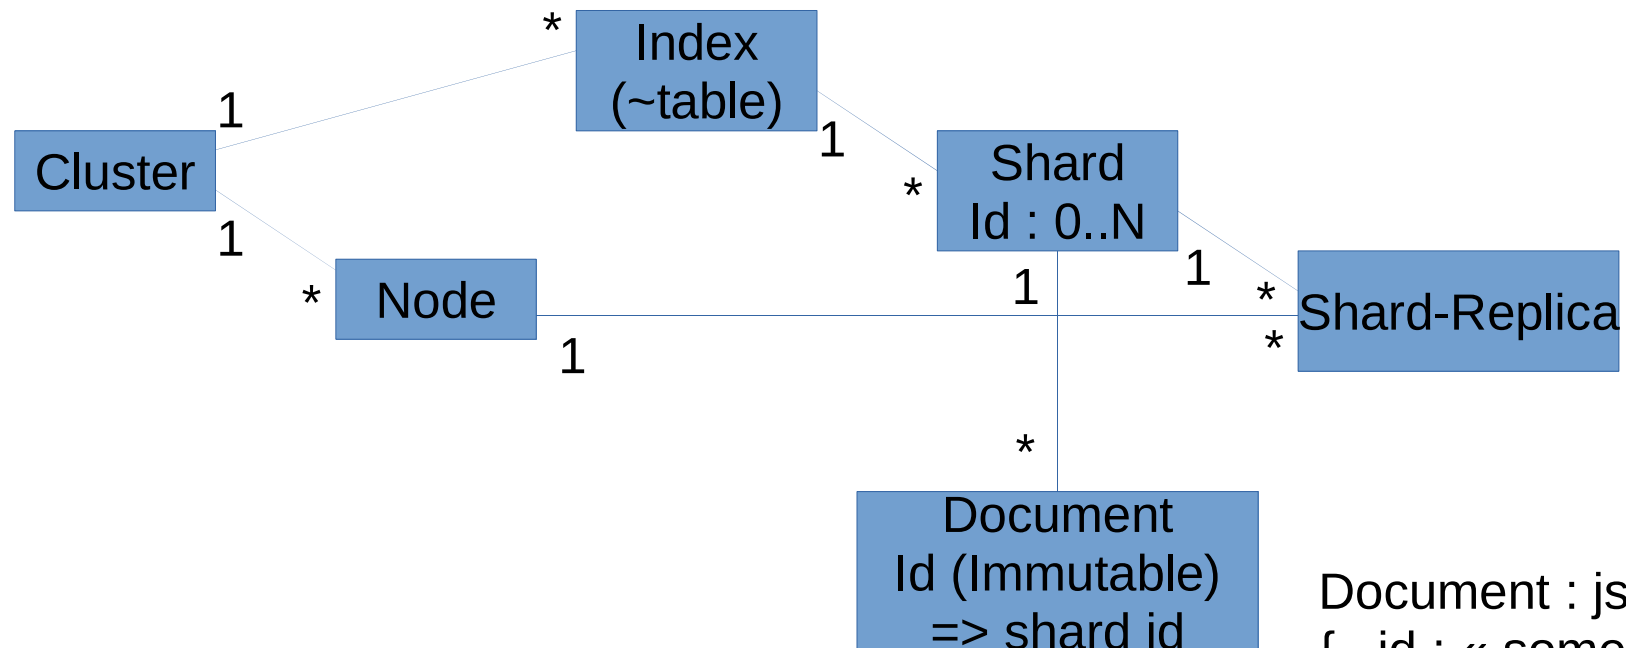
3 Examples & Comparisons for Sharding + Master/Replica



Example 1/3 : Elasticsearch



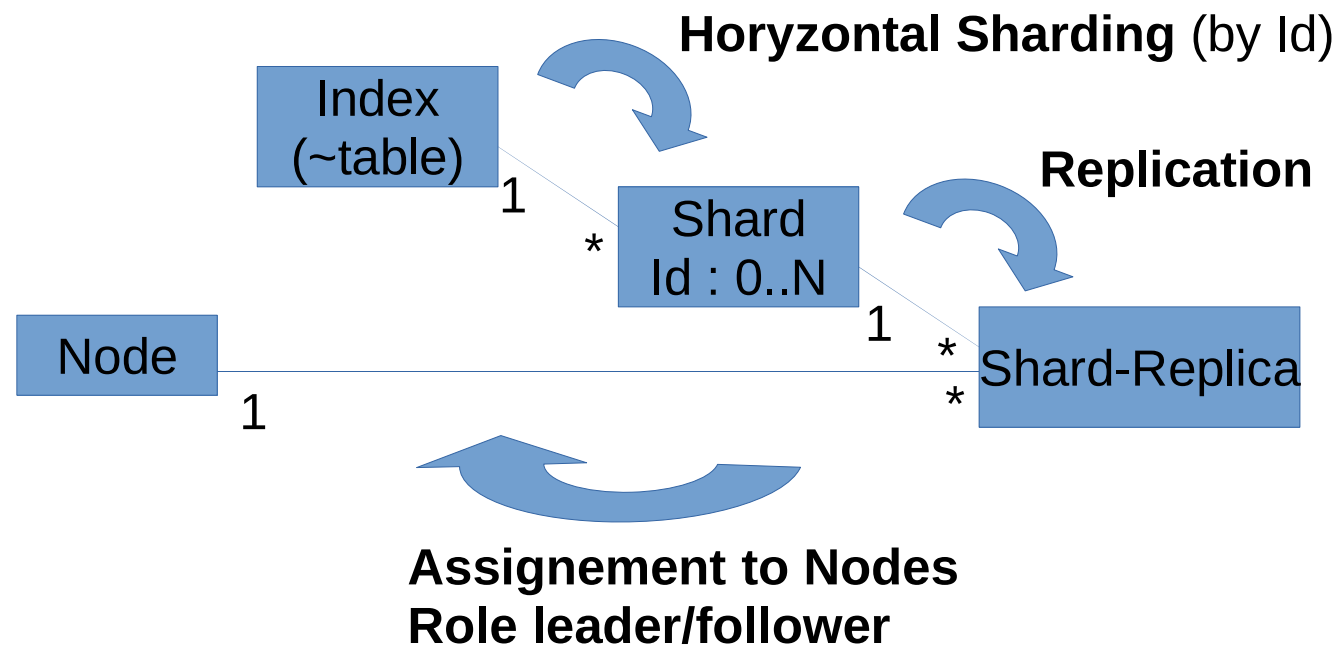
ElasticSearch ... UML model



Document : json text
{
 _id : « some-id1 »,
 field1 : 123,
 field2 : { subField : [« a »] }
}

ElasticSearch ...

Zooming Relations



HDFS



Hadoop

Overview

Datanodes

Datanode Volume Failures

Snapshot

Startup Progress

Utilities

Datanode Information

✓ In service

● Down

✂ Decommissioned

⌚ Decommissioned & dead

In operation

Show

25

entries

Search:

Node	Http Address	Last contact	Last Block Report	Capacity	Blocks	Block pool used	Version
✓ node-1:50010 (10.128.7.122:50010)	node-1:50075	4s	14m	18.74 GB	320	348.61 MB (1.82%)	2.7.3.2.6.4.0-91
✓ node-3:50010 (10.128.6.1:50010)	node-3:50075	3s	52106m	18.74 GB	1	486.1 MB (2.53%)	2.7.3.2.6.4.0-91
✓ node-4:50010 (10.129.2.1:50010)	node-4:50075	4s	29m	8.73 GB	473	596.96 MB (6.67%)	2.7.3.2.6.4.0-91
✓ node-5:50010 (10.131.2.1:50010)	node-5:50075	4s	76m	8.73 GB	715	1.32 GB (15.1%)	2.7.3.2.6.4.0-91
✓ node-7:50010 (10.128.0.1:50010)	node-7:50075	4s	52106m	8.73 GB	0	833.62 MB (9.32%)	2.7.3.2.6.4.0-91

Showing 1 to 5 of 5 entries

Previous

1

Next

File information - 6e584bed53234094b2c2e1cc1b443b63

Download

Block information - Block 0

Block ID: 1373741865

Block Pool ID: BP-1905800902-192.136.114.25-1494702271527

Generation Stamp: 1195

Size: 6053

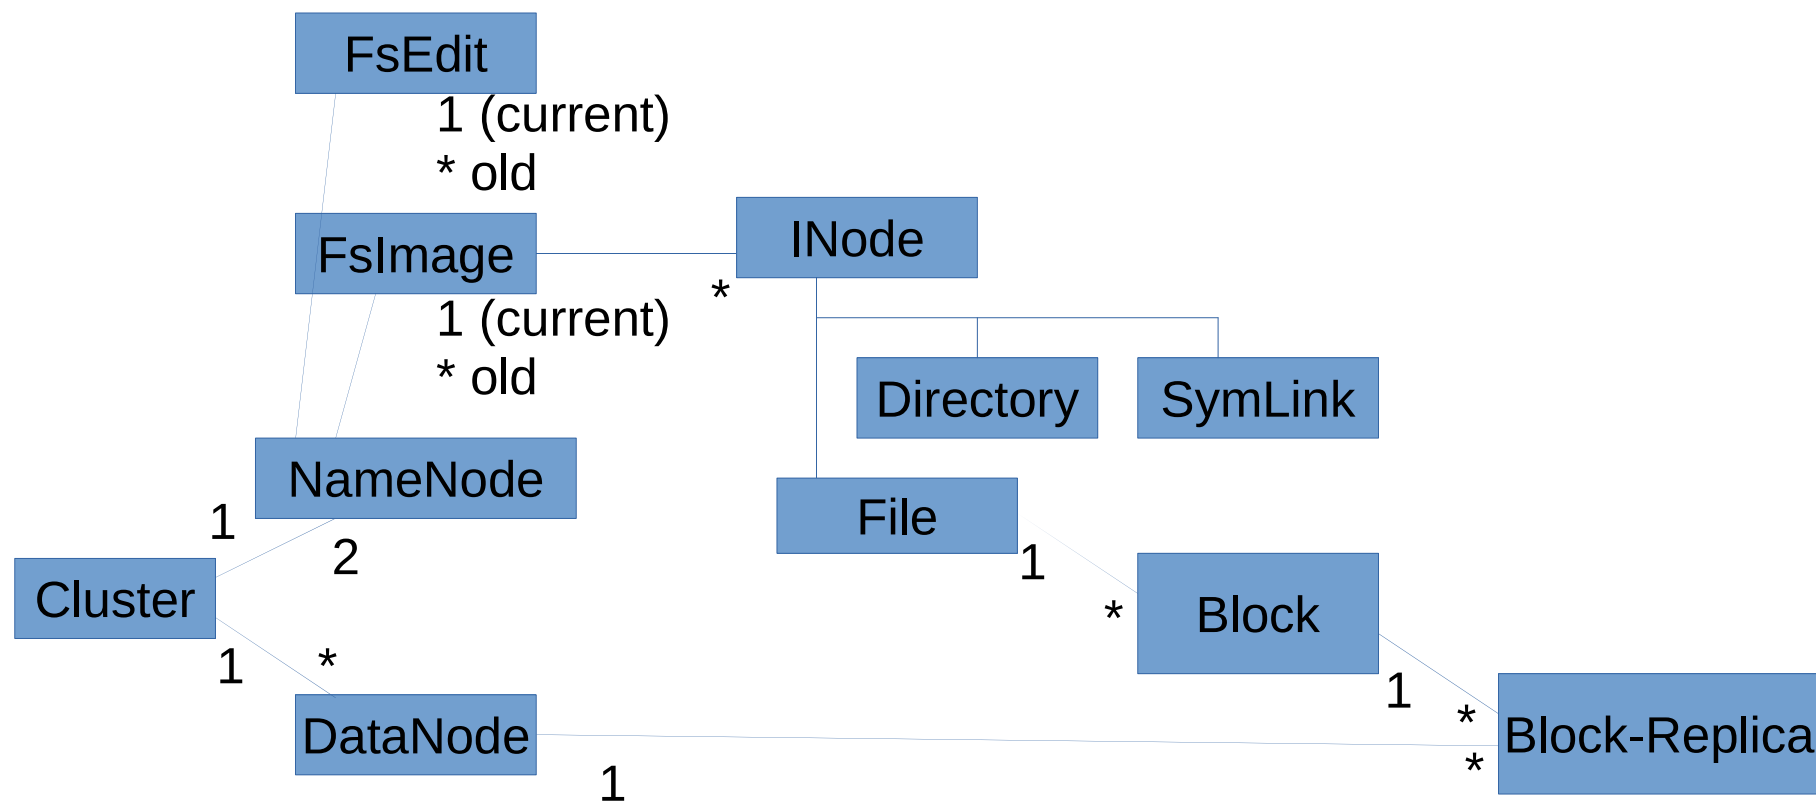
Availability:

- node-1:50010 (10.128.7.122:50010)
- node-3:50010 (10.128.6.1:50010)
- node-4:50010 (10.129.2.1:50010)

Close

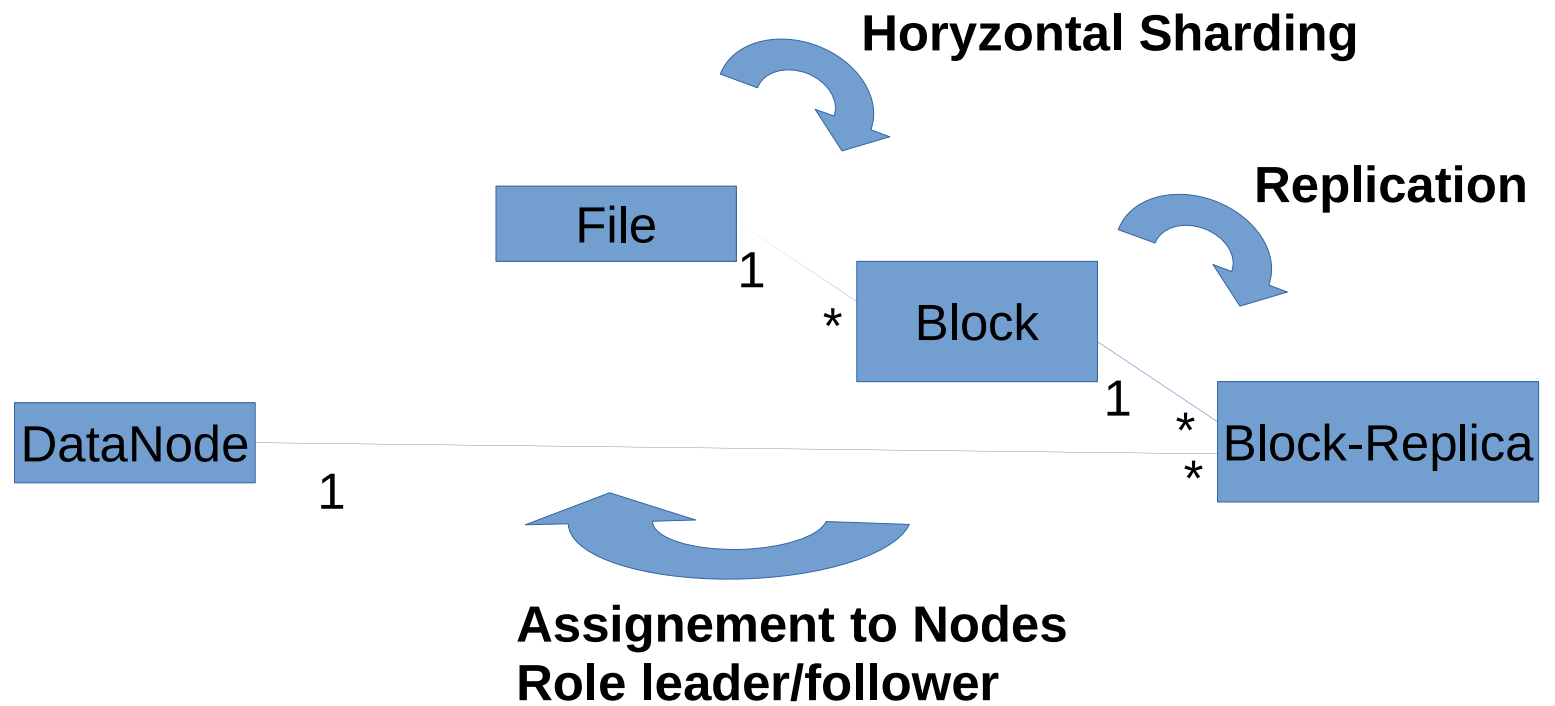


HDFS ... UML Model



File: binary large object « blob »
= metadata
(owner, group, chmod, acl, ..)
+ list of blocks

HDFS ... Zoomin Relations



HBase



Master hdp1

Region Servers

ServerName	Start time	Version	Requests Per Second	Num. Region Servers
hdp1.586025.1556711447905	Wed Feb 20 20:13:47 EST 2019	1.2.0-cdh5.7.6	8	5
hdp1.586025.1556711447460	Wed Feb 20 20:20:13:47 EST 2019	1.2.0-cdh5.7.6	8	4
hdp1.586025.1556711447460	Wed Feb 20 20:20:13:47 EST 2019	1.2.0-cdh5.7.6	8	3
Total			8	12

Tables

Namespace	Table Name	Online Regions	Offline Regions	Failed Regions	Split Regions	Other Regions	Description
default	test	1	0	0	0	0	'test', (NAME => 'cf')
default	test2	1	0	0	0	0	'test2', (NAME => 'id'), (NAME => 'temp')

HBase ... UML Model

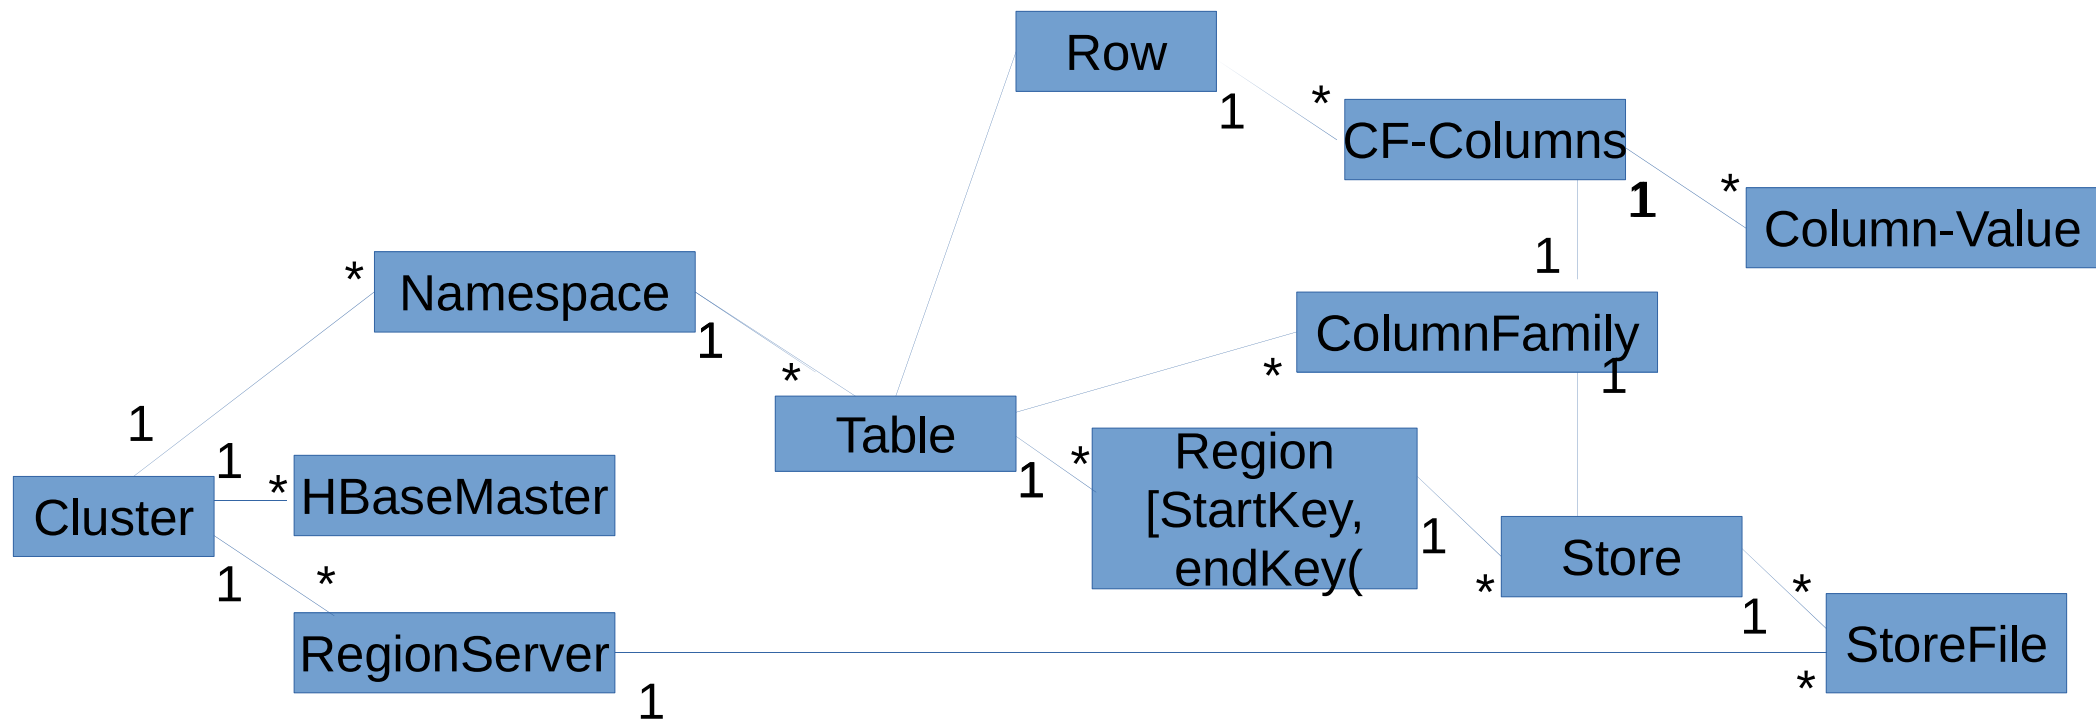


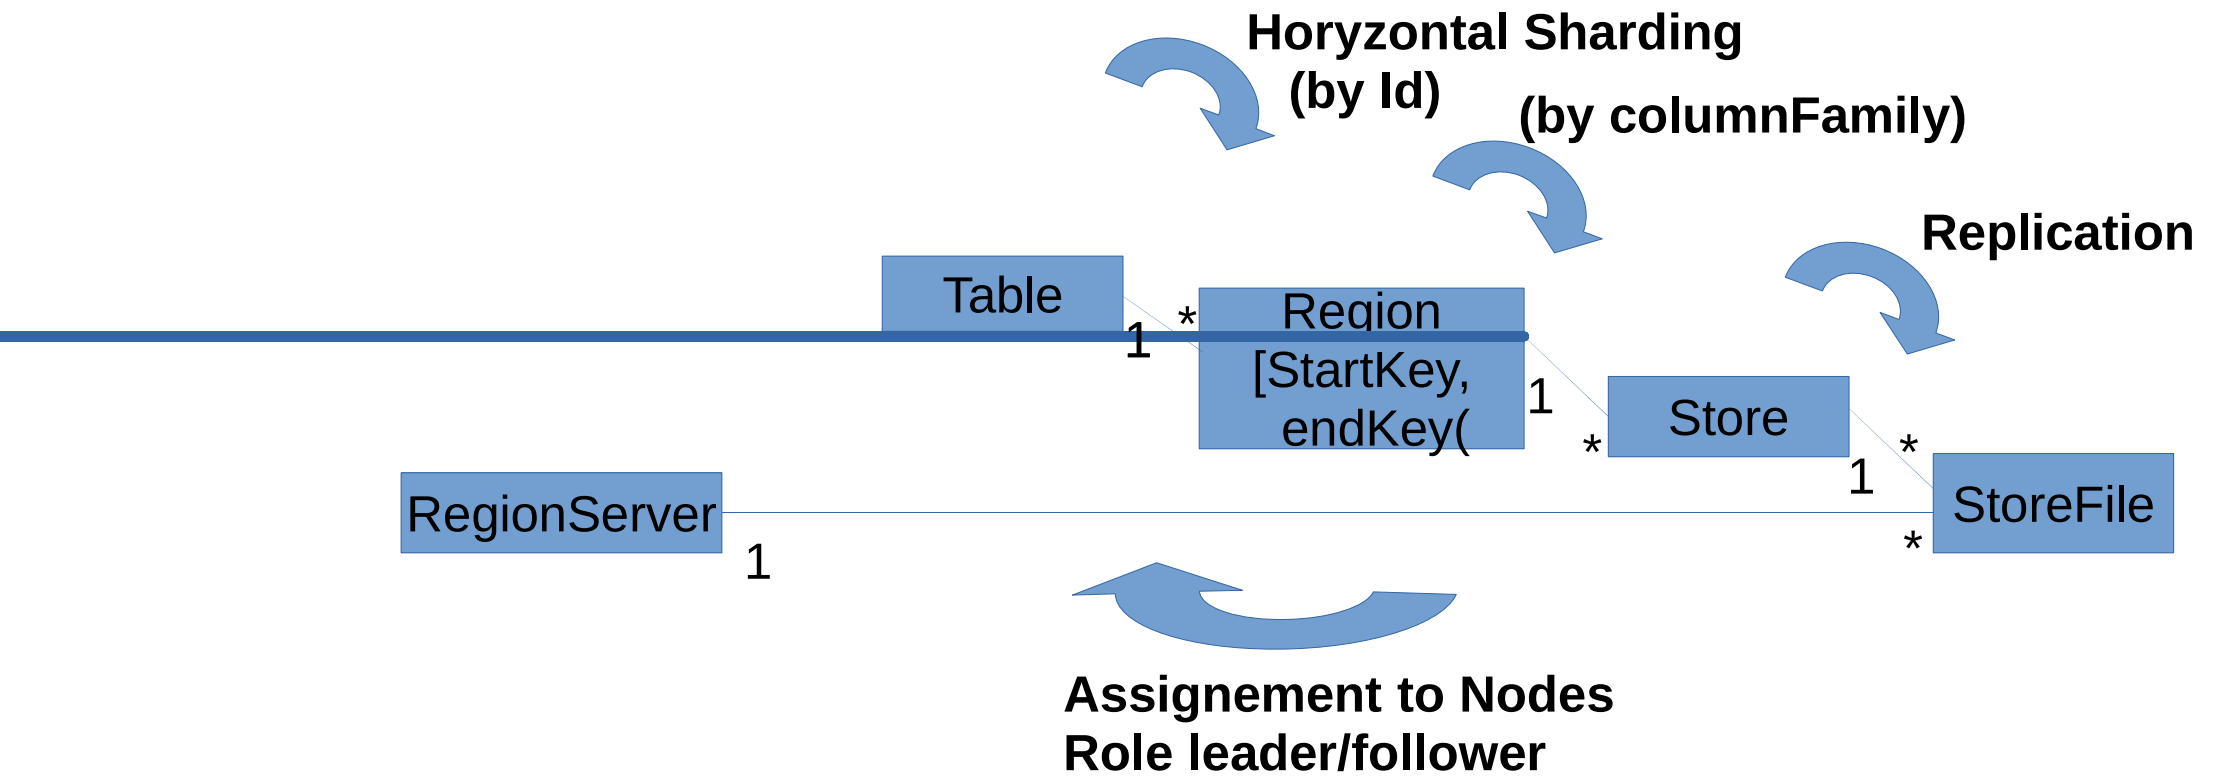
Table = Map<byte[], Map<byte[], byte[]>>

RowKey

ColumnFamily:Column

BlobValue

HBase ... Zooming Relations



Questions ?