Datalake File Format: Parquet

arnaud.nauwynck@gmail.com

Course Esilv 2026

This document:
https://github.com/Arnaud-Nauwynck/presentations/pres-bigdata/Datalake-File-Format-Parquet

Outline

- Parquet Caracteristics
  - Structured, with Schema
  - Data – Metadata (footer)
  - Immutable
  - Columnar ... Column Pruning
  - Splittable ... spark Dataset Partitions
  - Compression
  - Encoding
  - Statistics, Bloom ... spark Predicate-Push-Down
  - Optimize write once, read many
    ... spark dataset.repartition().sortWithinPartition().write

Parquet is a Structured Format
Strongly Typed, Schema-based

File Formats

Unstructured

Text

```
text line 1 \n
text line 2 \n
```

Csv

```
Col1;Col2;Col3\n
a1;b1;c1\n
a2;b2;c2\n
```

Semi-Structured

Json

```
{"a":"a1","b":"b1"} \n
{"a":"a2","b":"b2"} \n
```

Xml

```
<elt>
  <a>a1</a>
  <b>b1</b>
```

Structured

Serialization Structured

Avro,Thrift,Protobuf

```
schema:XX,
value:0101010101
```

Columnar Structured

Orc, Parquet

Structured : struct<>, array<>, map<>

Scalar Value

(= terminal element in grammar)

= primitive data-type

String
boolean
int
double
Date
…

Composite Value

struct<a:Type1, b:Type2, ...>

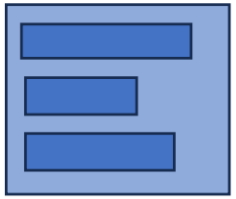

array<ElementType>

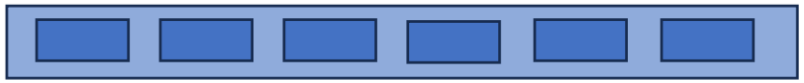

map<KeyType,ValueType>

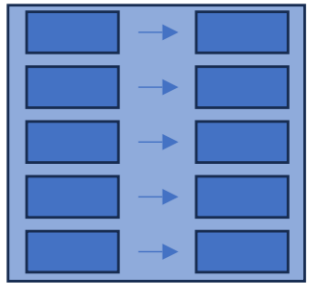

Type constraint = Schema

semi-structured
schema is optional

{ JSON } → JSON Schema (optional)

<xml> → XSD Schema (optional)

structured
schema is strict

.avro → .avsc = avro schema
included in avro datafile format header
implicit in avro client-server exchanges

thrift
protobuf

.parquet → included in parquet file footer (= metadata)

Parquet SDK / Converter / ObjectModel

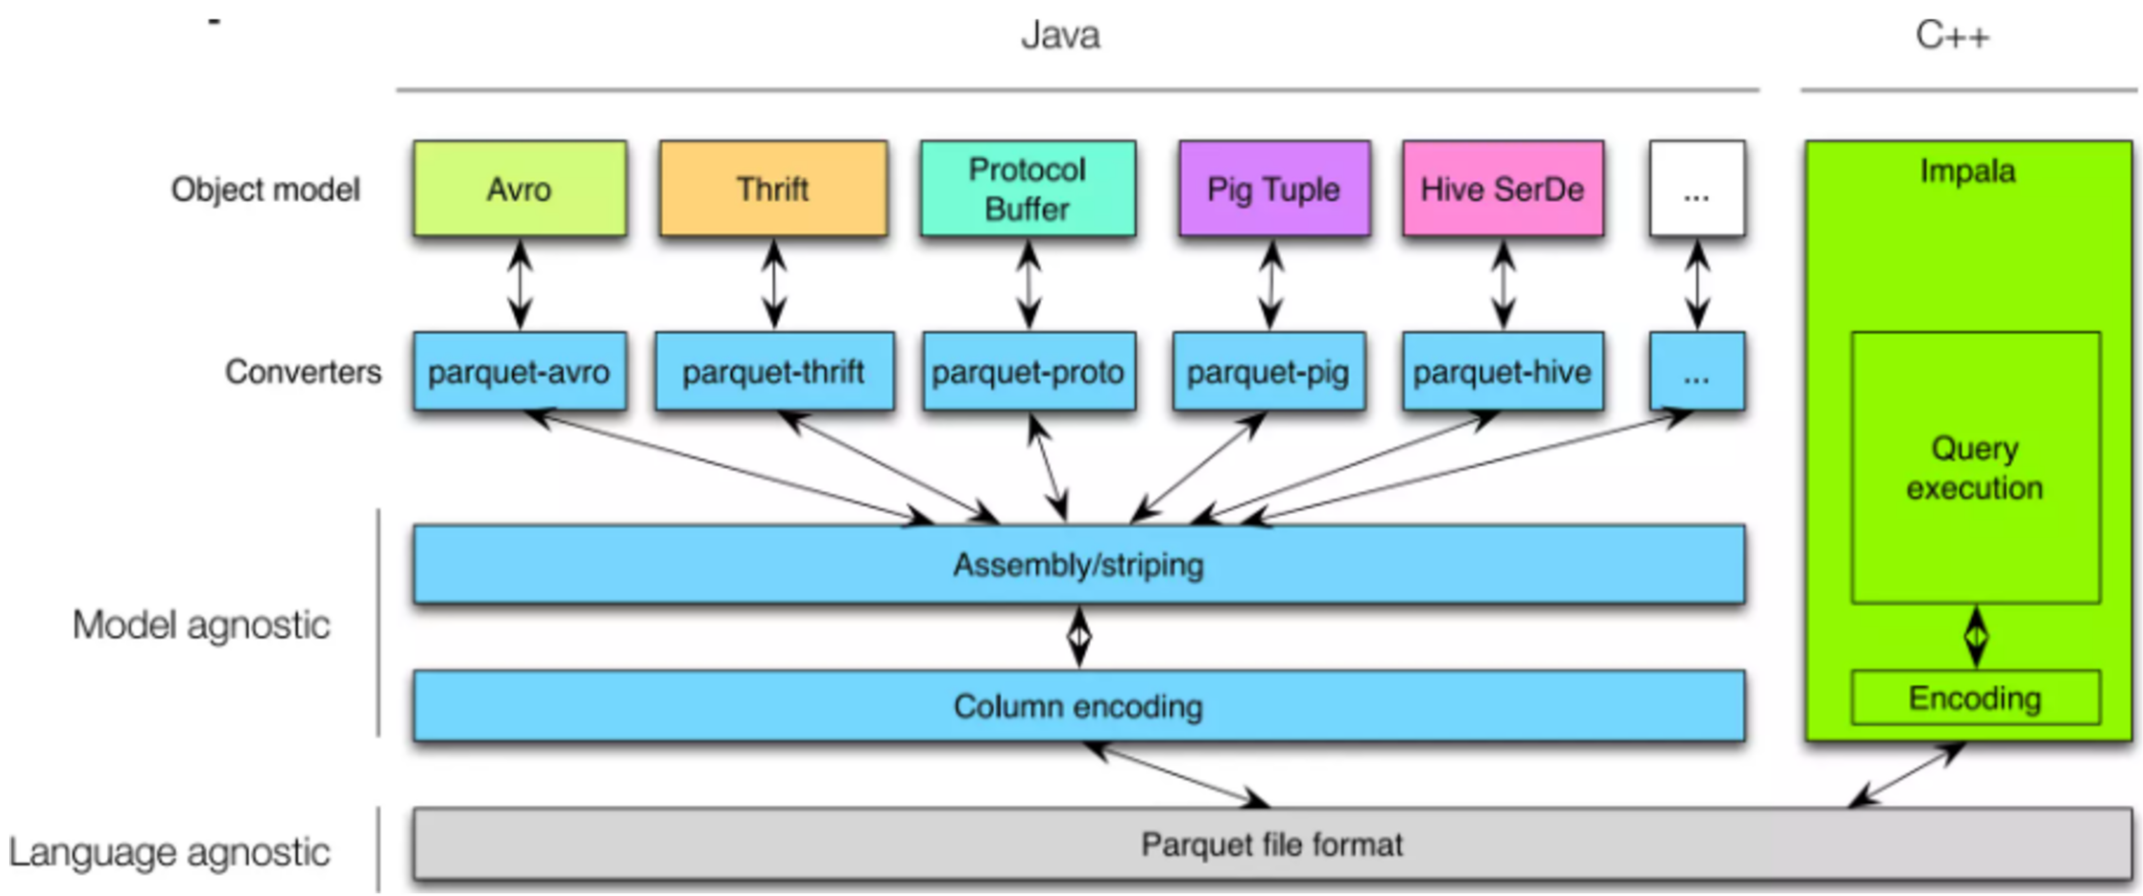

Parquet separates Data and MetaData

Parquet Schema : in metadata footer

data part
(usually N x blocks of 128 Mo)

footer: metadata part
(usually <= 50 ko)

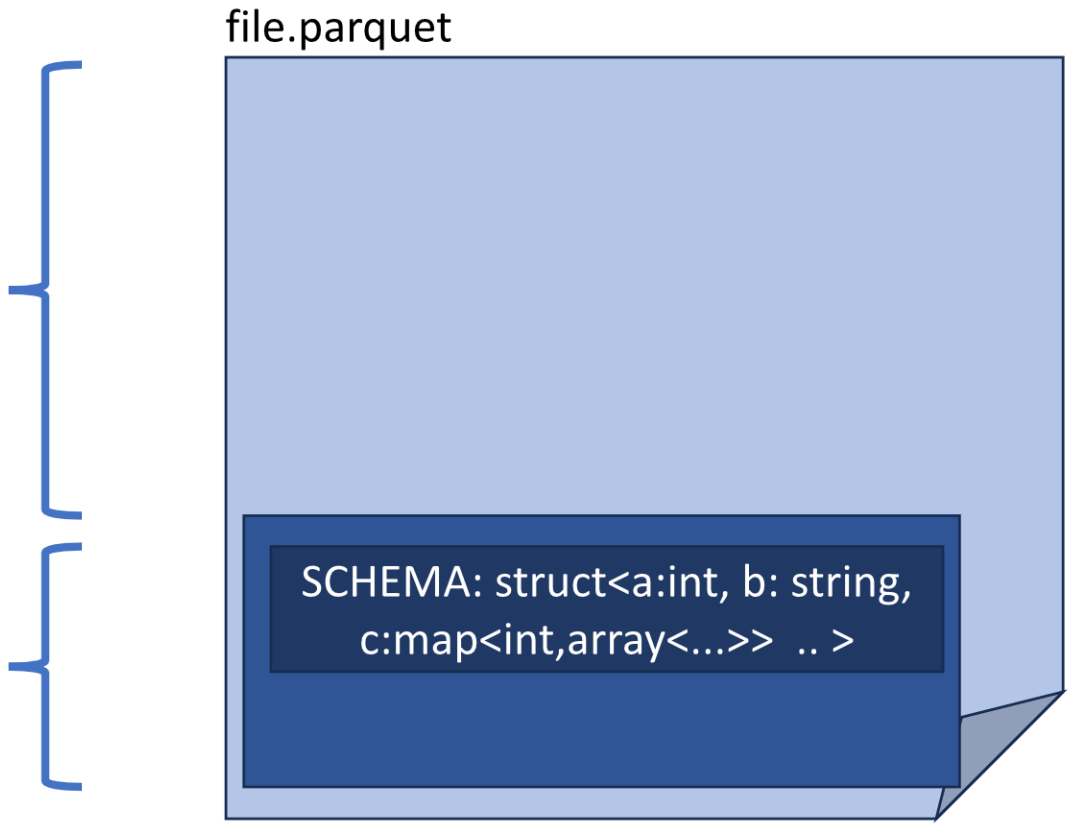

Parquet for Small Files ?

data part can be small
(even empty)

parquet files
ALWAYS have an
uncompressable footer,
containing SCHEMA

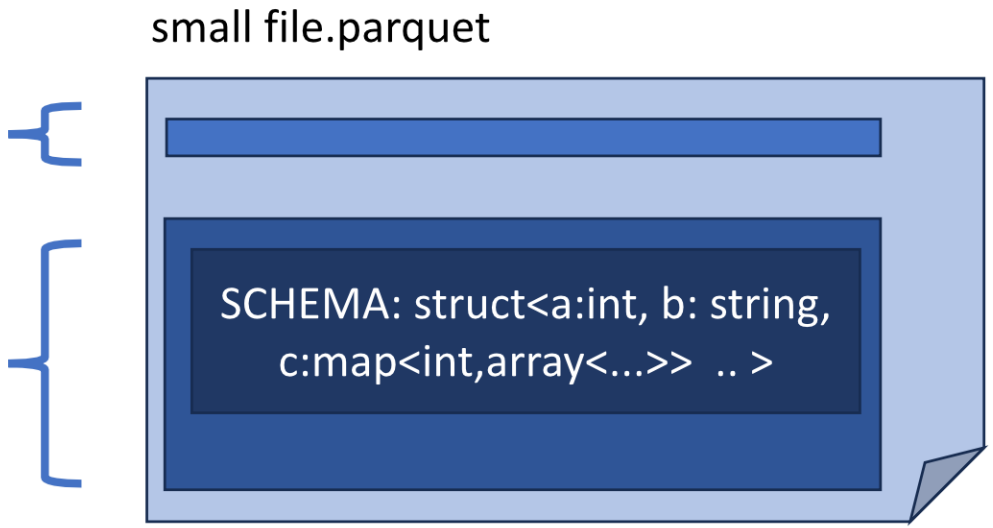

Parquet for small data is NOT efficient
bad ratio of "data / metadata" size

Read Parquet footer only

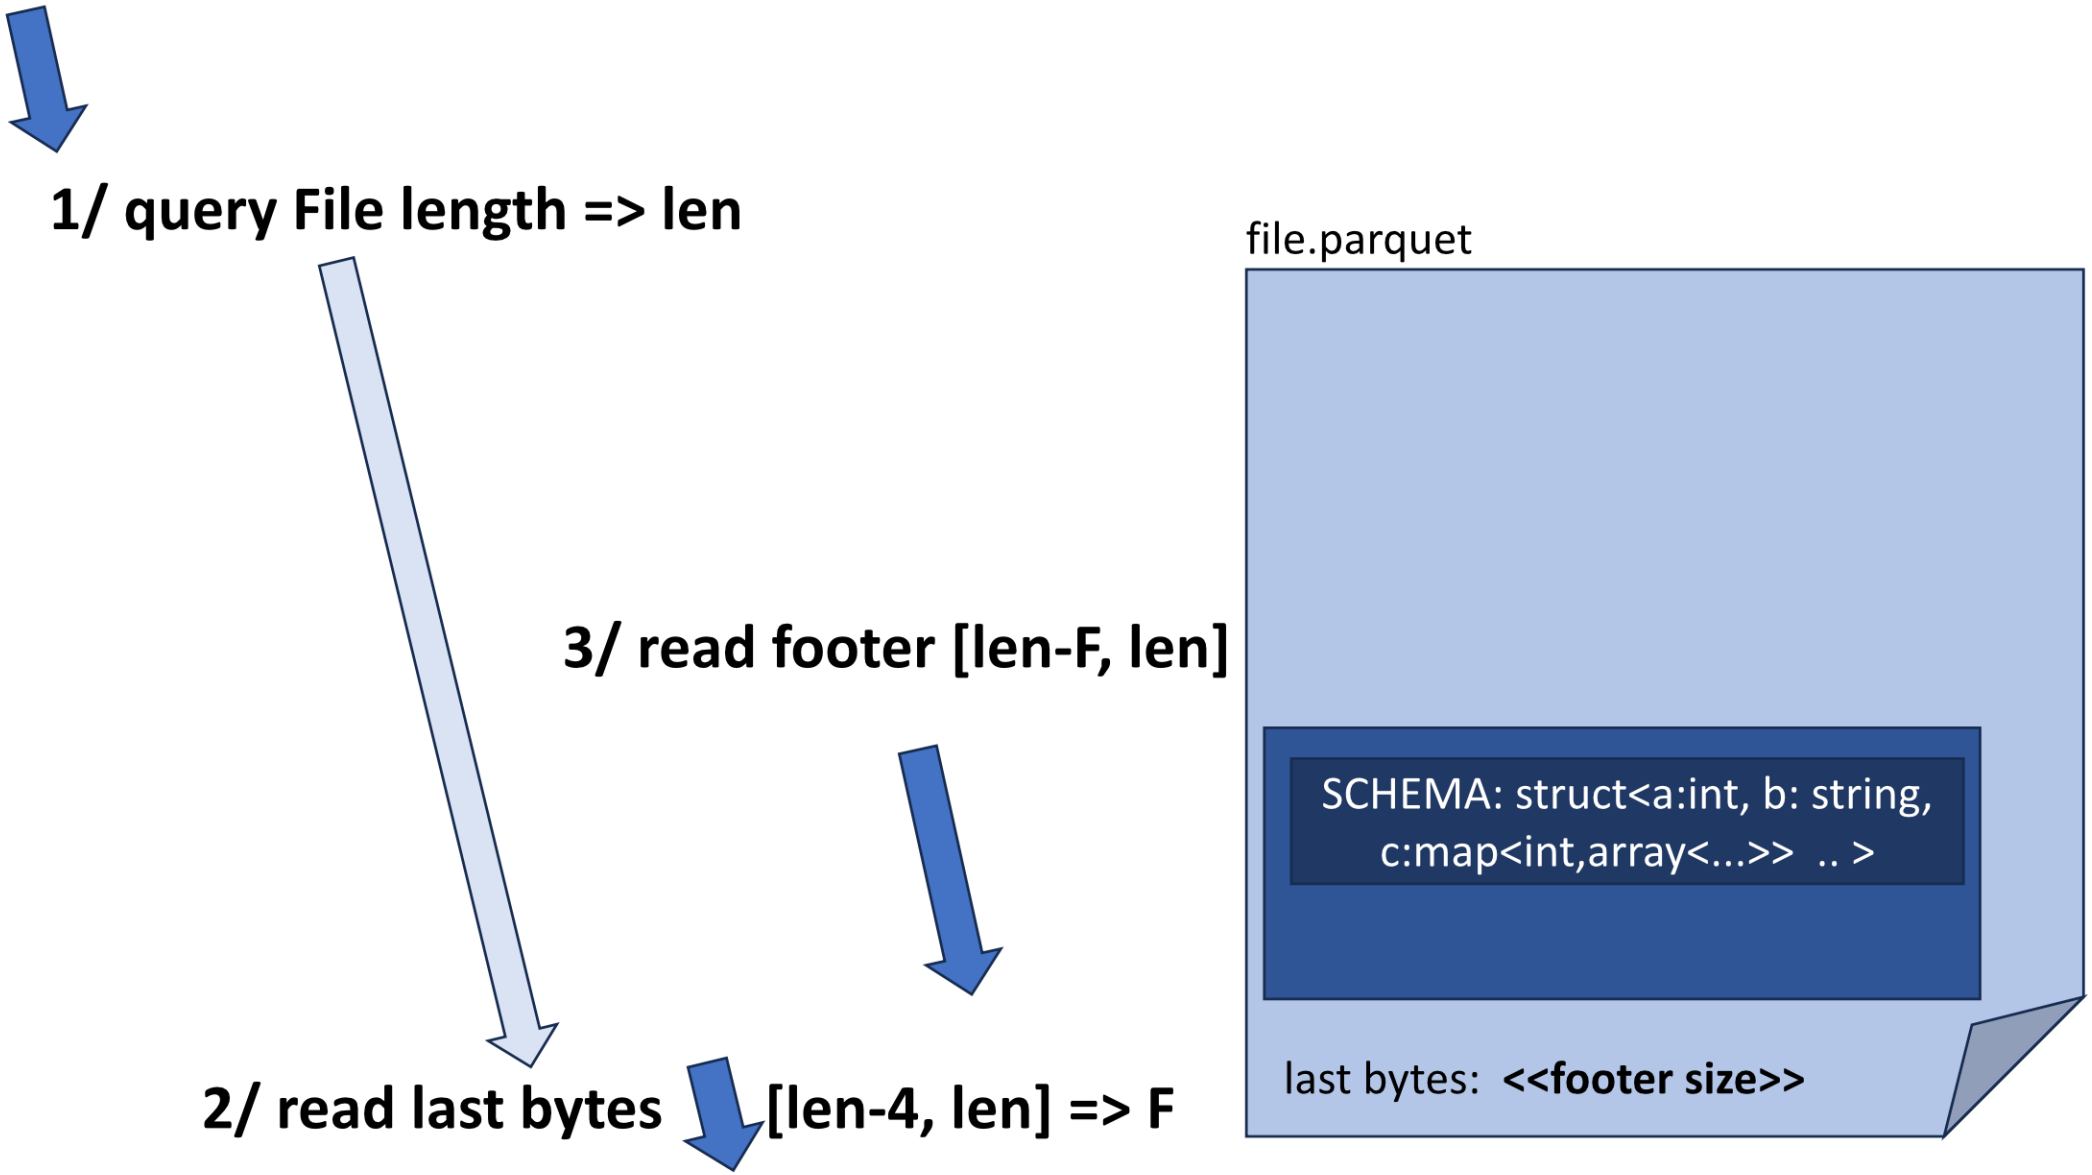

Why Footer instead of Header?

For reader : not a big overhead, ONLY 3 calls to read

For writer : MUCH more practical to "stream" write Nx rows,
keep in-memory only few metadata,
to flush write at end

Schema Evolution

- Adding columns: old files read as NULL
- Renaming columns: seen as drop + add
- Mismatched types: fail, or widened if compatible

file_v1
id, name

file_v2
id, name, email

file_v3
id, full_name, email

File	id	name	email	full_name
file_v1	1	Ann	NULL	NULL
file_v2	2	Bob	bob@x.com	NULL
file_v3	3	NULL	cy@x.com	Cy C.

Type conflict: age is int in v1, string in v4 → read fails, or type widened

mergeSchema

Default = NO mergeSchema

only 1 file schema read

= FASTER

MergeSchema =true

spark.read.option("mergeSchema", "true").load(..)

ALL file schemas read, then merged

= SLOWER, but exhaustive, support evolution

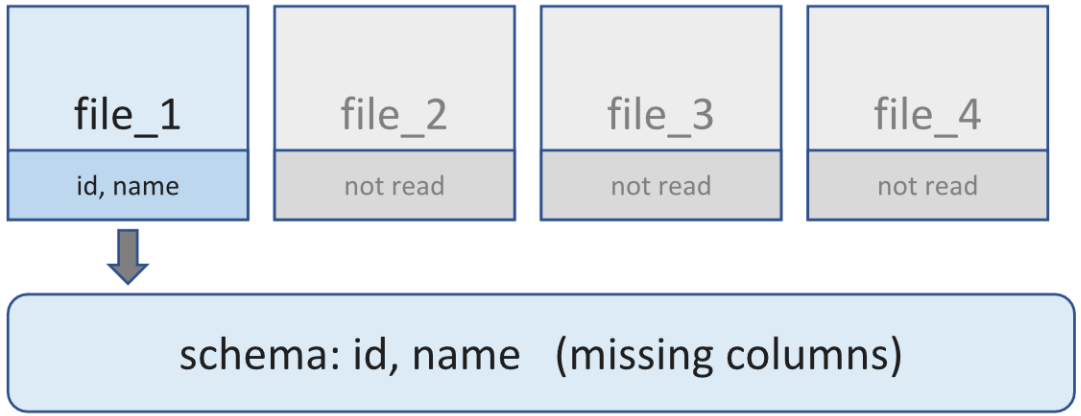

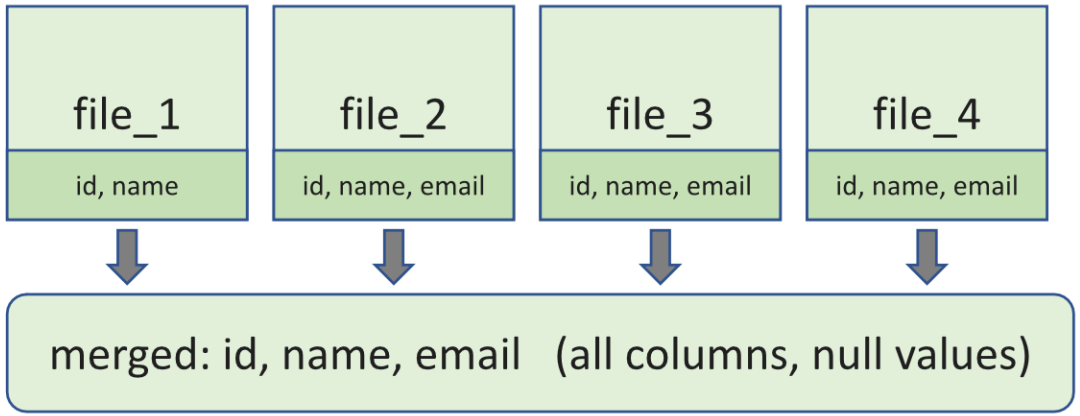

Parquet Files Schema Coercion from HiveMetastore Table DDL

Expected schema (Hive Metastore)

```
CREATE TABLE sales.users (
  id INT, name STRING, email STRING, age INT
) STORED AS PARQUET;
```

Actual file footers (Parquet)

file_1	file_2	file_3
id INT, name STRING	id INT, name STRING, email STRING, age STRING	id INT, name STRING, email STRING, age INT, city STRING

Match columns by name, not position · missing column → NULL · extra column → ignored
compatible type → cast (age "42" → 42)

id	name	email	age
1	Ann	NULL	NULL
2	Bob	bob@x.com	42 (cast)
3	Cy	cy@x.com	35 (city ignored)

file_4: age = "abc" → cast fails (read error)

Parquet is Immutable

Once Written (footer flushed), Data is Immutable = NO in-place Update

Update a row in-place: impossible

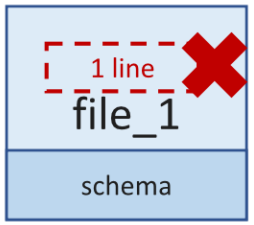

Update = copy the full file

1. copy + update 1 line

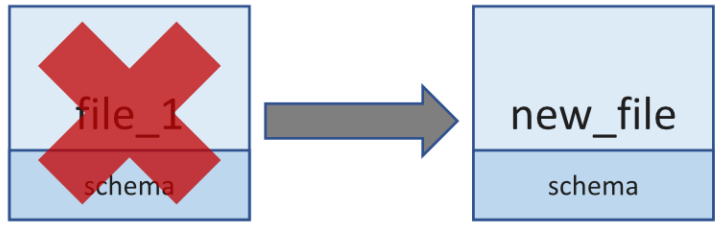

2. Delete old file_1

3. Create new file_2

Versionning Parquet in Legacy “Partition” Directories Hive Table

Hive table “path/day={yyyy-MM-dd}/*.parquet”
append = add a new file in the day folder

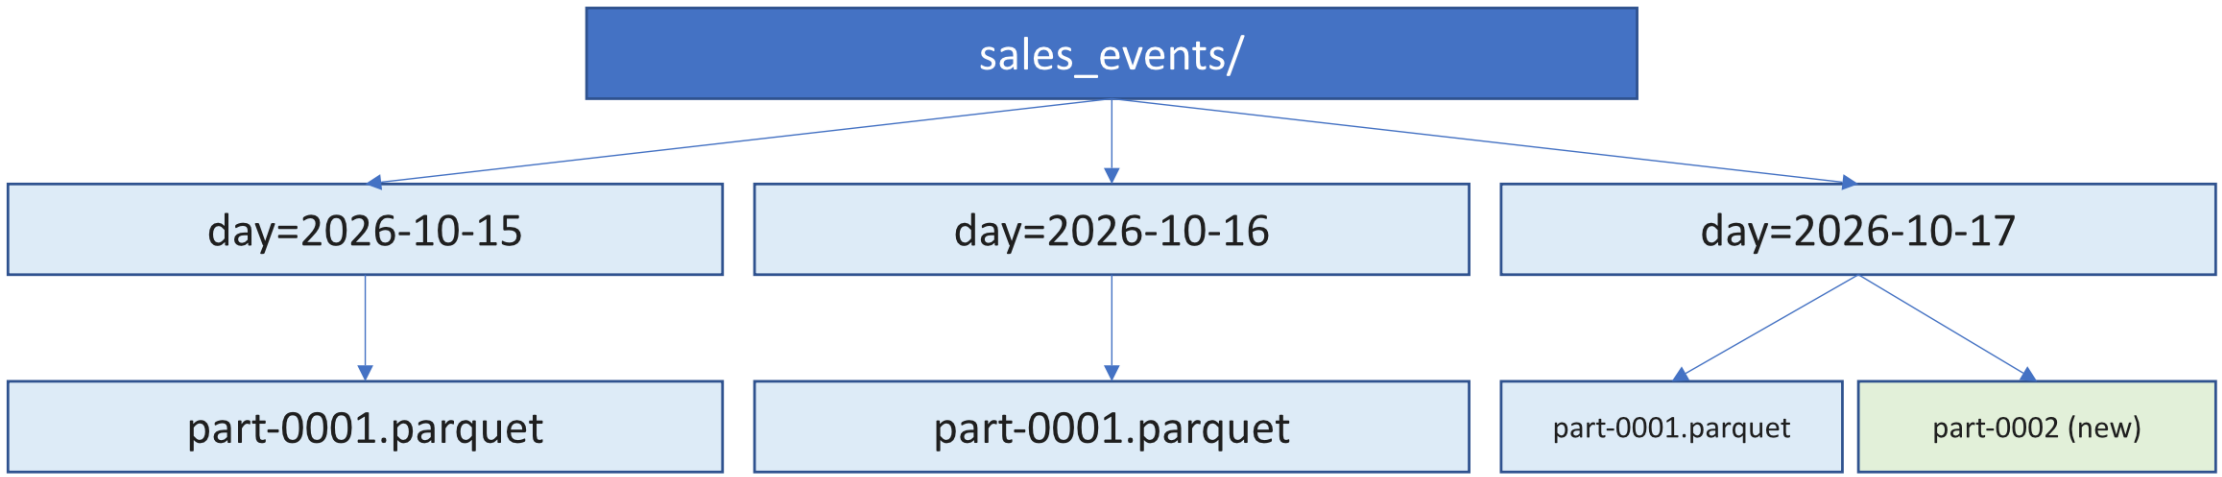

Existing files are untouched: the new file is simply added to today's folder.

Versionning Data:
Parquet within Iceberg

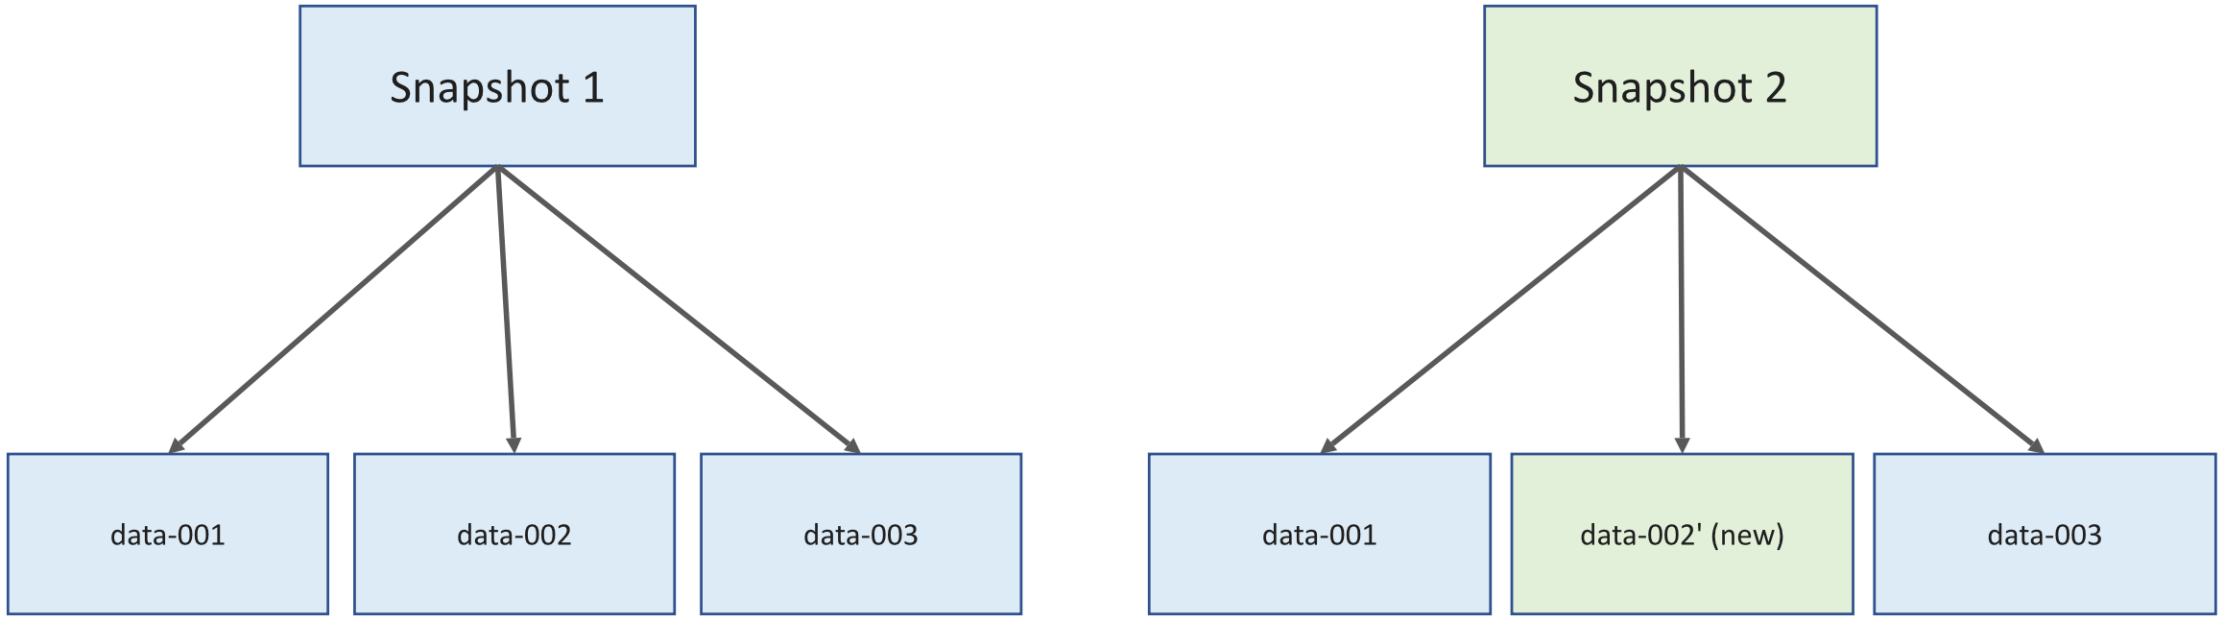

Parquet files are immutable: **Iceberg swaps the pointer, not the bytes.**
Old data-002 is no longer referenced from “current”.
Still visible “as-of” from previous version

Iceberg : Versionning System (like “Git”) over Parquet

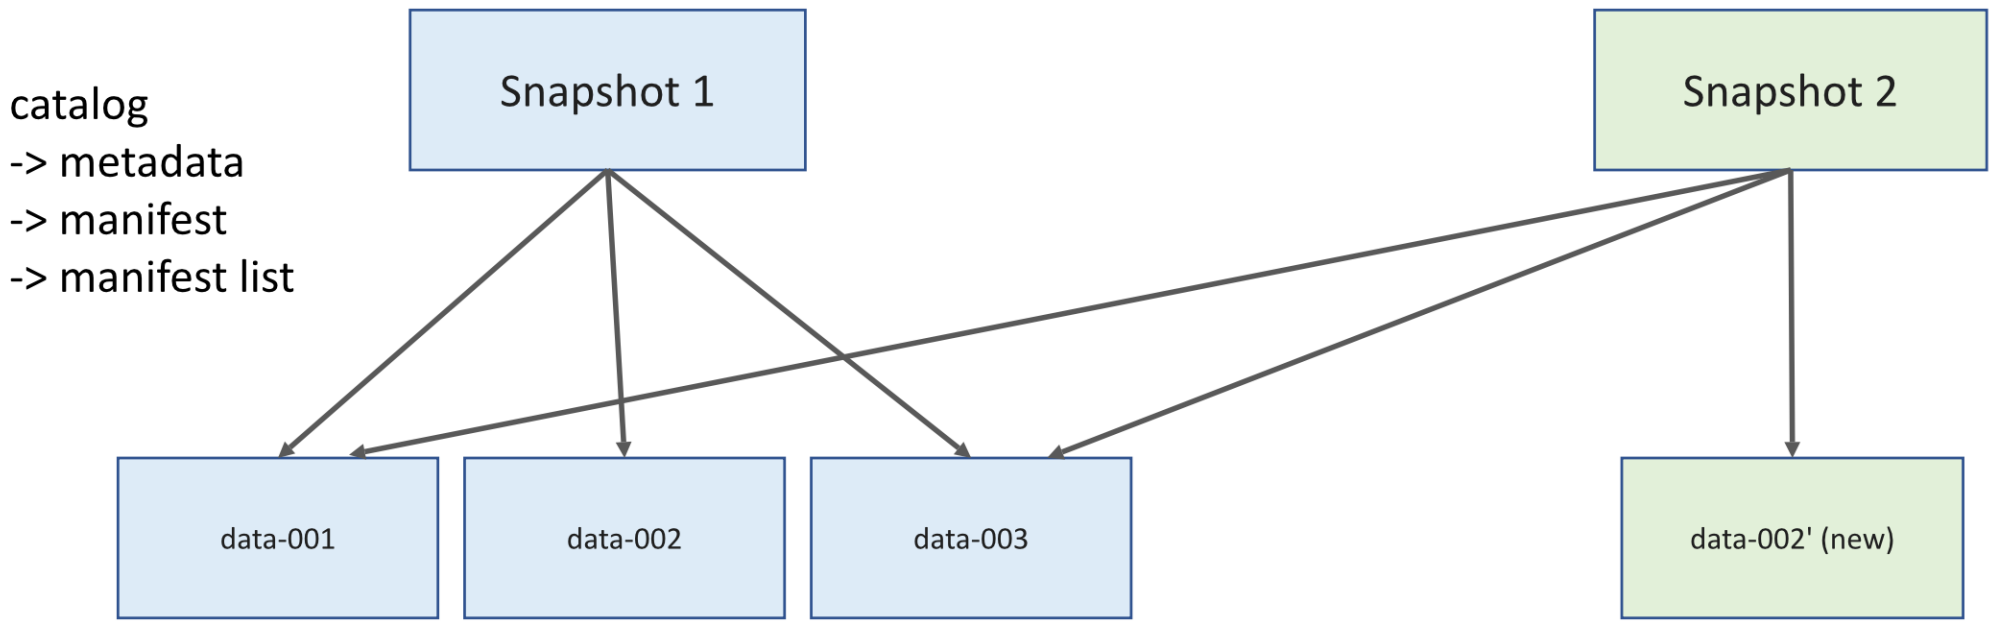

more at …. **https://iceberg.apache.org/**

Parquet is Columnar

Parquet : Columnar File Format

Logical view: rows - columns

a1	b1	c1	d1
a2	b2	c2	d2
a3	b3	c3	d3
a4	b4	c4	d4

On Disk Row Serialization: like CSV, JSON, AVRO,

On Disk: **Columnar**

Columnar
=> better memory aligned,
Vectorized CPU pipeline

```
struct A {

  boolean f1;   // <= 1 bit (on 1 byte)
                // in-memory padding 3 bytes
  int     f2;  // 4 bytes, aligned on multiple of 4
                //  in-memory padding 4 bytes
  long    f3  // 8 bytes
}
```

array[struct<..>] <=> **array[boolean]**
array[int]
array[long]

Vectorized Reader **~9x Faster**

Vectorized code … fewer "If", "Loop", "calls"

Better CPU Pipeline

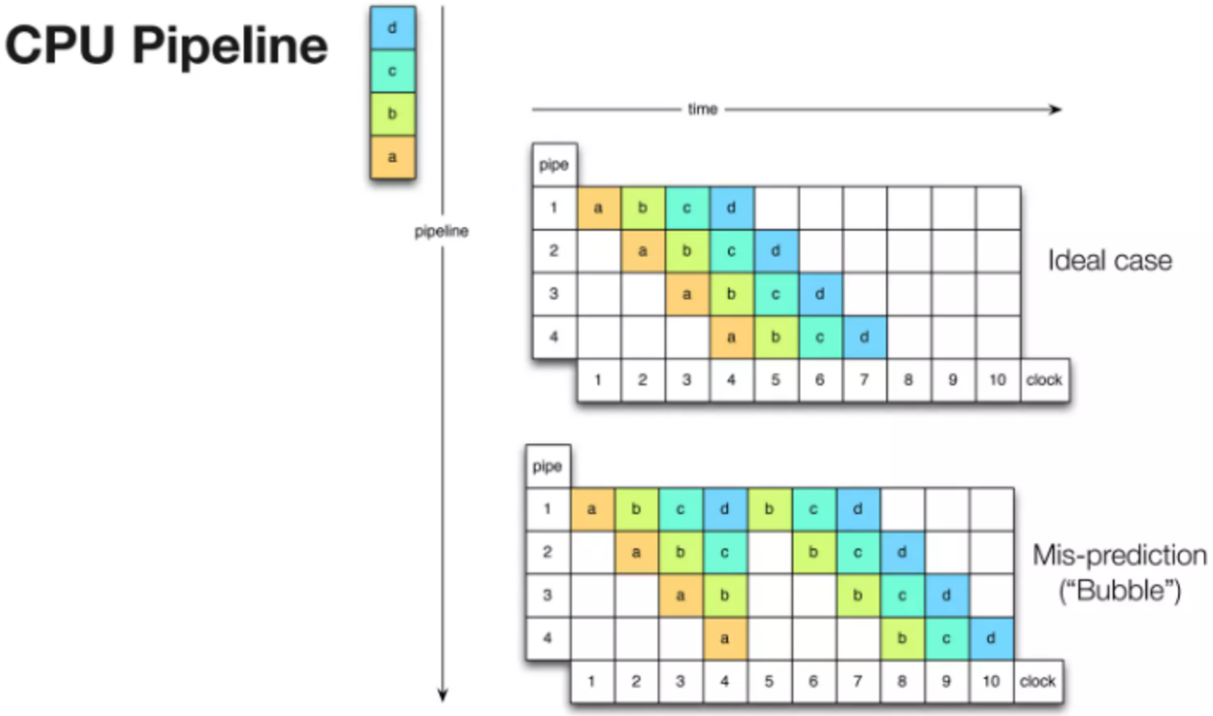

Better Memory Cache

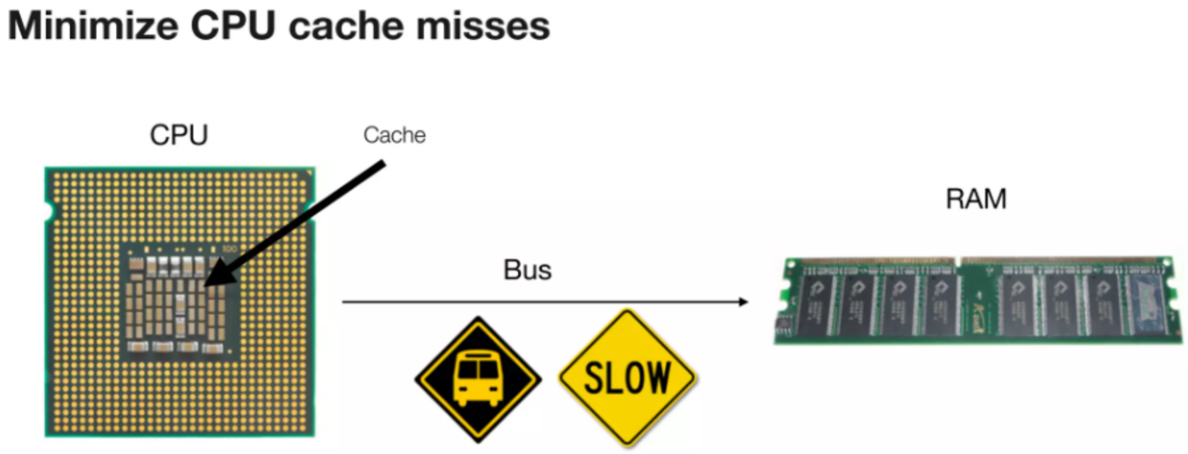

Column Pruning Optimization

SELECT b, d -- ONLY 2 columns
FROM table
WHERE ..

Logical view: rows - columns

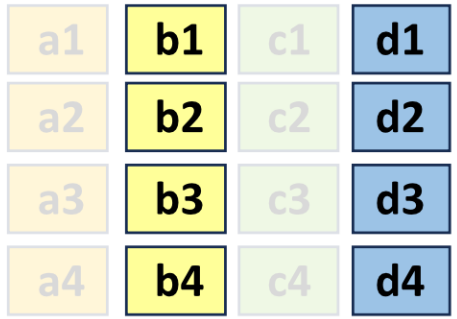

On Disk READ: columnar

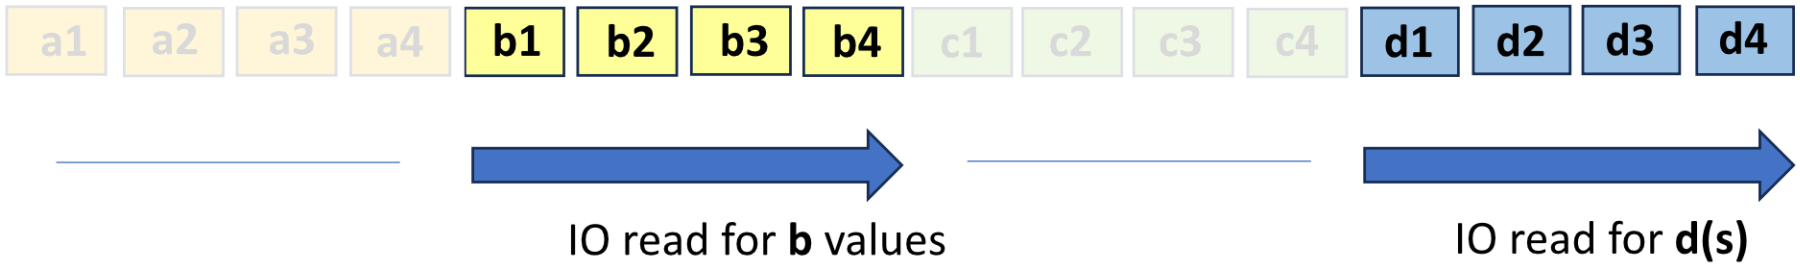

Compare IO Reads with JSON, CSV, Avro, ..

SELECT b, d -- ONLY 2 columns
FROM table
WHERE ..

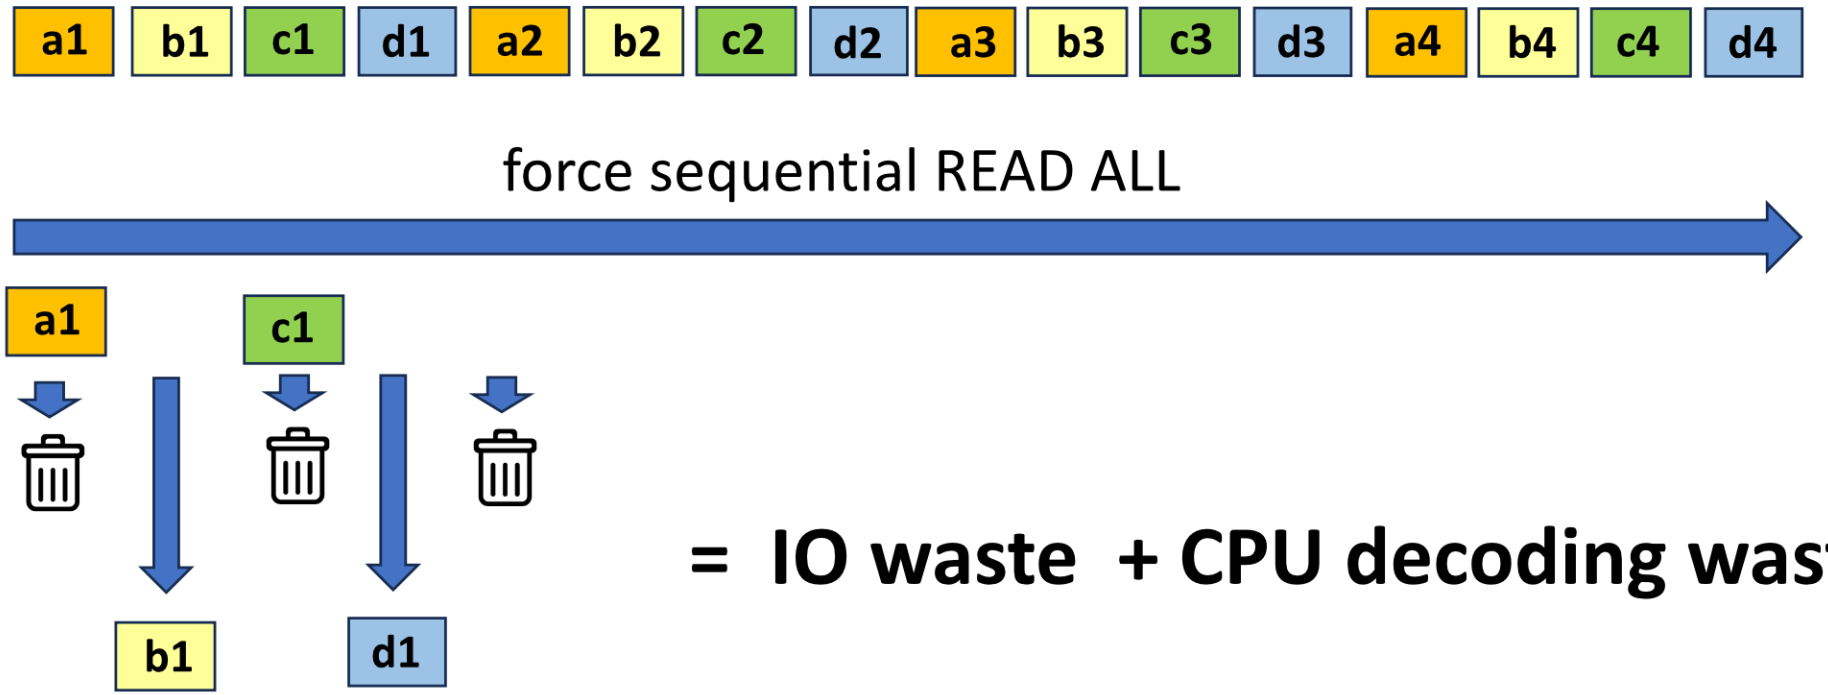

= IO waste + CPU decoding waste

Parquet is Splittable

Parquet split rows by RowGroups

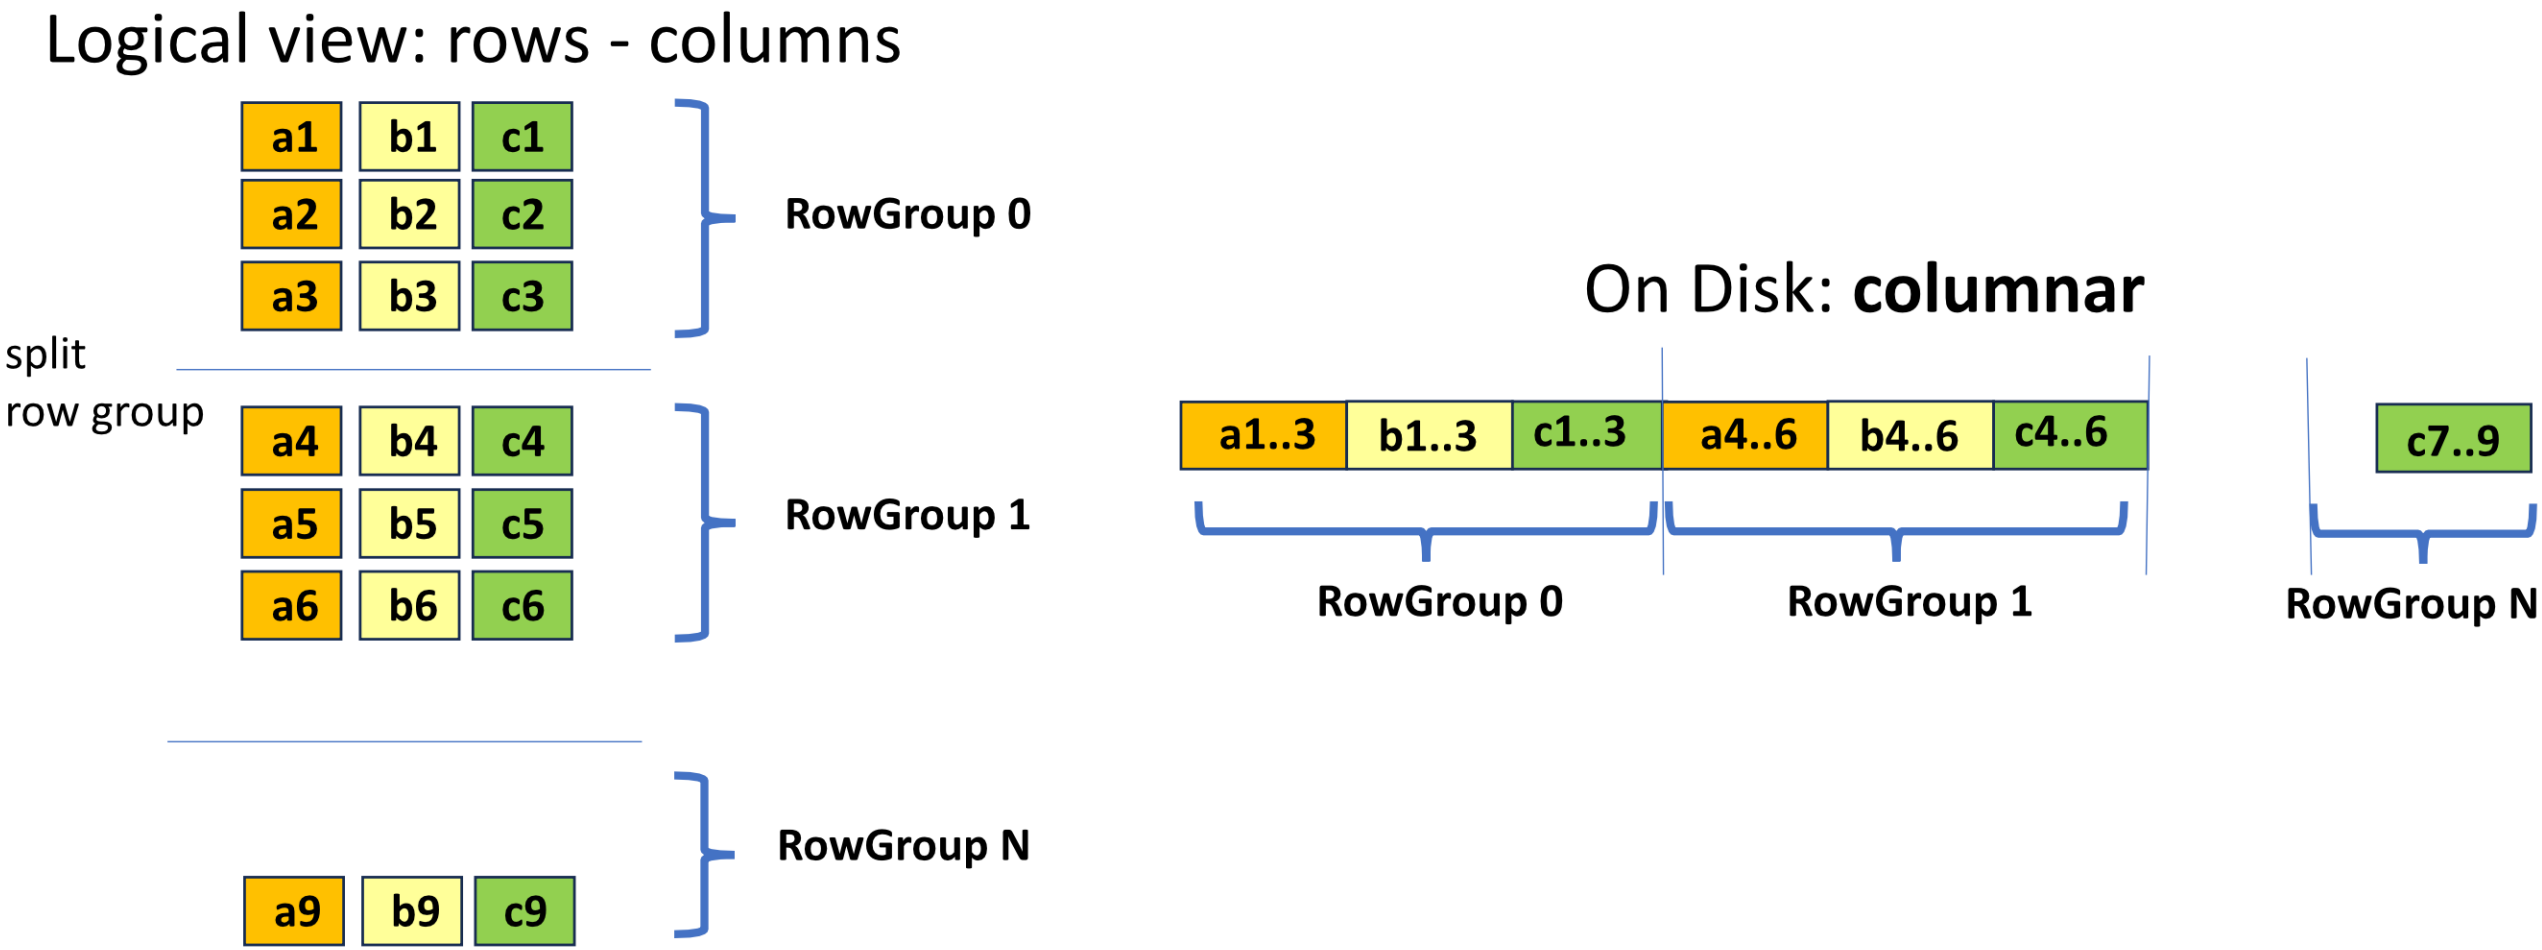

RowGroup ~ Spark Partition ~ Executor Thread

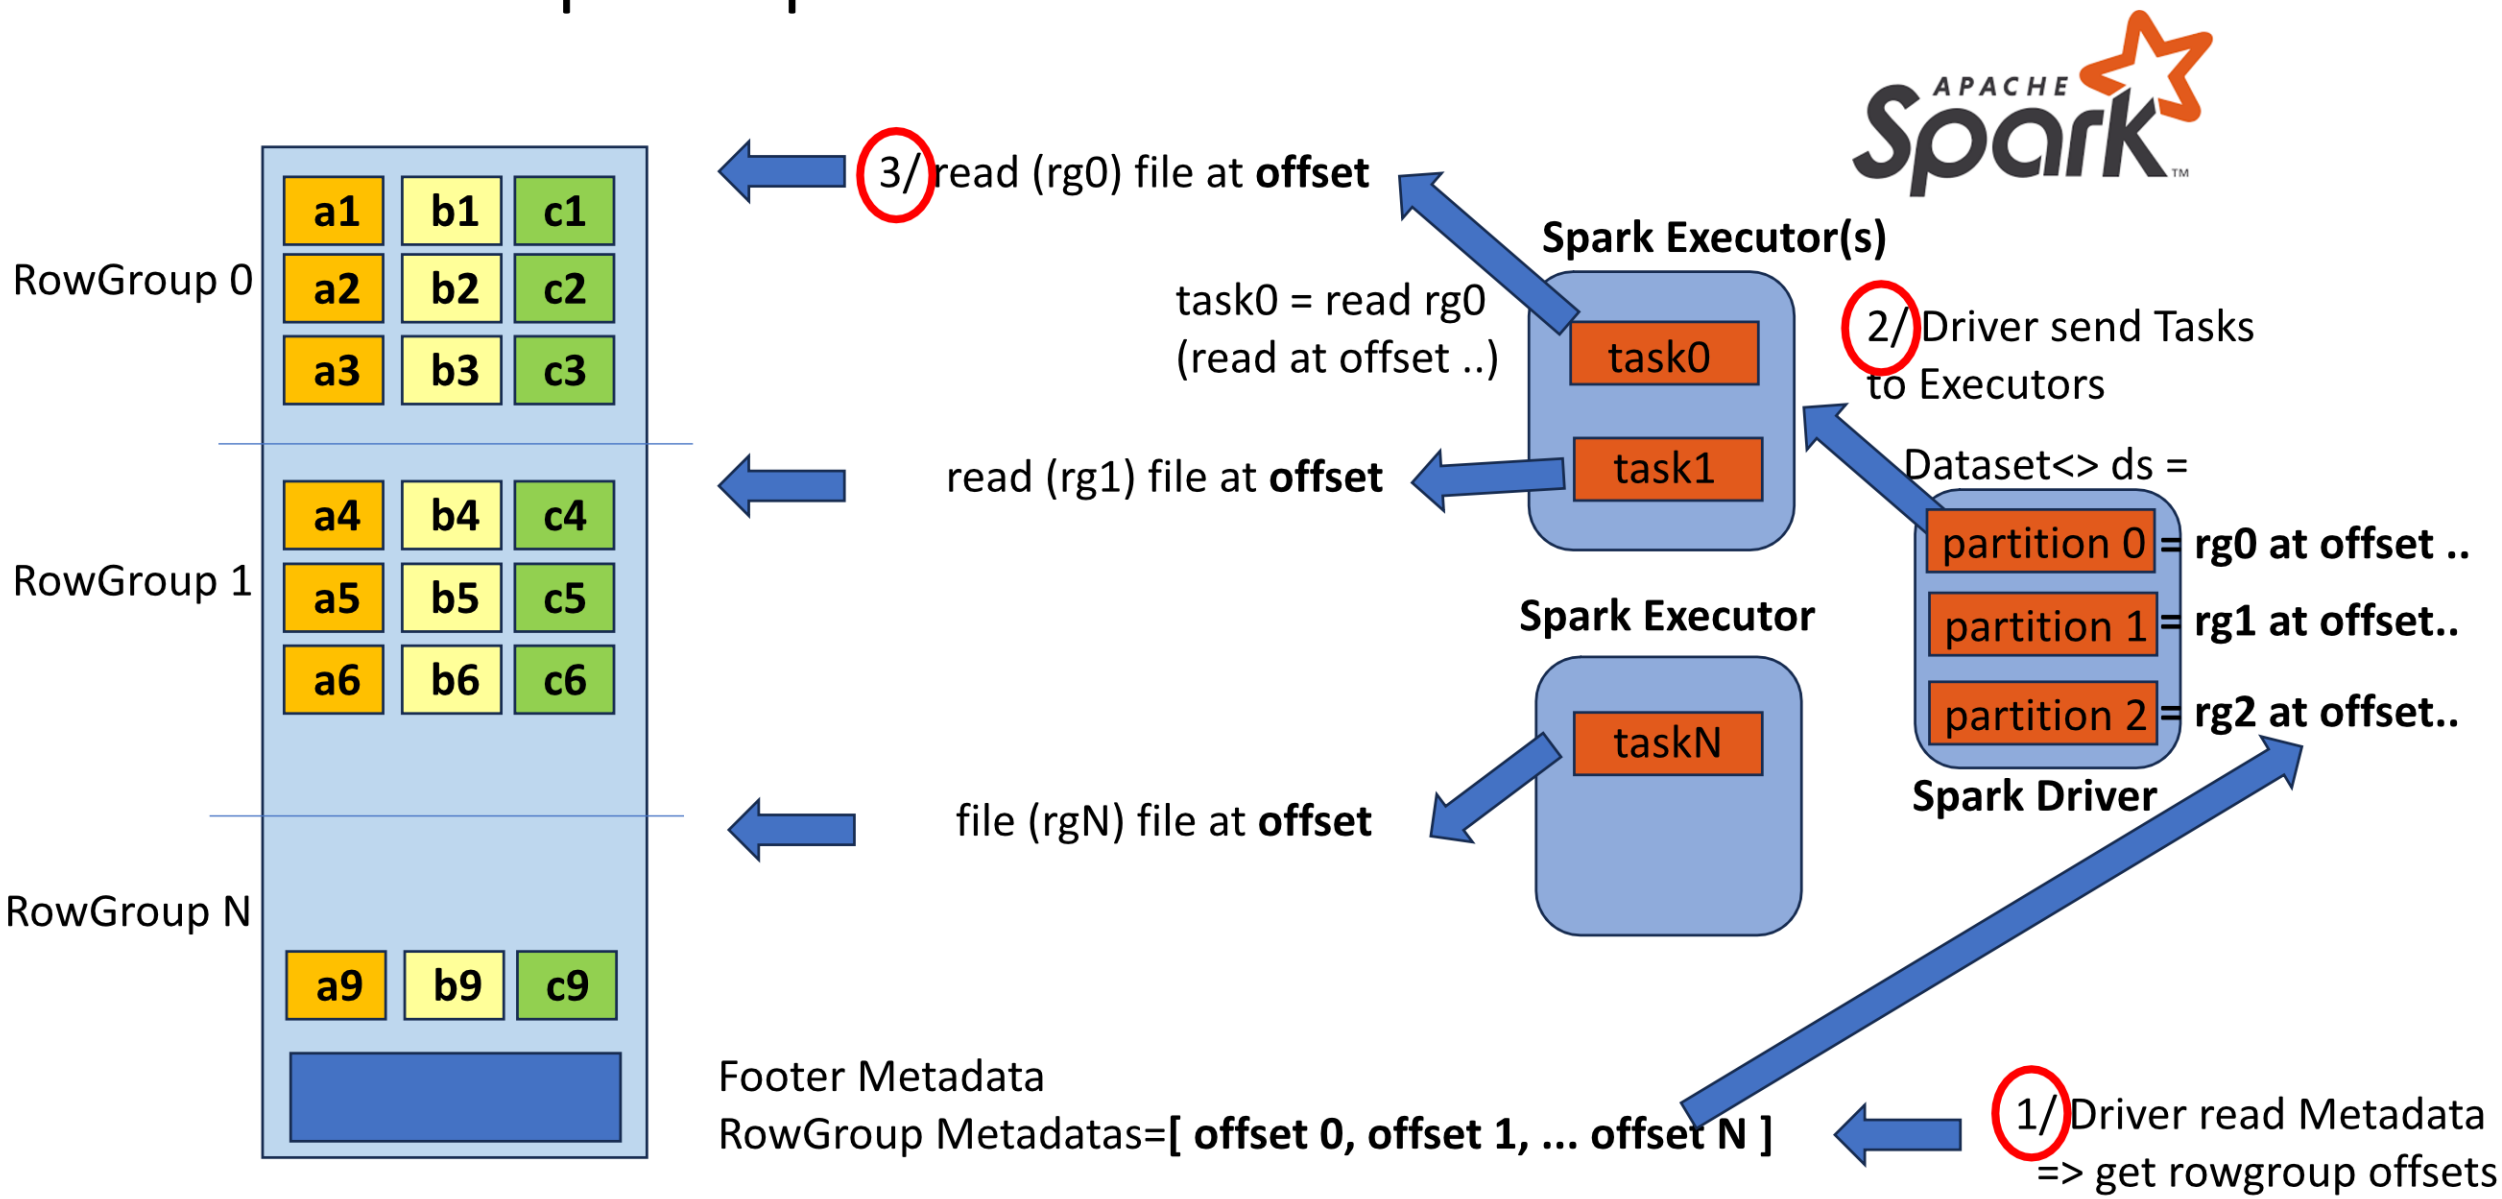

1 Parquet RowGroup(s)
-> default to 1 Spark DataSet partition

by default,

parquet.block.size = 128 Mo

=

spark.files.maxPartitionBytes = 128 Mo

=

spark.sql.files.maxPartitionBytes = 128 Mo

N Parquet RowGroups
-> FileSplit to P (<=N) Spark Partitions

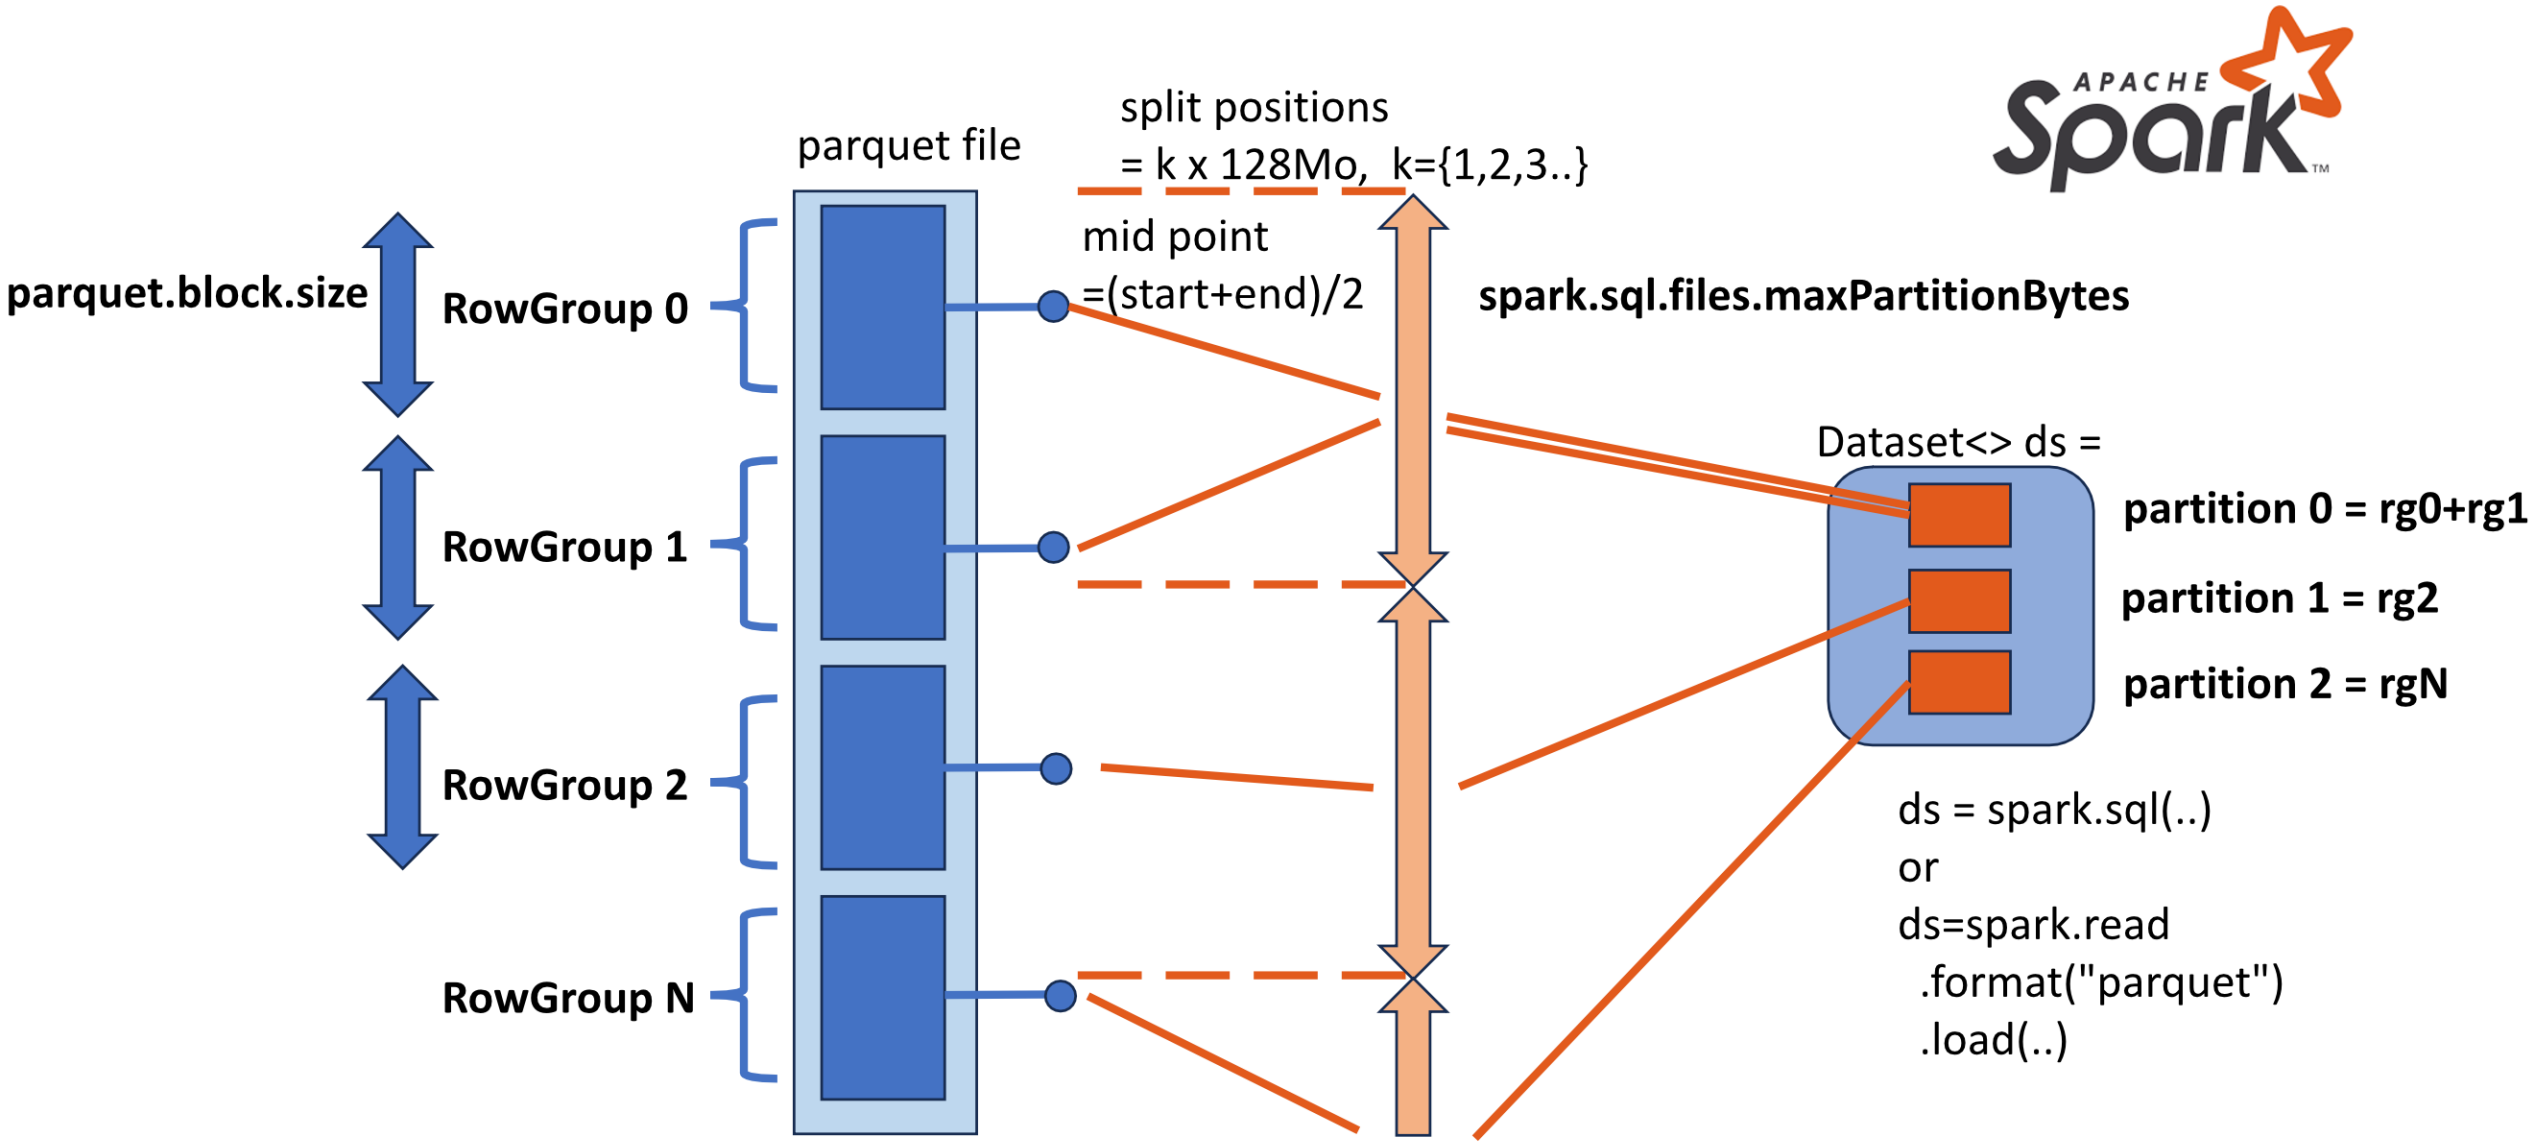

{Texts,CSV,Json} formats FileSplit

each spark executor reader Thread
  at split start => ignore first chars until '\n'
  at split end => read extra chars until '\n

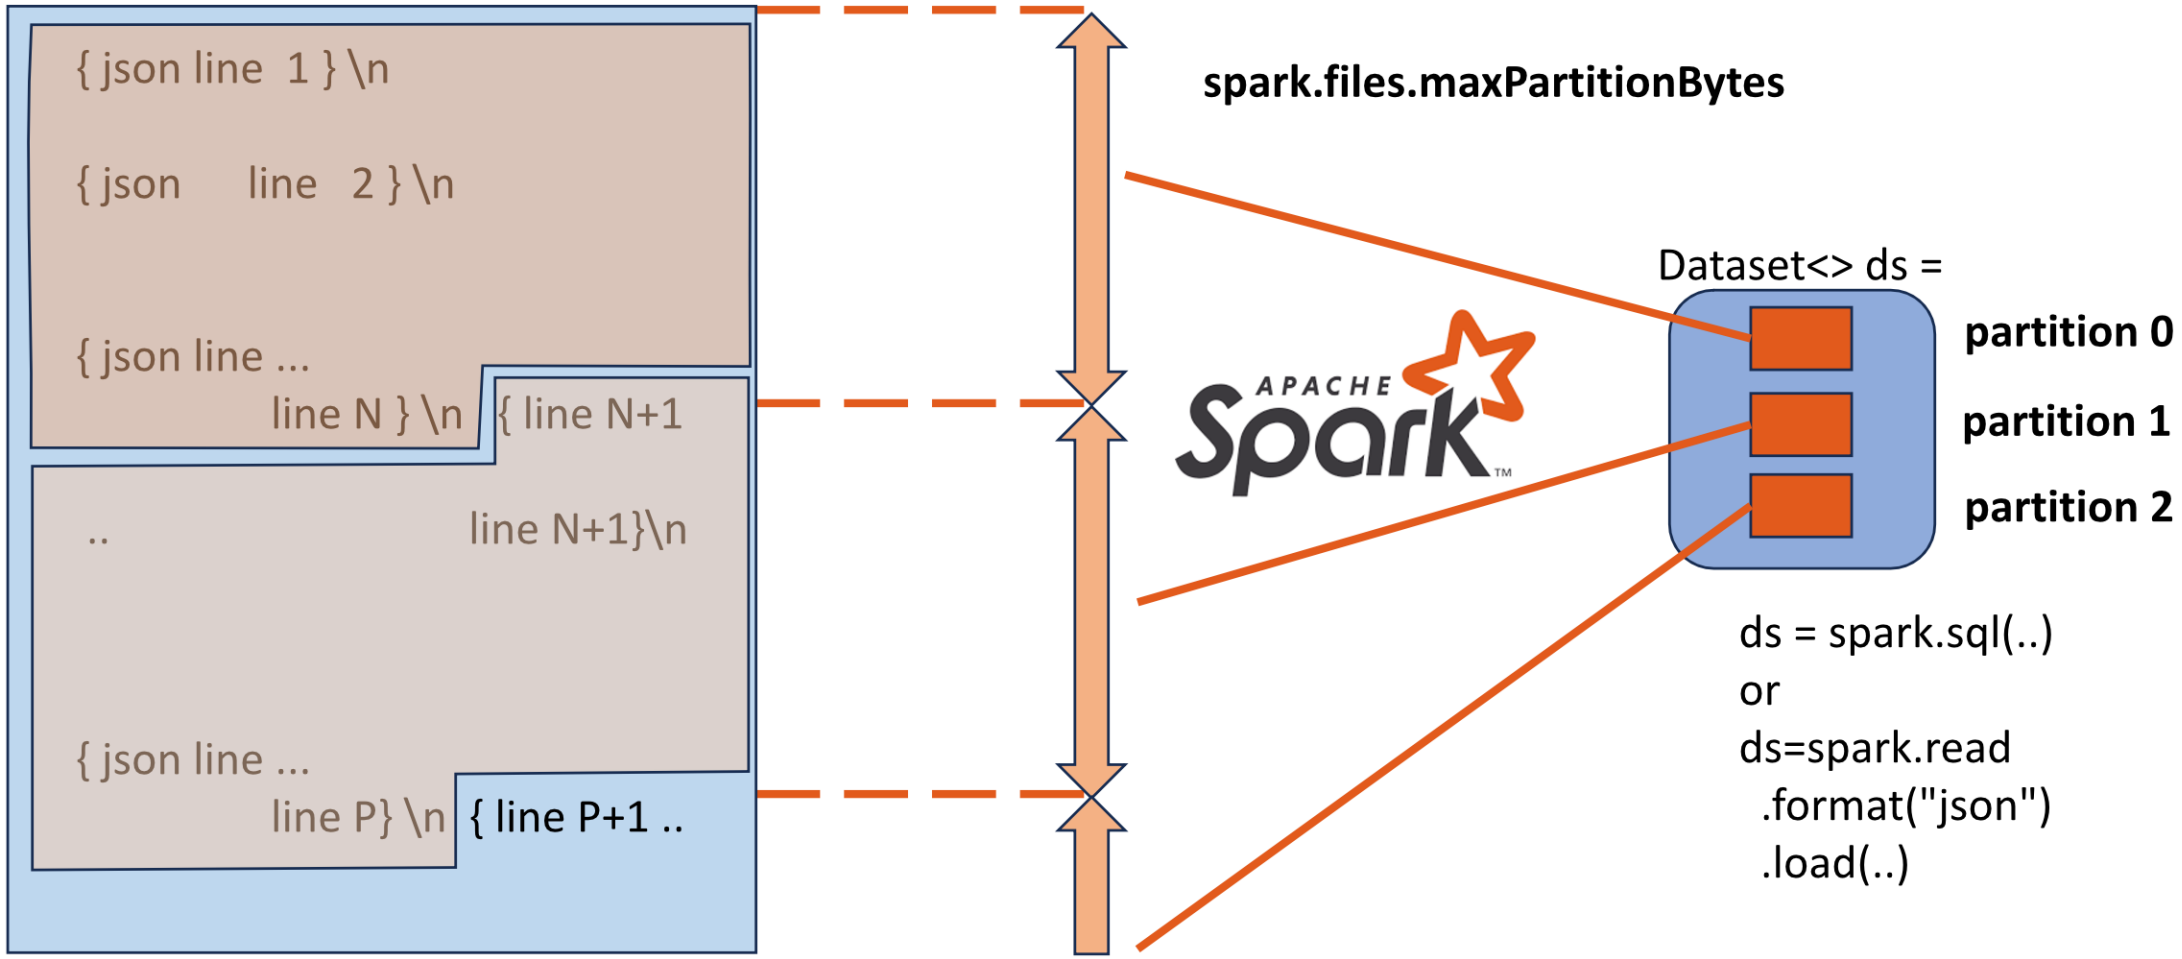

Example : Reading 1 CSV of 3.2 GB
=> 26 splits = 25 (~128 MB) + 1 very small

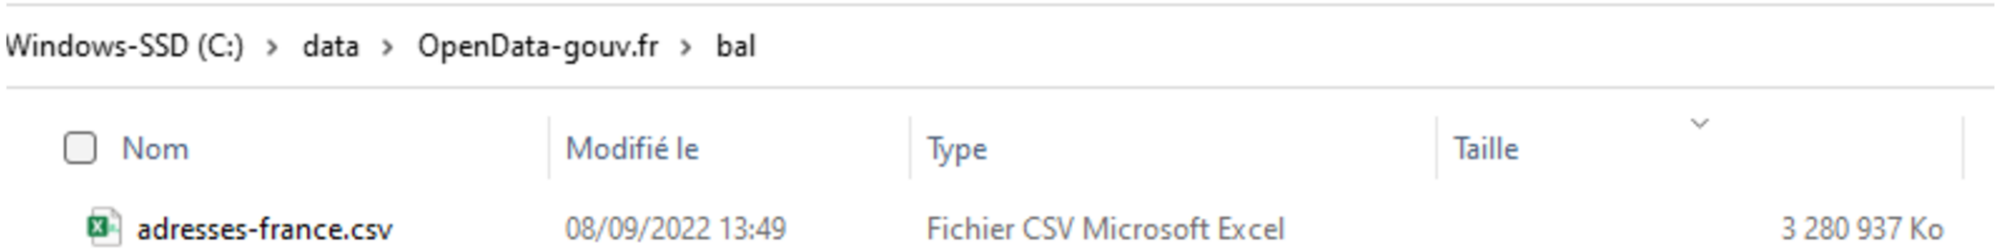

3 280 937 / (128*1024) = 25.03

```
scala> val ds = spark.read.format("csv").option("delimiter",";").load("C:/data/OpenData-gouv.fr/bal")
val ds: org.apache.spark.sql.DataFrame = [_c0: string, _c1: string ... 17 more fields]

scala> ds.toJavaRDD.getNumPartitions
val res0: Int = 26
```

Parquet is Compressible { snappy | gz }

*.gz are NOT Splittable !

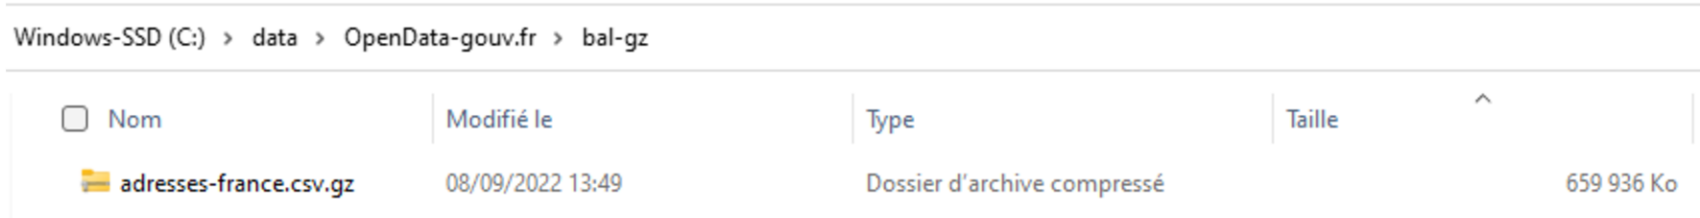

CSV file was 3.2 Go, now 660 Mo in .gz

but NOT Splittable => 1 spark partition, reading (CPU intensive) by 1 Thread only !

~8 times slower on a 8 cores PC

```
scala> val csvGzDs = spark.read.format("csv").option("delimiter",";").load("C:/data/OpenData-gouv.fr/bal-gz")
23/11/05 20:17:12 WARN ZlibFactory: Failed to load/initialize native-zlib library
val csvGzDs: org.apache.spark.sql.DataFrame = [_c0: string, _c1: string ... 17 more fields]

scala> csvGzDs.count()
[Stage 2:>                                                          (0 + 1) / 1]
```

Compressions Algorithms

data

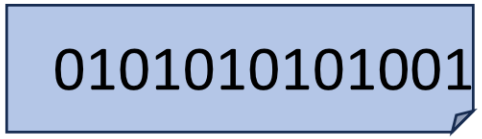

.snappy

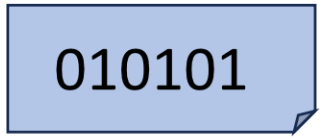

fast compression/decompression
(focus on speed)

.gz (gzip)

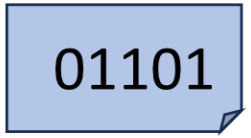

slower-compression,
better compression ratio
(focus on size)

.lz4

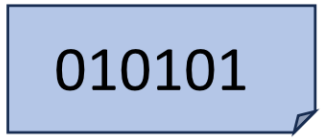

compromise between
fast / compression ratio

PARQUET.{snappy | gzip}

Logical view: rows - columns

a1	b1	c1	d1
a2	b2	c2	d2
a3	b3	c3	d3
a4	b4	c4	d4

On Disk: **Columnar NO compression**

On Disk: **Columnar parquet.{snappy | gzip }**

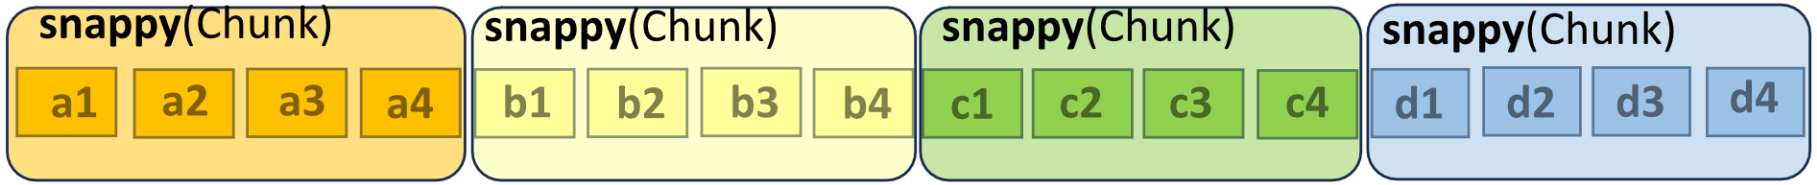

every "Chunk" of column data are compressed INDEPENDENTLY
=> file is STILL Splitteable

Parquet use Encodings

Run-length encoding (RLE)

a
a
a
bb
bb
bb
bb
bb
bb
c
c
c
c

repeated **x3** times

repeated **x6** times

repeated **x4** times

Var Length Encoder:

3 a
6 bb
4 c

Size of Adding "Constant" Column

adding a column with only "0" for Billions of rows

=> take only ~100 bytes per RowGroup

Dictionary Encoding

Manchester City,
Arsenal,
Manchester City,
FC Barcelone,
Arsenal,
Newcastle,
Manchester United,
Newcastle,
FC Barcelone,
Arsenal,
Manchester United,
...

Dictionary Encoder:

Distinct Dictionary Values:

1=Manchester City
2=Arsenal
3=FC Barcelone
4=Newcastle
5..

Value Indexes:

1, 2, 1, 3, 4, 5 ..

Dictionary Size Limit

By default, Dictionary size = max 1 Mega (per RowGroup - Column Chunk)

When using

Huge RowGroup (> 128Mo) => less Dictionaries used

Small RowGroup (32M, 64Mo) => more Dictionaries

Parameters:

parquet.enable.dictionary=true (default)

parquet.dictionary.page.size=1M

Delta Encoding

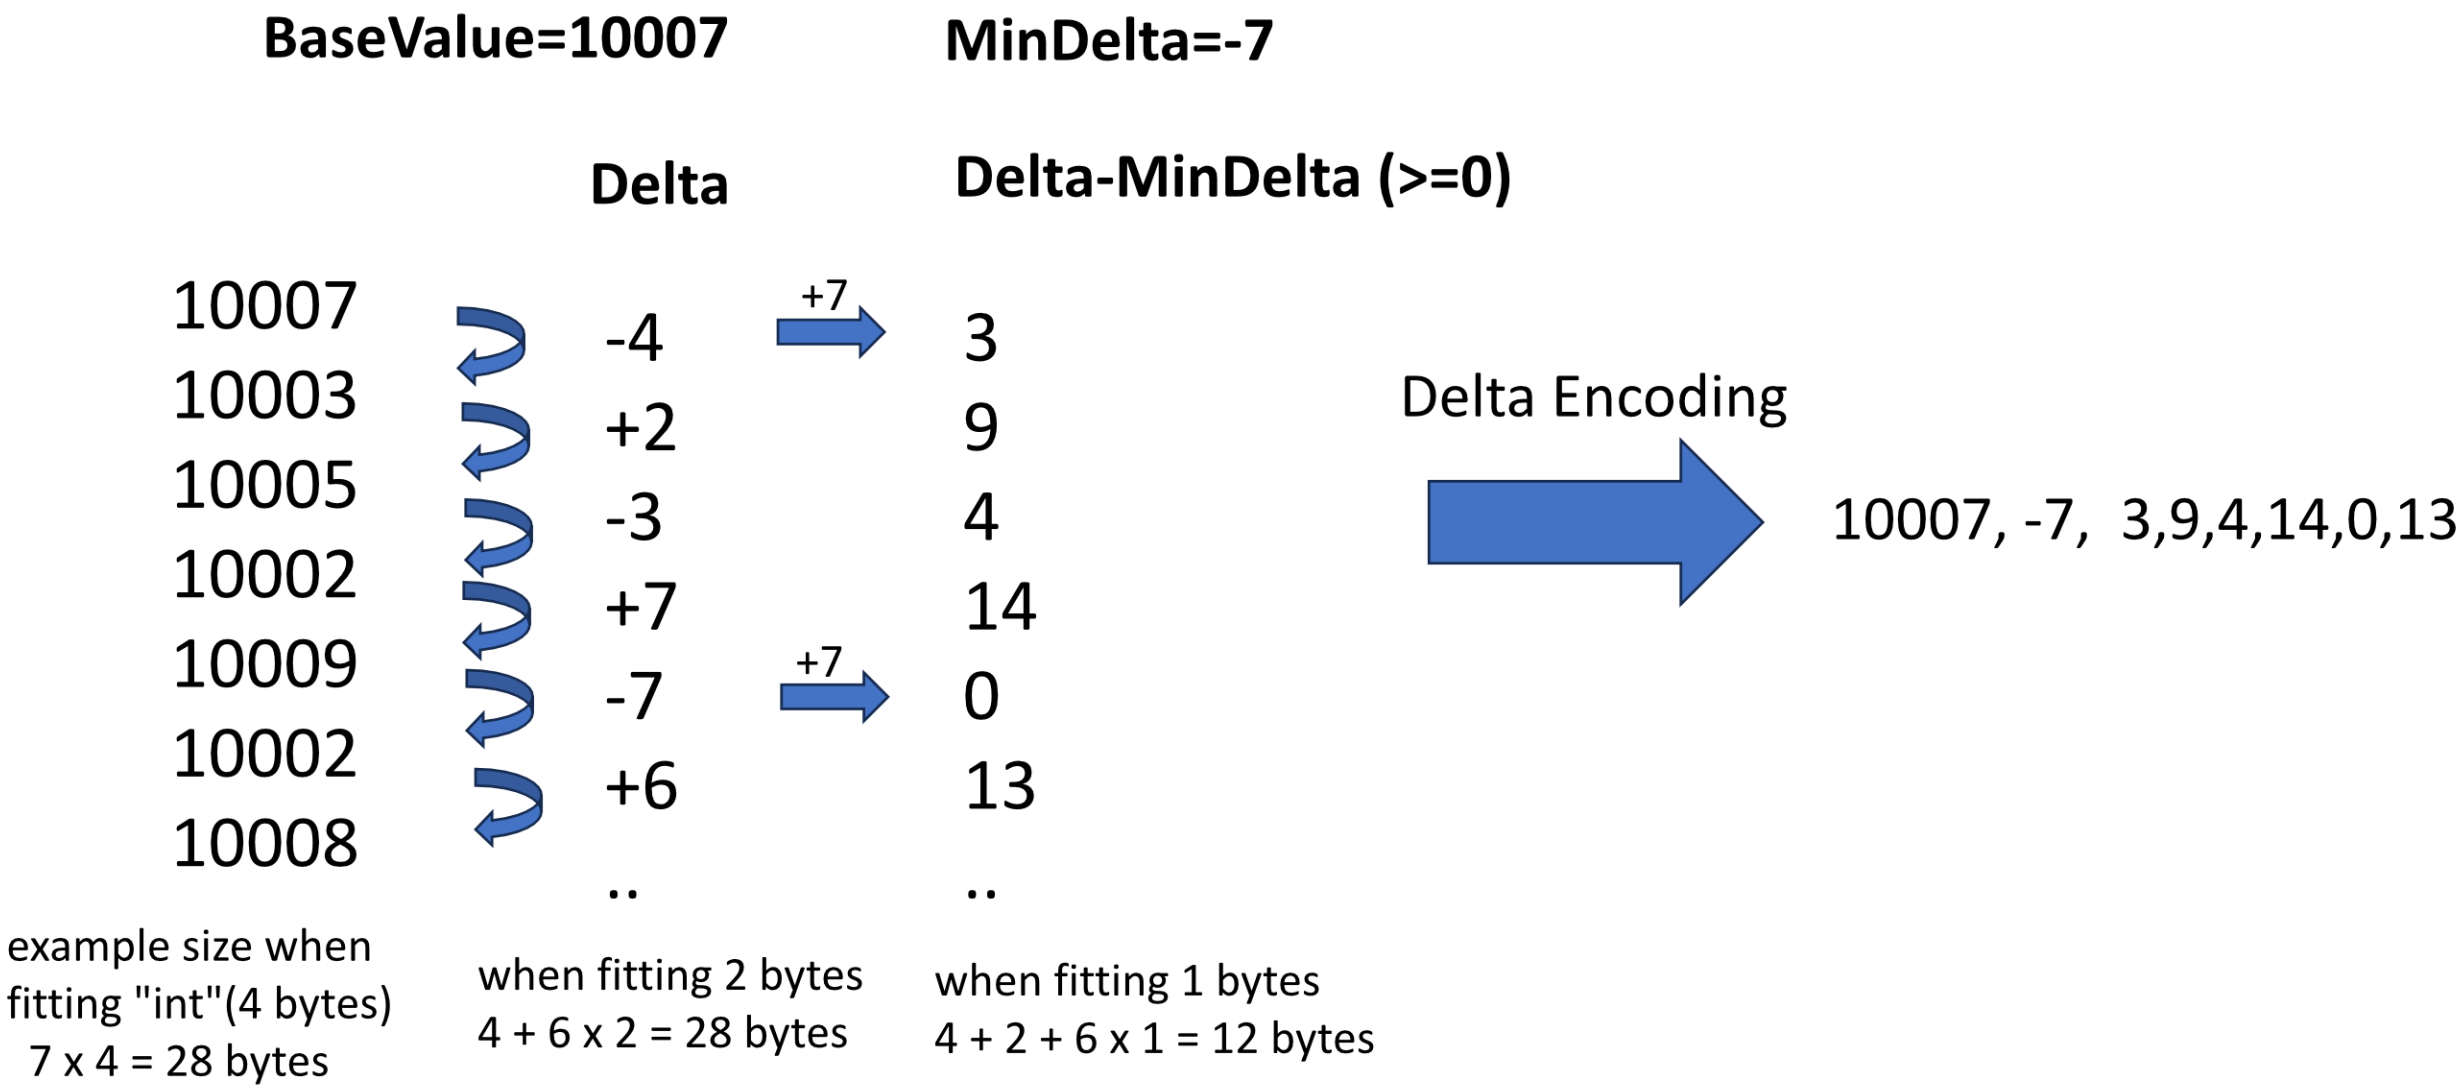

Low-Level Values Encoding… using Thrift « TCompactProtocol »

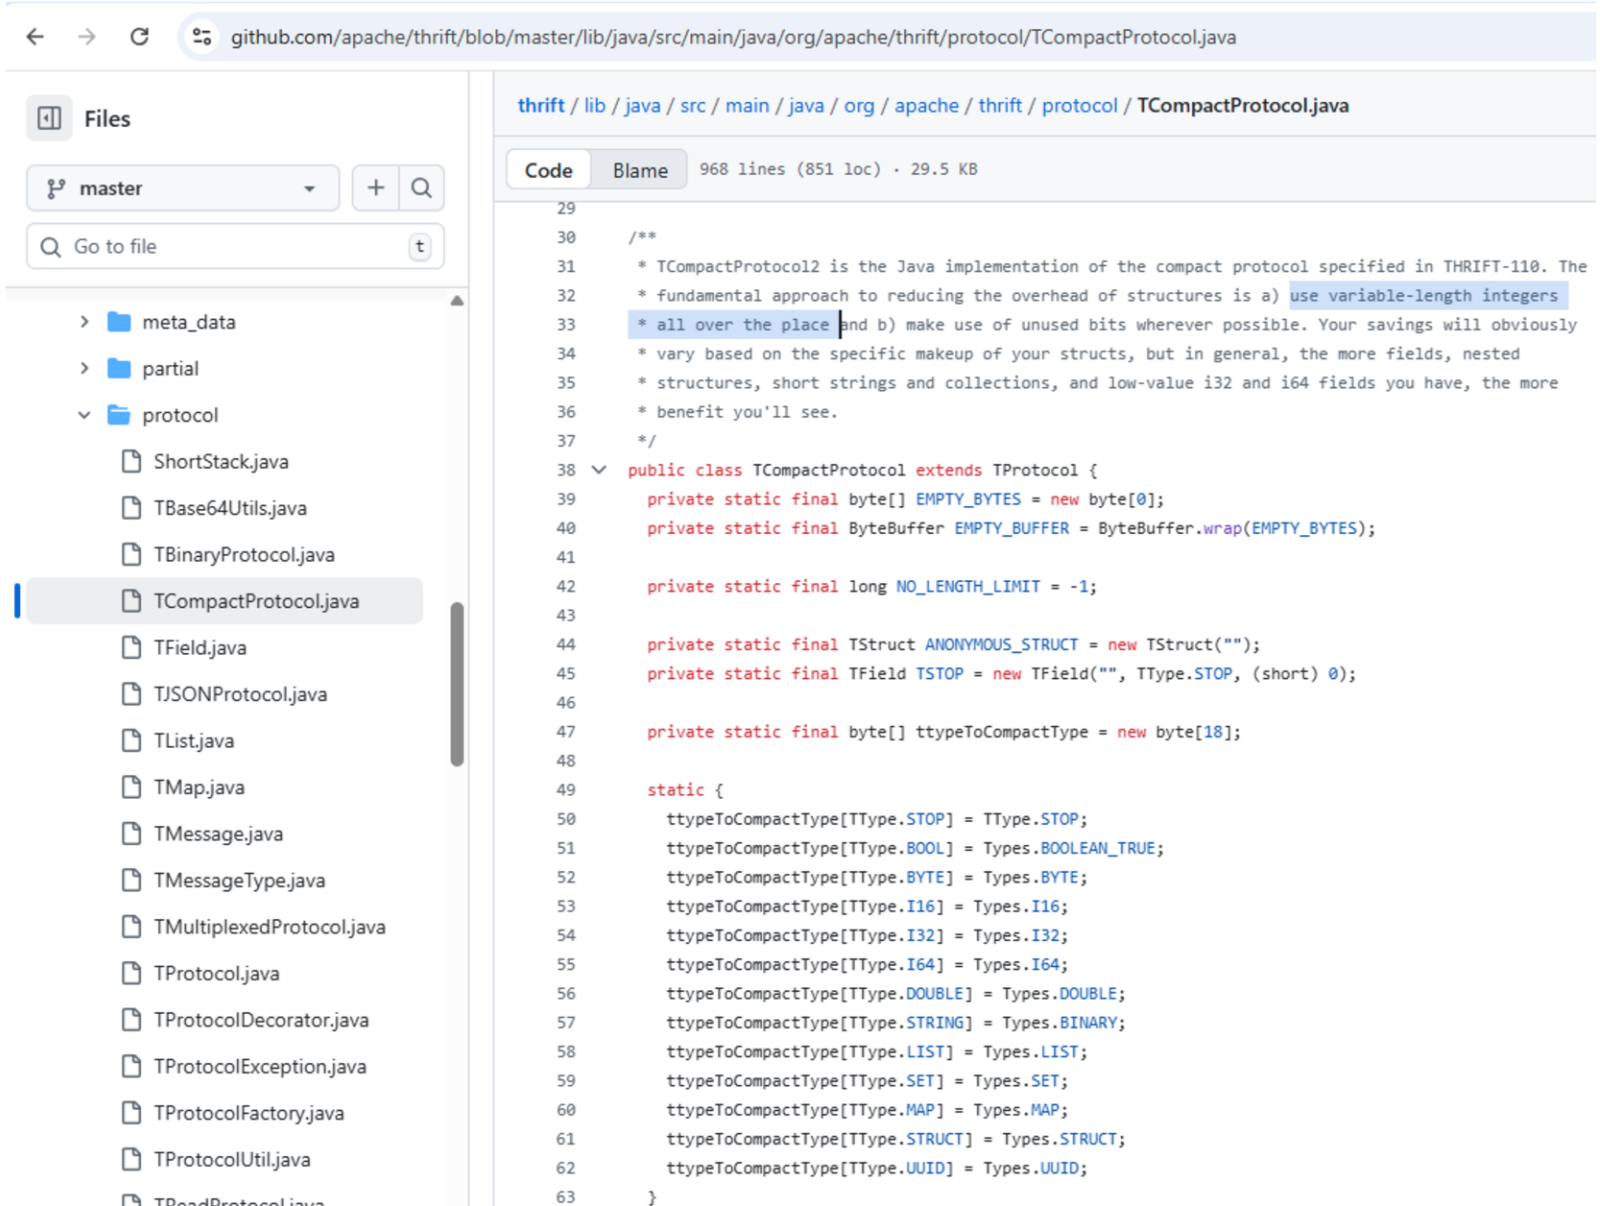


```
public interface TReadProtocol {

  TMessage readMessageBegin();
  void readMessageEnd();

  TStruct readStructBegin();
  void readStructEnd();

  TField readFieldBegin();
  void readFieldEnd();

  TMap readMapBegin();
  void readMapEnd();

  TList readListBegin();
  void readListEnd();

  TSet readSetBegin();
  void readSetEnd();

  boolean readBool();
  byte readByte();
  short readI16();
  int readI32();
  long readI64();
  UUID readUuid();
  double readDouble();
  String readString();
  ByteBuffer readBinary();
}
```

Variable Length for Int

Int32 (4 bytes) => encoded between 1 to 5 bytes
Int64 (8 bytes) => encoded between 1 to 10 bytes

Variable Length for Negative Sign ? ..

+1 = 0000...0001 → small varint
-1 = 1111...1111 → huge varint (10 bytes if naïvely encoded as unsigned!)

ZigZag encoding:

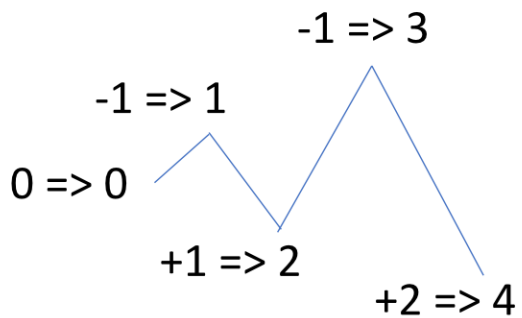

Example for int64 …
values in [-128,127] are encoded on 1 byte, then [-256,+255] on 2 bytes, …

Encoding Int = first convert to ZigZag encoding

thrift / lib / java / src / main / java / org / apache / thrift / protocol / TCompactProtocol.java

Code | Blame | 968 lines (851 loc) · 29.5 KB

```
303       public void writeBool(boolean b) throws TException {
}

/** Write an I16 as a zigzag varint. */
@Override
public void writeI16(short i16) throws TException {
  writeVarint32(intToZigZag(i16));
}

/** Write an i32 as a zigzag varint. */
@Override
public void writeI32(int i32) throws TException {
  writeVarint32(intToZigZag(i32));
}

/** Write an i64 as a zigzag varint. */
@Override
public void writeI64(long i64) throws TException {
  writeVarint64(longToZigzag(i64));
}
```

```
/**
 * Convert l into a zigzag long. This allows negative numbers to be represented compactly as a
 * varint.
 */
private long longToZigzag(long l) {
  return (l << 1) ^ (l >> 63);
}

/**
 * Convert n into a zigzag int. This allows negative numbers to be represented compactly as a
 * varint.
 */
private int intToZigZag(int n) {
  return (n << 1) ^ (n >> 31);
}
```

High-Level Encoding of Nested Structures (Map,Struct,List) => using **Definition** + **Repetition Levels**

Inspired from Google Dremel paper

https://static.googleusercontent.com/media/research.google.com/en/pubs/archive/36632.pdf

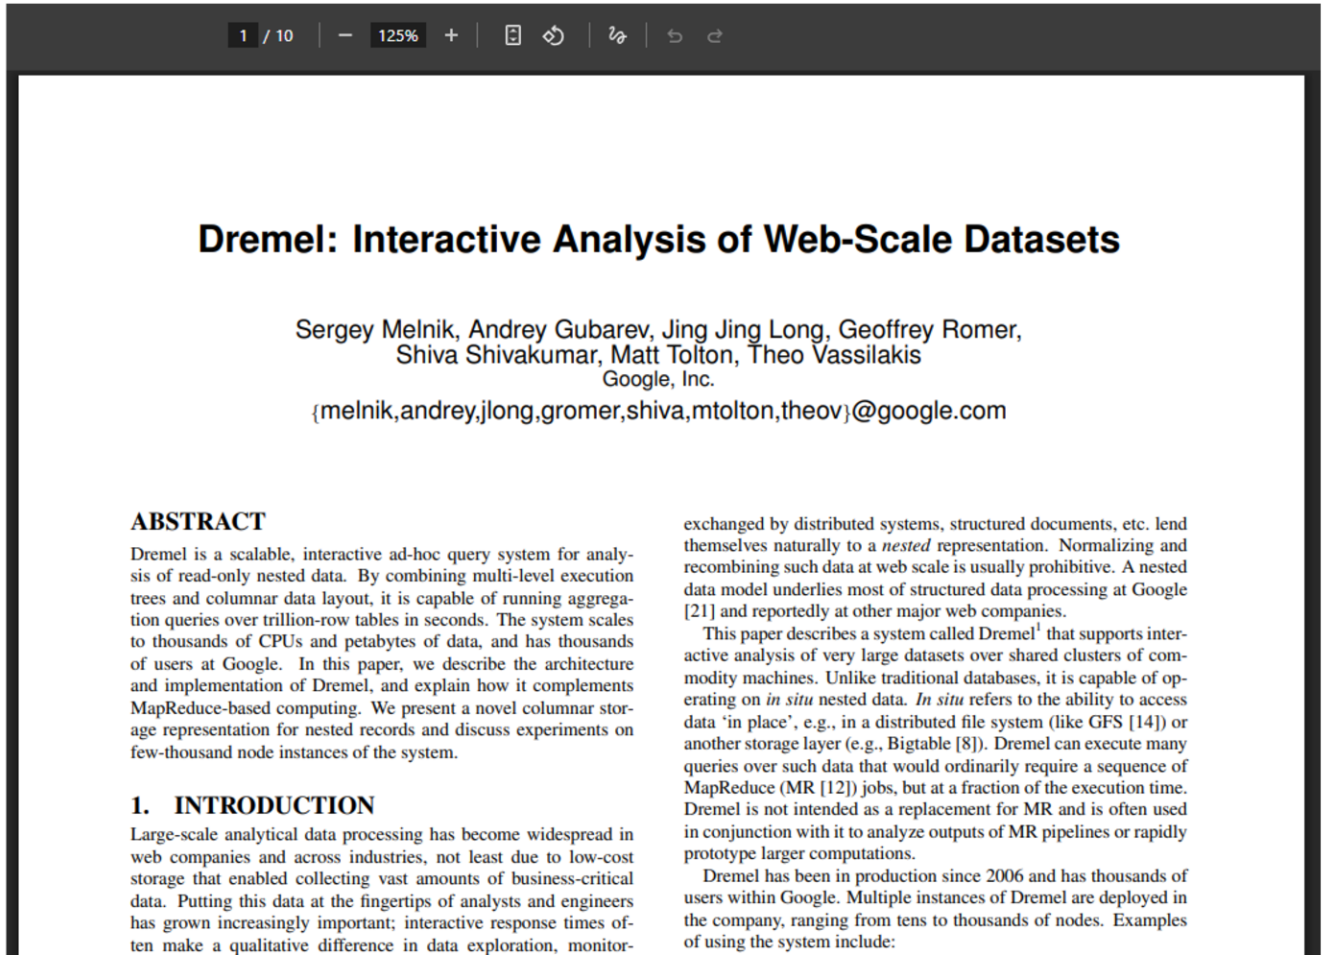

1 / 10 — 125%

Dremel: Interactive Analysis of Web-Scale Datasets

Sergey Melnik, Andrey Gubarev, Jing Jing Long, Geoffrey Romer,
Shiva Shivakumar, Matt Tolton, Theo Vassilakis
Google, Inc.
{melnik,andrey,jlong,gromer,shiva,mtolton,theov}@google.com

ABSTRACT

Dremel is a scalable, interactive ad-hoc query system for analysis of read-only nested data. By combining multi-level execution trees and columnar data layout, it is capable of running aggregation queries over trillion-row tables in seconds. The system scales to thousands of CPUs and petabytes of data, and has thousands of users at Google. In this paper, we describe the architecture and implementation of Dremel, and explain how it complements MapReduce-based computing. We present a novel columnar storage representation for nested records and discuss experiments on few-thousand node instances of the system.

1. INTRODUCTION

Large-scale analytical data processing has become widespread in web companies and across industries, not least due to low-cost storage that enabled collecting vast amounts of business-critical data. Putting this data at the fingertips of analysts and engineers has grown increasingly important; interactive response times often make a qualitative difference in data exploration, monitor-

exchanged by distributed systems, structured documents, etc. lend themselves naturally to a *nested* representation. Normalizing and recombining such data at web scale is usually prohibitive. A nested data model underlies most of structured data processing at Google [21] and reportedly at other major web companies.

This paper describes a system called Dremel[1] that supports interactive analysis of very large datasets over shared clusters of commodity machines. Unlike traditional databases, it is capable of operating on *in situ* nested data. *In situ* refers to the ability to access data 'in place', e.g., in a distributed file system (like GFS [14]) or another storage layer (e.g., Bigtable [8]). Dremel can execute many queries over such data that would ordinarily require a sequence of MapReduce (MR [12]) jobs, but at a fraction of the execution time. Dremel is not intended as a replacement for MR and is often used in conjunction with it to analyze outputs of MR pipelines or rapidly prototype larger computations.

Dremel has been in production since 2006 and has thousands of users within Google. Multiple instances of Dremel are deployed in the company, ranging from tens to thousands of nodes. Examples of using the system include:

Record-Oriented => stored as Path.Columns (R,D,V)

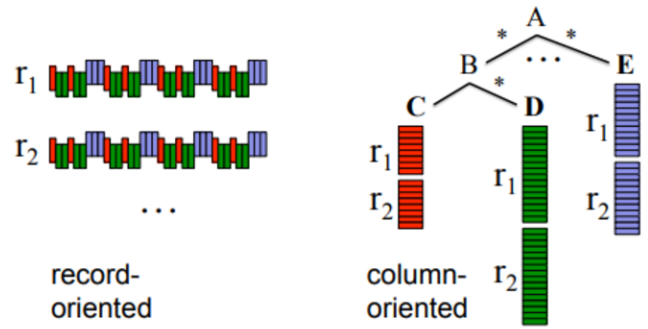

Figure 1: Record-wise vs. columnar representation of nested data

```
DocId: 10                    r1
Links
  Forward: 20
  Forward: 40
  Forward: 60
Name
  Language
    Code: 'en-us'
    Country: 'us'
  Language
    Code: 'en'
  Url: 'http://A'
Name
  Url: 'http://B'
Name
  Language
    Code: 'en-gb'
    Country: 'gb'
```

```
message Document {
  required int64 DocId;
  optional group Links {
    repeated int64 Backward;
    repeated int64 Forward; }
  repeated group Name {
    repeated group Language {
      required string Code;
      optional string Country; }
    optional string Url; }}
```

```
DocId: 20          r2
Links
  Backward: 10
  Backward: 30
  Forward:  80
Name
  Url: 'http://C'
```

Figure 2: Two sample nested records and their schema

DocId

value	r	d
10	0	0
20	0	0

Name.Url

value	r	d
http://A	0	2
http://B	1	2
NULL	1	1
http://C	0	2

Links.Forward

value	r	d
20	0	2
40	1	2
60	1	2
80	0	2

Links.Backward

value	r	d
NULL	0	1
10	0	2
30	1	2

Name.Language.Code

value	r	d
en-us	0	2
en	2	2
NULL	1	1
en-gb	1	2
NULL	0	1

Name.Language.Country

value	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
NULL	0	1

Figure 3: Column-striped representation of the sample data in Figure 2, showing repetition levels (r) and definition levels (d)

(R,D,V) triplets … complex but super efficient

Repetition
Definition
Value

Integer Values in column « R » and « D » are mostly compressed using RLE

The algorithm is complex and can be ignored by end-users of Parquet

The Good News = it works extremely well !!

Parquet can search/compress at any « path » inside nested Records

Parquet use Statistics

Column Statistics: min/max Value per RowGroup

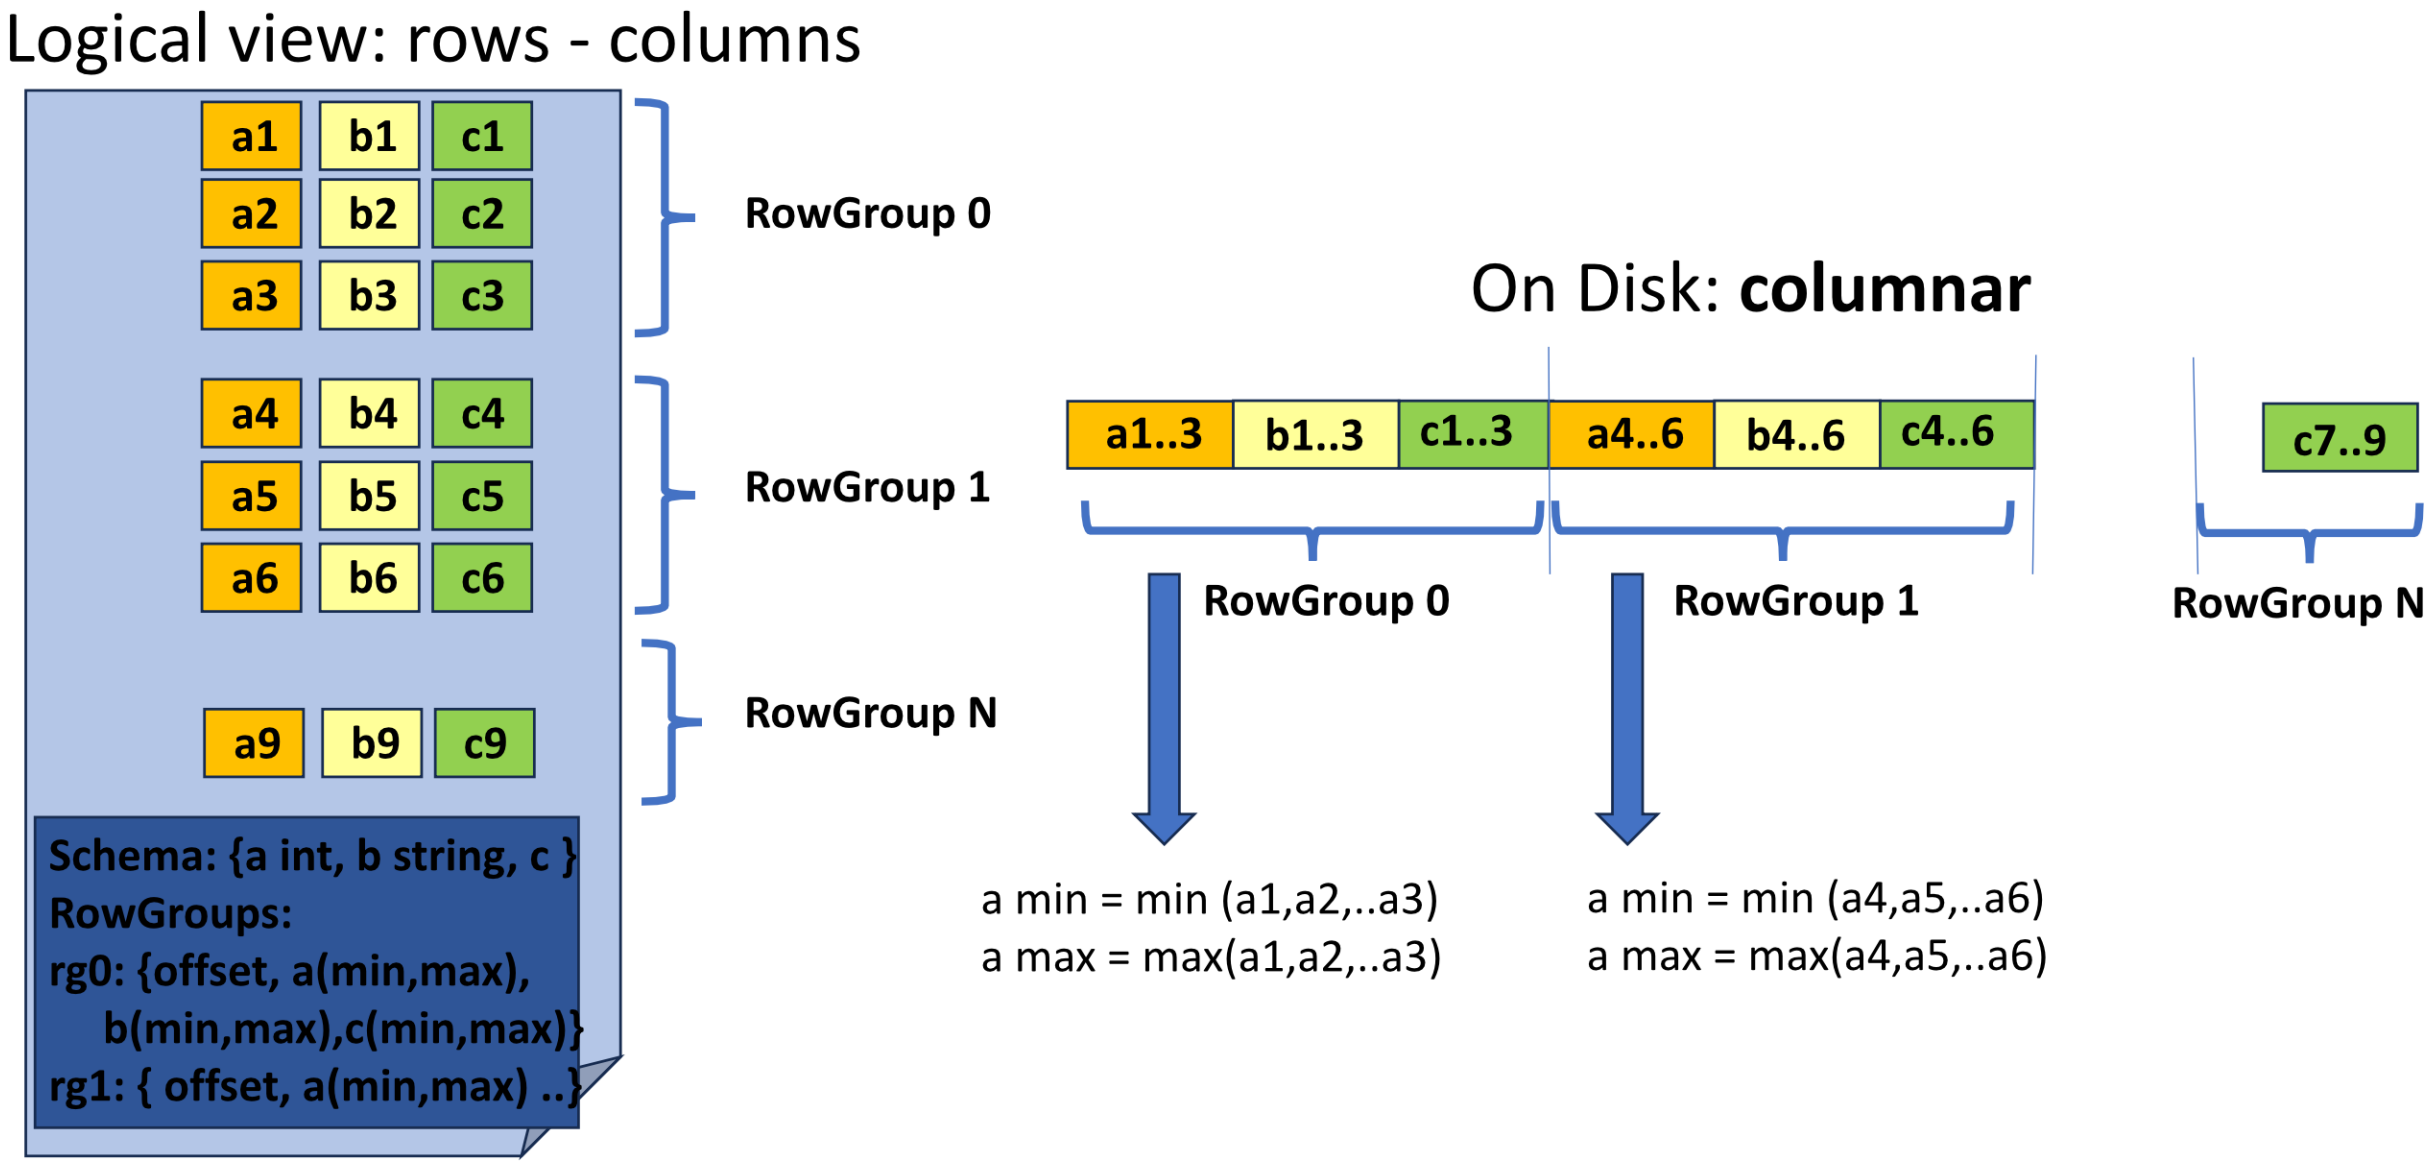

Reading Metadata
=> read schema + offset + statistics

data part
(usually N x blocks of 128 Mo)

footer: metadata part
(usually <= 50 ko)

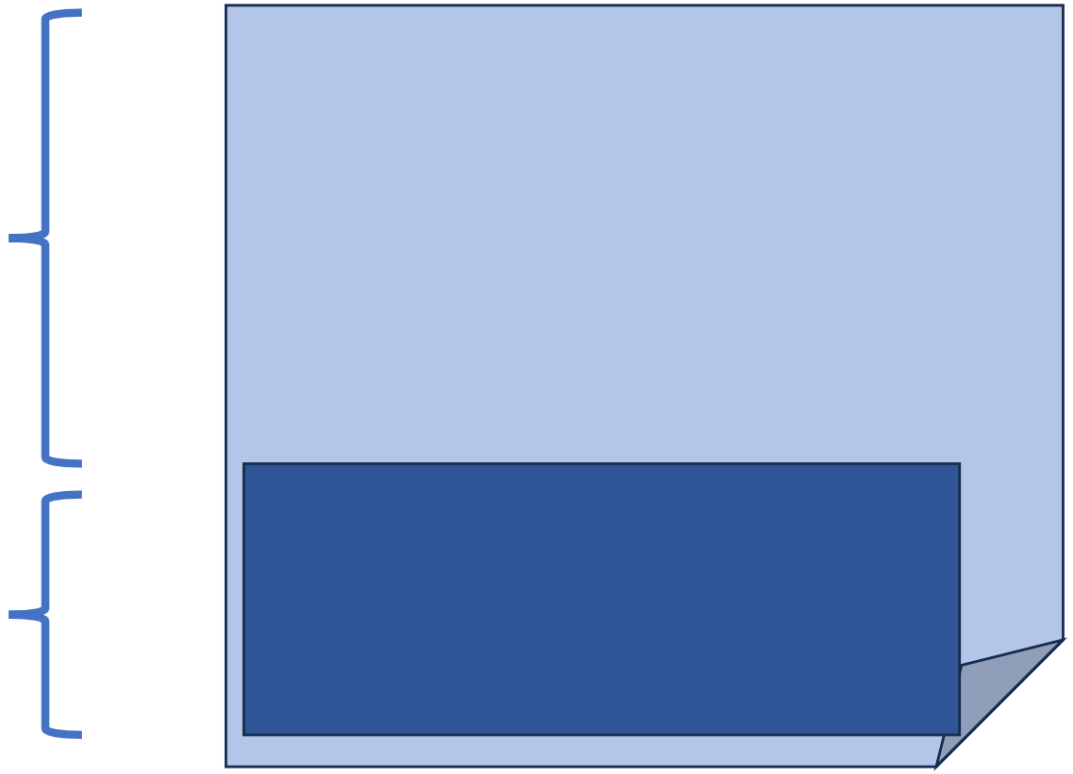

Parquet can use Bloom Filter

Bloom Filter

is W in the set {x,y,z} ?

hash(W) = 0000100...10100

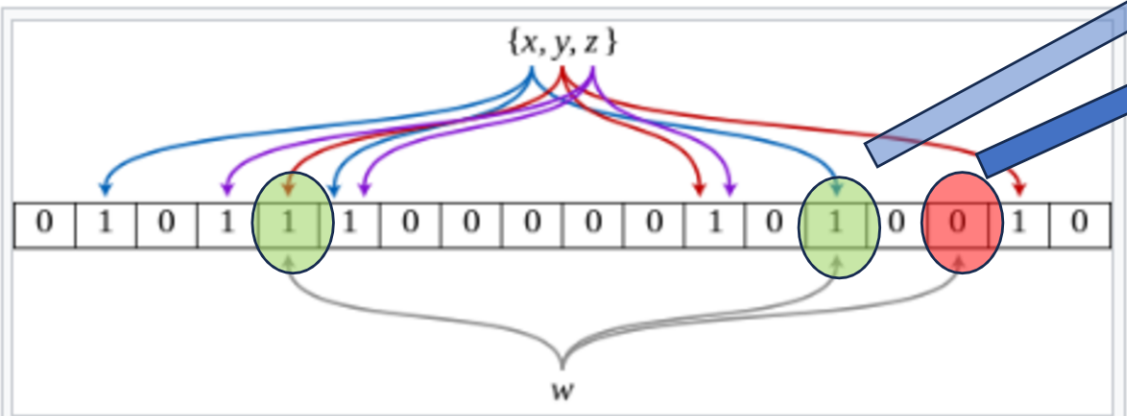

An example of a Bloom filter, representing the set $\{x, y, z\}$. The colored arrows show the positions in the bit array that each set element is mapped to. The element w is not in the set $\{x, y, z\}$, because it hashes to one bit-array position containing 0. For this figure, $m = 18$ and $k = 3$.

Probably Yes ...

Firm NO

https://en.wikipedia.org/wiki/Bloom_filter

Parquet Predicate-Push-Down

Push-Down :
from Spark-> Parquet Lib-> IO Storage

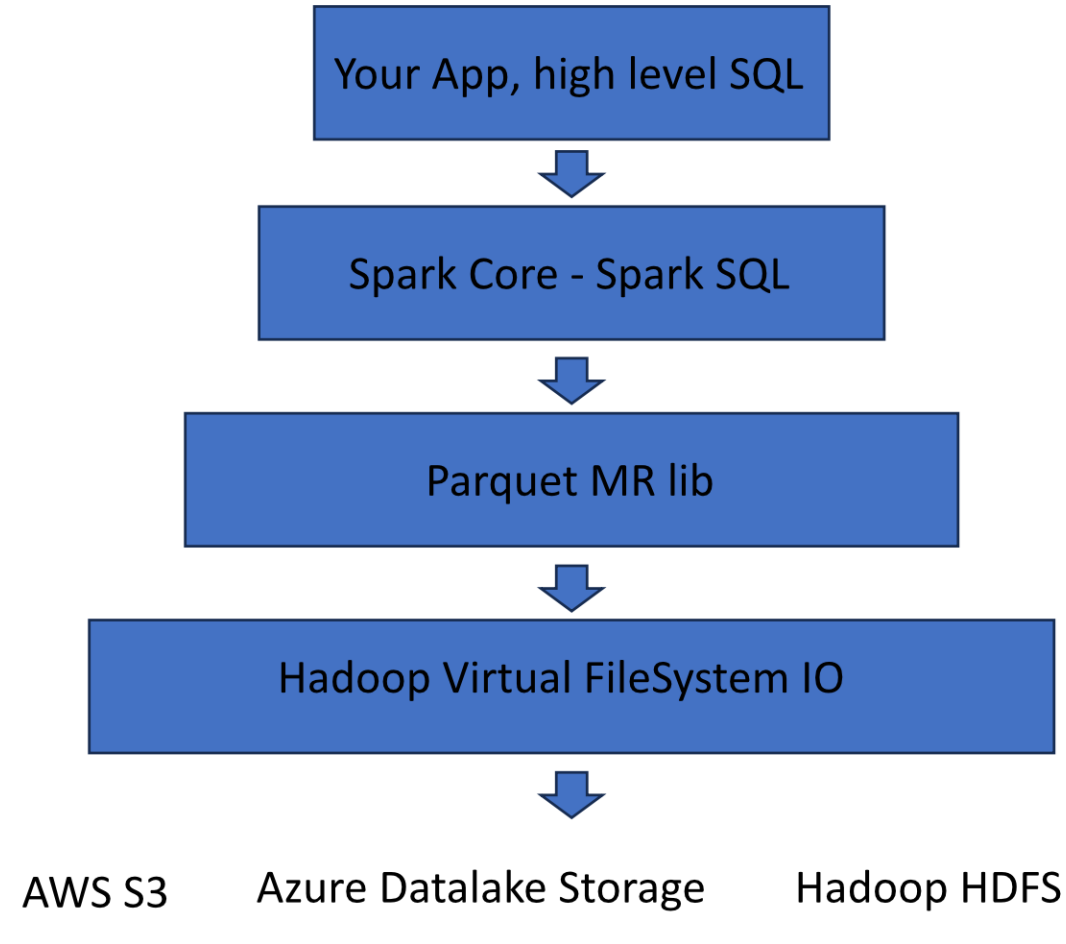

Statistics for skip read RowGroup
SELECT .. WHERE column=value

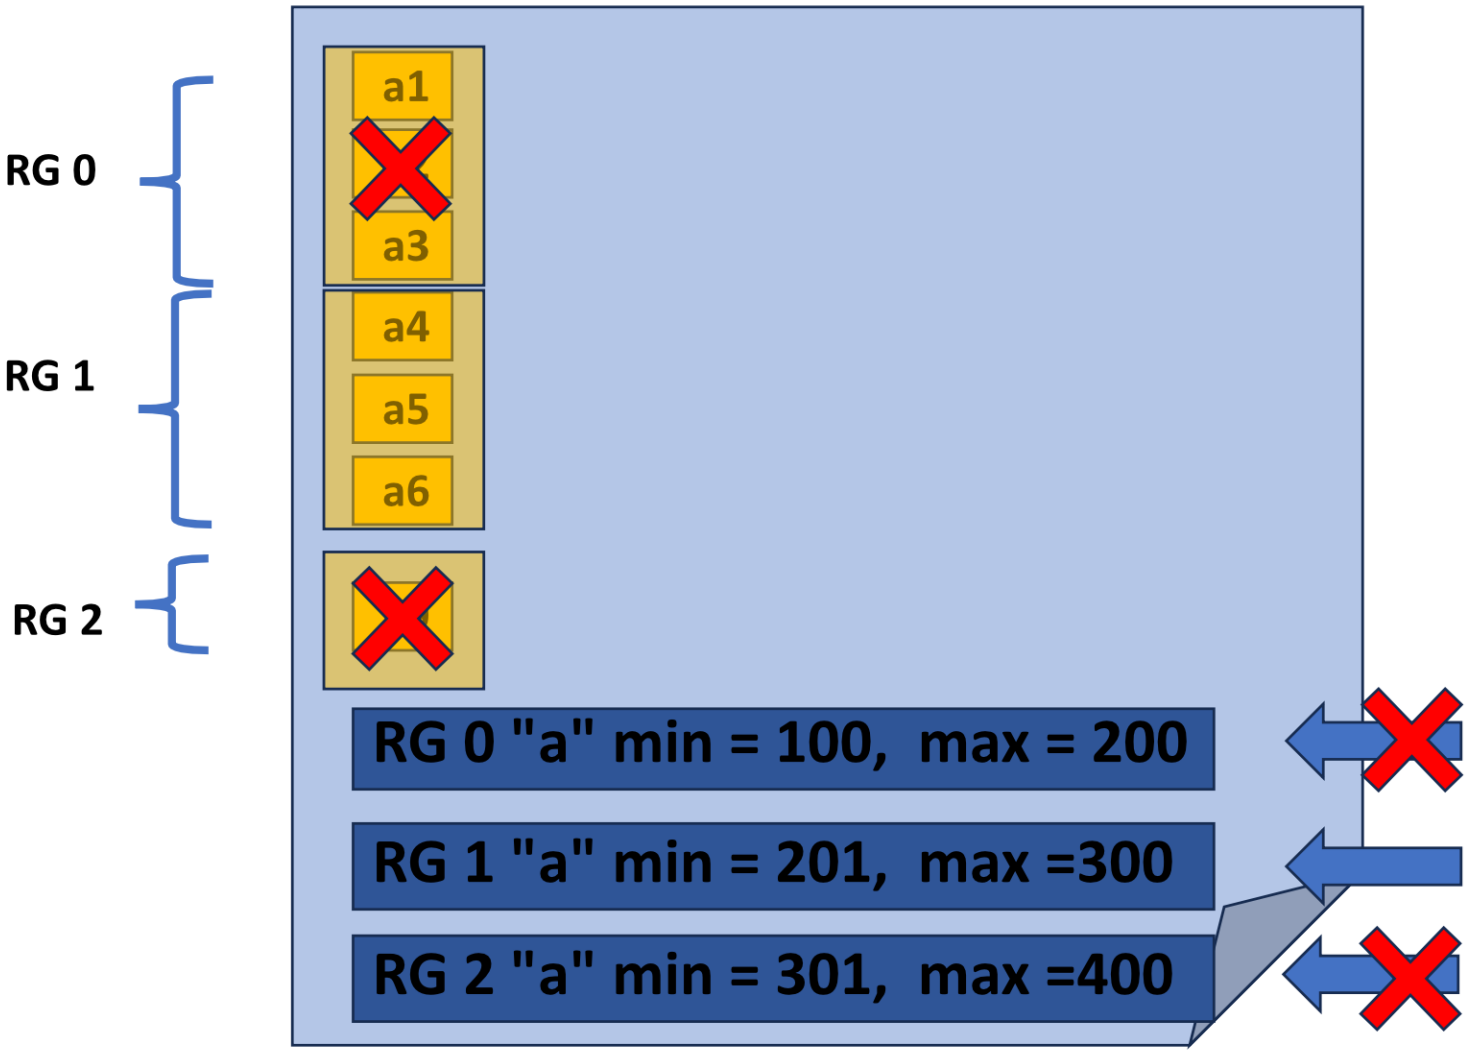

SELECT ...
WHERE a = 250

NO match max < a

OK min <= a <= max (maybe false positive)

NO match a < min

Column Pruning
+ RowGroup Pruning(PPD) = minimal Reads

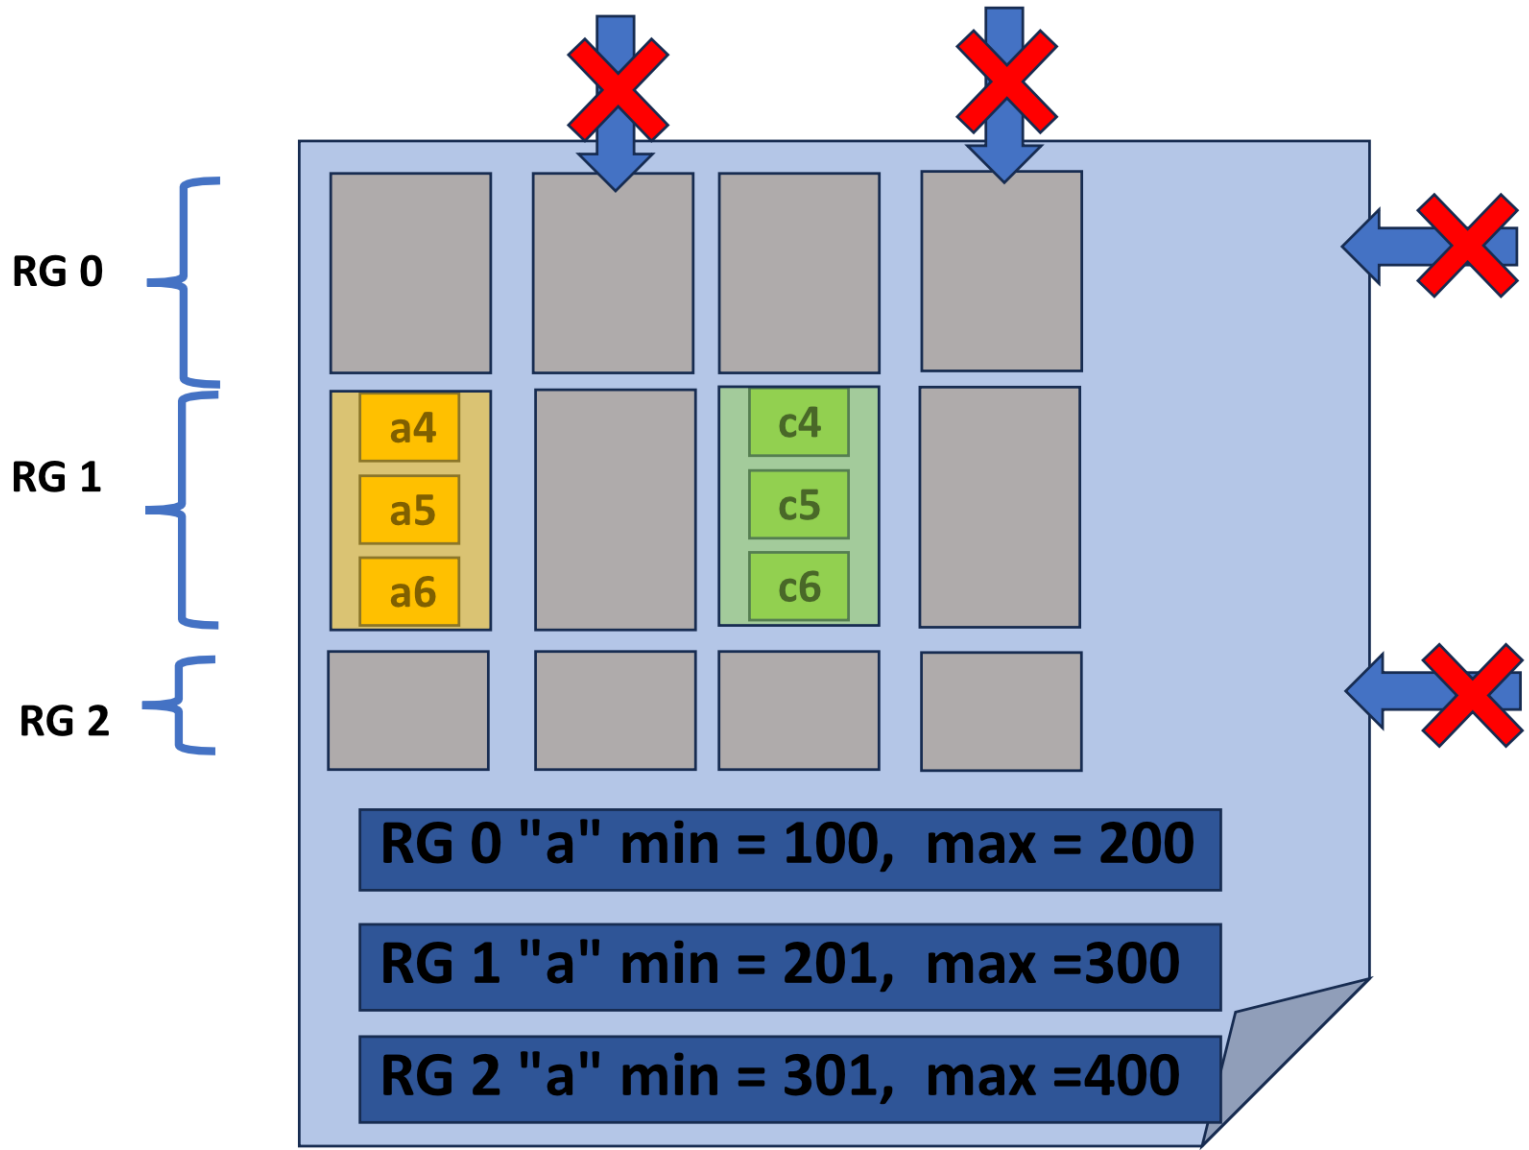

SELECT c
WHERE a = 250

Badly sorted files
=> Bad min/max statistics => False positives

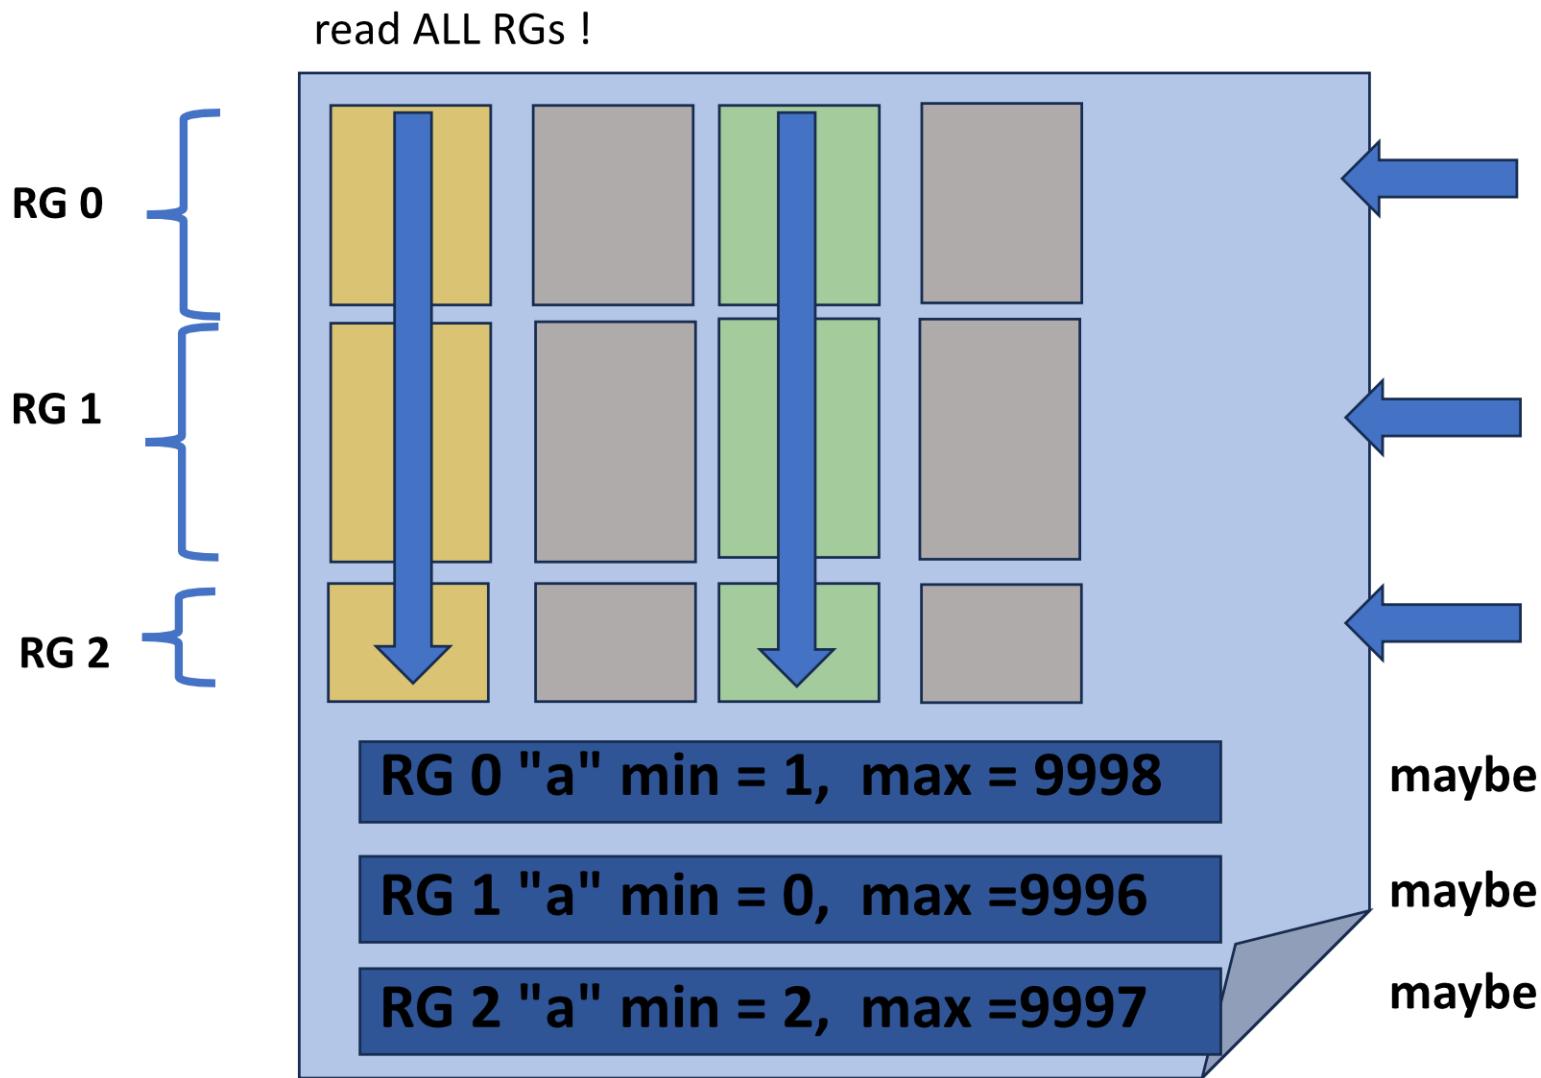

Optim Write Once / Read Many

when writing Parquet files ... think how file might be read later !

spend CPU when writing to save CPU / IO later reading

.repartion(…) // avoid too many small files
.sortWithinPartitions(...) // sort within each files

```
dataset
    .repartition(nRepartitionCount)
    // or  .repartition("col1", nRepartitionHash)

    .sortWithinPartitions("colA", "colB")

    .write
    .option("parquet.block.size", 32*MEGA)

    .format("parquet")
    .save("file://some-dir")
    // or
    // .format("hive").insertInto("hiveDb.hiveTableName")
```

avoid many small files
dataset.repartition(1)
or .coalesce(1)

dataset
.repartition(nRepartitionCount) // or .coalesce(nRepartitionCount)

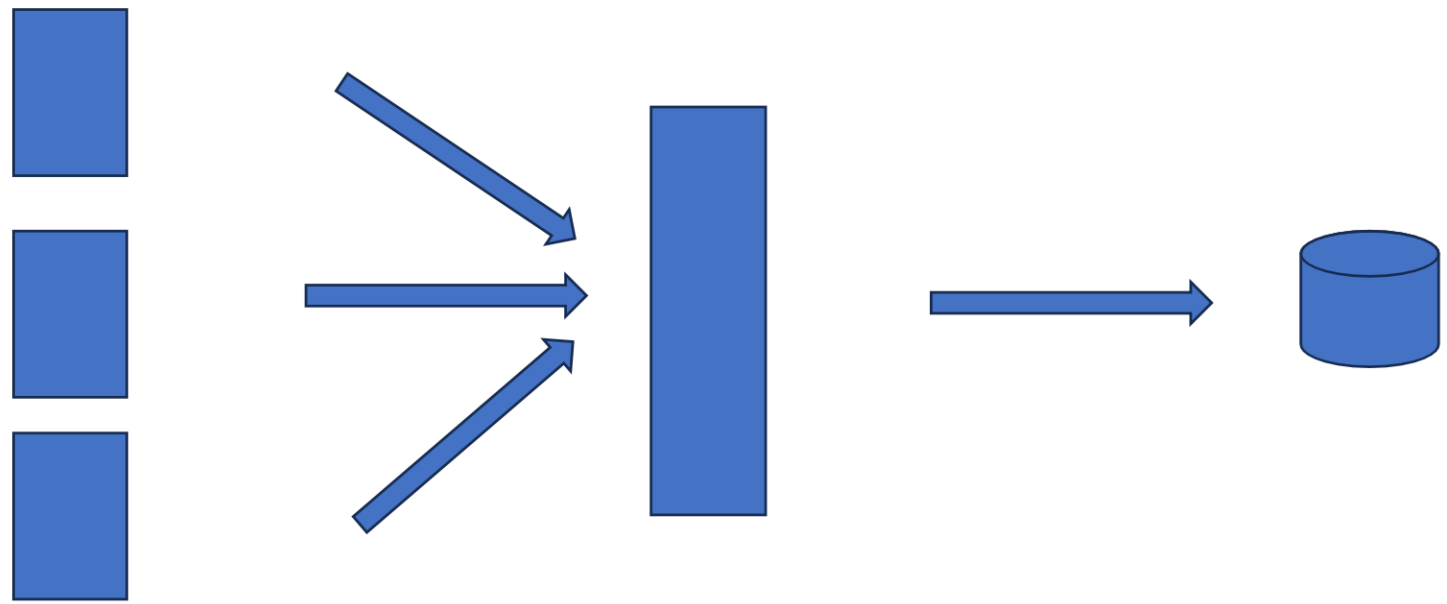

N spark in-memory Partitions
distributed over N executors

1 spark in-memory Partition
(should fit in spark.executor.memory !!)

write 1 parquet file

Does not fit in memory.. compromise to .repartition(smallN)

typically 200 partitions

default
spark.shuffle.size=200

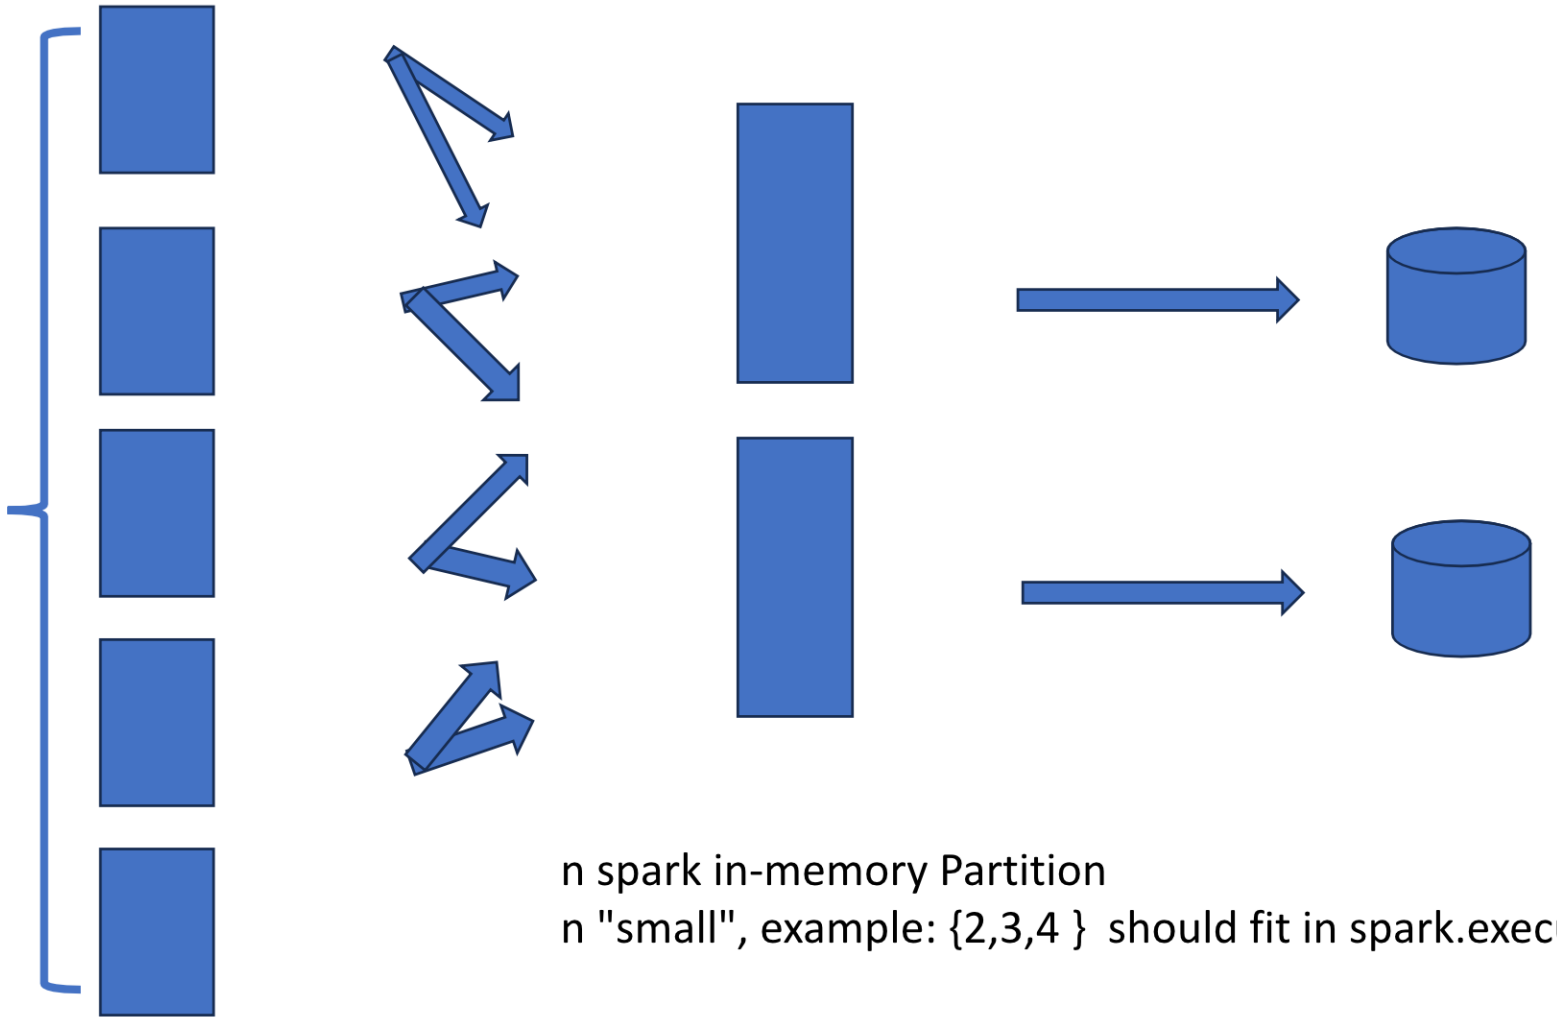

n spark in-memory Partition
n "small", example: {2,3,4 } should fit in spark.executor.memory !!

.repartition("column")

typical usage:
column "a" has FEW distinct values { 1, 2, 3, 4 }

ds = dataset.repartition("a")

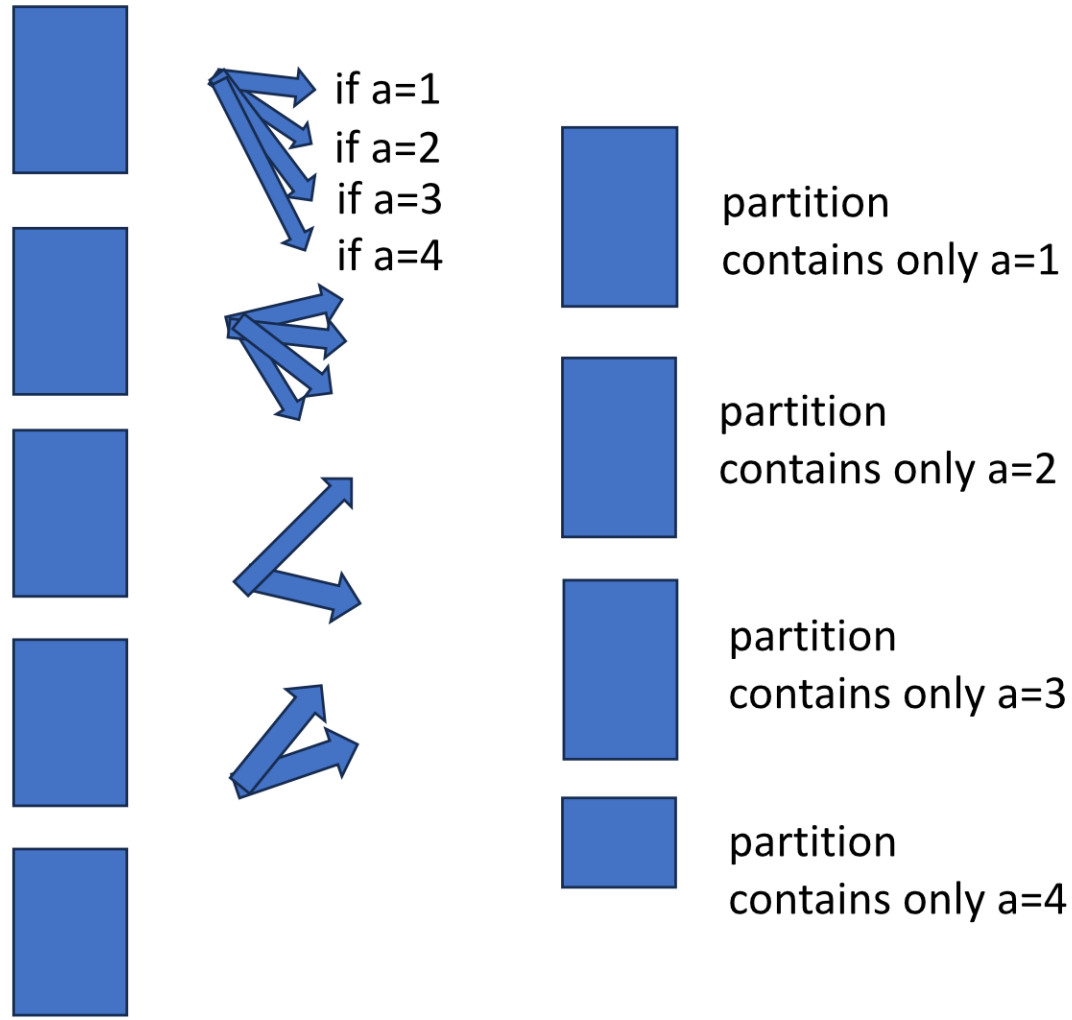

.repartition("column", nHashCount)

example ... having many distinct values

col "a" values in [1, 2, 3, ... 500000]

=> h = hash(a) % 15 in [0, 1, .. 14]

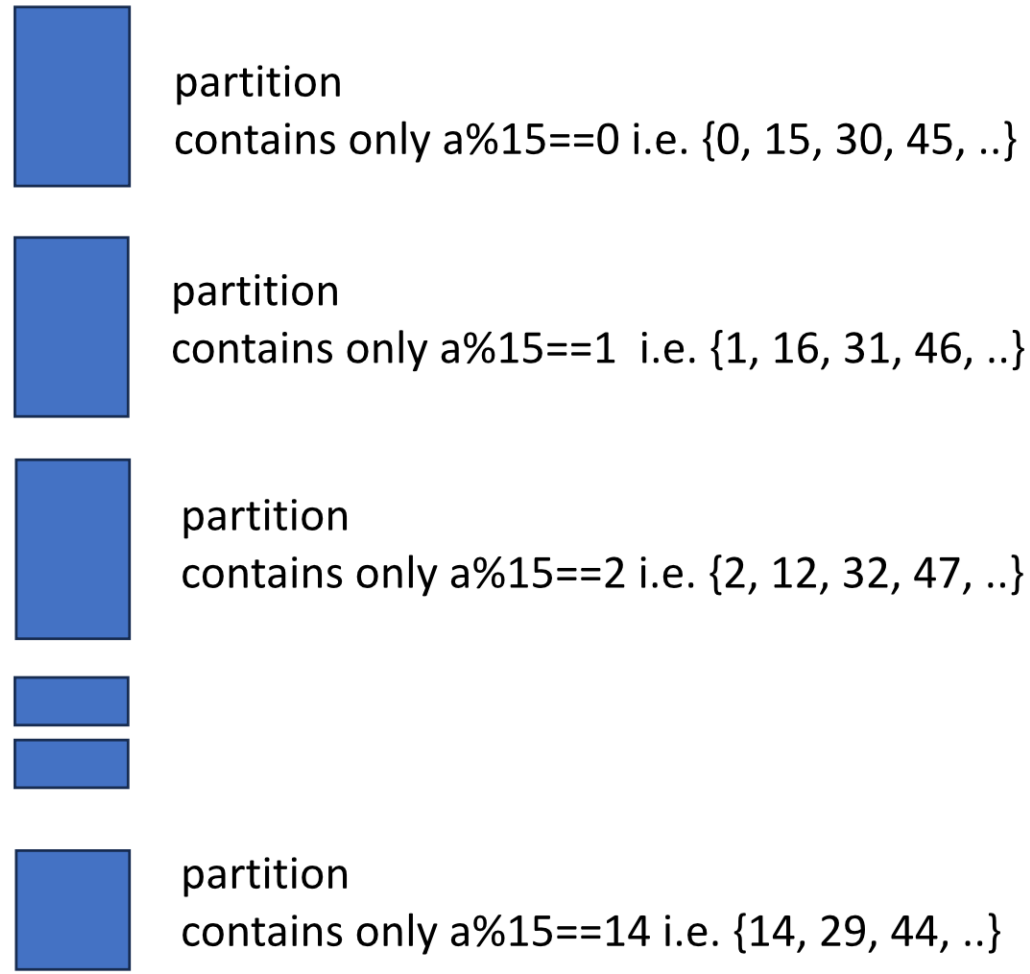

.sortWithinPartitions("a", "b", ...)

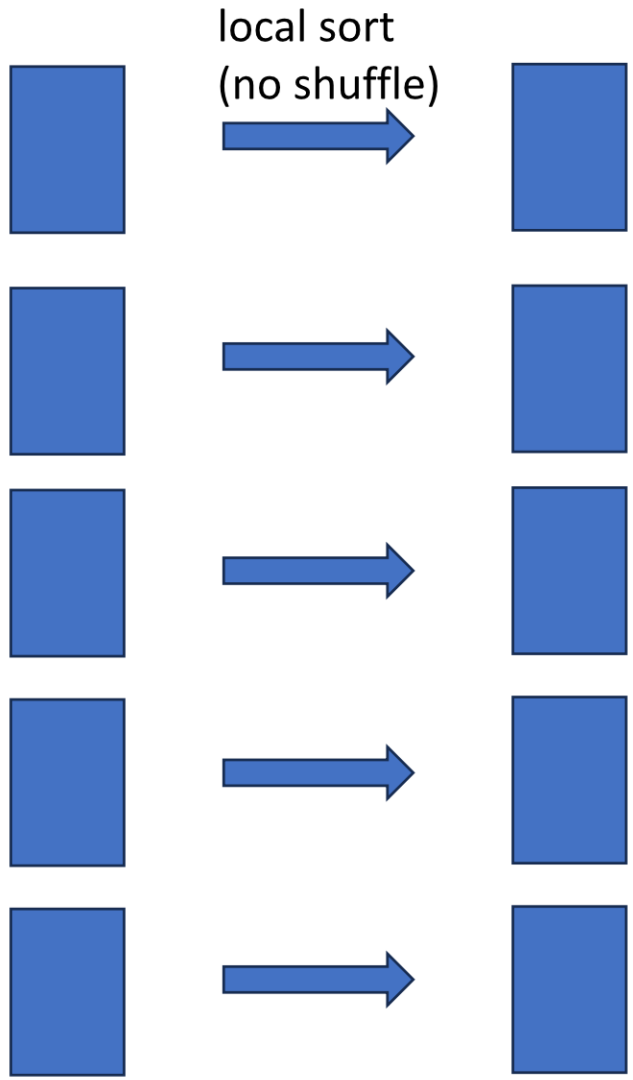

.sortWithinPartitions
=> RowGroups stats more "compact"

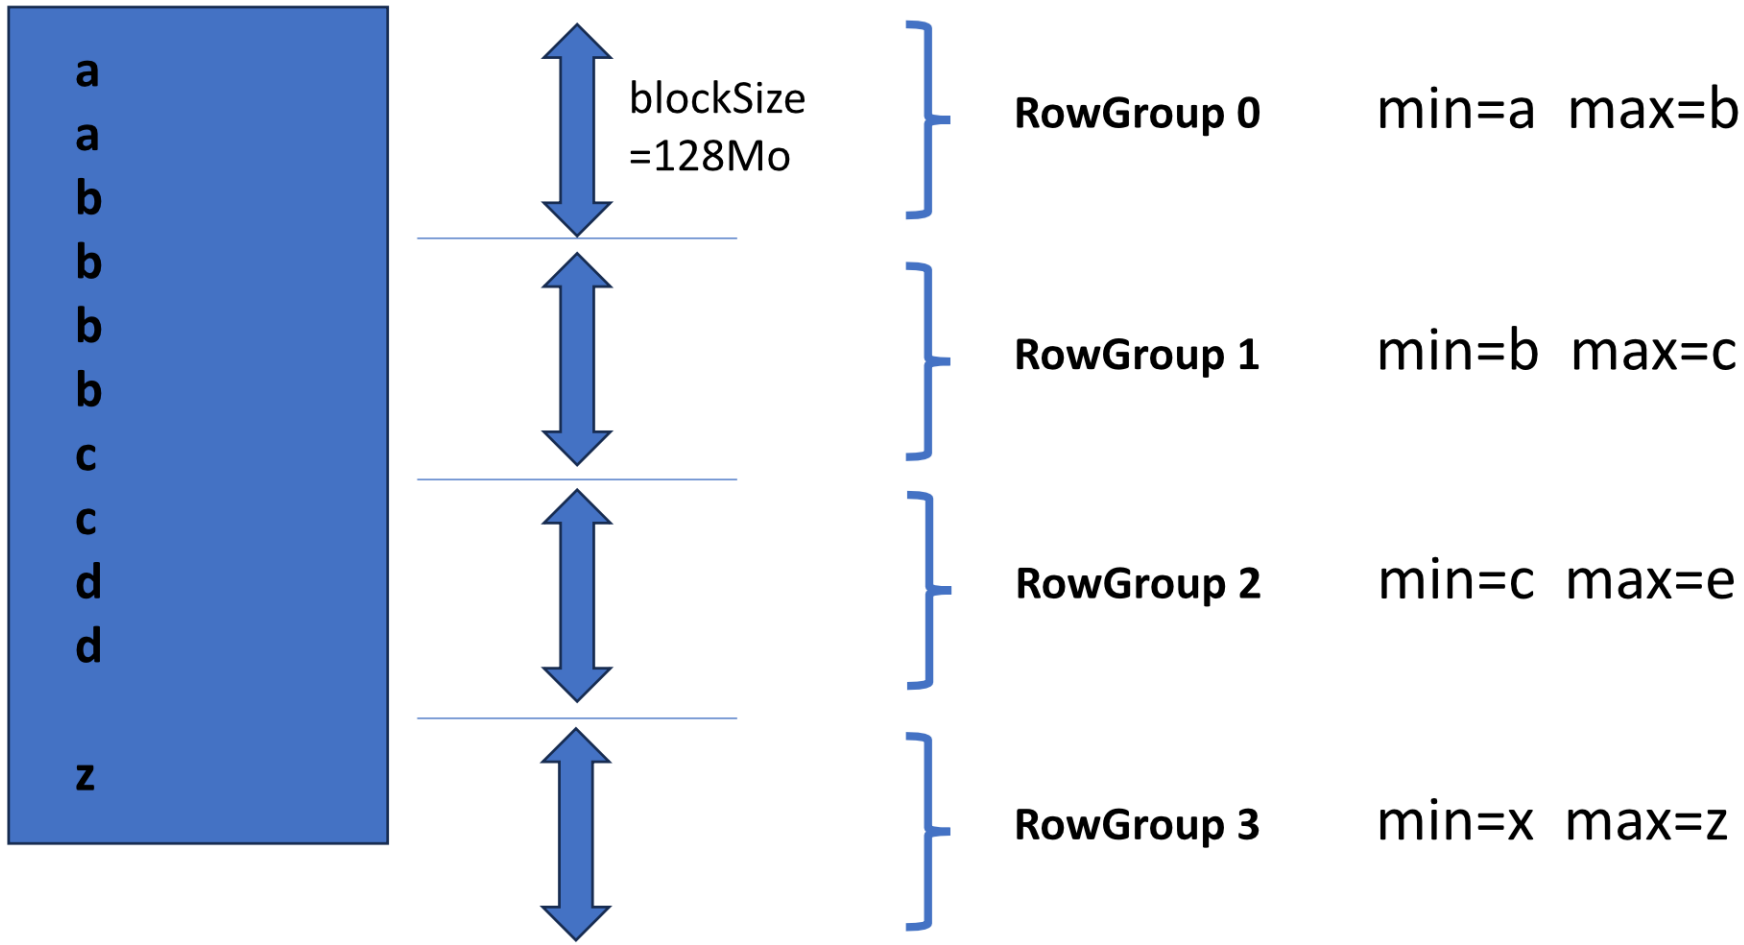

see also Dataset.write .bucketBy(..)

explain spark Dataset write bucketBy

In Apache Spark, the `bucketBy` method is used to organize data into buckets (or partitions) based on the values of specific columns. This improves performance for certain operations like joins, aggregations, and lookups, particularly when the same bucketing is used across datasets.

What is Bucketing?

Bucketing is a technique to partition data into a fixed number of buckets based on the hash value of one or more columns. Unlike partitioning by directory structure, bucketing determines which bucket a record belongs to by applying a hash function to the column values.

Dataset's `write.bucketBy`

The `bucketBy` method in Spark is part of the `DataFrameWriter` API and is used to specify bucketing when writing data. It allows you to save a dataset in a bucketed format.

Syntax

java Copy code

```java
Dataset<Row> dataset = ...;

// Write data with bucketing
dataset.write()
       .bucketBy(numBuckets, "column1", "column2") // Define the number of buckets and buc
       .sortBy("column1")                          // (Optional) Specify sort order within
       .format("parquet")                          // Specify file format (e.g., parquet, O
       .save("/path/to/output");
```

Parameters

1. `numBuckets`:
   - The number of buckets to divide the data into.
   - This is a fixed number and cannot be changed later.
2. `column1`, `column2`, ...:
   - The columns based on which the data will be hashed and distributed into buckets.

Analyzing Parquet Files, using "parquet-cli"

Analyzing Parquet File, using "parquet-cli"

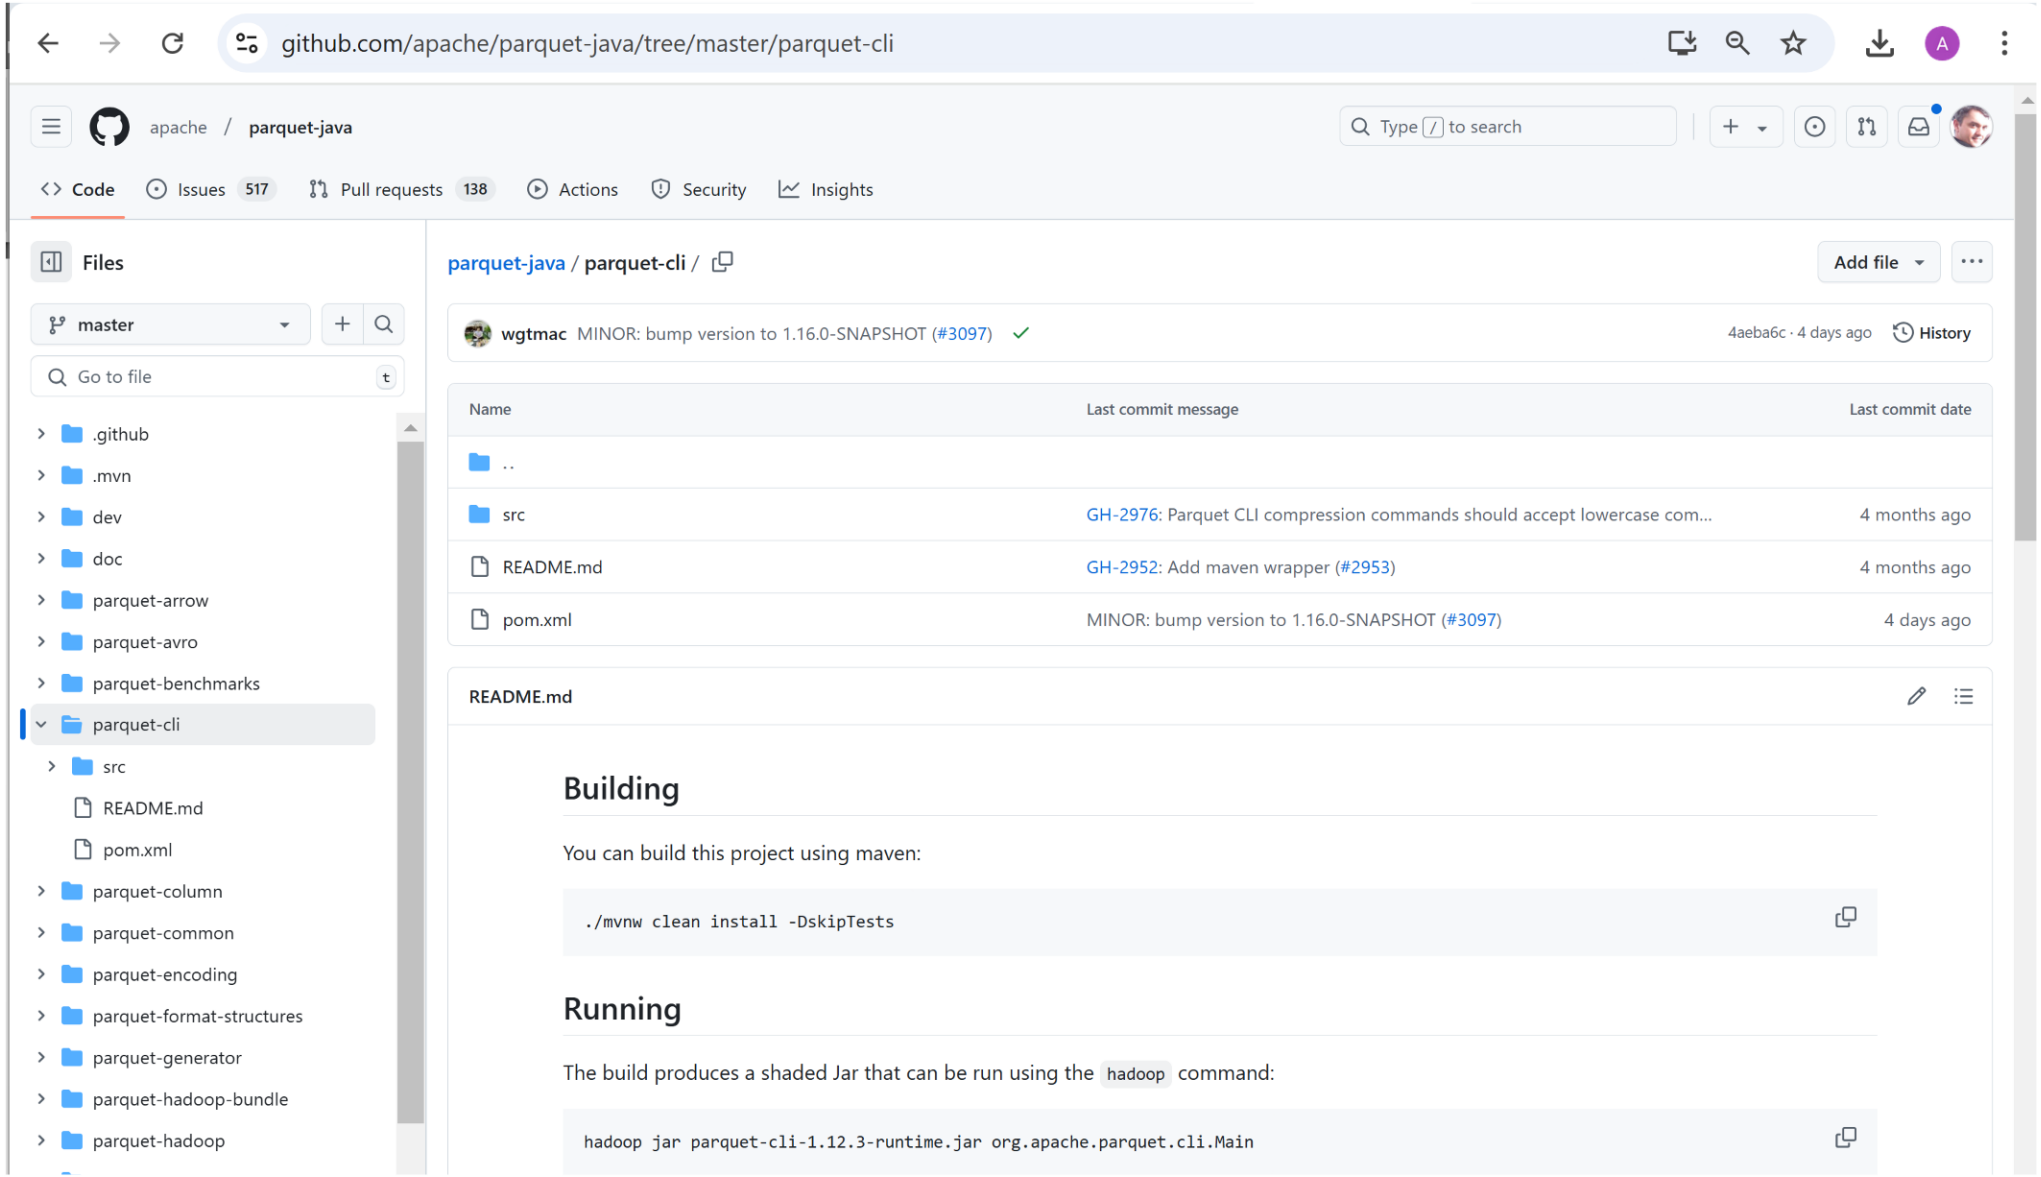

Search jar in maven repo: https://search.maven.org

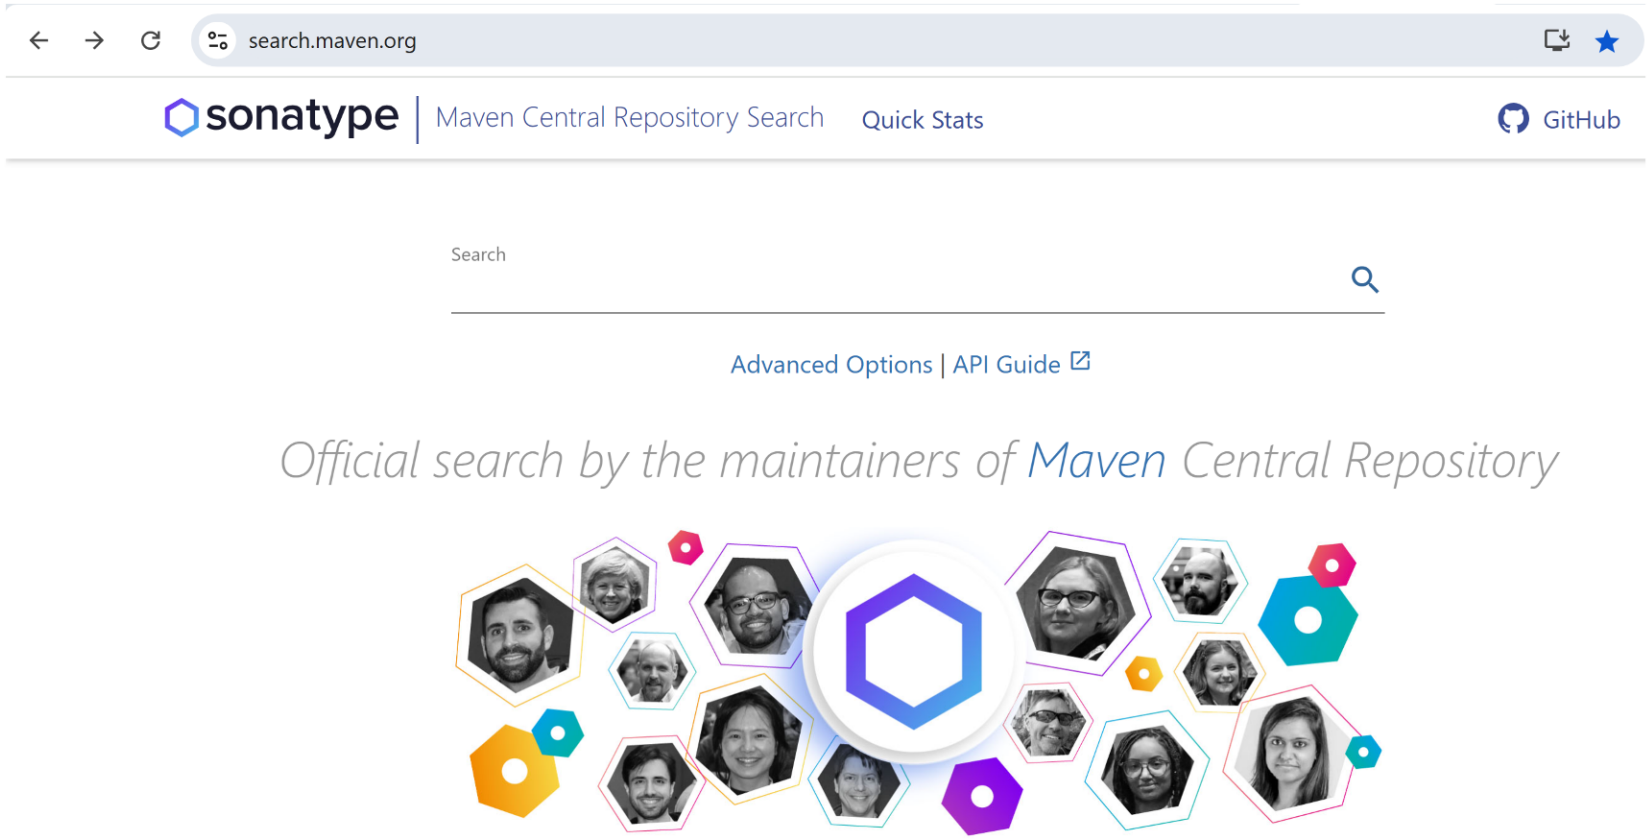

type search "a:parquet-cli" g:org.apache.parquet

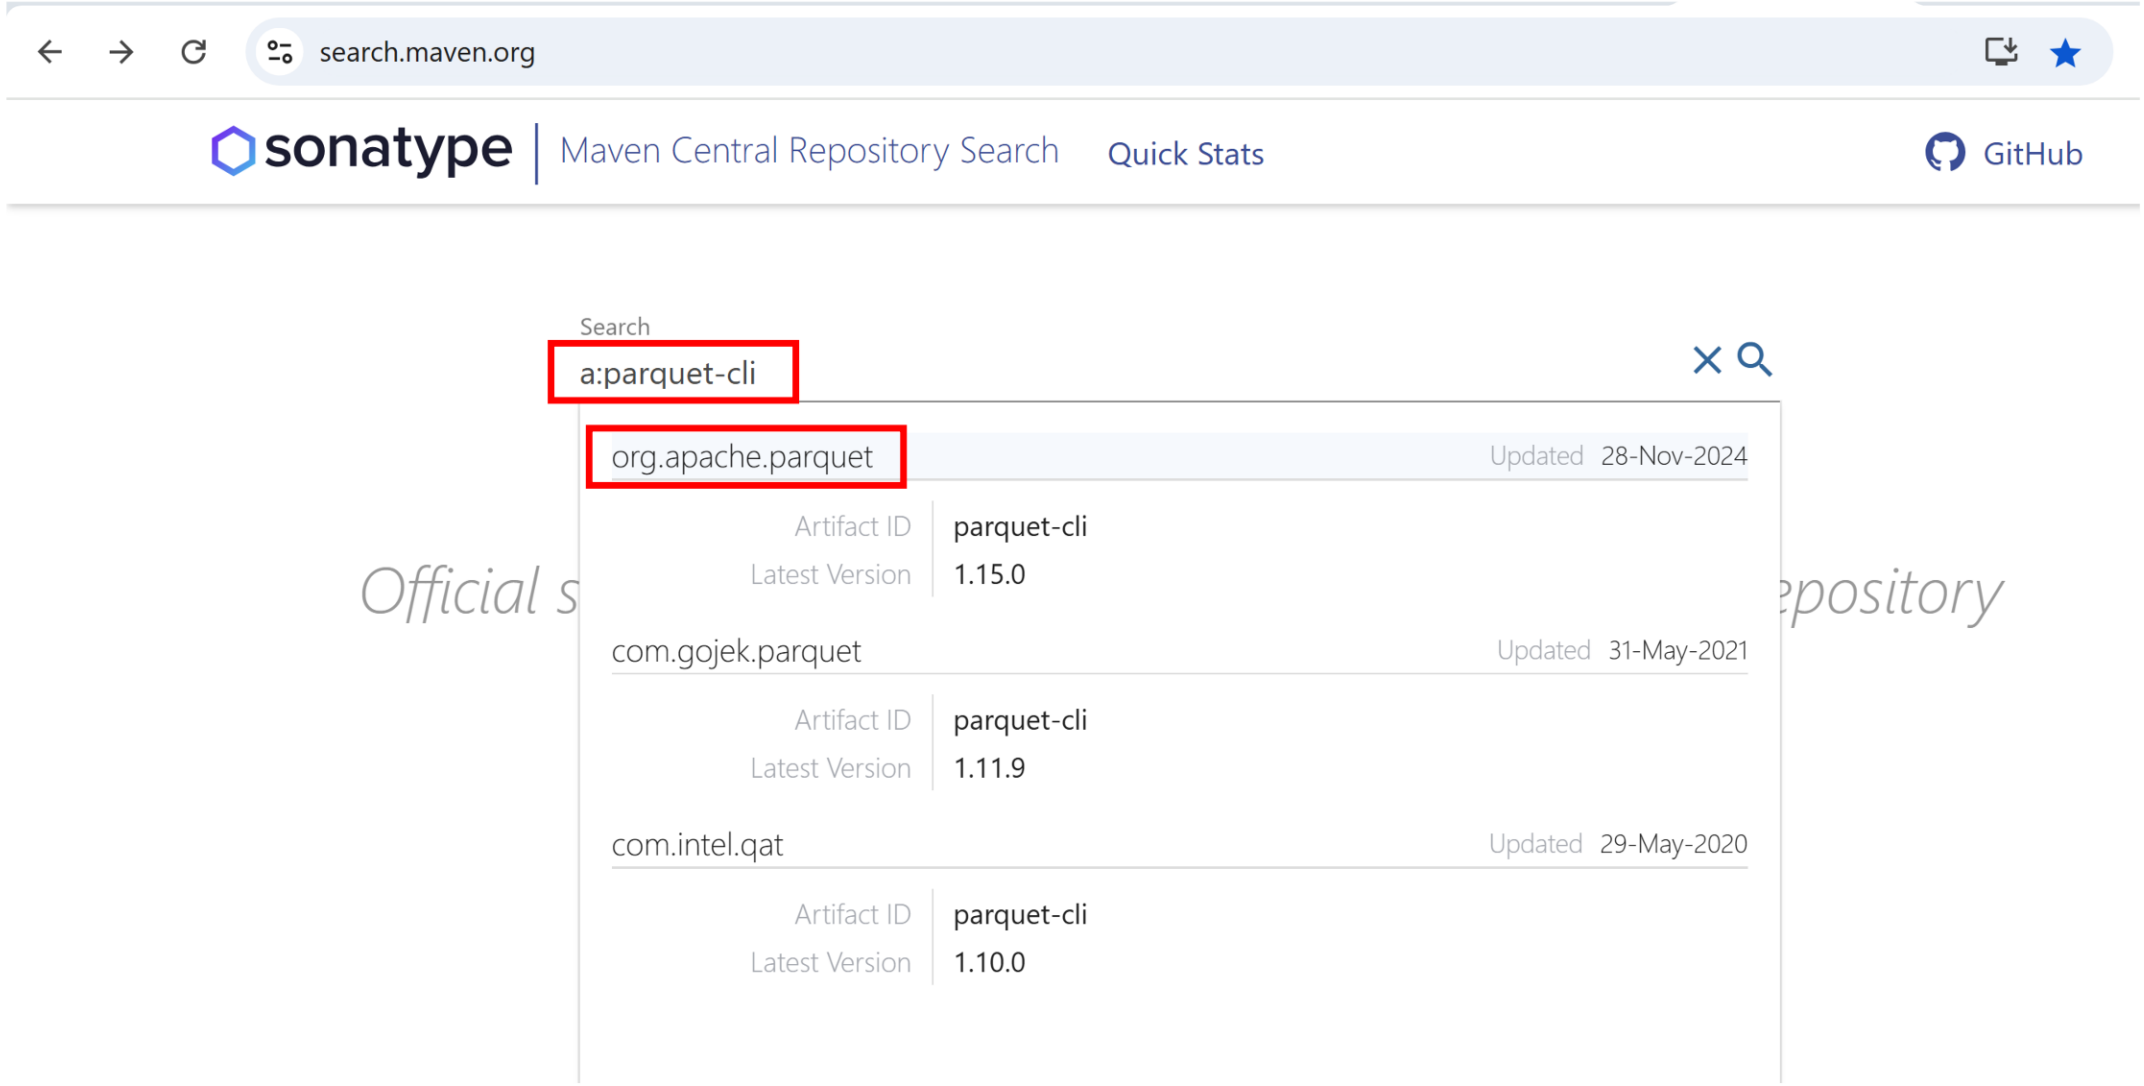

Downloading "parquet-cli" runtime.jar for same version 1.13.0 (same as spark/jars/*.jar)

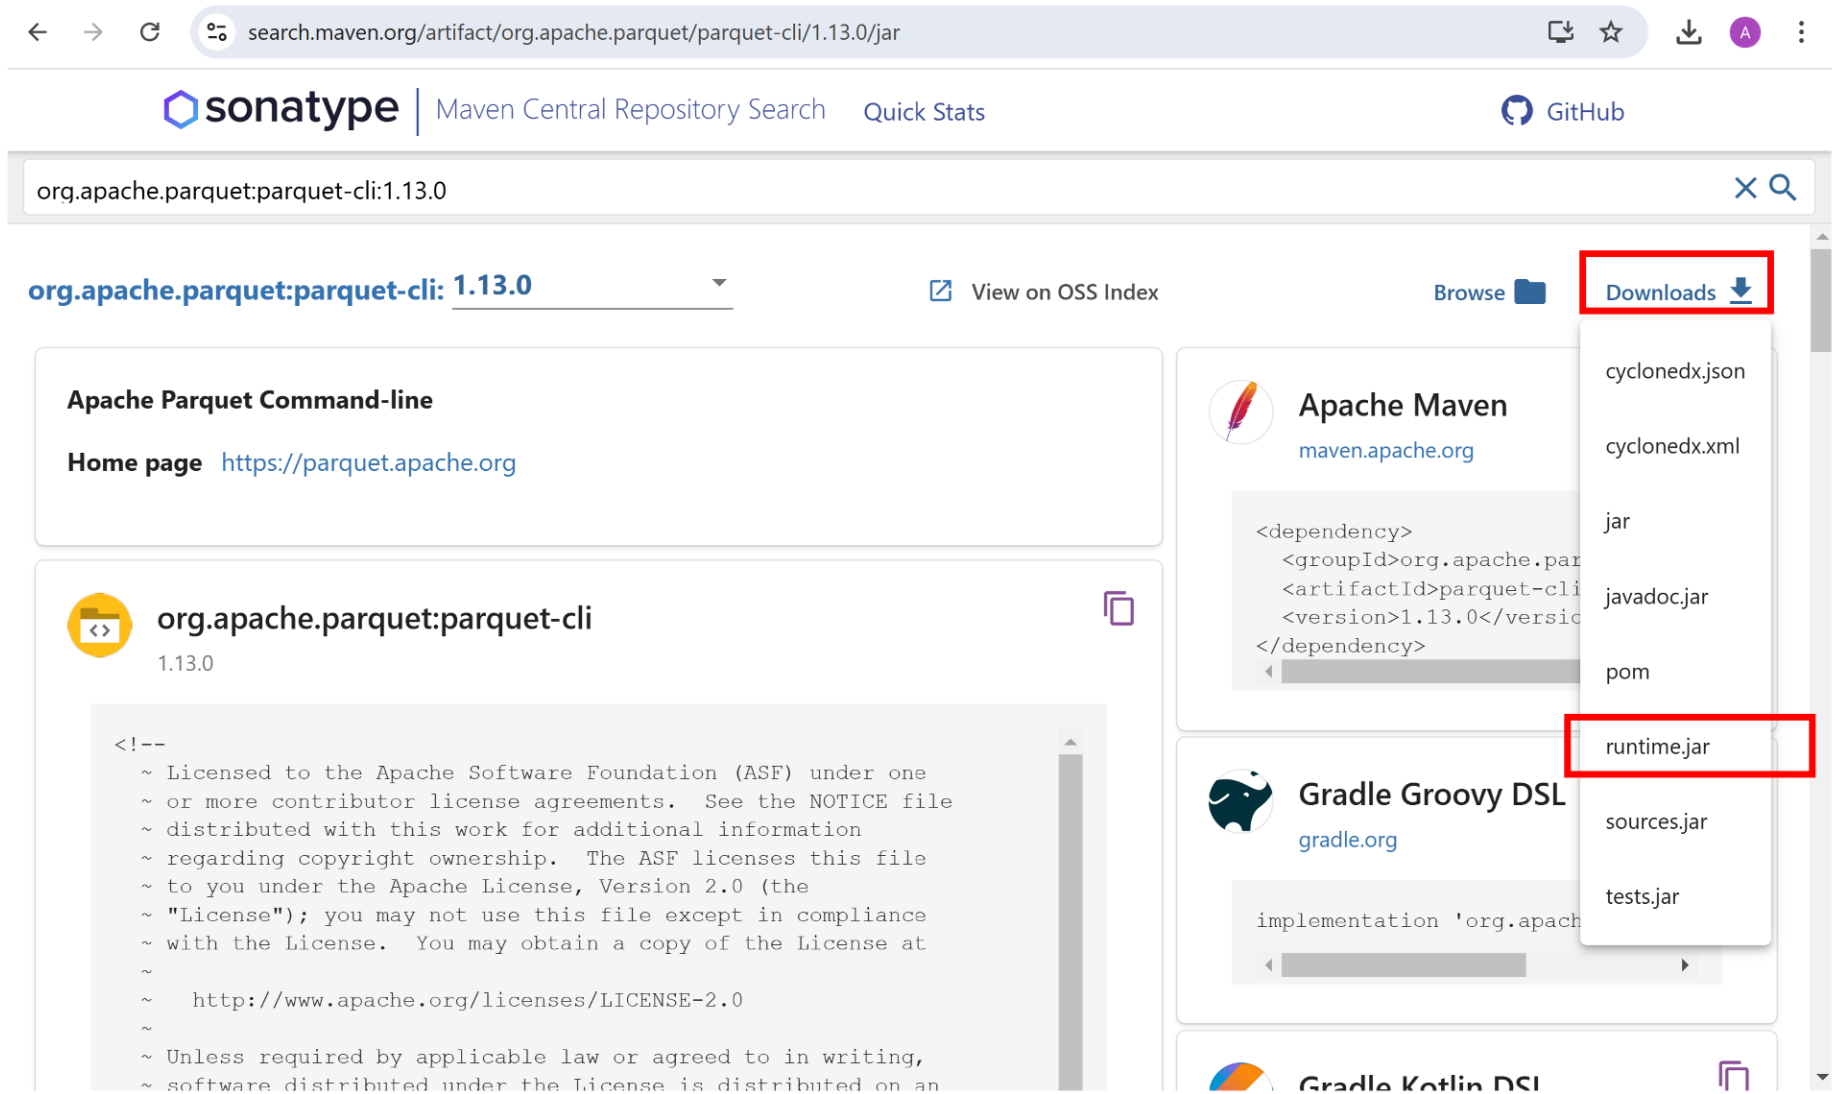

Launching parquet-cli

java -jar parquet-cli-runtime.jar help

(using shaded jar, otherwise may need to add hadoop classpath)

```
C:\apps\spark>java -jar parquet-cli-1.15.0-SNAPSHOT-runtime.jar help

Usage: parquet [options] [command] [command options]

  Options:

    -v, --verbose, --debug
        Print extra debugging information

  Commands:

    help
        Retrieves details on the functions of other commands
    meta
        Print a Parquet file's metadata
    pages
        Print page summaries for a Parquet file
    dictionary
        Print dictionaries for a Parquet column
    check-stats
        Check Parquet files for corrupt page and column stats (PARQUET-251)
    schema
        Print the Avro schema for a file
    csv-schema
        Build a schema from a CSV data sample
    convert-csv
```

parquet-cli meta [1/2]

java -jar parquet-cli-runtime.jar meta file.parquet

```
C:\apps\spark>java -jar parquet-cli-1.15.0-SNAPSHOT-runtime.jar meta  C:\apps\spark\spark
t

File path:  C:\apps\spark\spark-warehouse\db1.db\addr\part-c000.snappy.parquet
Created by: parquet-mr version 1.13.1 (build db4183109d5b734ec5930d870cdae161e408ddba)
Properties:
                   org.apache.spark.version: 3.5.0
  org.apache.spark.sql.parquet.row.metadata: {"type":"struct","fields":[{"name":"uid_adre
e":"cle_interop","type":"string","nullable":true,"metadata":{}},{"name":"commune_insee","
mmune_nom","type":"string","nullable":true,"metadata":{}},{"name":"commune_deleguee_insee
:"commune_deleguee_nom","type":"string","nullable":true,"metadata":{}},{"name":"voie_nom"
lieudit_complement_nom","type":"string","nullable":true,"metadata":{}},{"name":"numero","
uffixe","type":"string","nullable":true,"metadata":{}},{"name":"position","type":"string"
e","nullable":true,"metadata":{}},{"name":"y","type":"double","nullable":true,"metadata":
ata":{}},{"name":"lat","type":"double","nullable":true,"metadata":{}},{"name":"cad_parcel
e":"source","type":"string","nullable":true,"metadata":{}},{"name":"certification_commune
:"date_der_maj","type":"date","nullable":false,"metadata":{}}]}
Schema:
message spark_schema {
  optional binary uid_adresse (STRING);
  optional binary cle_interop (STRING);
  optional binary commune_insee (STRING);
  optional binary commune_nom (STRING);
  optional int32 commune_deleguee_insee;
  optional binary commune_deleguee_nom (STRING);
  optional binary voie_nom (STRING);
  optional binary lieudit_complement_nom (STRING);
  optional int32 numero;
  optional binary suffixe (STRING);
  optional binary position (STRING);
  optional double x;
```

parquet-cli meta [2/2]

```
Row group 0:  count: 219962  71,92 B records  start: 4  total(compressed): 15,086 MB total(uncompressed):33,177 MB
--------------------------------------------------------------------------------
                         type      encodings count     avg size   nulls   min / max
uid_adresse              BINARY    S _ R_ F  219962    34,79 B    47399   " @a:00003b17-45be-48c1-aa..." / "Argence""
cle_interop              BINARY    S   _     219962    5,47 B     0       "01001_0005_00026" / "54"
commune_insee            BINARY    S _ R     219962    0,01 B     1       "01001" / "01350"
commune_nom              BINARY    S _ R     219962    0,02 B     0       "Ambléon" / "Évosges"
commune_deleguee_insee   INT32     S _ R     219962    0,02 B     202278  "1015" / "1442"
commune_deleguee_nom     BINARY    S _ R     219962    0,02 B     202277  "6518380.18" / "Étrez"
voie_nom                 BINARY    S _ R     219962    1,44 B     0       """le château""" / "Îlot Grammont"
lieudit_complement_nom   BINARY    S _ R     219962    0,09 B     197279  """le château""" / "Étrez"
numero                   INT32     S _ R     219962    1,45 B     1       "0" / "99999"
suffixe                  BINARY    S _ R     219962    0,11 B     206335  "1" / "z"
position                 BINARY    S _ R     219962    0,25 B     15376   "bâtiment" / "service technique"
x                        DOUBLE    S   _     219962    5,06 B     226     "-0.0" / "943078.82"
y                        DOUBLE    S   _     219962    5,37 B     227     "6505390.74" / "6603083.59"
long                     DOUBLE    S   _     219962    7,99 B     227     "4.730471" / "6.163086"
lat                      DOUBLE    S   _     219962    7,64 B     227     "45.617249" / "46.506743"
cad_parcelles            BINARY    S _ R     219962    1,87 B     167172  "010010000A1033" / "01380000_E0291"
source                   BINARY    S _ R     219962    0,06 B     2266    "arcep" / "inconnue"
certification_commune    INT32     S _ R     219962    0,03 B     1       "0" / "1"
date_der_maj             INT32     S _ R     219962    0,23 B     0       "1901-01-01" / "2024-12-05"

Row group 1:  count: 215938  72,72 B records  start: 15819187  total(compressed): 14,976 MB total(uncompressed):33,489 MB
--------------------------------------------------------------------------------
                         type      encodings count     avg size   nulls   min / max
uid_adresse              BINARY    S   _     215938    36,79 B    39793   " @a:00003774-58fc-46a7-b1..." / " @v:ffe769c6-a8b
cle_interop              BINARY    S   _     215938    5,19 B     0       "01350_0005_00707" / "02546_rcnpc4_00054"
commune_insee            BINARY    S _ R     215938    0,02 B     0       "01350" / "02546"
commune_nom              BINARY    S _ R     215938    0,04 B     0       "Abbécourt" / "Évergnicourt"
commune_deleguee_insee   INT32     S   _     215938    0,01 B     208894  "1059" / "2811"
commune_deleguee_nom     BINARY    S   _     215938    0,03 B     208894  "Anizy-le-Château" / "Virieu-le-Petit"
```

parquet-cli meta | grep "Row group"

java -jar parquet-cli-runtime.jar meta file.parquet > meta.txt

grep "Row group" meta.txt | head

Row group 0: count: 219962 71,92 B records start: 4 total(compressed): 15,086 MB total(uncompressed):33,177 MB
Row group 1: count: 215938 72,72 B records start: 15819187 total(compressed): 14,976 MB total(uncompressed):33,489 MB
Row group 2: count: 213992 73,68 B records start: 31523114 total(compressed): 15,036 MB total(uncompressed):33,821 MB
Row group 3: count: 233199 70,29 B records start: 47289138 total(compressed): 15,631 MB total(uncompressed):34,232 MB
Row group 4: count: 231558 69,04 B records start: 63679653 total(compressed): 15,246 MB total(uncompressed):33,383 MB
Row group 5: count: 219962 74,89 B records start: 79665824 total(compressed): 15,710 MB total(uncompressed):34,531 MB
Row group 6: count: 218409 73,91 B records start: 96138854 total(compressed): 15,394 MB total(uncompressed):33,957 MB
Row group 7: count: 199907 75,02 B records start: 112281114 total(compressed): 14,301 MB total(uncompressed):31,914 MB
Row group 8: count: 218409 69,87 B records start: 127277278 total(compressed): 14,552 MB total(uncompressed):34,174 MB
Row group 9: count: 272568 65,80 B records start: 142536439 total(compressed): 17,104 MB total(uncompressed):35,616 MB

grep "Row group" meta.txt | tail

Row group 108: count: 252150 64,55 B records start: 1719335249 total(compressed): 15,522 MB total(uncompressed):34,027 MB
Row group 109: count: 215240 71,68 B records start: 1735611436 total(compressed): 14,715 MB total(uncompressed):34,014 MB
Row group 110: count: 219962 68,48 B records start: 1751040857 total(compressed): 14,365 MB total(uncompressed):32,978 MB
Row group 111: count: 219962 68,18 B records start: 1766103960 total(compressed): 14,302 MB total(uncompressed):32,914 MB
Row group 112: count: 234887 64,52 B records start: 1781100645 total(compressed): 14,454 MB total(uncompressed):33,355 MB
Row group 113: count: 236624 66,44 B records start: 1796256273 total(compressed): 14,993 MB total(uncompressed):37,495 MB
Row group 114: count: 234063 65,95 B records start: 1811977242 total(compressed): 14,722 MB total(uncompressed):33,626 MB
Row group 115: count: 213992 73,42 B records start: 1827414634 total(compressed): 14,983 MB total(uncompressed):34,526 MB
Row group 116: count: 230100 69,58 B records start: 1843125866 total(compressed): 15,269 MB total(uncompressed):36,039 MB
Row group 117: count: 152925 67,36 B records start: 1859136956 total(compressed): 9,823 MB total(uncompressed):26,622 MB

parquet-cli column-size

java -jar parquet-cli-runtime.jar column-size file.parquet

```
C:\apps\spark>java -jar parquet-cli-1.15.0-SNAPSHOT-runtime.jar column-size  C:\apps\spark\spa
commune_deleguee_insee-> Size In Bytes: 431641 Size In Ratio: 2.3089352E-4
commune_insee-> Size In Bytes: 322609 Size In Ratio: 1.725701E-4
numero-> Size In Bytes: 26951190 Size In Ratio: 0.014416738
voie_nom-> Size In Bytes: 35412828 Size In Ratio: 0.01894304
uid_adresse-> Size In Bytes: 882820344 Size In Ratio: 0.4722385
date_der_maj-> Size In Bytes: 7501158 Size In Ratio: 0.0040125214
cle_interop-> Size In Bytes: 139359234 Size In Ratio: 0.07454608
source-> Size In Bytes: 1142012 Size In Ratio: 6.108854E-4
long-> Size In Bytes: 213866608 Size In Ratio: 0.11440159
commune_deleguee_nom-> Size In Bytes: 836453 Size In Ratio: 4.4743568E-4
suffixe-> Size In Bytes: 3721170 Size In Ratio: 0.0019905292
cad_parcelles-> Size In Bytes: 66126109 Size In Ratio: 0.035372198
certification_commune-> Size In Bytes: 673339 Size In Ratio: 3.601827E-4
x-> Size In Bytes: 133896306 Size In Ratio: 0.07162385
lieudit_complement_nom-> Size In Bytes: 2886263 Size In Ratio: 0.0015439206
y-> Size In Bytes: 143779166 Size In Ratio: 0.07691039
commune_nom-> Size In Bytes: 549711 Size In Ratio: 2.9405157E-4
position-> Size In Bytes: 4716951 Size In Ratio: 0.0025231927
lat-> Size In Bytes: 204444396 Size In Ratio: 0.109361455
```

Column Sizes ... sorting

cat column-size
 | sed 's/-> Size In Bytes//g' | sed 's/ Size In Ratio//g' | sed 's/: /;/g' > column-size.csv

Column	Size	Ratio	Size in Mo
uid_adresse	882820344	47,22%	841,9
long	213866608	11,44%	204,0
lat	204444396	10,94%	195,0
y	143779166	7,69%	137,1
cle_interop	139359234	7,45%	132,9
x	133896306	7,16%	127,7
cad_parcelles	66126109	3,54%	63,1
voie_nom	35412828	1,89%	33,8
numero	26951190	1,44%	25,7
date_der_maj	7501158	0,40%	7,2
position	4716951	0,25%	4,5
suffixe	3721170	0,20%	3,5
lieudit_complement_nom	2886263	0,15%	2,8
source	1142012	0,06%	1,1
commune_deleguee_nom	836453	0,04%	0,8
certification_commune	673339	0,04%	0,6
commune_nom	549711	0,03%	0,5
commune_deleguee_insee	431641	0,02%	0,4
commune_insee	322609	0,02%	0,3
total	1869437488	100,00%	1782,8

Parquet Column Sizes Summary

zipcode ("commun_insee") : 0.3 Mo
+
city name ("commun_nom") : 0.5 Mo
+
street name ("voie_nom") : 33.8 Mo
+
numero : 25.7 Mo

=> 60 Mega (3% of file)

+ longitude, latitude : 204 Mo + 195 Mo
=> 459 Mo

most of the column size is

uid_address : 840 Mo (47.2%)
+
x, y ... redundant with longitude,lattitude
+
cle_interop : 132 Mo (7.4%)

... we will see parquet "column-pruning" in action : skipping unnecessary columns !

More Graphical

Web UI tools ?

Parquet Parser library in JavaScript : « hyparquet »

github.com/hyparam/hyparquet

README | MIT license

Dependency free since 2023!

What is hyparquet?

Hyparquet is a lightweight, dependency-free, pure JavaScript library for parsing Apache Parquet files. Apache Parquet is a popular columnar storage format that is widely used in data engineering, data science, and machine learning applications for efficiently storing and processing large datasets.

Hyparquet aims to be the world's most compliant parquet parser. And it runs in the browser.

Parquet Viewer

Try hyparquet online: Drag and drop your parquet file onto hyperparam.app to view it directly in your browser. This service is powered by hyparquet's in-browser capabilities.

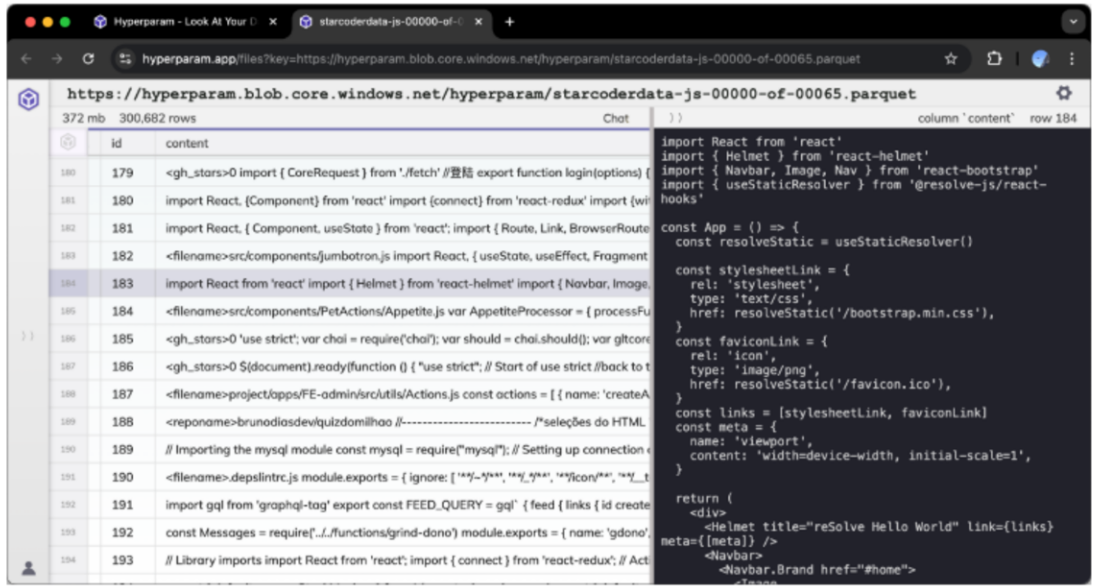

Parquet Metadata Visualizer

https://julien.ledem.net/experiment/parquet-visualizer.html

created by the creator of Parquet: Julien Ledem

Project source code in github:

https://github.com/julienledem/parquet-metadata-vizualizer

https://github.com/julienledem/parquet-metadata-vizualizer

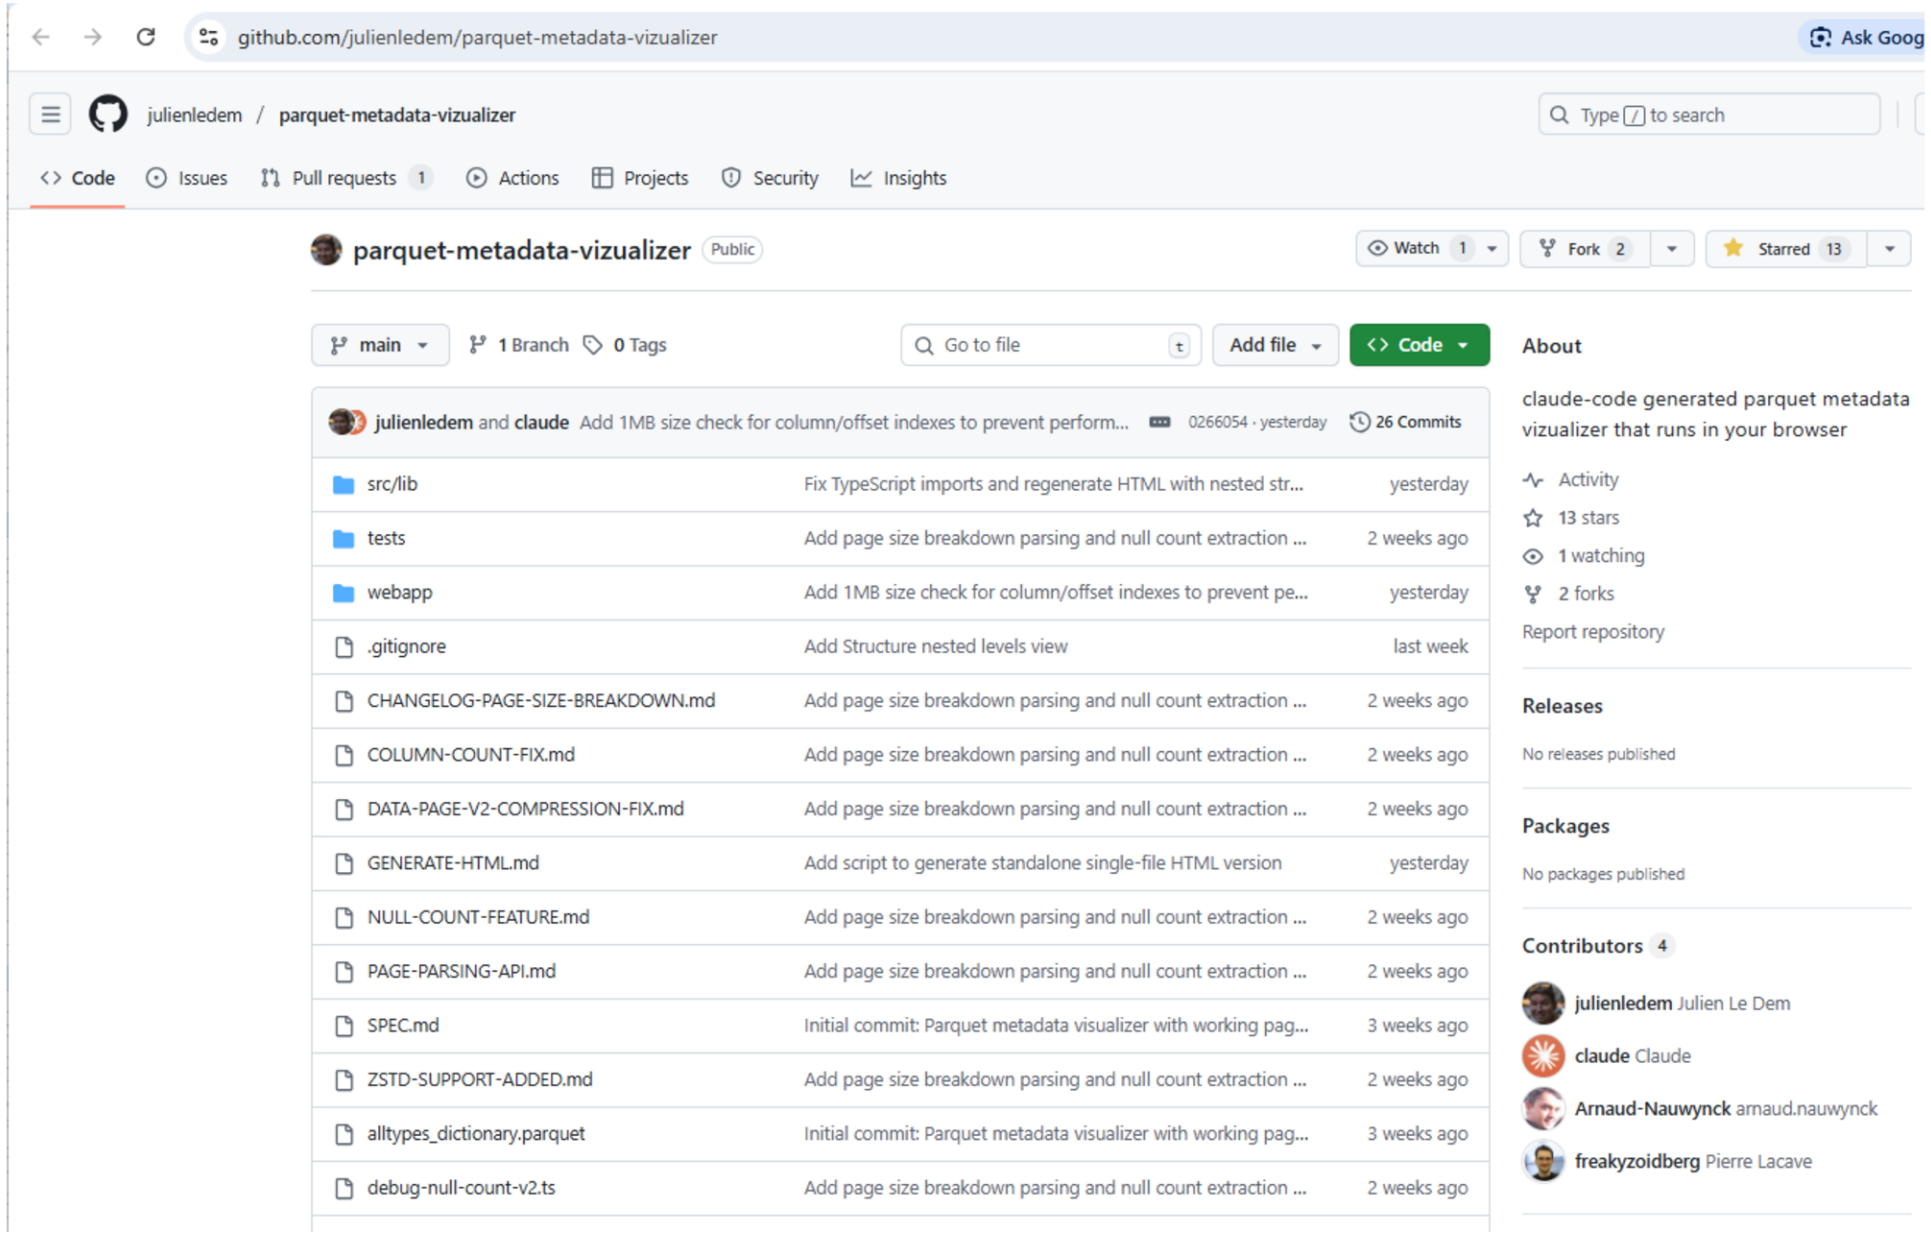

Parquet Metadata Viewer

https://julien.ledem.net/experiment/parquet-visualizer.html

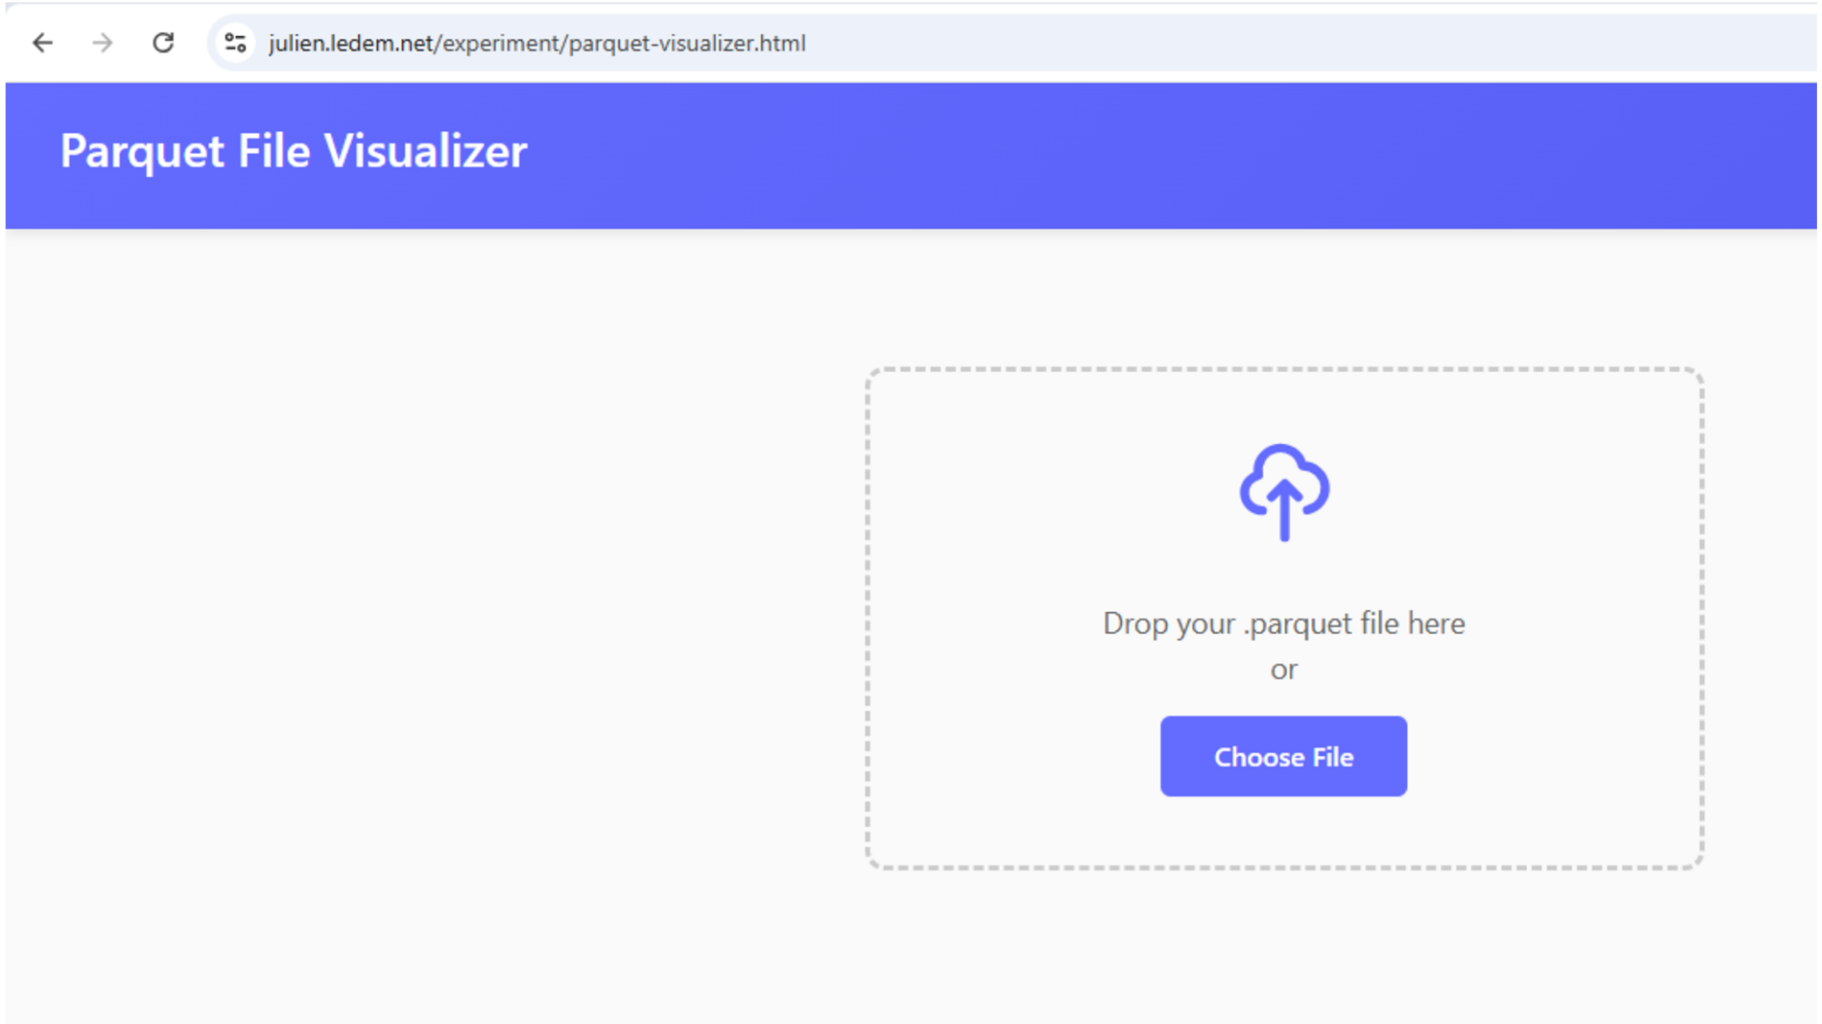

Parquet Metadata Viewer

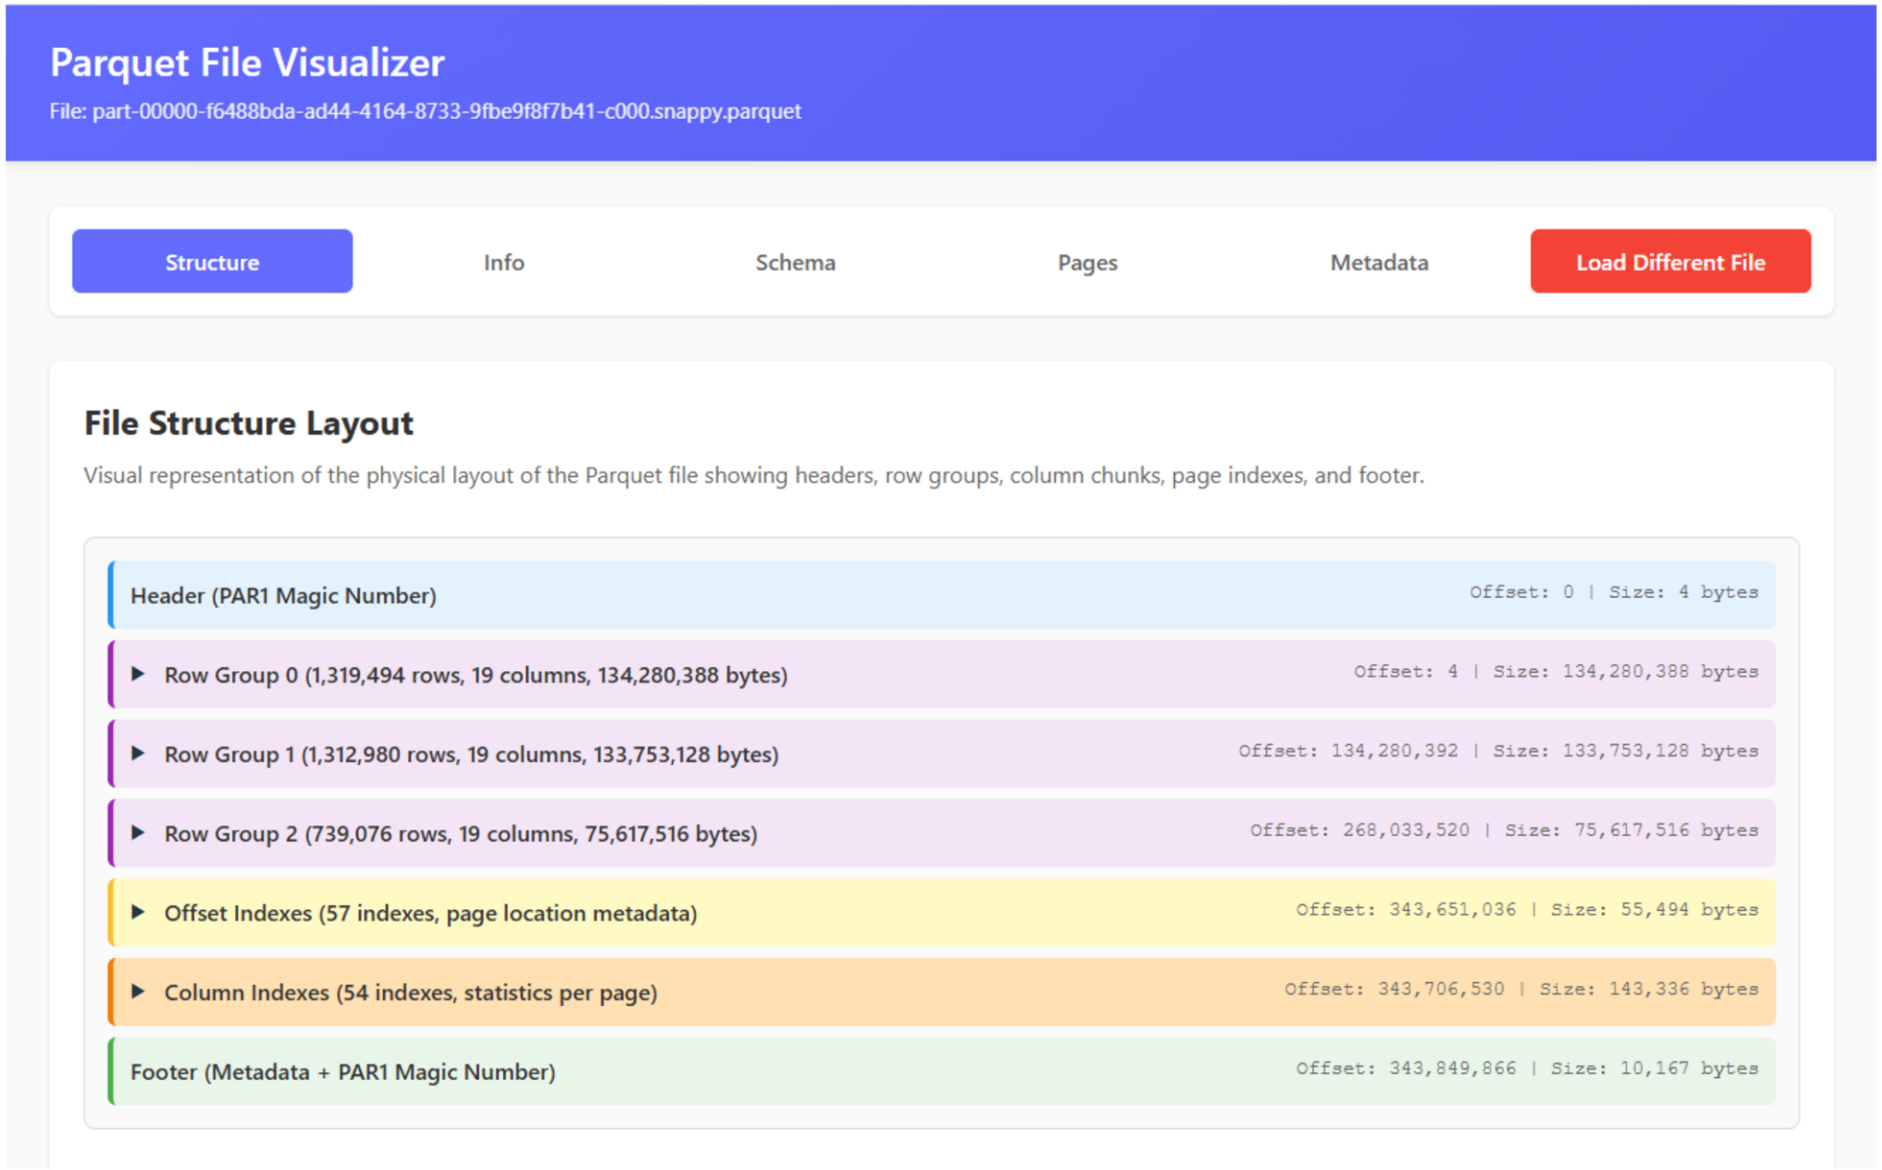

RowGroup- ColumnChunk

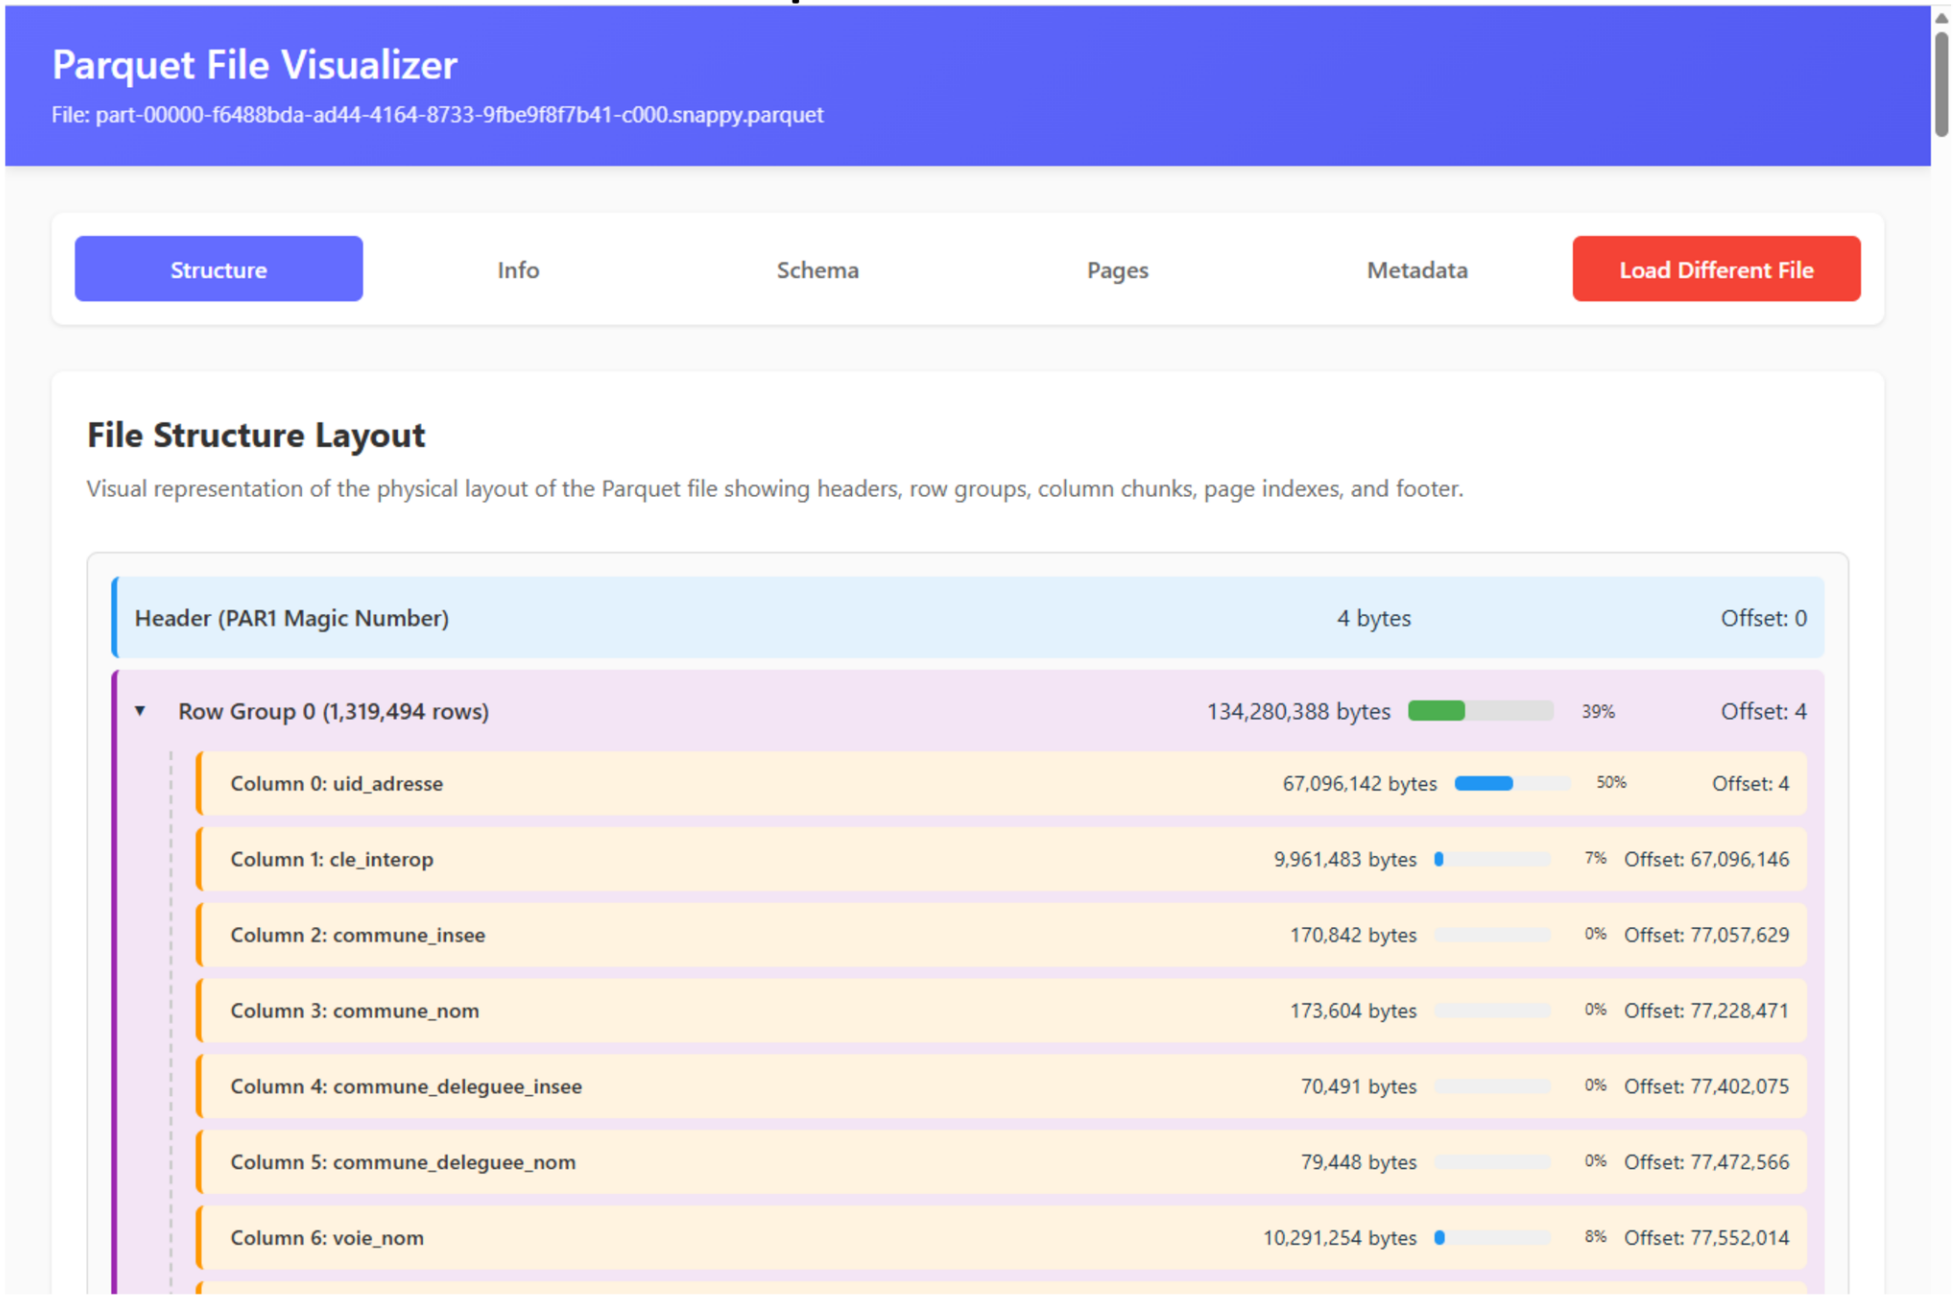

ColumnChunk- Pages

COMPRESSED SIZE:

Avg: 115,781 bytes

Min: 115,781 | Max: 115,781

UNCOMPRESSED SIZE:

Avg: 205,332 bytes

Min: 205,332 | Max: 205,332

COMPRESSION RATIO:

`1.77x`

AVG VALUES/PAGE:

`13,441`

COMPRESSED SIZE:

Avg: 849 bytes

Min: 461 | Max: 1,262

UNCOMPRESSED SIZE:

Avg: 846 bytes

Min: 456 | Max: 1,258

COMPRESSION RATIO:

`1.00x`

AVG VALUES/PAGE:

`19,992`

Page Size Histogram

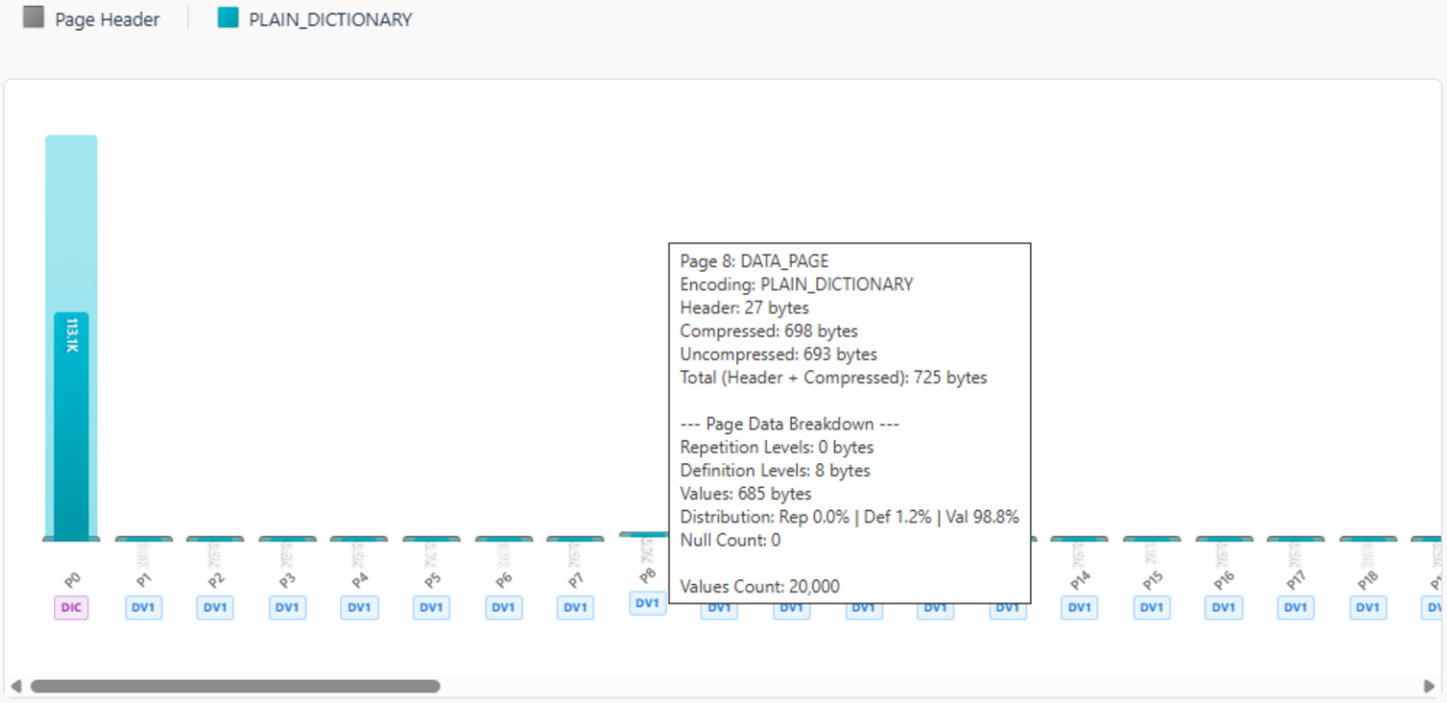

Offset Indexes, Column Indexes

A small project … welcoming PullRequest Contributions

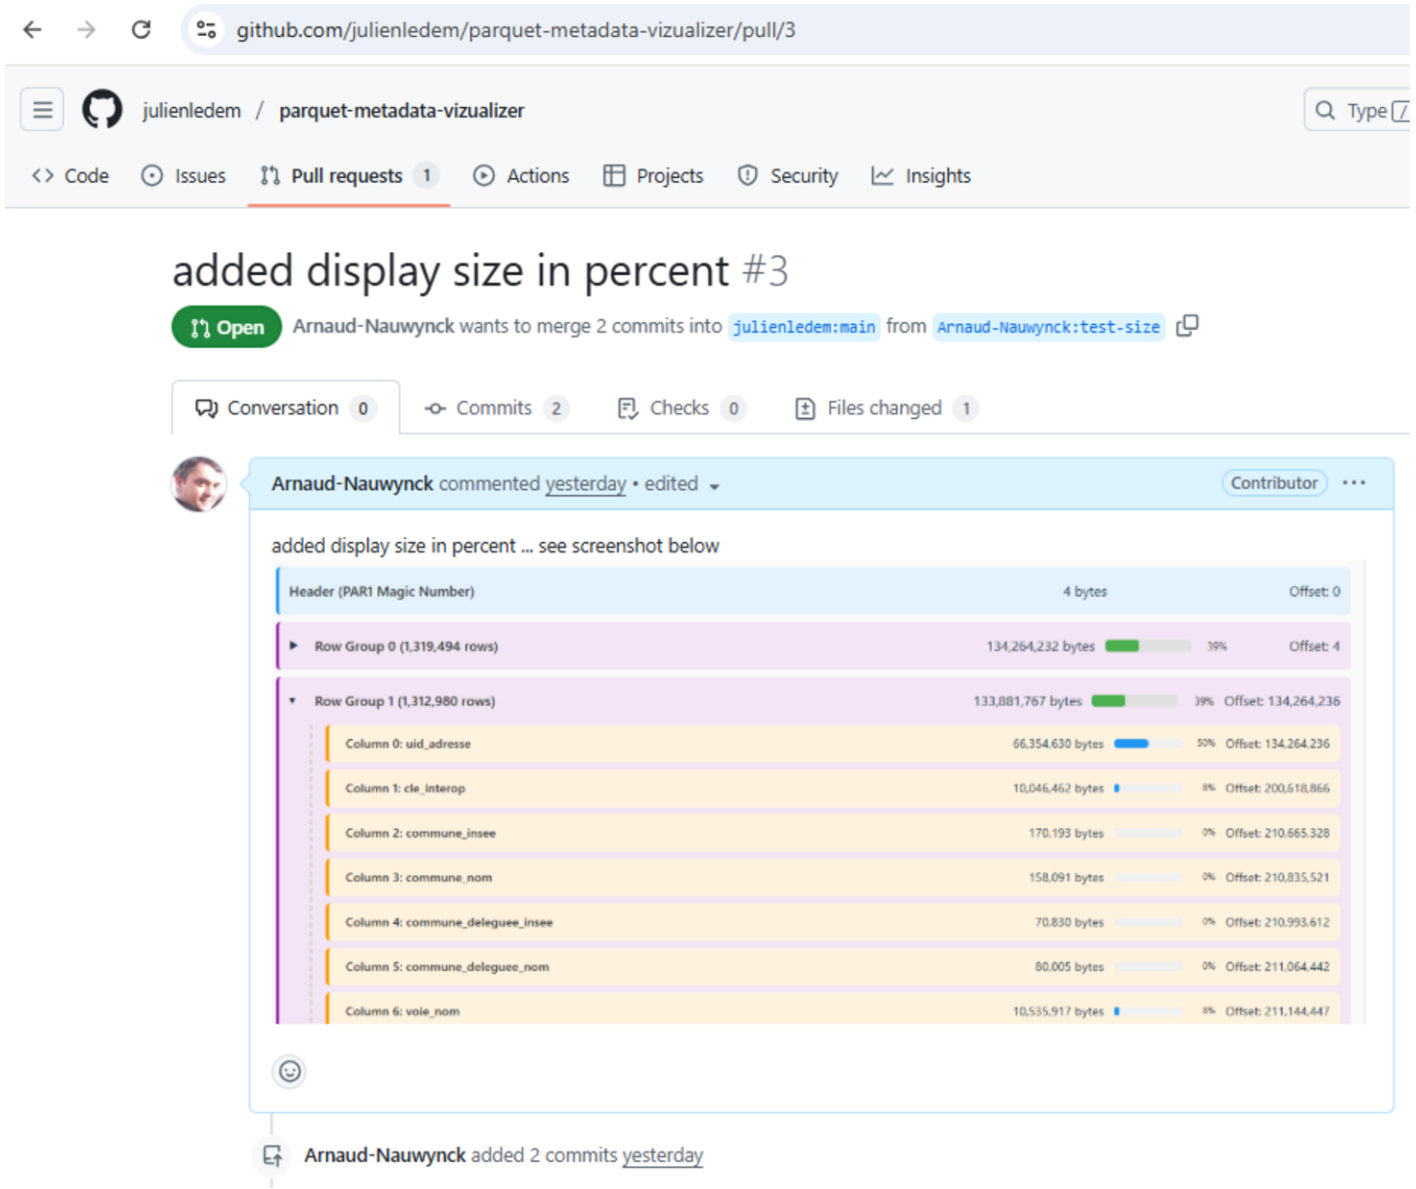

Project using « Vibe Coding » with Claude AI

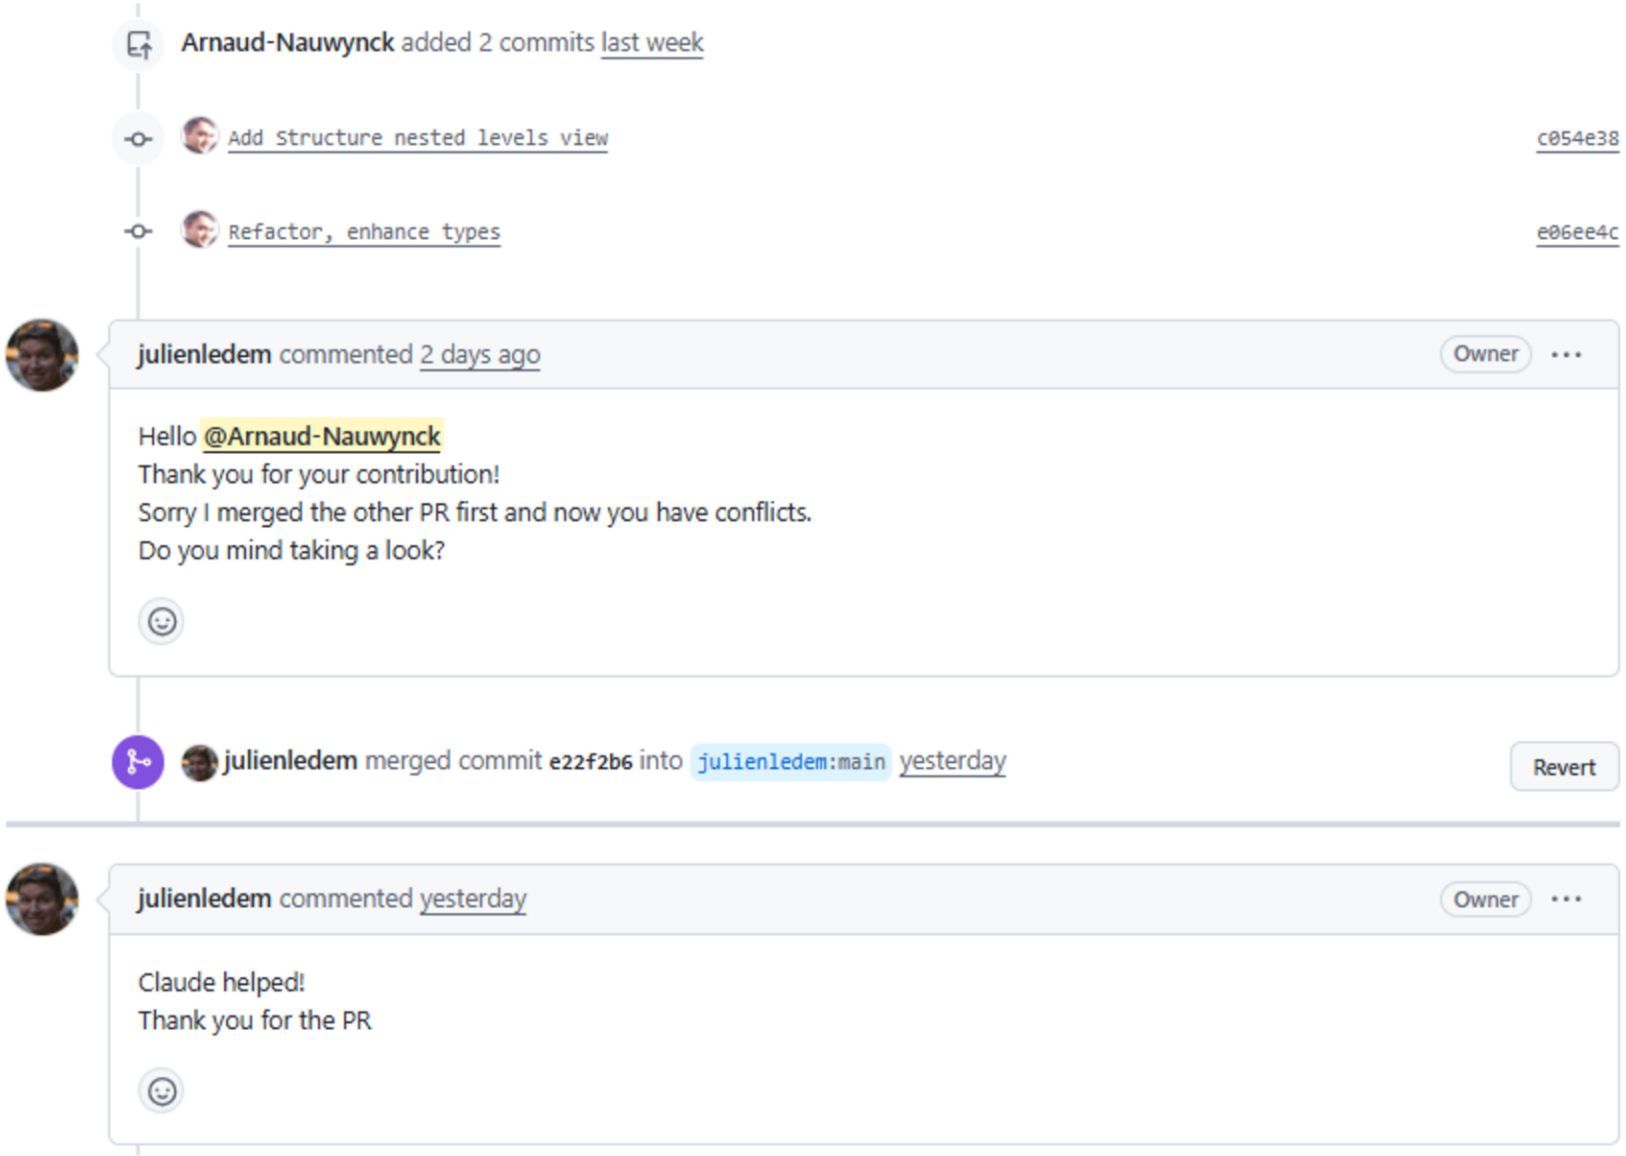

Work In Progress…

Many simple ideas still to implement !!

Example:

- Show more nested levels (per Page, per Encodings)
- show Predicate-Push-Down effects using min/max statistics
- Compare different Files for same data, but different write options used
- Scan multiple Files within 1 Directory, do statistics on File reads
- etc…

Conclusion

Parquet File Format is AMAZING

Spark is great using Parquet

Doing BigData processing = doing Spark + Parquet
with good .repartition(...)
.bucketBy(...)
.sortWithinPartitions(...)

Mind Map

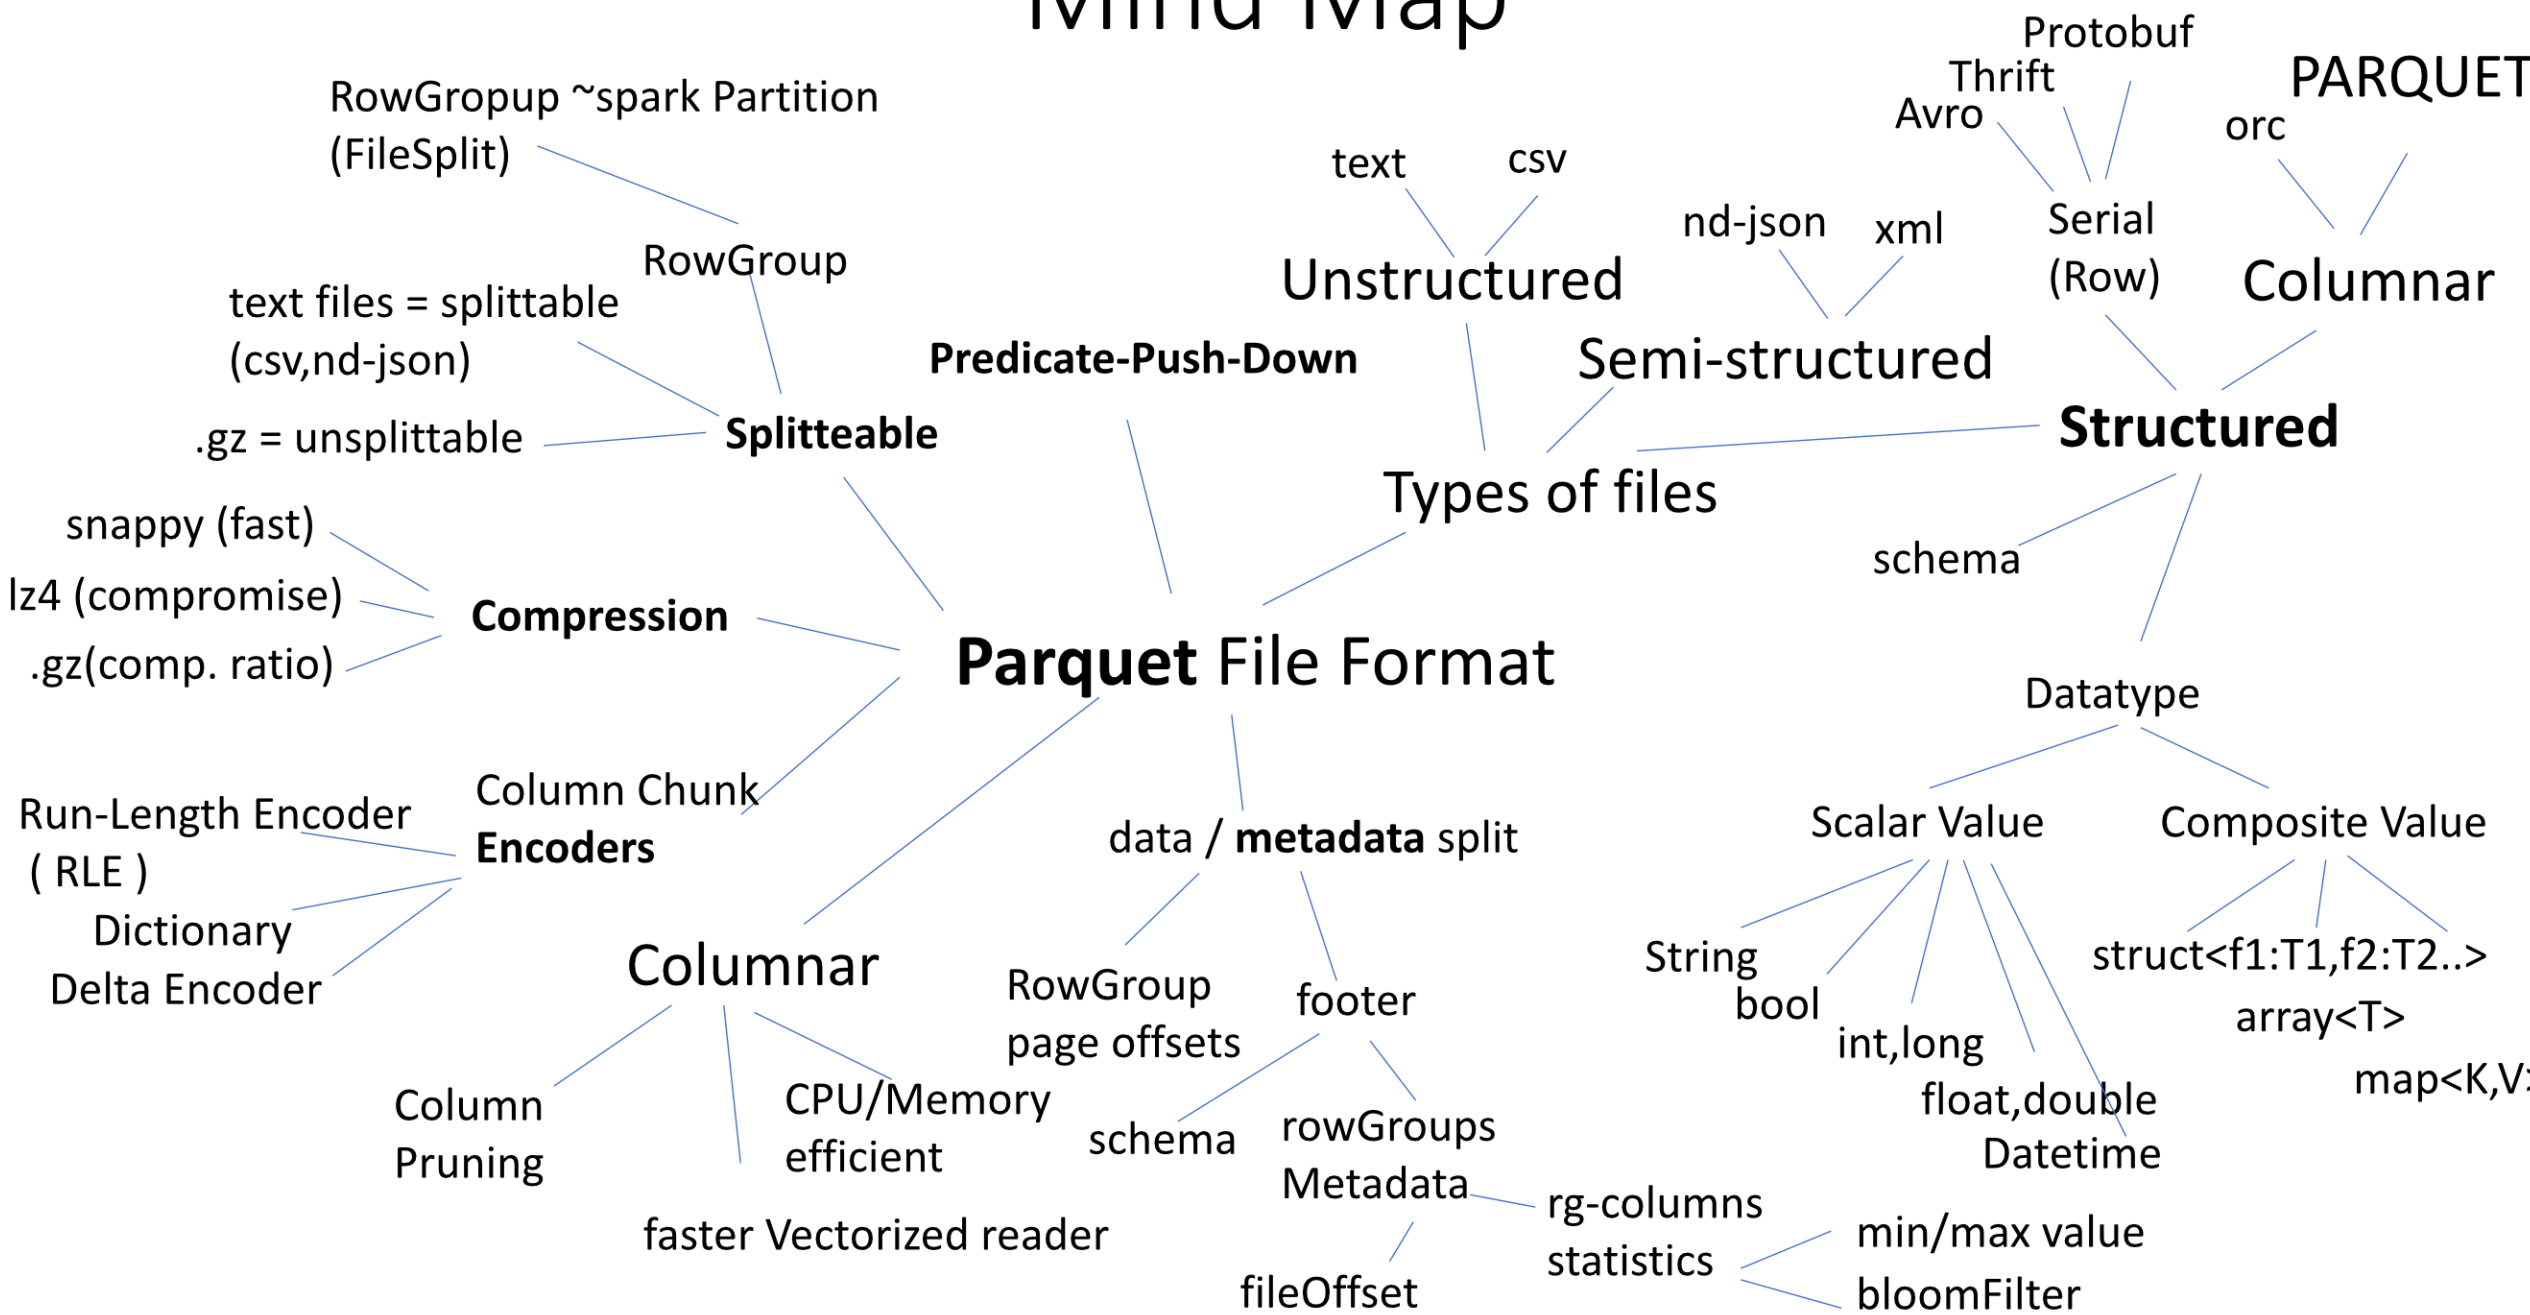
PARQUET
Protobuf
Thrift
Avro
orc
RowGropup ~spark Partition
(FileSplit)
text
csv
nd-json
xml
Serial
(Row)
Columnar
RowGroup
text files = splittable
(csv,nd-json)
Unstructured
Predicate-Push-Down
Semi-structured
.gz = unsplittable
Splitteable
Structured
Types of files
snappy (fast)
lz4 (compromise)
.gz(comp. ratio)
Compression
schema
Parquet File Format
Datatype
Column Chunk
Encoders
Run-Length Encoder
(RLE)
Dictionary
Delta Encoder
data / metadata split
Scalar Value
Composite Value
Columnar
String
bool
int,long
float,double
Datetime
struct<f1:T1,f2:T2..>
array<T>
map<K,V>
RowGroup
page offsets
footer
Column
Pruning
CPU/Memory
efficient
schema
rowGroups
Metadata
faster Vectorized reader
rg-columns
statistics
min/max value
bloomFilter
fileOffset

Questions?

arnaud.nauwynck@gmail.com