

Spark Introduction

course Esilv 2024

arnaud.nauwynck@gmail.com

this document:

[https://github.com/Arnaud-Nauwynck/presentations/
/pres-bigdata/9-spark-intro](https://github.com/Arnaud-Nauwynck/presentations/pres-bigdata/9-spark-intro)

Spark (Recent) History & Ancestors

MapReduce @Google

2002
@Google

2004 Google
Paper published

2014 Google
No more used of MapReduce



2006 @Yahoo
Hadoop implementation

2012 Yarn (v2)

2008 Apache Open-Source


2021 MapReduce bashing
... HDFS & Hadoop also
HortonWork bought by Cloudera
HDInsight @Azure...very bad choice
.. To be abandoned

 **MPI**
1995 Message Passing Interface



2010 Spark paper

2013 Apache top-level



2015
Kubernetes



2020
Spark on K8s

Simple => Many Specific Systems => Unified



« **Simple** » ecosystem
(verbose inefficient &
complex java code)

« **Bazard** » ecosystem
(**Too MANY TOO** **S P E C I F I C**
redundant, complexes)

“Unified” ecosystem
Simple
+ **extensible modules**

Spark = « Unified Engine »


[Download](#)[Libraries ▾](#)[Documentation ▾](#)[Examples](#)[Community ▾](#)[Developers ▾](#)[Apache Software Foundation ▾](#)

Unified engine for large-scale data analytics

[GET STARTED](#)

What is Apache Spark™?


Apache Spark™ is a multi-language engine for executing data engineering, data science, and machine learning on single-node machines or clusters.



Multi Purposes – Multi Languages


Simple.
Fast.
Scalable.
Unified.

Key features




Batch/streaming data

Unify the processing of your data in batches and real-time streaming, using your preferred language: Python, SQL, Scala, Java or R.




SQL analytics

Execute fast, distributed ANSI SQL queries for dashboarding and ad-hoc reporting. Runs faster than most data warehouses.



Data science at scale

Perform Exploratory Data Analysis (EDA) on petabyte-scale data without having to resort to downsampling



Machine learning

Train machine learning algorithms on a laptop and use the same code to scale to fault-tolerant clusters of thousands of machines.

Python

SQL



Scala

Java

R

Spark-Core + ...

Structured
Data



Spark SQL

Spark
Streaming

Spark MLlib

Spark
GraphX



Modules



DataSource Connectors
(Hadoop API)



Cluster Manager



Langages Support

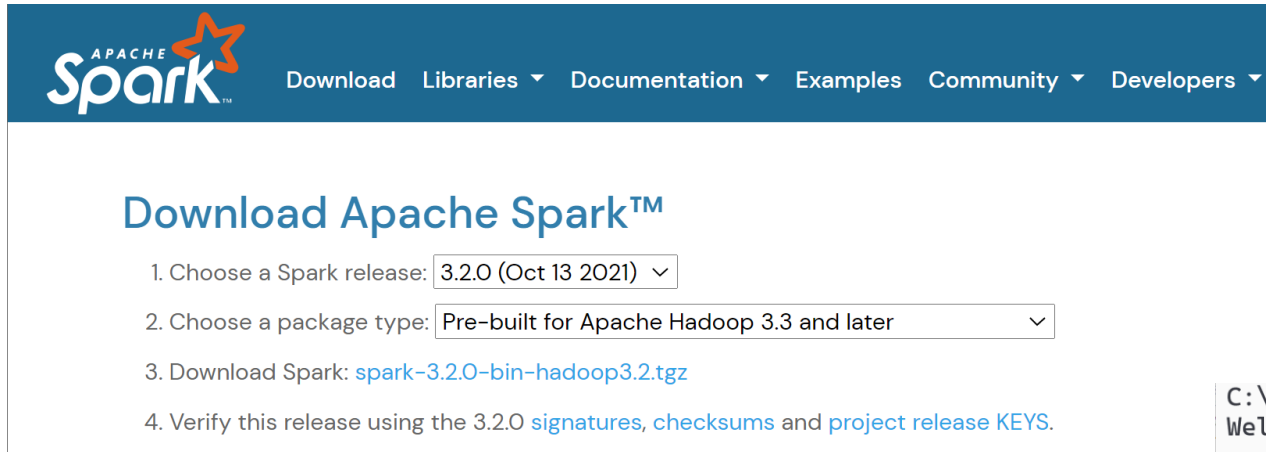


Getting Started

1/ Download

2/ Unzip + add to PATH

3/ launch



```
C:> bin\spark-shell
```

(or spark-submit, or spark-sql)

```
C:\Users\arnaud>spark-shell
Welcome to

      /_--\  /_--\  /_--\  /_--\  /_--\
     /  \  /  \  /  \  /  \  /  \
    /    \  /    \  /    \  /    \  /    \
   /      \  /      \  /      \  /      \  /      \
  /        \  /        \  /        \  /        \  /        \
 /          \  /          \  /          \  /          \  /          \
/            \  /            \  /            \  /            \  /            \
 \            /  \            /  \            /  \            /  \            /
  \          /    \          /    \          /    \          /    \          /
   \        /      \        /      \        /      \        /      \        /
    \      /        \      /        \      /        \      /        \      /
     \    /          \    /          \    /          \    /          \    /
      \  /            \  /            \  /            \  /            \  /
       \_/              \_/              \_/              \_/              \_/

version 3.5.0

Using Scala version 2.13.8 (OpenJDK 64-Bit Server VM, Java 20.0.1)
Type in expressions to have them evaluated.
Type :help for more information.
Spark context Web UI available at http://DesktopArnaud:4040
Spark context available as 'sc' (master = local[*], app id = local-1734184851521).
Spark session available as 'spark'.

scala>
```

```
scala>
```

```
scala> println(« Hello spark »);
```

launch spark-shell from terminal



```
C:\Users\arnaud>spark-shell
Welcome to
```

version 3.5.0

Using Scala version 2.13.8 (OpenJDK 64-Bit Server VM, Java 20.0.1)

Type in expressions to have them evaluated.

```
Type :help for more information.
```

Spark context Web UI available at <http://DesktopArnaud:4040>

```
Spark context available as 'sc' (master = local[*], app id = local-1734184851521).
```

Spark session available as 'spark'.

```
scala> println("Hello spark")
```

Hello spark

```
scala>
```

inside spark
type scala code



Spark-shell> SCALA code

:paste

```
scala> :paste
// Entering paste mode (ctrl-D to finish)

for(i <- 0 to 5) {
  println(s"Hello ${i}")
}
|
```

Ctrl-D

```
// Exiting paste mode, now interpreting.

Hello 0
Hello 1
Hello 2
Hello 3
Hello 4
Hello 5

scala>
```

spark-shell> help

```
scala> :help
```

All commands can be abbreviated, e.g., :he instead of :help.

:help [command]	print this summary or command-specific help
:completions <string>	output completions for the given string
:imports [name name ...]	show import history, identifying sources of names
:implicit [-v]	show the implicits in scope
:javap <path class>	disassemble a file or class name
:line <id> <line>	place line(s) at the end of history
:load <path>	interpret lines in a file
:paste [-raw] [path]	enter paste mode or paste a file
:power	enable power user mode
:quit	exit the REPL
:replay [options]	reset the REPL and replay all previous commands
:require <path>	add a jar to the classpath
:reset [options]	reset the REPL to its initial state, forgetting all session entries
:save <path>	save replayable session to a file
:sh <command line>	run a shell command (result is implicitly => List[String])
:settings <options>	update compiler options, if possible; see reset
:silent	disable/enable automatic printing of results
:type [-v] <expr>	display the type of an expression without evaluating it
:kind [-v] <type>	display the kind of a type. see also :help kind
:warnings	show the suppressed warnings from the most recent line which had any

Useful default key bindings:

TAB	code completion
CTRL-ALT-T	show type at cursor, hit again to show code with types/implicit inferred.

```
scala>
```

ds = spark.createDataSet(..)

```
scala> val data = Array(1, 2, 3, 4, 5)
data: Array[Int] = Array(1, 2, 3, 4, 5)
```

```
scala> val ds = spark.createDataset(data)
ds: org.apache.spark.sql.Dataset[Int] = [value: int]

scala> ds.reduce((a, b) => a+b)
res4: Int = 15
```

ds = spark.read.textFile(..)

loremIpsum.txt

Lorem Ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

```
scala> val ds=spark.read.textFile("c:/data/loremIpsum.txt")  
ds: org.apache.spark.sql.Dataset[String] = [value: string]
```

```
scala> ds.count // count lines  
res1: Long = 4
```

dataset.show ()

default show(20 /*line*/, true /*truncate*/)

```
scala> val ds = spark.read.textFile("c:/data/loremIpsum.txt")
val ds: org.apache.spark.sql.Dataset[String] = [value: string]
```

```
scala> ds.show()
+-----+
|              value|
+-----+
|Lorem Ipsum dolor...|
|Ut enim ad minim ...|
|Duis aute irure d...|
|Excepteur sint oc...|
+-----+
```

```
scala> ds.show(2, false)
```

```
+-----+
|value|
+-----+
|Lorem Ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.|
|Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.|
+-----+
```

only showing top 2 rows

```
scala>
```

Words Count : flatMap(...).count

```
scala> val wordDs = ds.flatMap(line =>
  line.replaceAll("[,;.:!?]", " ")
    .replaceAll(" ", " ")
    .split(" ")
)
words: .. Dataset[String] = [value: string]
```

```
scala> wordDs.count
res3: Long = 69
```

```
scala> wordDs.show(10, false)
```

```
scala> wordDs.count()
val res5: Long = 69

scala> wordDs.show(10)
+-----+
|      value      |
+-----+
|      Lorem      |
|      Ipsum      |
|      dolor      |
|      sit        |
|      amet       |
|consectetur      |
|adipiscing       |
|      elit       |
|      sed        |
|      do         |
+-----+
only showing top 10 rows
```

Words Count ... local[*] Debug in Java IDE

The screenshot shows a Java IDE with a Spark application running in debug mode. The main window displays the source code of `SparkWordsCountAppMain.java`, which counts words in a file. The left sidebar shows a list of threads, including the main thread and various daemon threads. The right sidebar shows the 'Variables' window with a table of current variables and their values. The bottom console shows the output of the application, including the word count result.

Variables Window:

Name	Value
no method return value	
args	String[0] (id=38)
sparkConf	SparkConf (id=39)
spark	SparkSession (id=40)
loremIpsum	Dataset<T> (id=41)
lineCount	4
wordDs	Dataset<T> (id=42)
wordCount	69

Console Output:

```
11:06:20 INFO fr.an.tests.testspark.SparkWordsCountAppMain: words count:69
```


spark-submit --class <<MainClass>> <<JarFile>>


```
c:\arn\devPerso\test-snippets\test-spark (master -> origin)
λ spark-submit --class fr.an.tests.testspark.SparkWordsCountAppMain target/tests-spark-0.0.1-SNAPSHOT.jar
22/01/06 11:14:11 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
22/01/06 11:14:11 INFO SparkContext: Running Spark version 3.1.1
```

... skipped 1000 lines LOG ...


```
22/01/06 11:14:22 INFO DAGScheduler: Submitting ResultStage 5 (MapPartitionsRDD[28] at count at SparkWordsCountAppMain.java:51),
  which has no missing parents
22/01/06 11:14:22 INFO MemoryStore: Block broadcast_8 stored as values in memory (estimated size 10.1 KiB, free 413.6 MiB)
22/01/06 11:14:22 INFO MemoryStore: Block broadcast_8_piece0 stored as bytes in memory (estimated size 5.0 KiB, free 413.6 MiB)
22/01/06 11:14:22 INFO BlockManagerInfo: Added broadcast_8_piece0 in memory on DESKTOP-2EGCC8R:64427 (size: 5.0 KiB, free: 413.9
  MiB)
22/01/06 11:14:22 INFO SparkContext: Created broadcast 8 from broadcast at DAGScheduler.scala:1383
22/01/06 11:14:22 INFO DAGScheduler: Submitting 1 missing tasks from ResultStage 5 (MapPartitionsRDD[28] at count at SparkWordsC
  ountAppMain.java:51) (first 15 tasks are for partitions Vector(0))
22/01/06 11:14:22 INFO TaskSchedulerImpl: Adding task set 5.0 with 1 tasks resource profile 0
22/01/06 11:14:22 INFO TaskSetManager: Starting task 0.0 in stage 5.0 (TID 5) (DESKTOP-2EGCC8R, executor driver, partition 0, NO
  DE_LOCAL, 4453 bytes) taskResourceAssignments Map()
22/01/06 11:14:22 INFO Executor: Running task 0.0 in stage 5.0 (TID 5)
22/01/06 11:14:22 INFO ShuffleBlockFetcherIterator: Getting 1 (60.0 B) non-empty blocks including 1 (60.0 B) local and 0 (0.0 B)
  host-local and 0 (0.0 B) remote blocks
22/01/06 11:14:22 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 4 ms
22/01/06 11:14:22 INFO Executor: Finished task 0.0 in stage 5.0 (TID 5). 2605 bytes result sent to driver
22/01/06 11:14:22 INFO TaskSetManager: Finished task 0.0 in stage 5.0 (TID 5) in 22 ms on DESKTOP-2EGCC8R (executor driver) (1/1
  )
22/01/06 11:14:22 INFO TaskSchedulerImpl: Removed TaskSet 5.0, whose tasks have all completed, from pool
22/01/06 11:14:22 INFO DAGScheduler: ResultStage 5 (count at SparkWordsCountAppMain.java:51) finished in 0,038 s
22/01/06 11:14:22 INFO DAGScheduler: Job 3 is finished. Cancelling potential speculative or zombie tasks for this job
22/01/06 11:14:22 INFO TaskSchedulerImpl: Killing all running tasks in stage 5: Stage finished
22/01/06 11:14:22 INFO DAGScheduler: Job 3 finished: count at SparkWordsCountAppMain.java:51, took 0,093309 s
22/01/06 11:14:22 INFO SparkWordsCountAppMain: words count:69
22/01/06 11:14:22 INFO SparkWordsCountAppMain: finished
22/01/06 11:14:22 INFO SparkUI: Stopped Spark web UI at http://DESKTOP-2EGCC8R:4040
22/01/06 11:14:22 INFO BlockManagerInfo: Removed broadcast_7_piece0 on DESKTOP-2EGCC8R:64427 in memory (size: 10.5 KiB, free: 41
  3.9 MiB)
22/01/06 11:14:22 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
```


resultDs = dataset.map(x -> {... return y; })

Source Dataset = List<X>




Mapping function (or lambda) :
x -> return y



dataset.map(func)

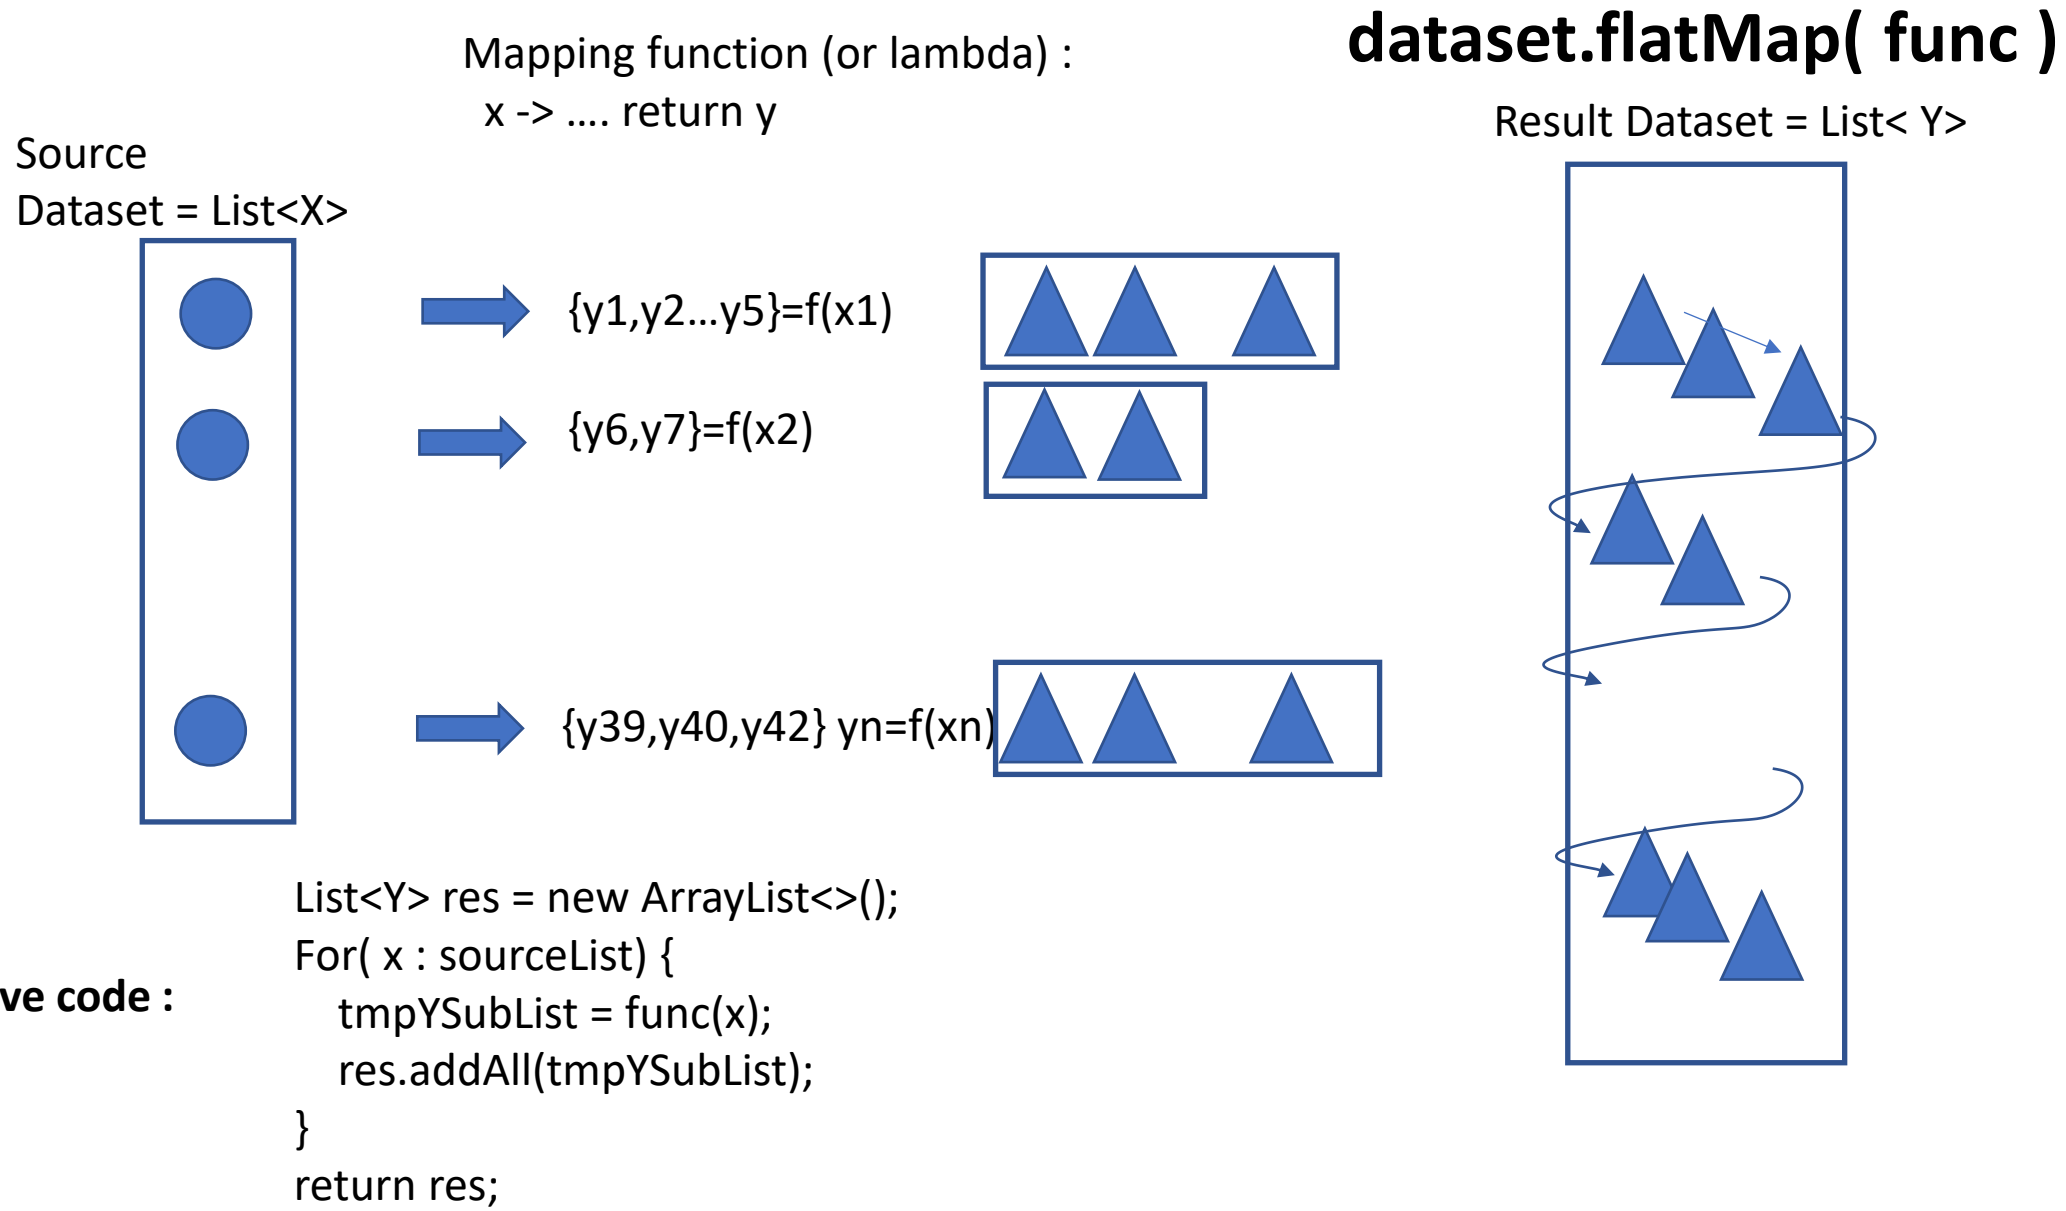
Result = List< Y>



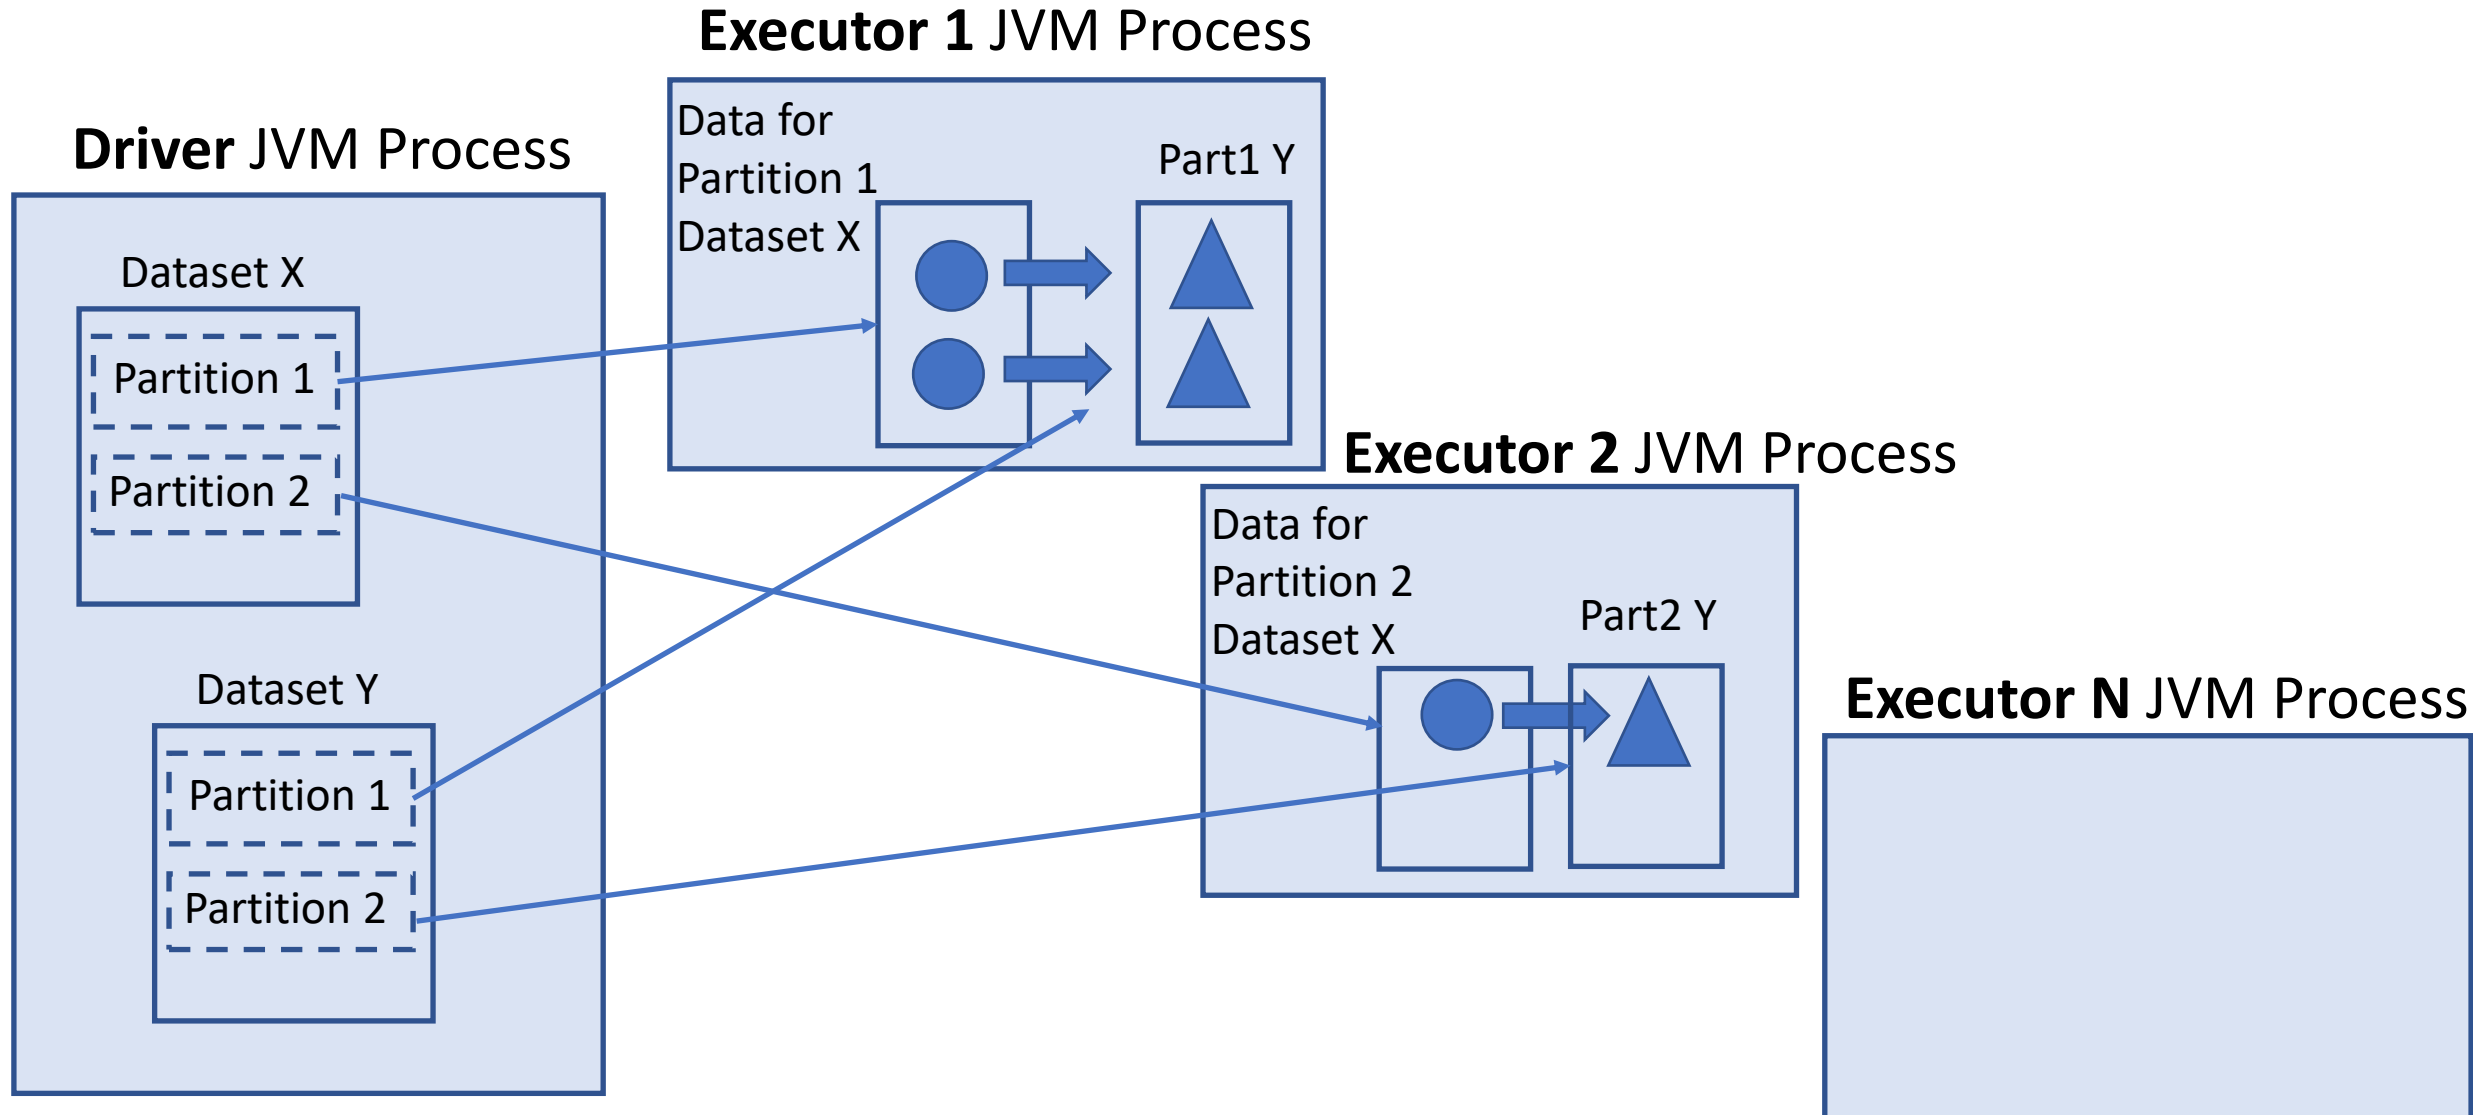
« Equivalent » imperative code :

```
List<Y> res = new ArrayList<>();  
For( x : sourceList) {  
    y = func(x);  
    res.add(y);  
}  
return res;
```

resultDs = dataset.flatMap(x -> {... return list<y>; })

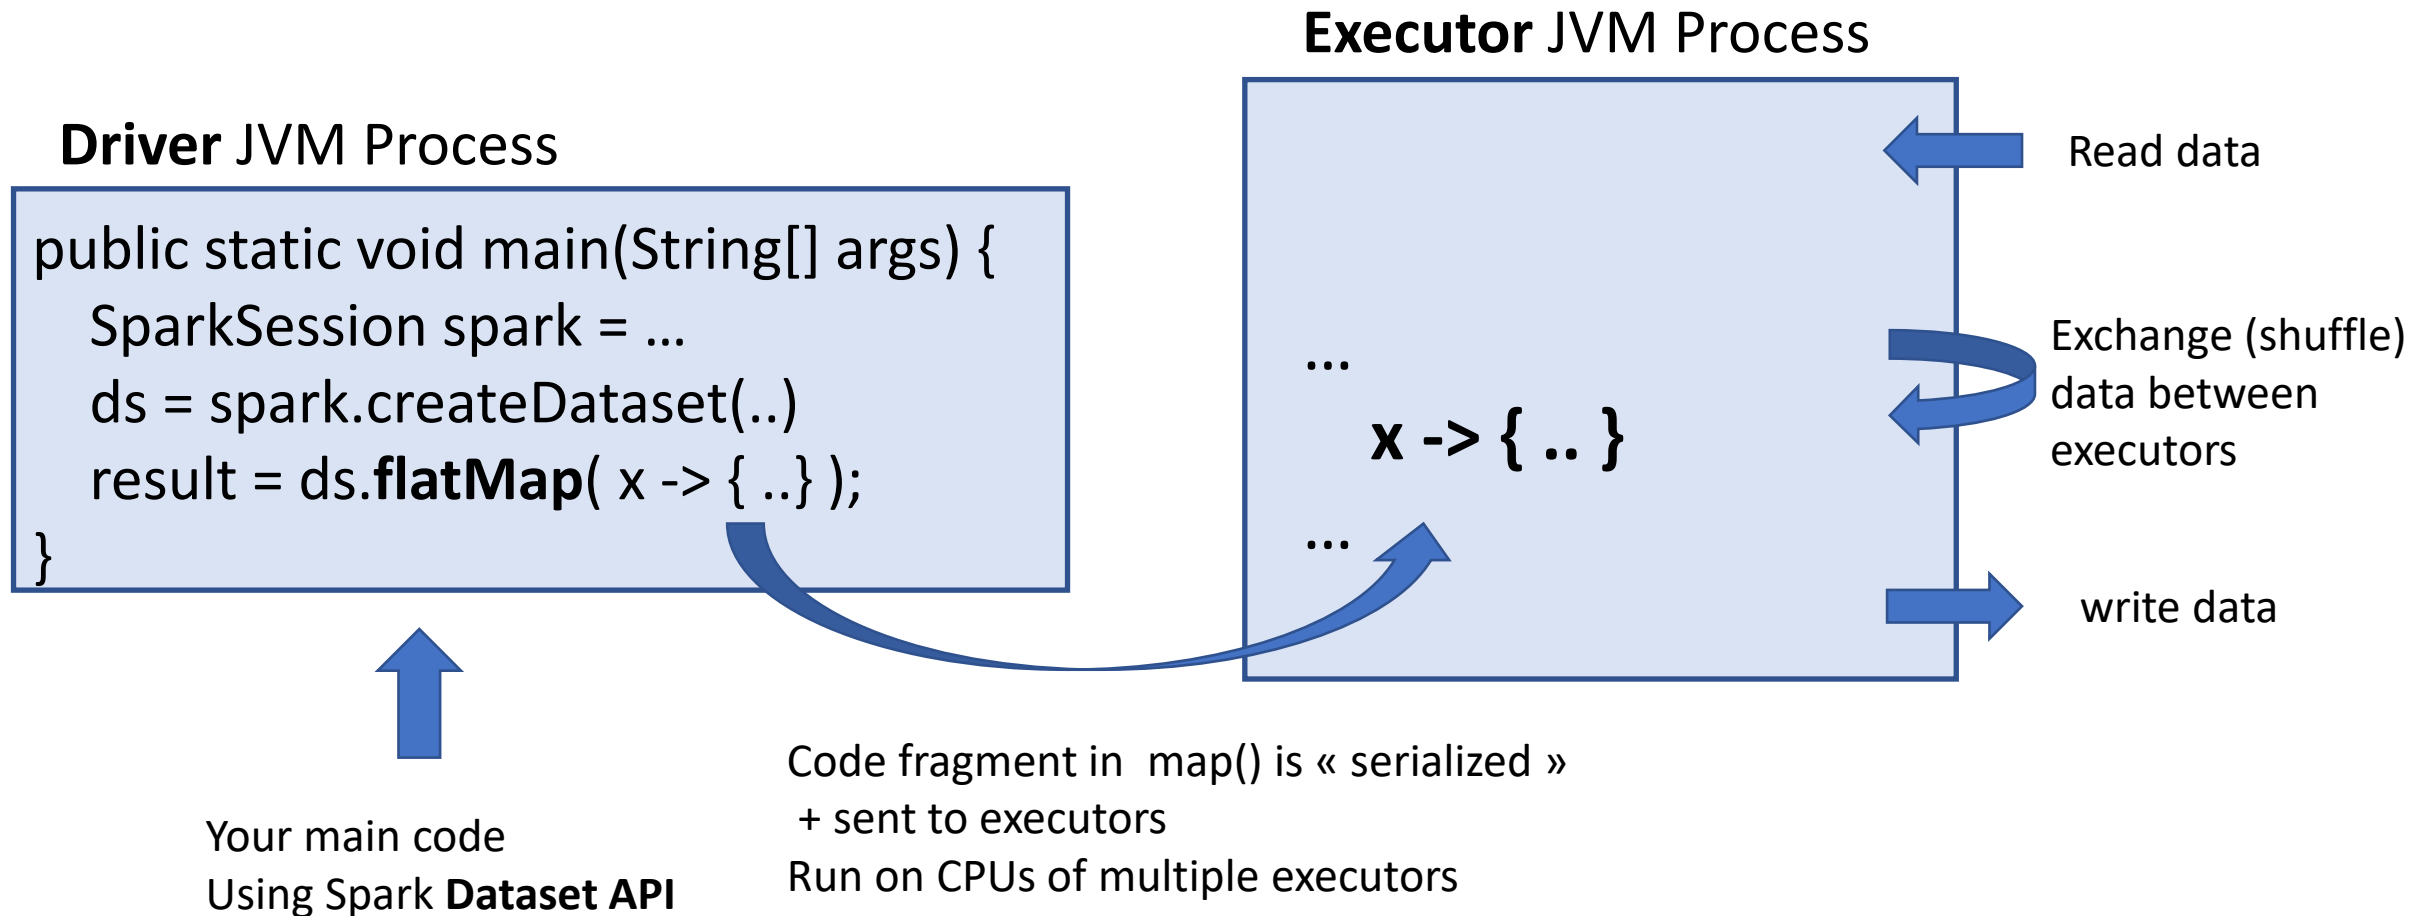


Dataset: « logical View » of « partitions » on Driver
... « data » allocated in-memory on Executors



Driver : Drives the main() program

Executors : Execute the functions on data



See the difference ?
Driver Api / Executor Engine
Code Logic / Data+Cpu Internal




Exchange / Shuffle / Repartition / Coalesce / Distribute Wide transformation / Reduce


Executor 1
JVM Process



Executor 2
JVM Process



Executor 3
JVM Process



Wide vs Narrow Transformation


Wide Transformation (Exchange) between

```
dataSet Y = x.repartition(3).shuffle(..) .reduce(..) .sortBy()
```

Executor 1

Executor 2

Executor 3



Network exchange

Between executors

... serialization/deserialization of byte data


Narrow Transformation

```
dataSet Y = x.map( func1 ) .flatMap(func2).filter( pred ) .first(10) .sortWithinPartition ( )
```

Partition Xi

+ corresp. Yi

on Executor 1/2/3



No data exchange

only in-memory pointer,
within same thread/process

... computation changes
but same logical partitioning

Things get complex...

« repartition() .flatMap() .explain() »

```
scala> loremIpsum.flatMap(line => line.replaceAll("[,;.:!?!]", " ").replaceAll(" ", " ").split(" ")).explain
```

```
== Physical Plan ==
```

```
*(1) SerializeFromObject [... AS value#116]
```

```
+ MapPartitions org.apache.spark.sql.Dataset$$Lambda$3735/14922651@1f3e32c, obj#115: java.lang.String
```

```
+ DeserializeToObject value#12.toString, obj#114: java.lang.String
```

```
+ FileScan text [value#12] Batched: false, DataFilters: [], Format: Text,
```

```
Location: InMemoryFileIndex[file:/c:/data/loremIpsum.txt], PartitionFilters: [], PushedFilters: [], ReadSchema: struct<value
```

```
scala> loremIpsum.repartition(3).flatMap(line => line.replaceAll("[,;.:!?!]", " ").replaceAll(" ", " ").split(" ")).explain
```

```
== Physical Plan ==
```

```
*(1) SerializeFromObject [staticinvoke(... ) AS value#112]
```

```
+ MapPartitions org.apache.spark.sql.Dataset$$Lambda$3735/14922651@18e4389, obj#111: java.lang.String
```

```
+ DeserializeToObject value#12.toString, obj#110: java.lang.String
```

```
+ Exchange RoundRobinPartitioning(3), REPARTITION_WITH_NUM, [id=#235]
```

```
+ FileScan text [value#12] Batched: false, DataFilters: [], Format: Text,
```

```
Location: InMemoryFileIndex[file:/c:/data/loremIpsum.txt], PartitionFilters: [], PushedFilters: [], ReadSchema: struct<valu
```


Confusing Questions at first Glance...

- 1/ RDD vs DataSet vs DataFrame ?
- 2/ meaning of Sql / Job / Task / Staging / Action ?
- 3/ Driver vs Executor ... where is executed my code ?
- 4/ Batch and Streaming api ?
- 5/ SQL or Code ? Functional API in java-scala-python ?
- 6/ Spark-shell / spark-sql / spark-submit / spark-thrift server / spark-history server ?
- 7/ Master = yarn/standalone/k8s + mode = Client vs Cluster vs ...
- 8/ use Spark-ui / Console / logs?
- 9/ Performance diagnostic?
DAG ? Metrics ? Shuffle ? SpillToDisk ? SkewedPartition?
- 10/ Optimize or add more resources ?