BigData – Spark – Processing

Spark cluster mode distribution & Tuning

Course Esilv 2026

arnaud.nauwynck@gmail.com

this document:

https://github.com/Arnaud-Nauwynck/presentations/pres-bigdata/14-Spark-ClusterMode

Outline

Reminder Spark-driver / spark-executors

Cluster Tuning Params

deployMode (local, spark://, yarn, k8s)

Other launchers (java embedded, livy, jupyter, pyspark ..)

Reminder

Spark-Core + …

Structured Data

PULSAR

kafka

Spark SQL

Spark Streaming

Spark MLlib

Spark GraphX

Modules

Amazon S3

Azure Data Lake Storage Gen2

DELTA LAKE

ICEBERG

DataSource Connectors (Hadoop API)

Cluster Manager

Standalone Cluster

Languages Support

Java

Scala

python

R

Reminder

1 Spark application

= 1 Spark-driver + N spark-executors

Reminder

Read N Files – assign Tasks to Executors

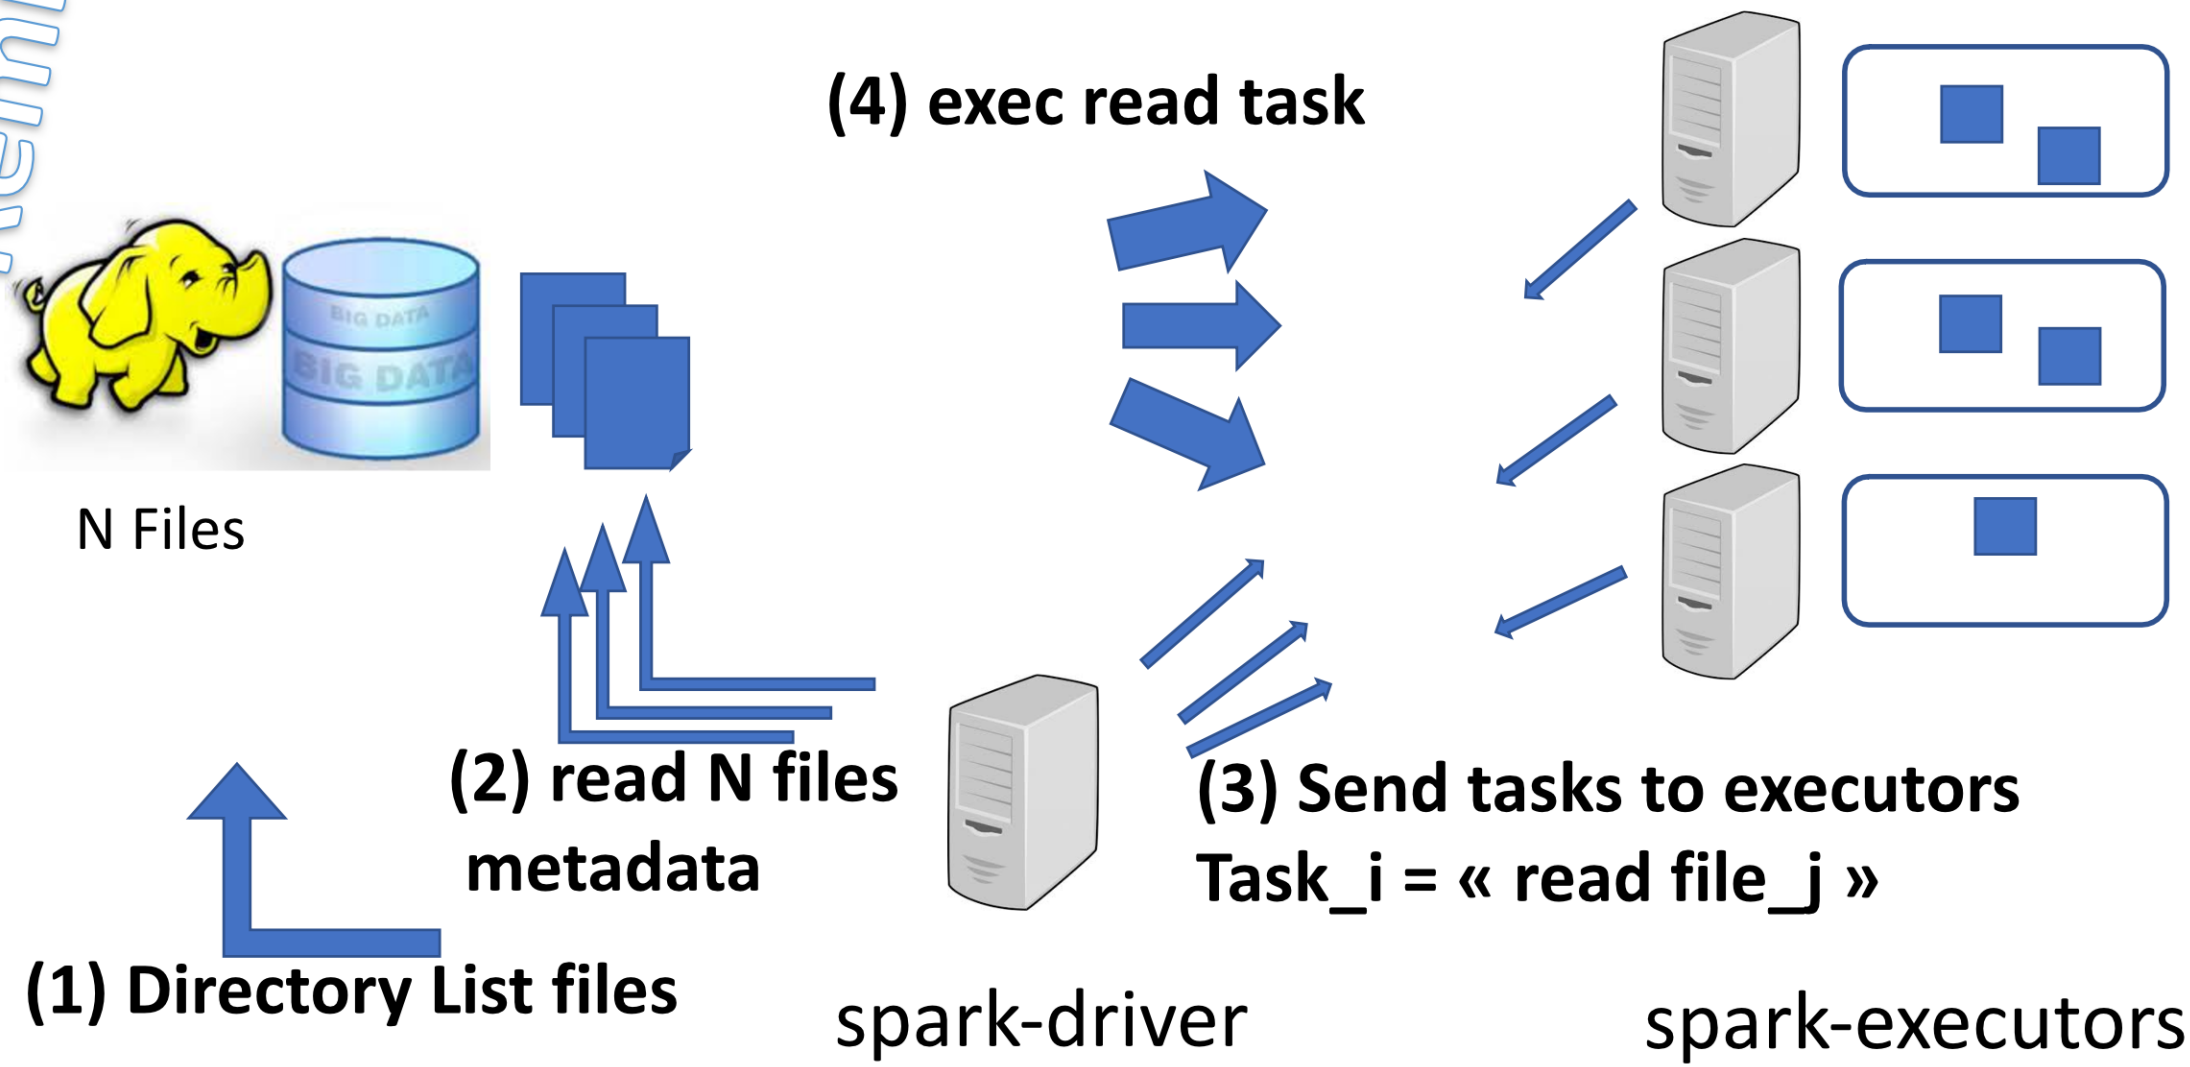

Reminder

Parallelizable per partition / available CPU
1 or several partitions per Executor
..1 or several Executors

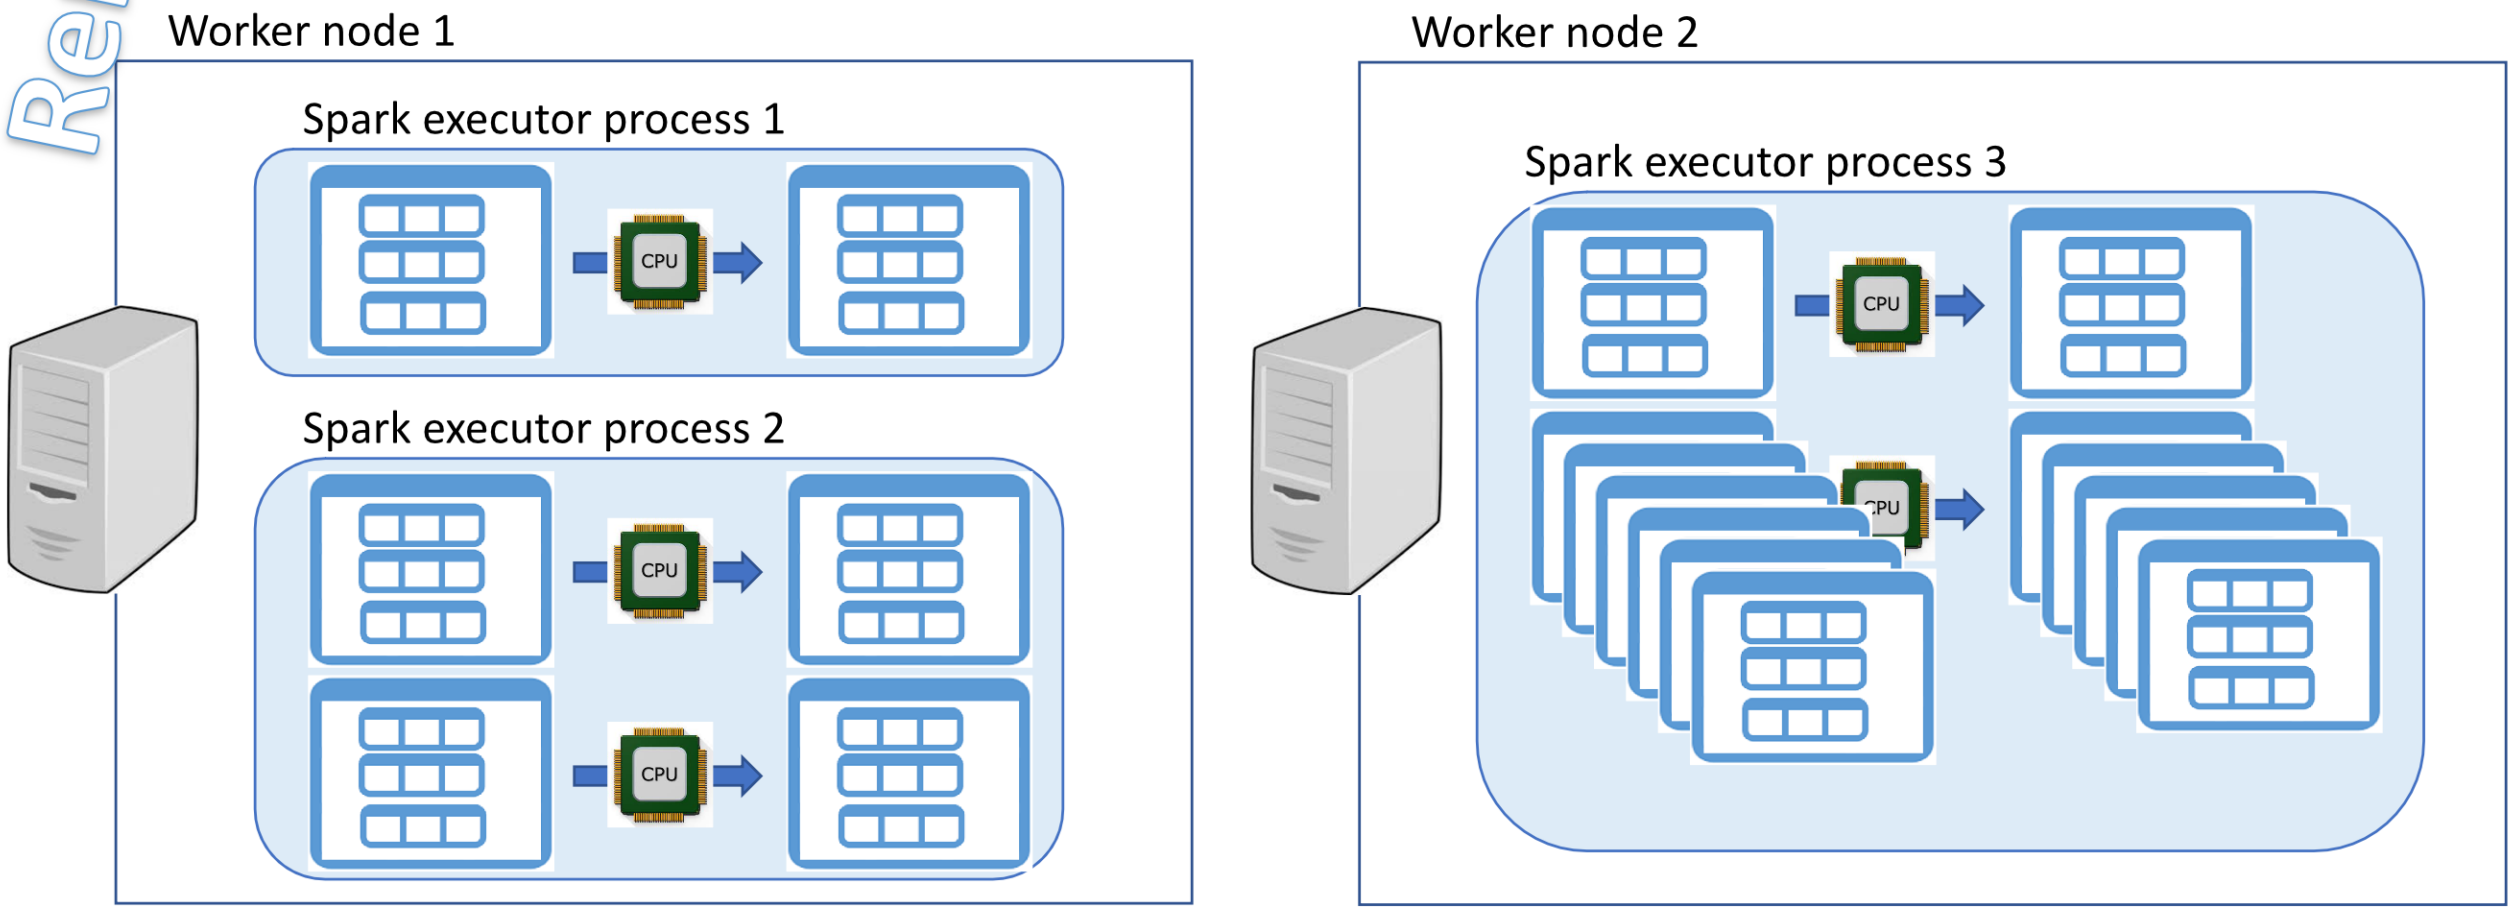

WholeStageCodeGen

Reminder

Program
= Dataset instructions

```
36  val sc = new SparkContext(args(0), "GroupBy Test",
37    System.getenv("SPARK_HOME"), SparkContext.jarOfClass(this.getClass))
38
39  val pairs1 = sc.parallelize(0 until numMappers, numMappers).flatMap { p =>
40    val ranGen = new Random
41    var arr1 = new Array[(Int, Array[Byte])](numKVPairs)
42    for (i <- 0 until numKVPairs) {
43      val byteArr = new Array[Byte](valSize)
44      ranGen.nextBytes(byteArr)
45      arr1(i) = (ranGen.nextInt(Int.MaxValue), byteArr)
46    }
47    arr1
```

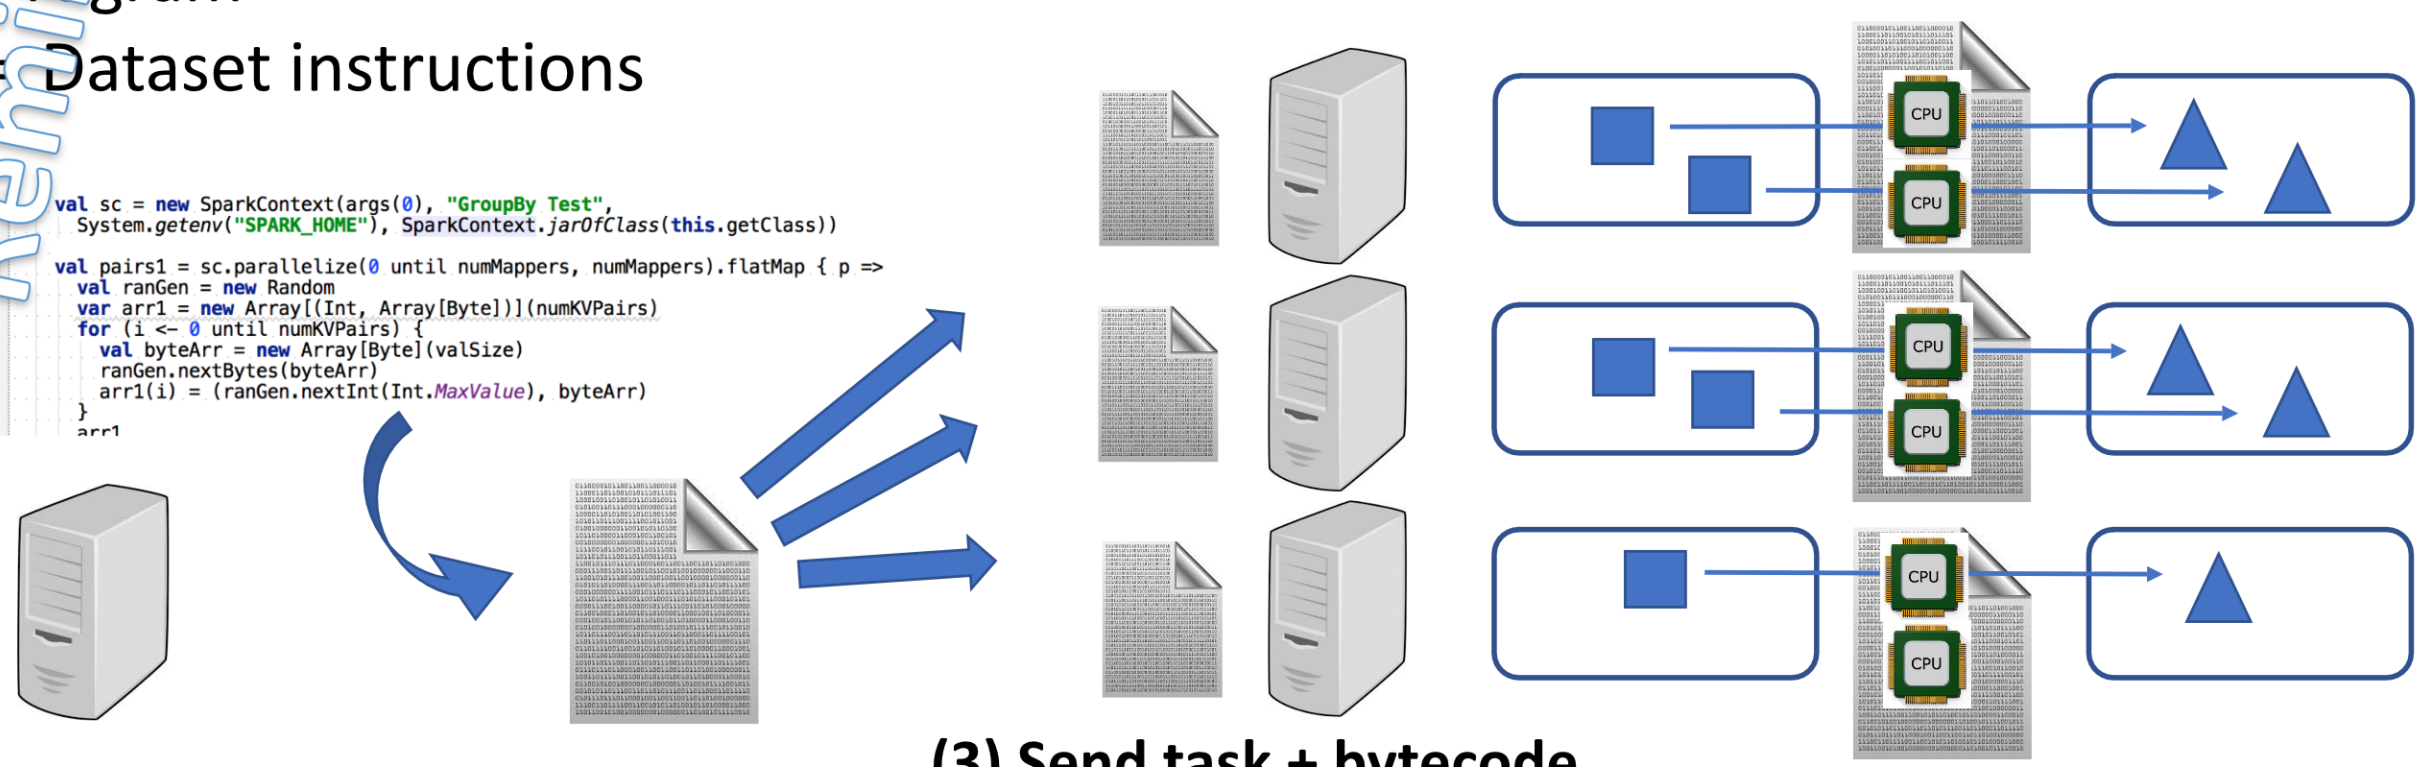

(1) Generate java code

(RDD Spark sub-class « WholeStageCodeGen$i »)

(2) Compile Bytecode

(3) Send task + bytecode to spark-executors

(4) Execute tasks

Reminder

Cluster Manager Extensions: TaskScheduler API

```
abstract class TaskScheduler {
  ..start(), stop(),

  submitTasks,cancelTasks,

  notify Host-Executor-Task changes
}
```

Reminder

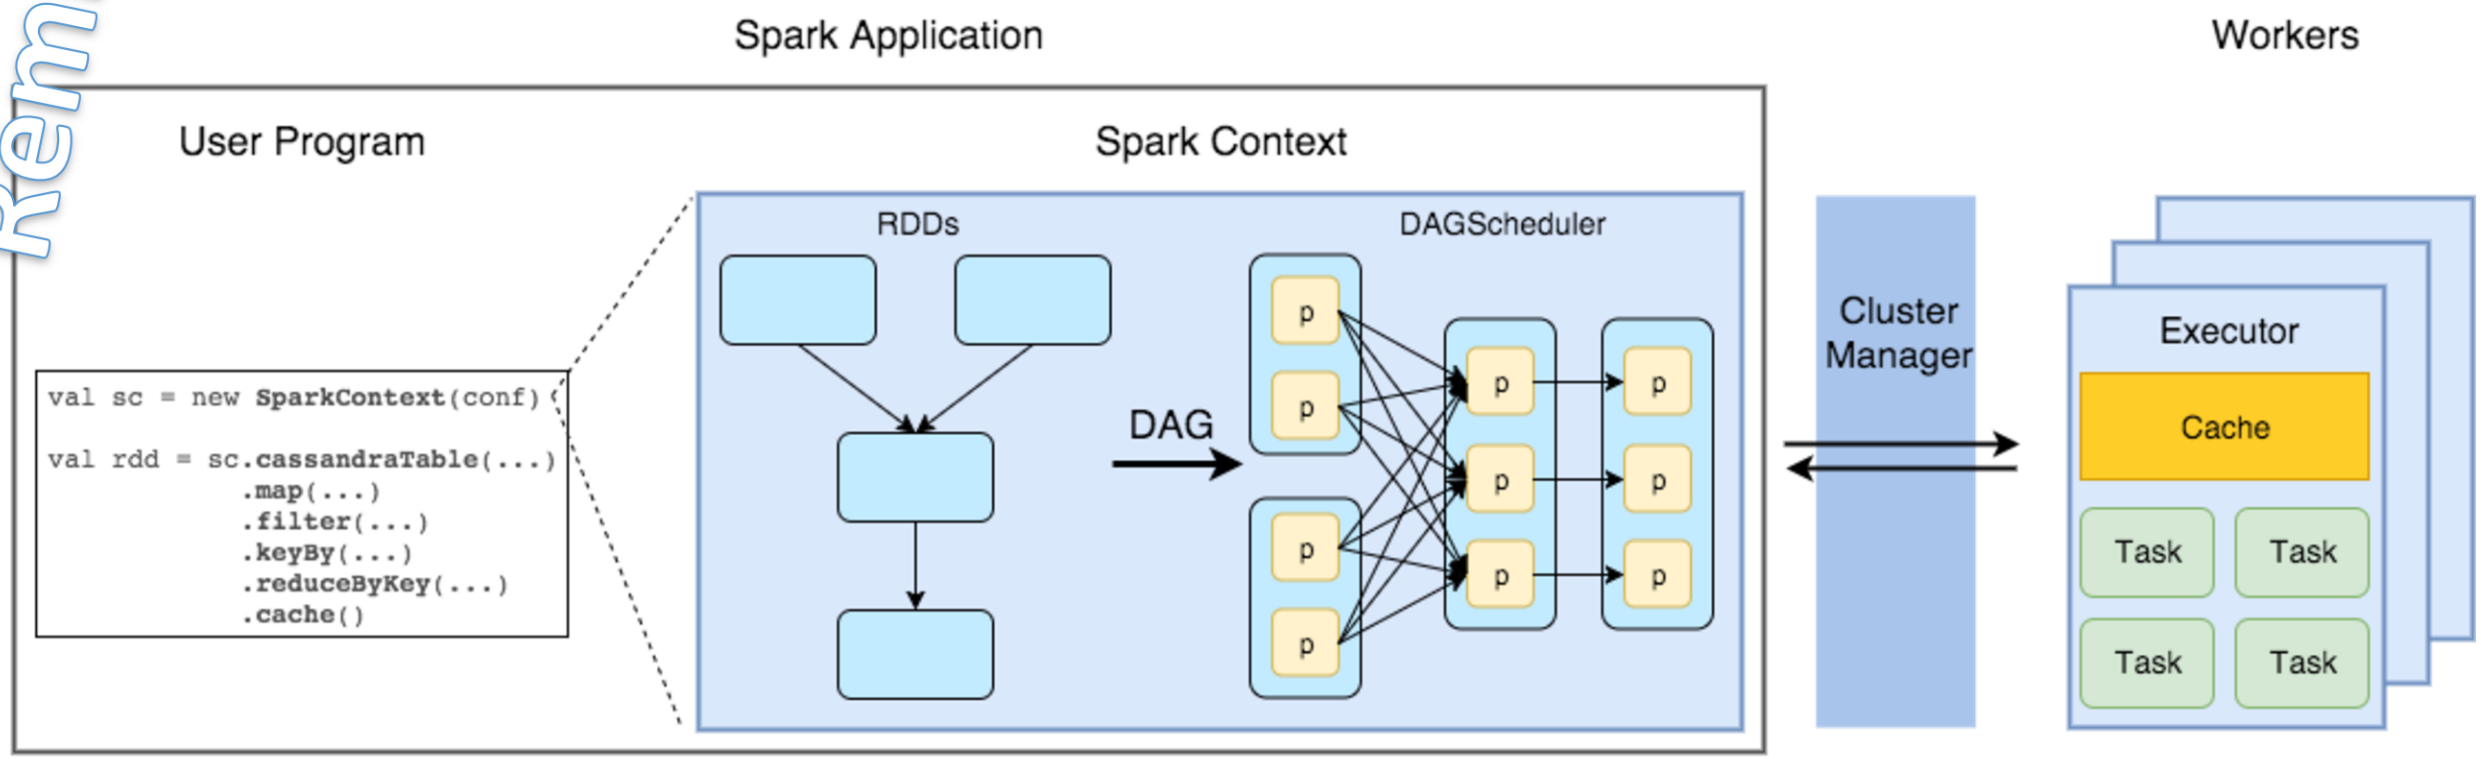
Spark Application
Workers
User Program
Spark Context
RDDs
DAGScheduler
DAG
val sc = new SparkContext(conf)
val rdd = sc.cassandraTable(...)
.map(...)
.filter(...)
.keyBy(...)
.reduceByKey(...)
.cache()
p
Cluster Manager
Executor
Cache
Task
Task
Task
Task

Reminder

CoarseGrainScheduler
… CoarseGrain Executor

Spark-driver
Implements Fault Tolerance+Distribution

Spark-executor
… internally called « CoarseGrainExecutor » main loop

CoarseGrain
partition
= unit or recomputation

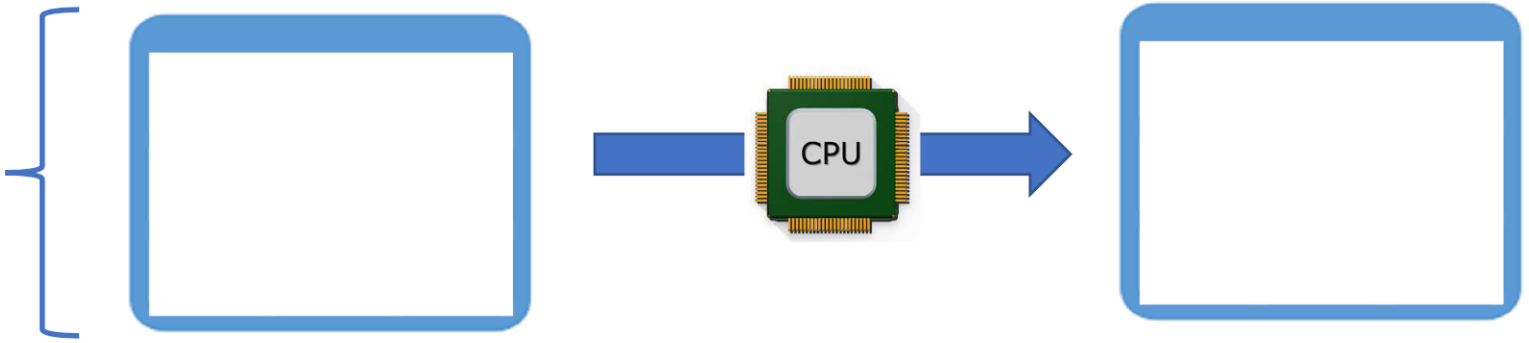

Outline

Reminder Spark-driver / spark-executors

Cluster Tuning Params

deployMode (local, spark://, yarn, k8s)

Other launchers (java embedded, livy, jupyter, pyspark ..)

Tuning Spark to use cluster power ?

Distribution / Partitions / Parallelism

Memory needs for Dataset

driver != executor

Dataset = Partitions metadata / data

1 Dataset
= N partitions

1 Partition
= p rows
(p1 != p2 != ..pN)

Row

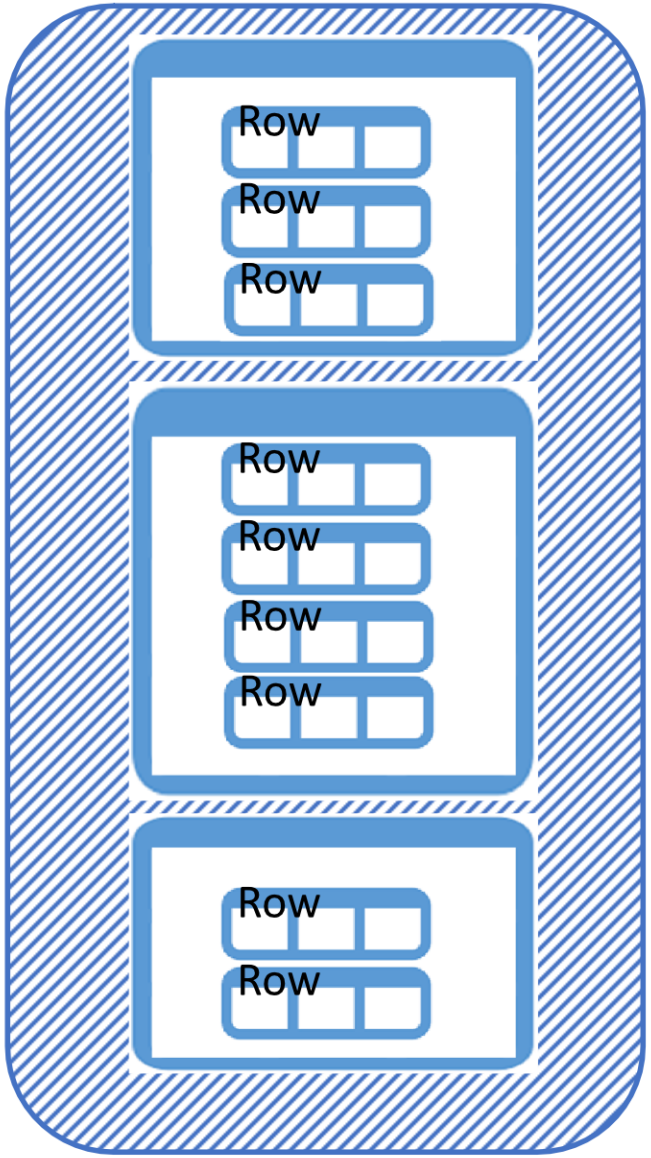

On Driver = metadata / on Executor = Data

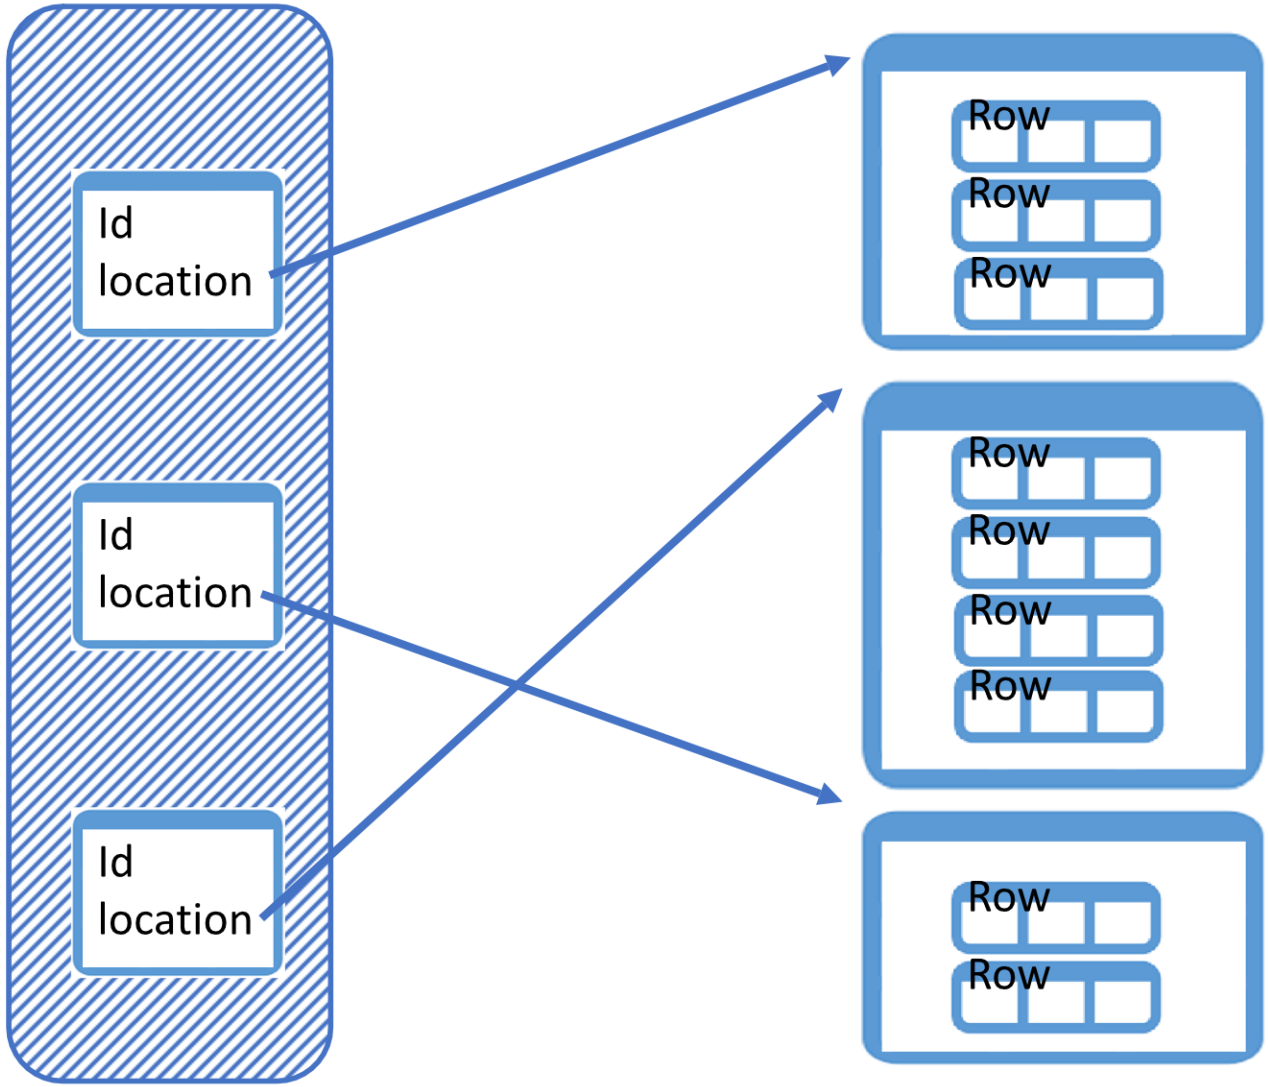

1 Dataset on spark-driver

N x Partitions (blocks) on spark-executors

Dataset … Partitions Distributed on Executors

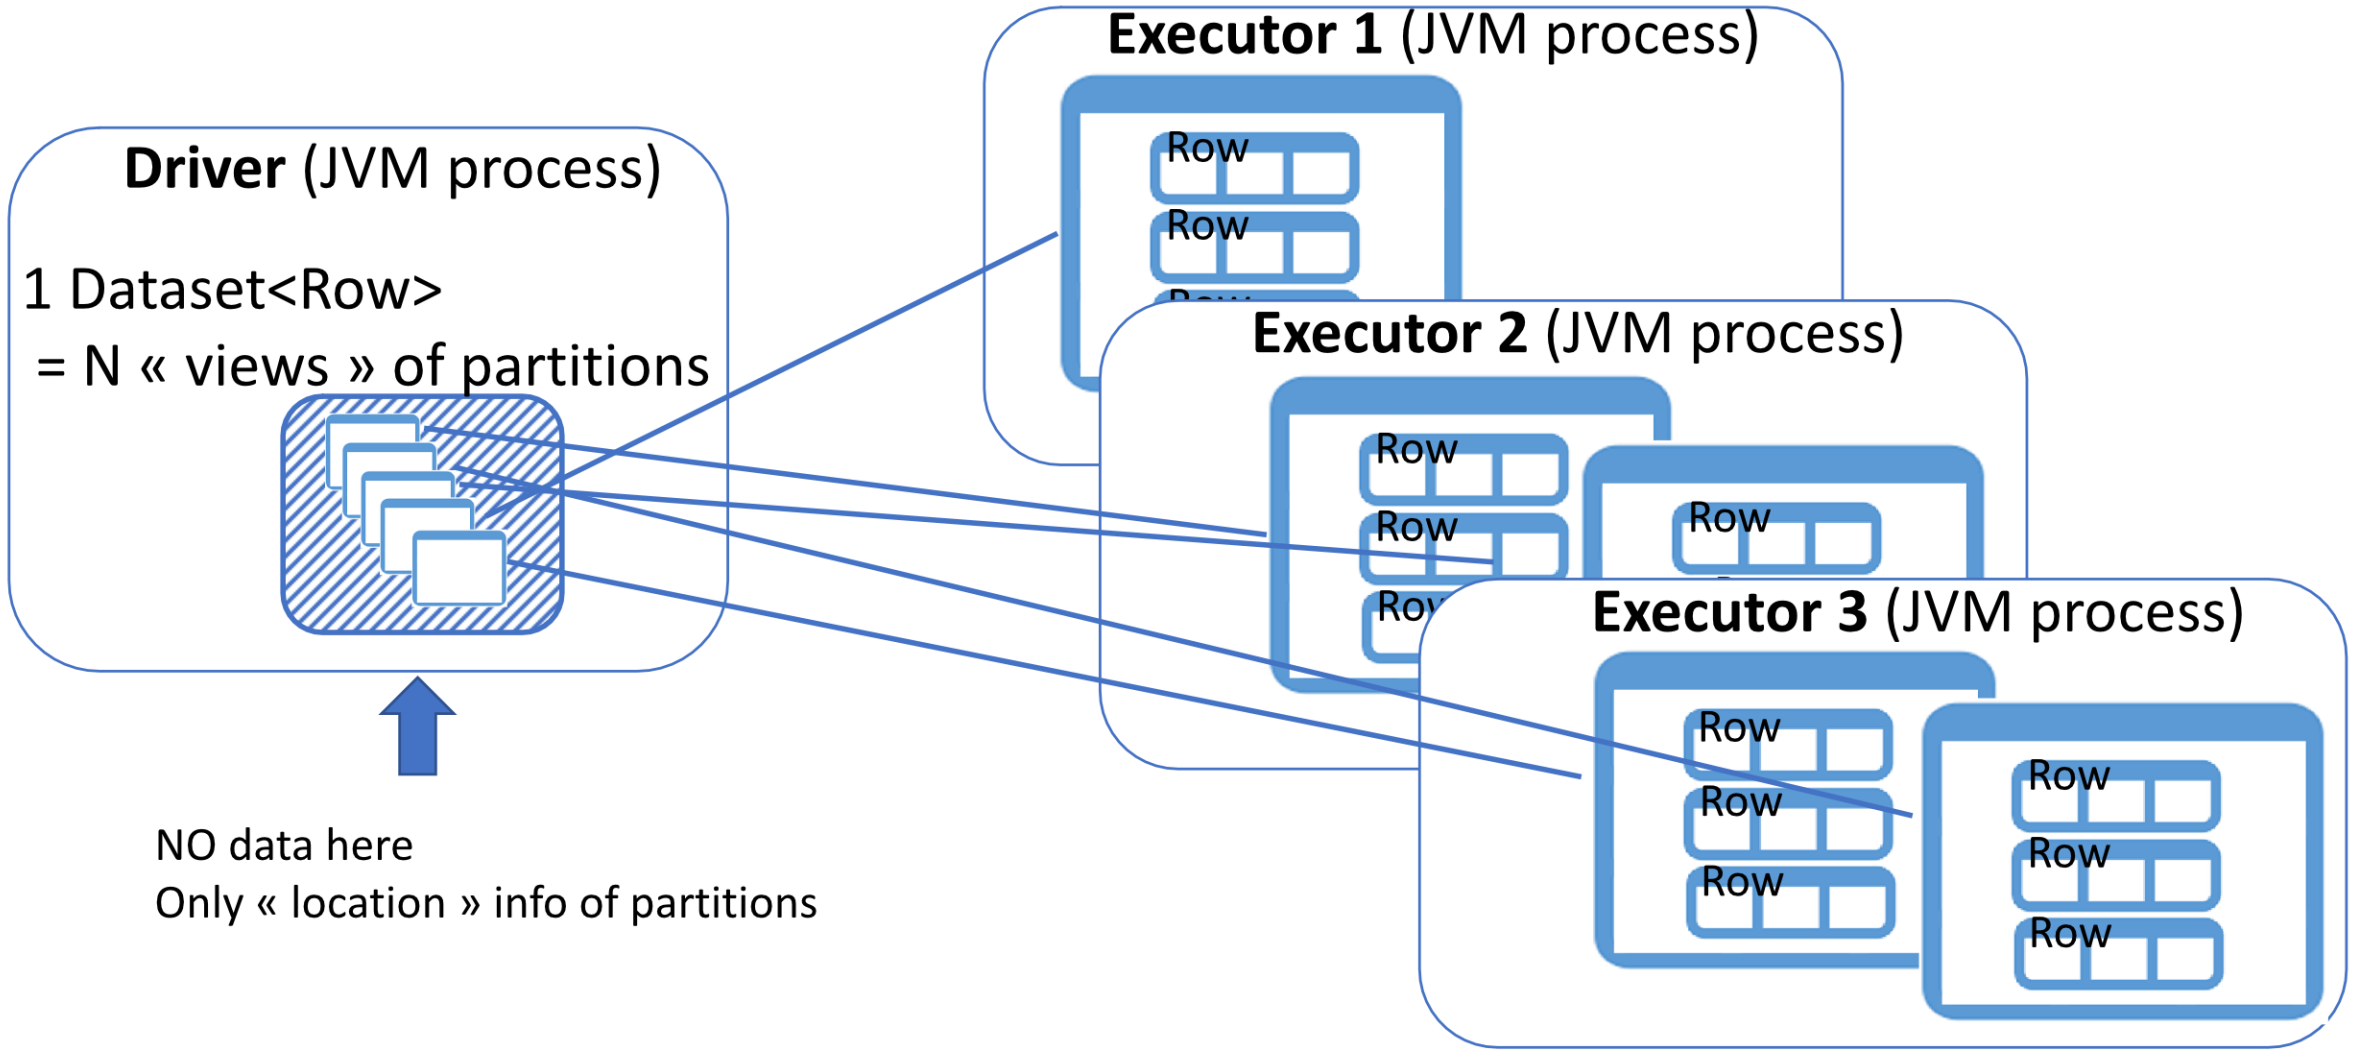

dataset = Spark.createDataset(inmemoryList)
… serialize (send) data to Executors

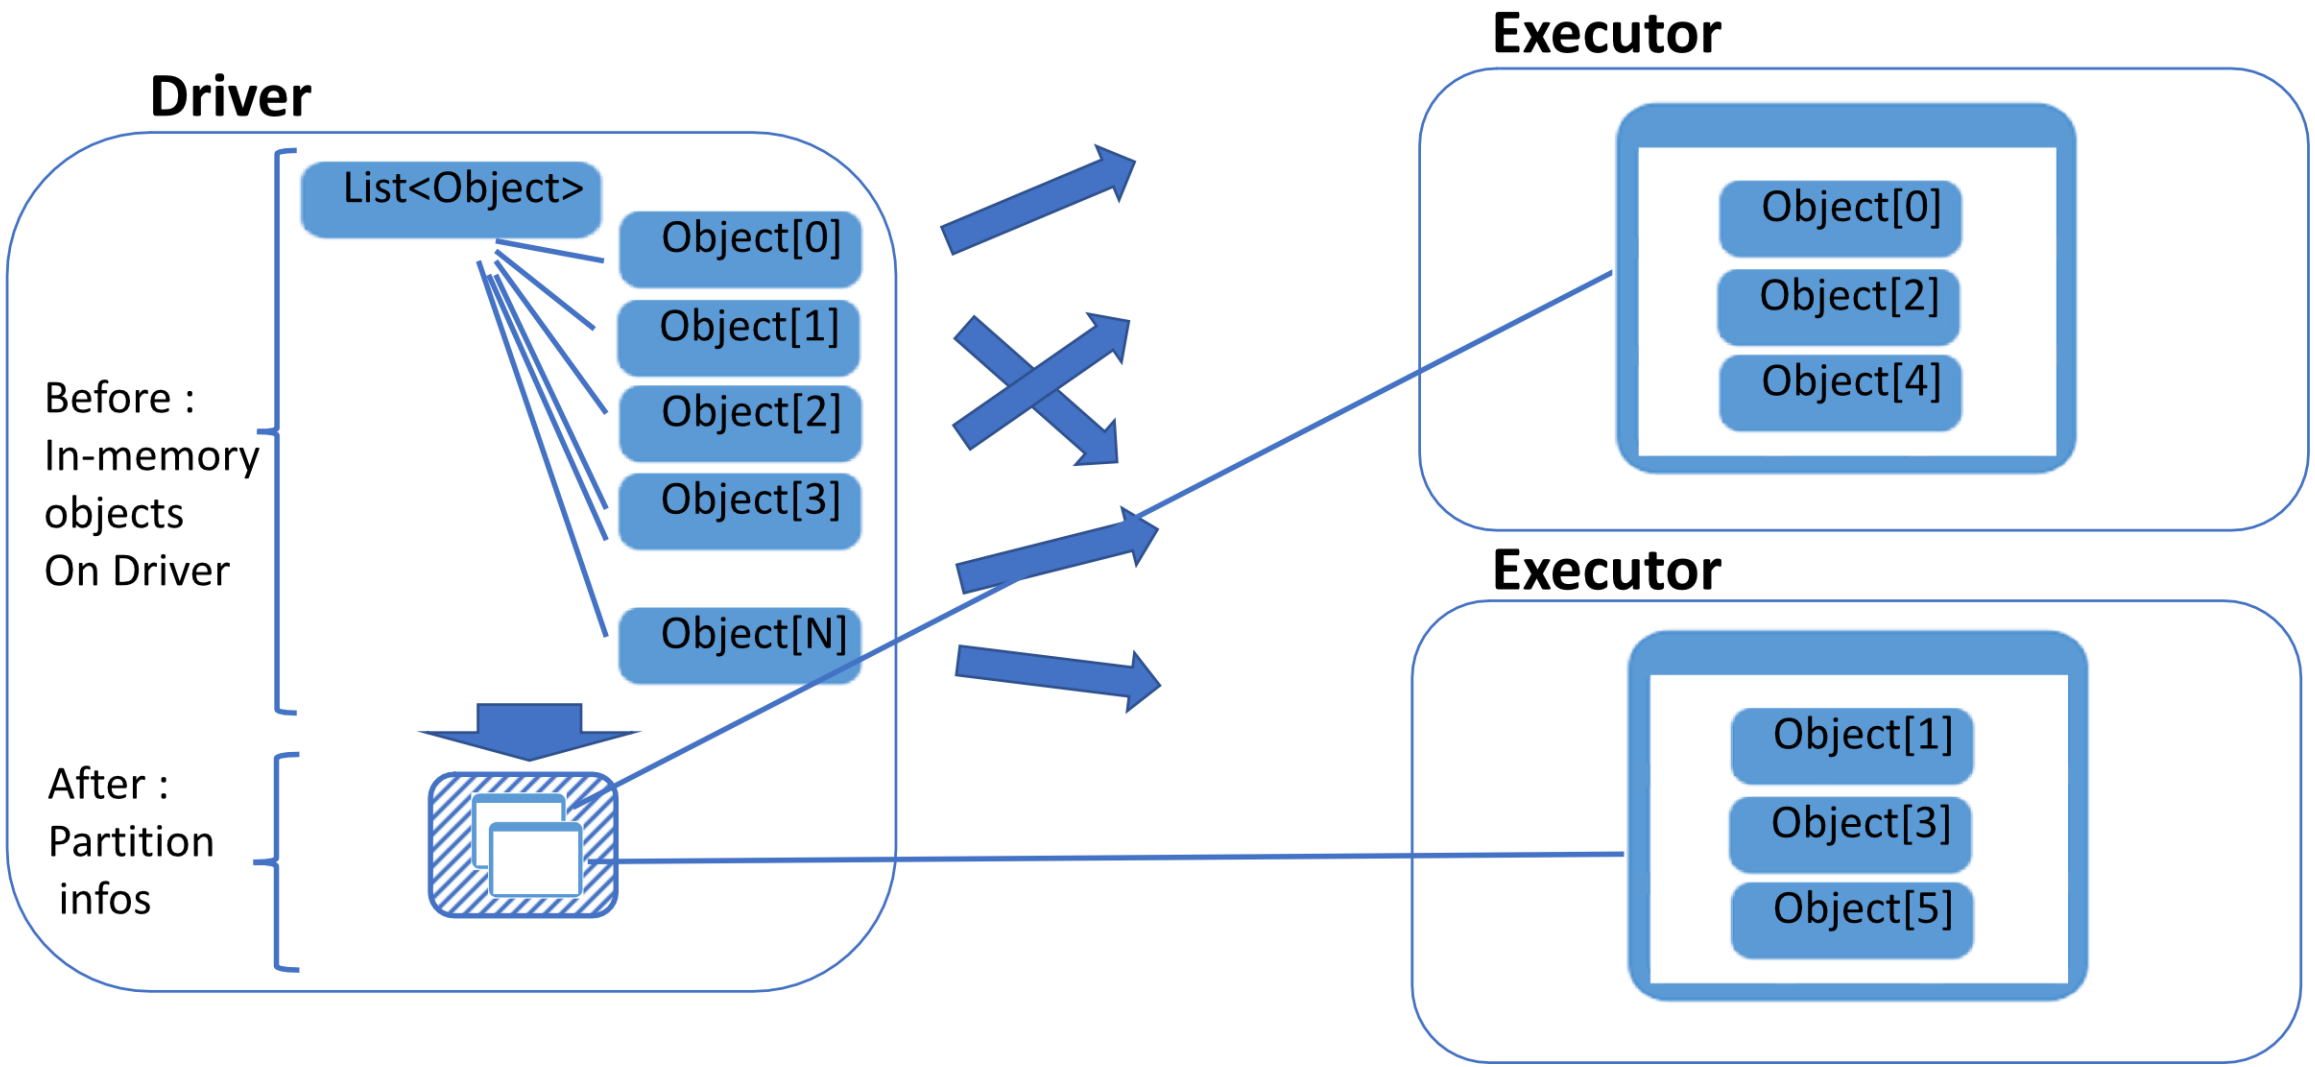

Memory needs on Driver << Executors

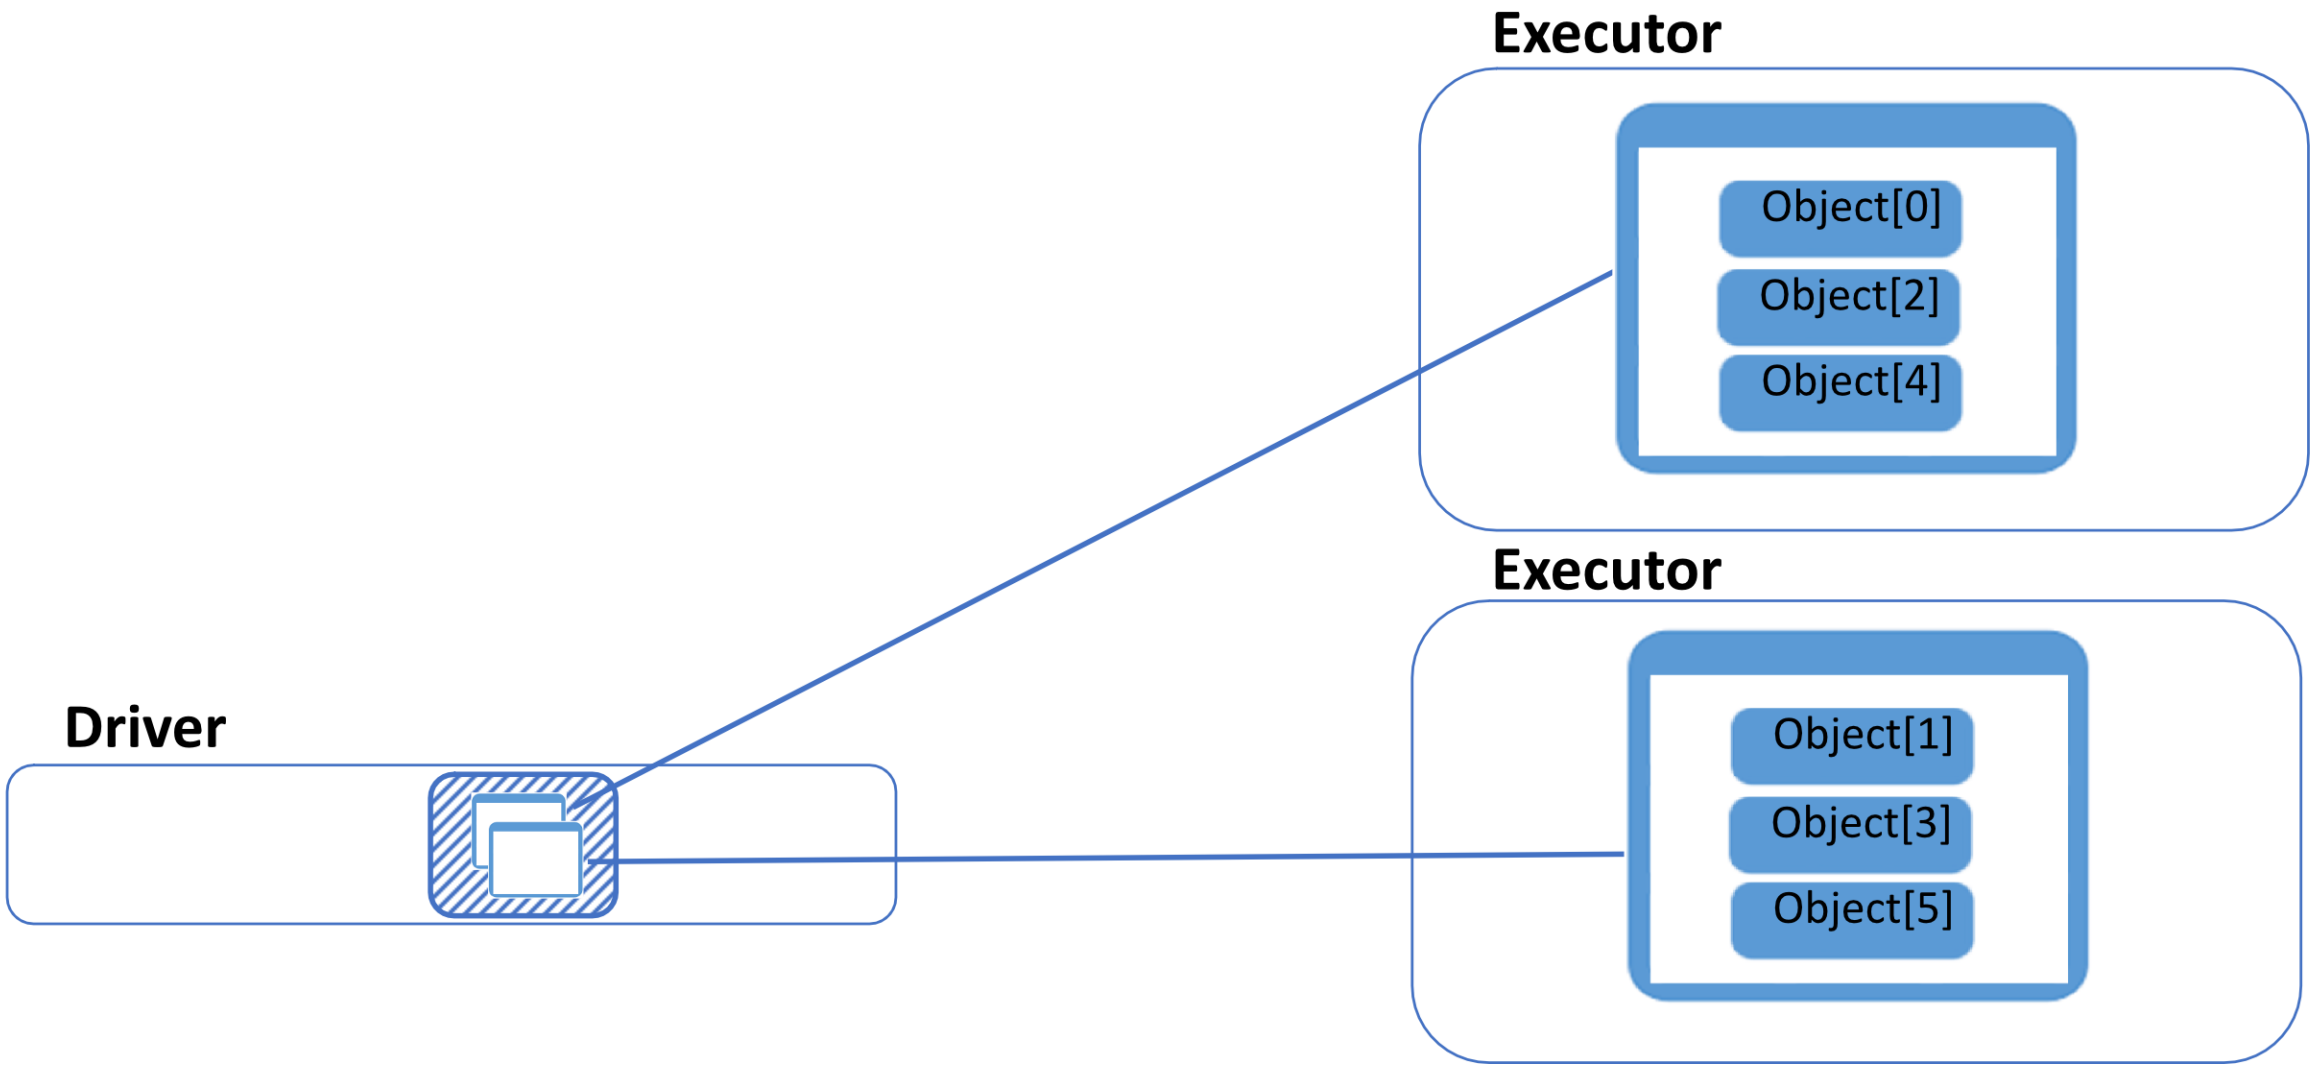

Spark Memory Arguments

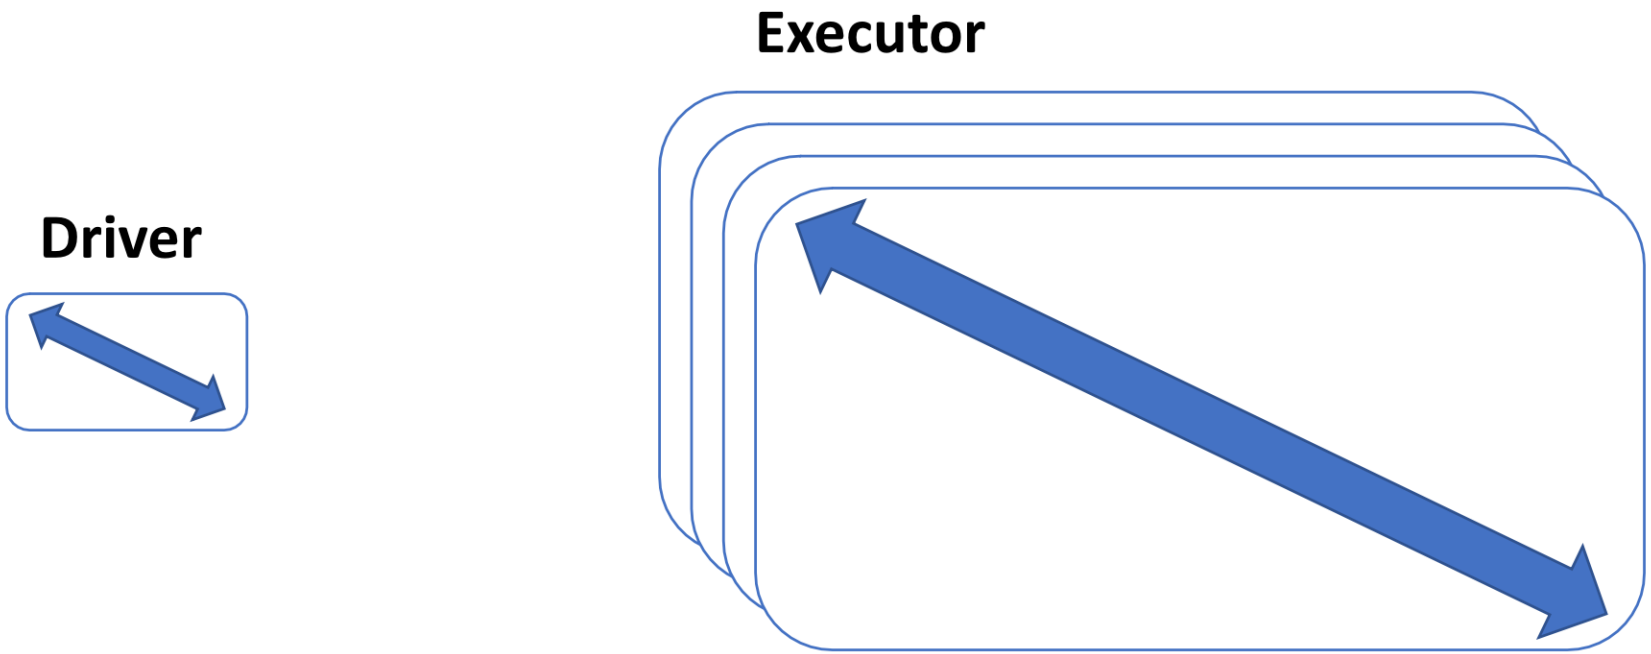

--driver-memory 1g

spark.driver.memory=1g

--executor-memory 50g

spark.executor.memory=50g

Spark Memory …
to {driver|executor} java –Xms .. –Xmx ..

`spark.driver.memory`	1g	Amount of memory to use for the driver process, i.e. where SparkContext is initialized, in the same format as JVM memory strings with a size unit suffix ("k", "m", "g" or "t") (e.g. `512m`, `2g`). *Note:* In client mode, this config must not be set through the `SparkConf` directly in your application, because the driver JVM has already started at that point. Instead, please set this through the `--driver-memory` command line option or in your default properties file.
`spark.executor.memory`	1g	Amount of memory to use per executor process, in the same format as JVM memory strings with a size unit suffix ("k", "m", "g" or "t") (e.g. `512m`, `2g`).

Use Cluster Resources :

Distributing N executors

Cluster Manager (SPI)

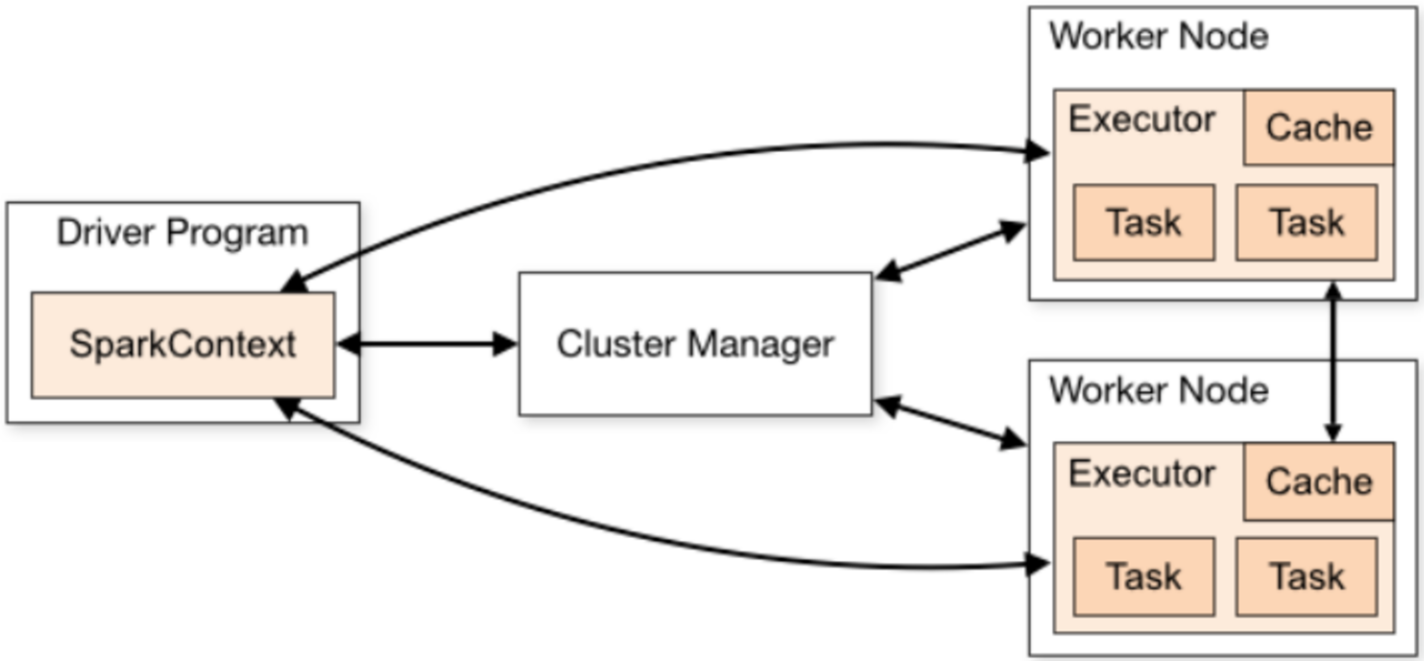

Executors .. JVM processes launched on Cluster

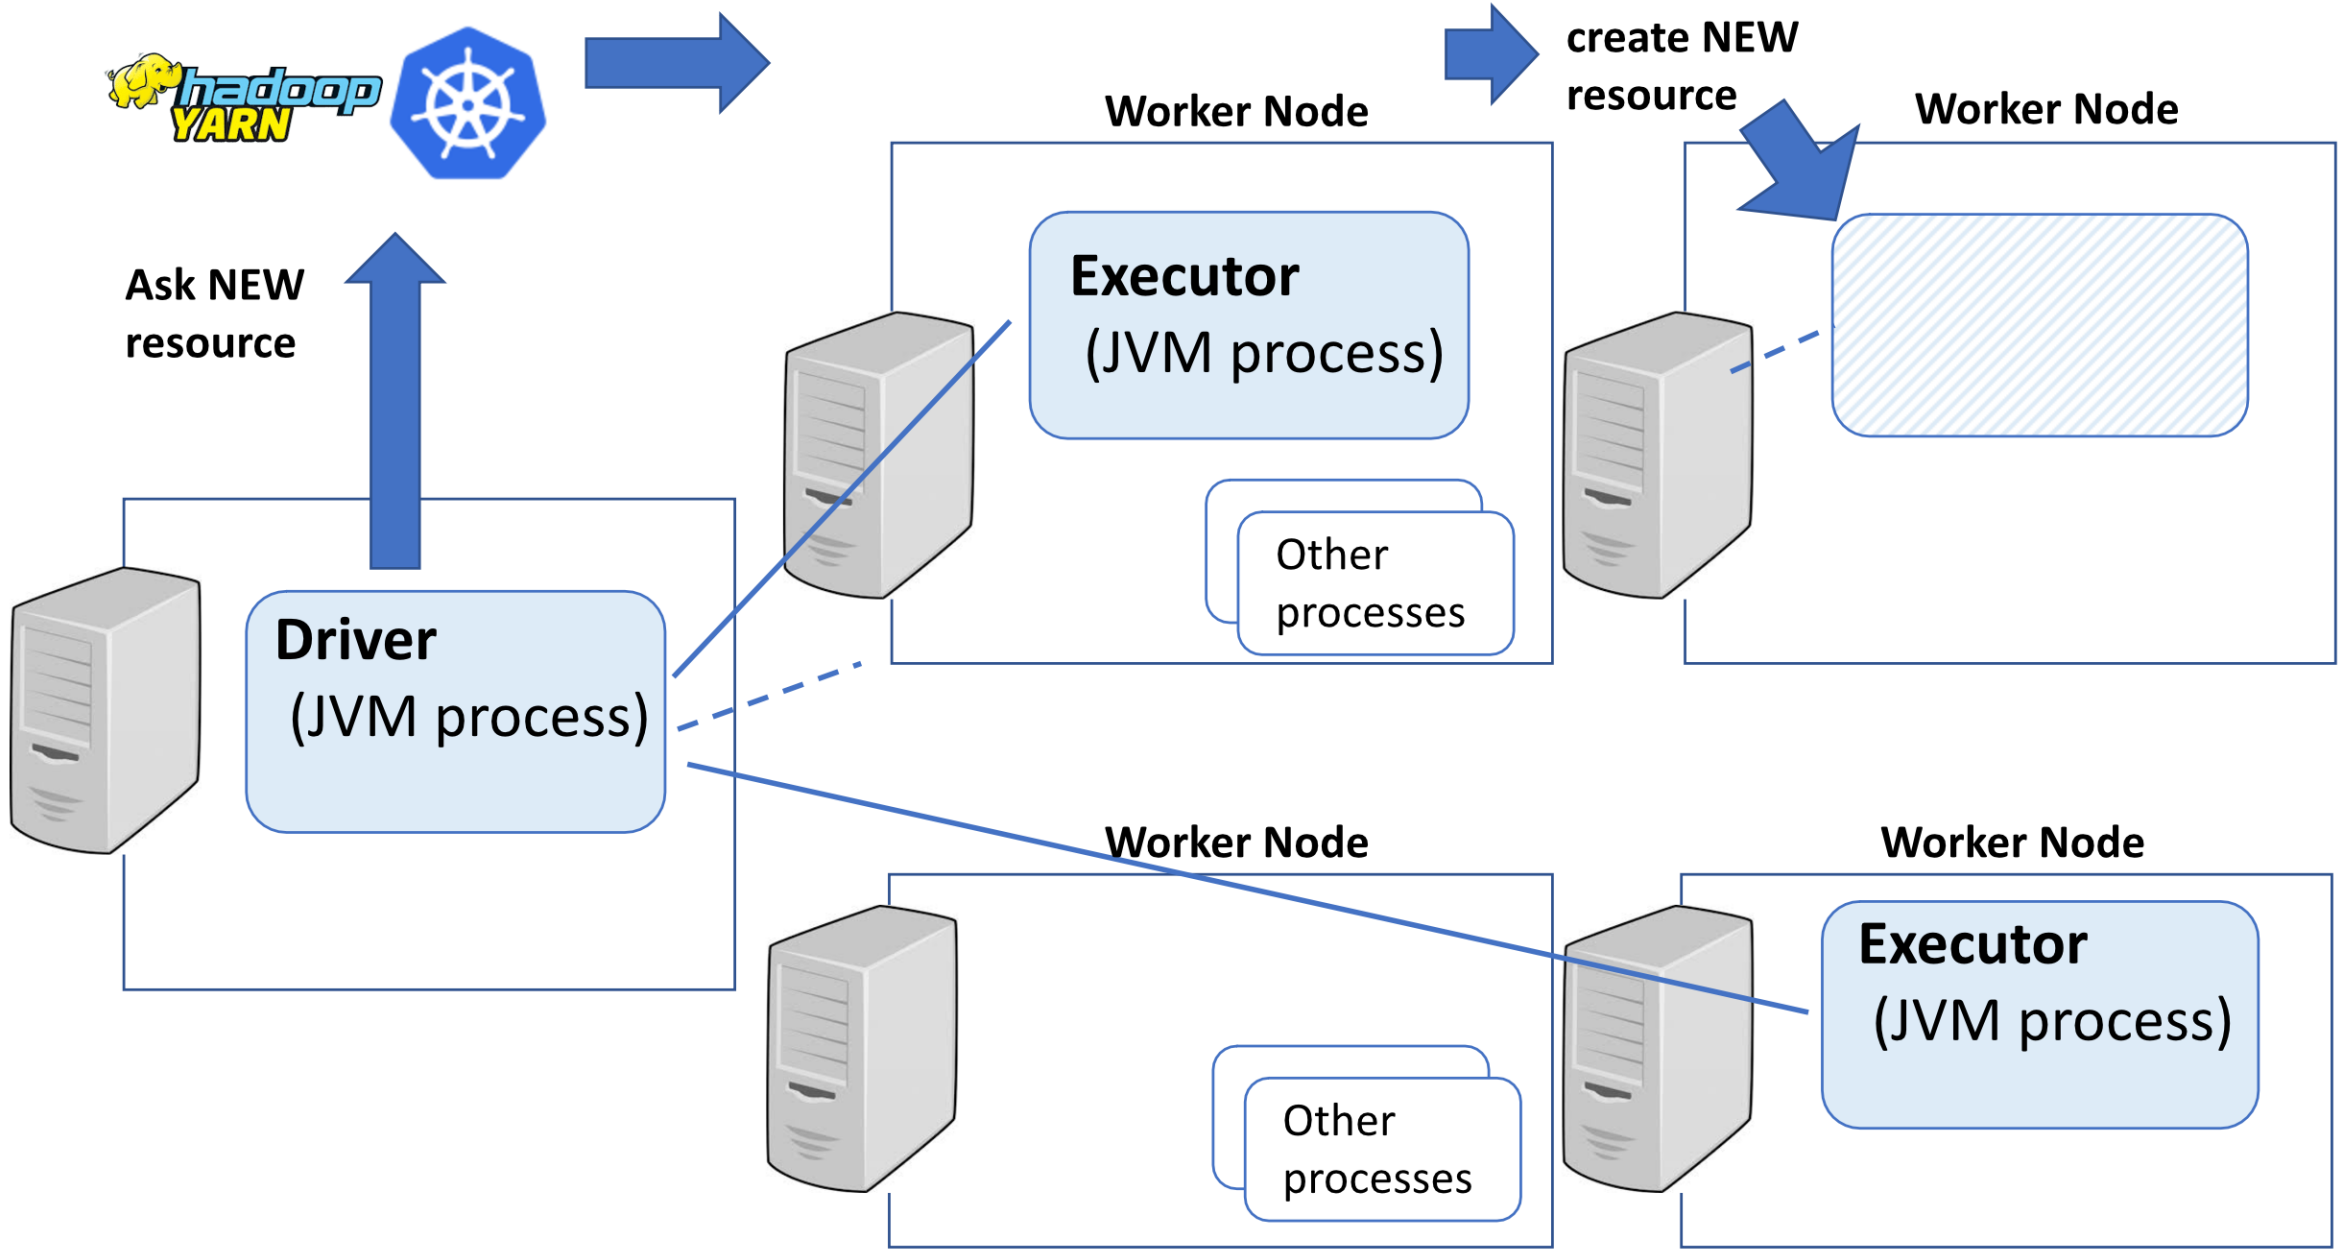

Driver … Launch spark-executors

${JAVA_HOME}/bin/java

JVM args

-server -Xmx30720m
-Djava.io.tmpdir=/yarn/local/usercache/<user> /appcache/<application_id>/<container_id>/tmp
-Dspark.history.ui.port=18080
-Dspark.driver.port=33029
-Dspark.yarn.app.container.log.dir /yarn/log/<application_id>/<container_id>/

Main class

org.apache.spark.executor.**CoarseGrainedExecutorBackend**

main(
String[] args)

--driver-url spark://CoarseGrainedScheduler@<driverhost>:33029
--executor-id 1
--hostname <executorhost>
--cores 5
--app-id <application_id>
--user-class-path file:/yarn/local/usercache/<user> /appcache/<application_id>/<container_id>/__app__.jar

Spark Cluster arguments

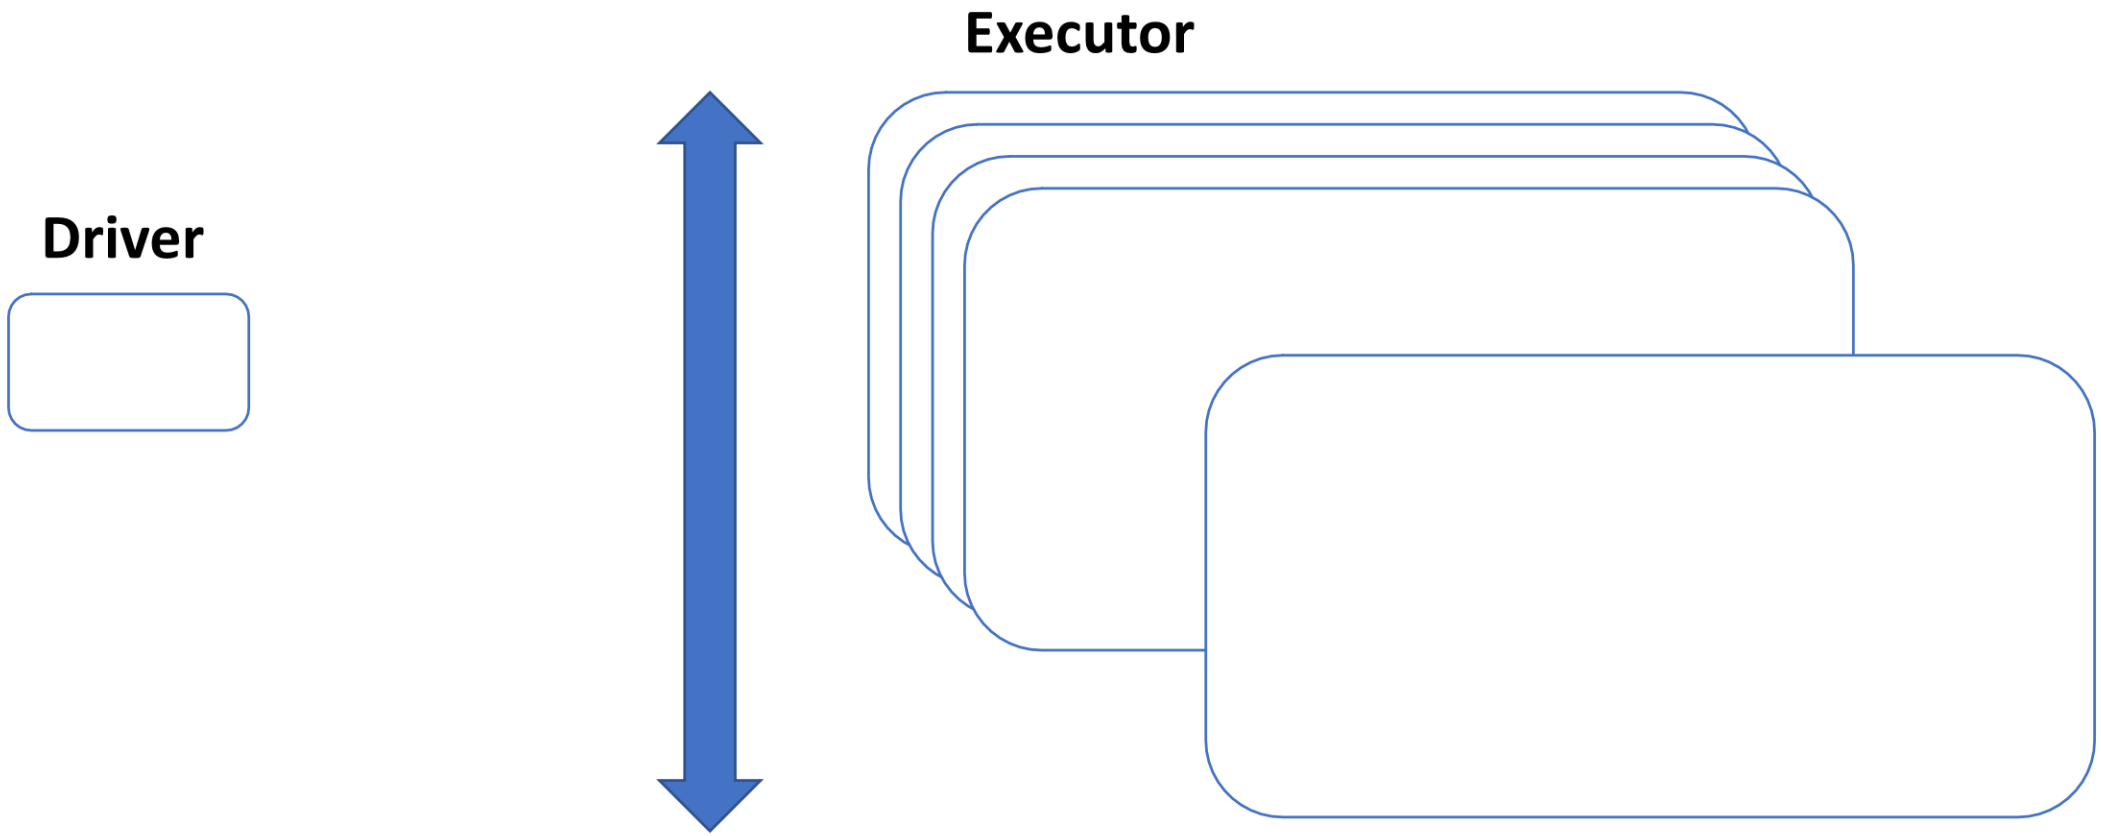

spark.executor.instances=5

spark.executor.instances
.. for Static allocation

`spark.executor.instances`	2	The number of executors for static allocation. With `spark.dynamicAllocation.enabled`, the initial set of executors will be at least this large.

.. See also Dynamic Allocation

Dynamic Allocation

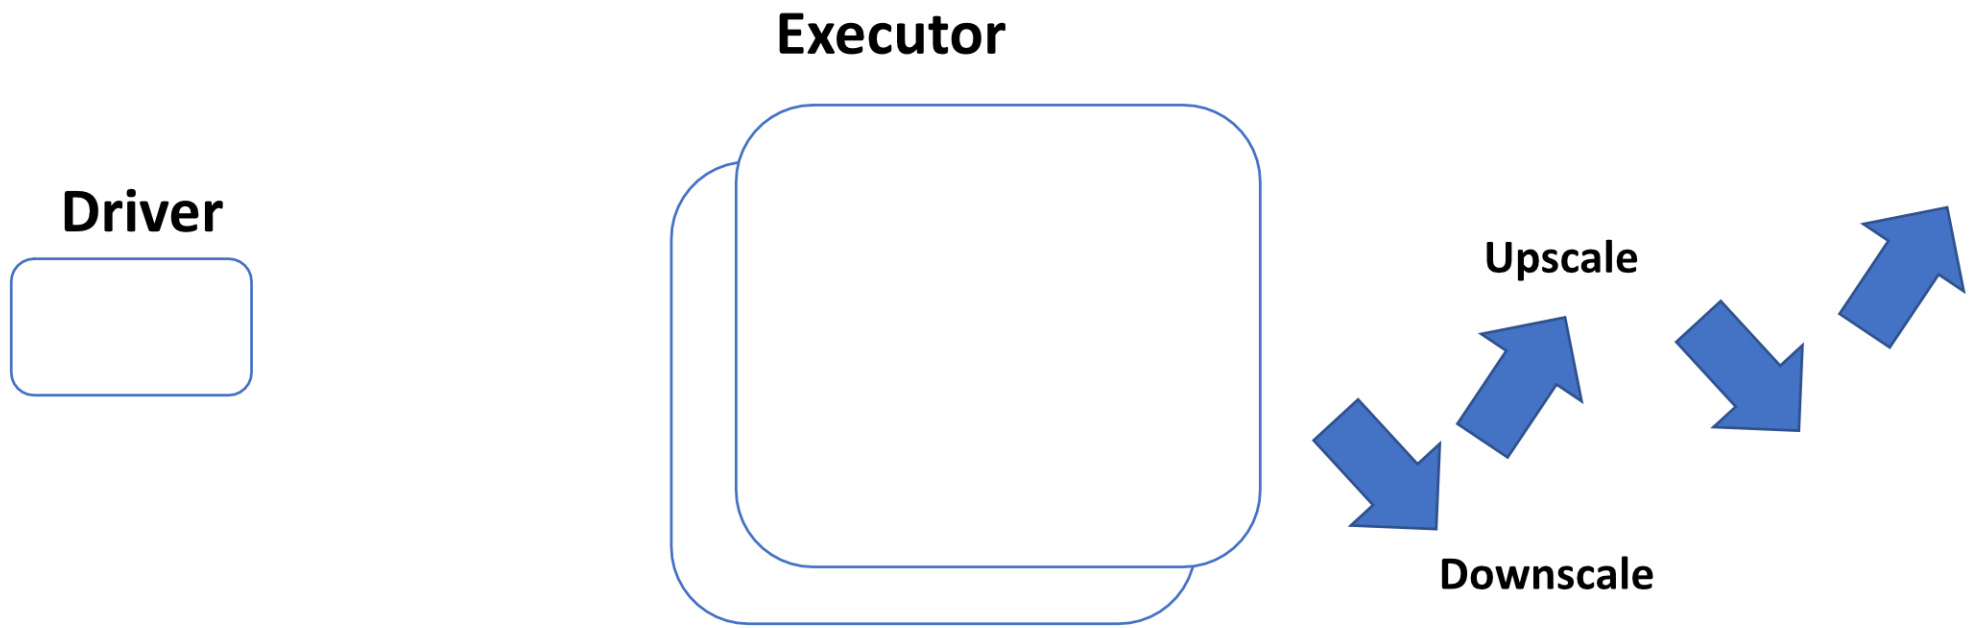

--conf spark.dynamic.allocation.enabled true

--conf spark.dynamic.allocation.{min | max | initialExecutors | executorIdleTimeout | … }

Dynamic Allocation arguments

Dynamic Allocation

Property Name	Default	Meaning	Since Version
`spark.dynamicAllocation.enabled`	false	Whether to use dynamic resource allocation, which scales the number of executors registered with this application up and down based on the workload. For more detail, see the description here. This requires `spark.shuffle.service.enabled` or `spark.dynamicAllocation.shuffleTracking.enabled` to be set. The following configurations are also relevant: `spark.dynamicAllocation.minExecutors`, `spark.dynamicAllocation.maxExecutors`, and `spark.dynamicAllocation.initialExecutors` `spark.dynamicAllocation.executorAllocationRatio`	1.2.0
`spark.dynamicAllocation.executorIdleTimeout`	60s	If dynamic allocation is enabled and an executor has been idle for more than this duration, the executor will be removed. For more detail, see this description.	1.2.0
`spark.dynamicAllocation.cachedExecutorIdleTimeout`	infinity	If dynamic allocation is enabled and an executor which has cached data blocks has been idle for more than this duration, the executor will be removed. For more details, see this description.	1.4.0
`spark.dynamicAllocation.initialExecutors`	`spark.dynamicAllocation.minExecutors`	Initial number of executors to run if dynamic allocation is enabled. If `--num-executors` (or `spark.executor.instances`) is set and larger than this value, it will be used as the initial number of executors.	1.3.0
`spark.dynamicAllocation.maxExecutors`	infinity	Upper bound for the number of executors if dynamic allocation is enabled.	1.2.0
`spark.dynamicAllocation.minExecutors`	0	Lower bound for the number of executors if dynamic allocation is enabled.	1.2.0
`spark.dynamicAllocation.executorAllocationRatio`	1	By default, the dynamic allocation will request enough executors to maximize the parallelism according to the number of tasks to process. While this minimizes the latency of the job, with small tasks this setting can waste a lot of resources due to executor allocation overhead, as some executor might not even do any work. This setting allows to set a ratio that will be used to reduce the number of executors w.r.t. full parallelism. Defaults to 1.0 to give maximum parallelism. 0.5 will divide the target number of executors by 2 The target number of	2.4.0

Dynamic Spark + Dynamic K8s Cloud Cluster

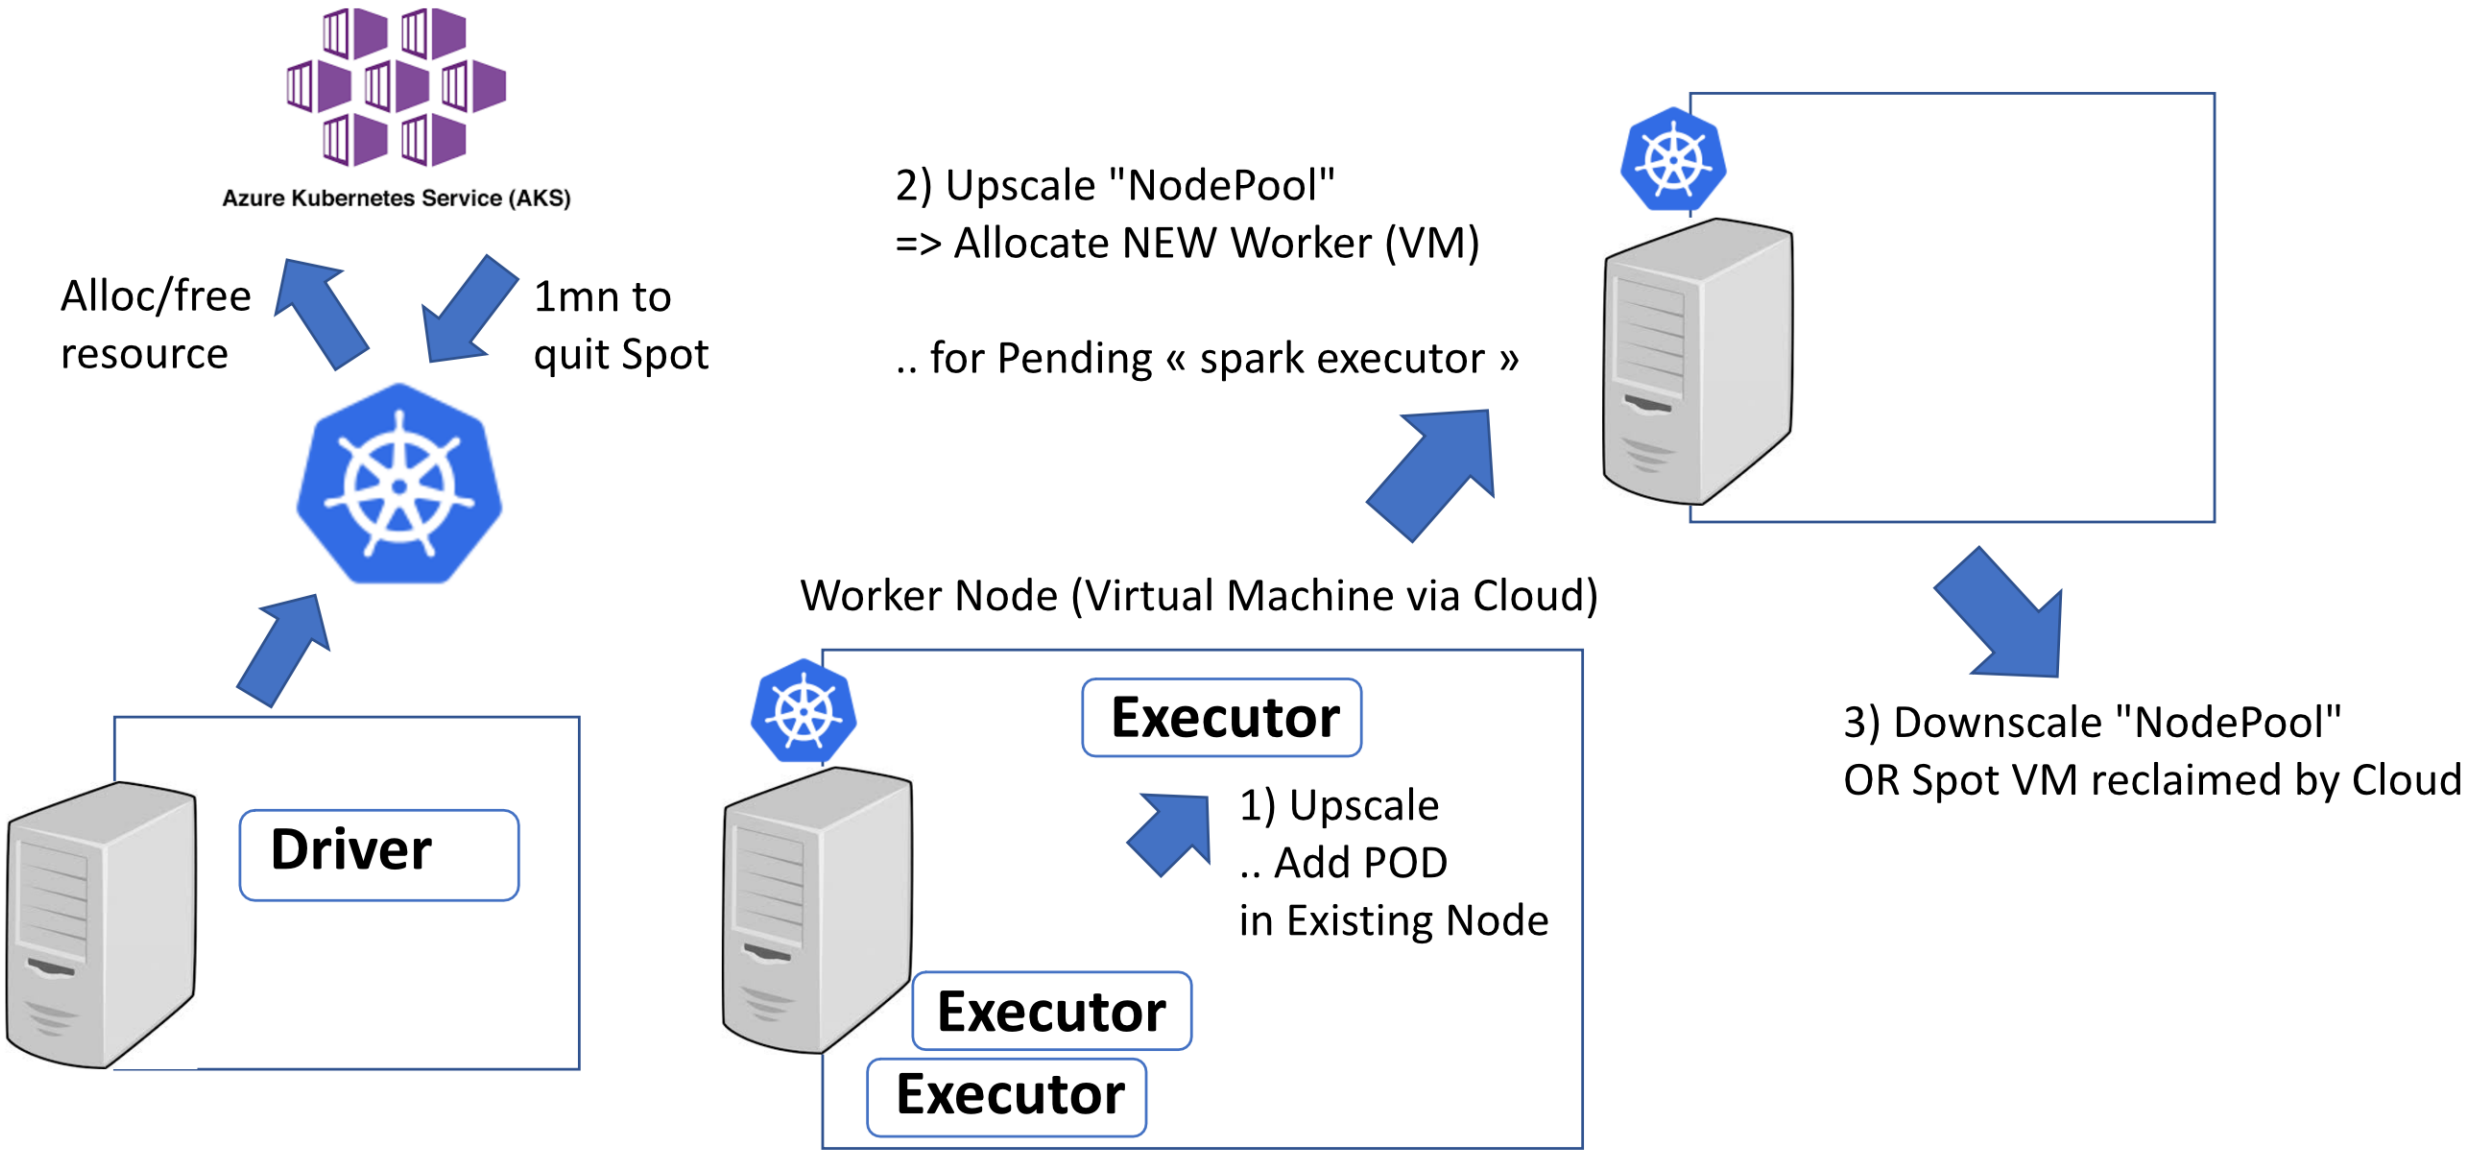

Spark dynamicAllocation
<==> K8s NodePool autoScaling

=> more (system) VMs from Cloud Provider

=> when full: NodePool "TriggerScaleUp"   (VM provisioning may take ~3mn)

=>   need more k8s resources

allocating new (applicative) "executor"

releasing (applicative) "executor"

=>   use less k8s resources

=> when empty (0 running pod on Node), VM could be evicted from NodePool

(example: configure with 15mn delay)

=> less (system) VMs from Cloud Provider

K8s NodePool, Affinity, Toleration, Spot VM ..

Node Pool 1

Small VMs … reserved
=> for long running services, spark Drivers

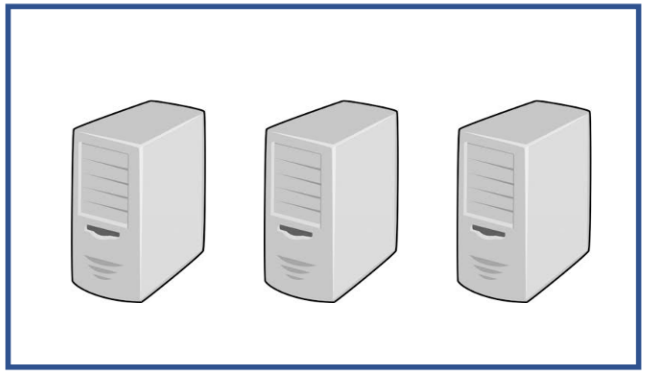

Node Pool 2

Big VMs … « SPOT »
(cheaper, may be reclaimed by cloud)
=> for spark Executors

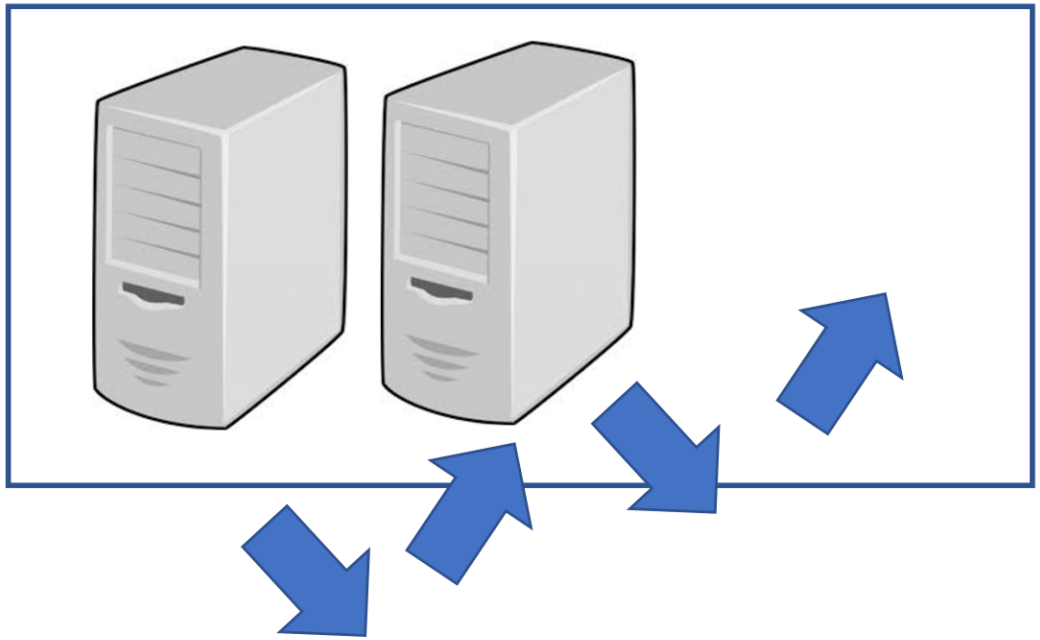

Use CPU Resource: 100% or sharing ?

CPU Cores <-> Threads <-> max Active Tasks

CPU

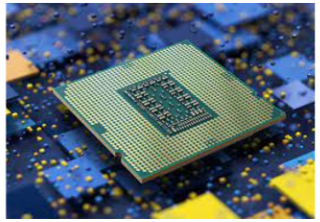

x Cores

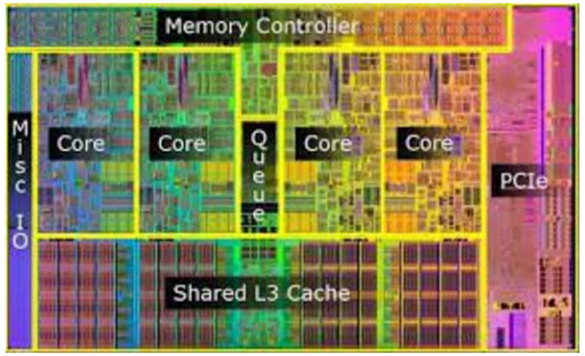

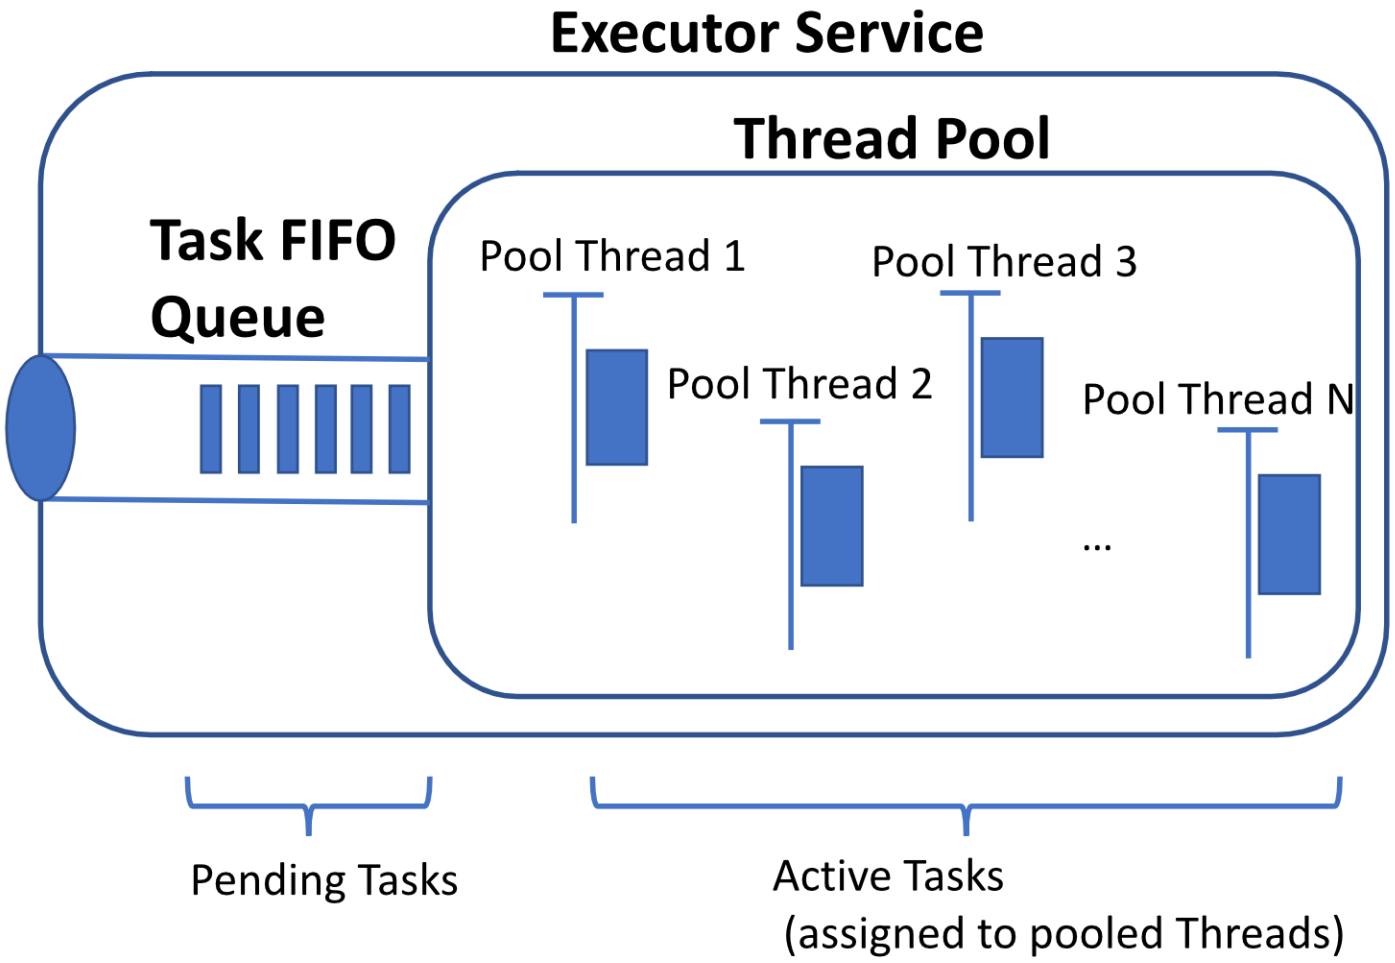

Executor Cores argument

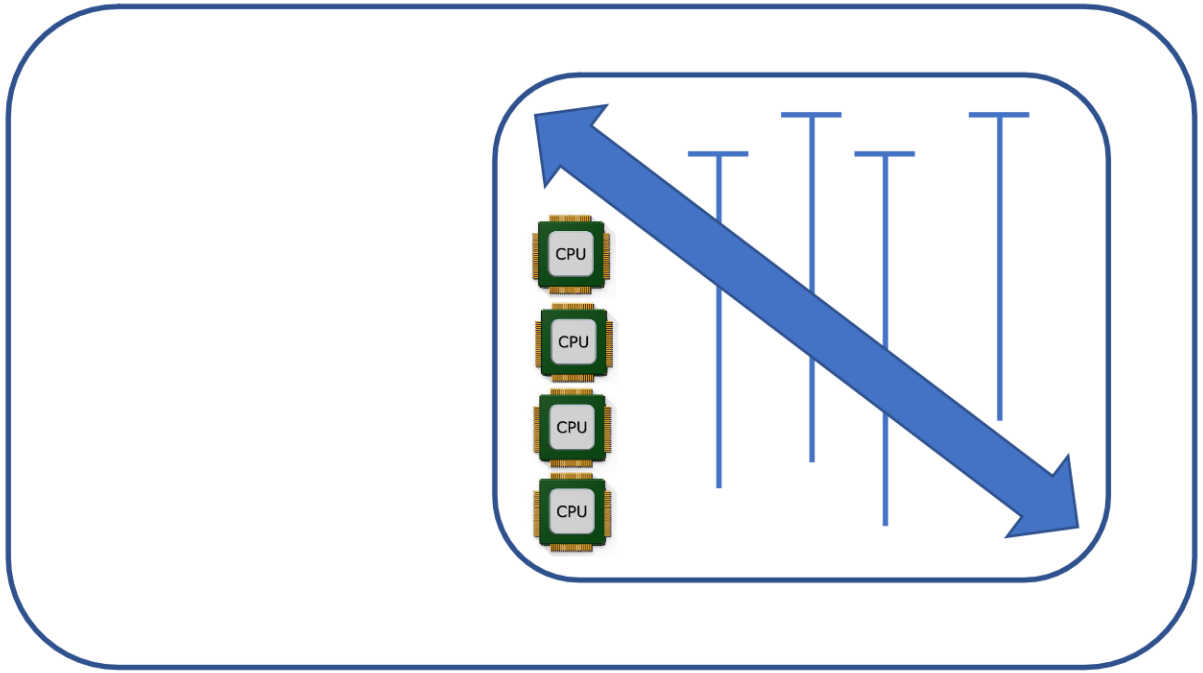

`spark.executor.cores`	1 in YARN mode, all the available cores on the worker in standalone and Mesos coarse-grained modes.	The number of cores to use on each executor. In standalone and Mesos coarse-grained modes, for more detail, see this description.

--executor-cores N

spark.executor.cores=N

Relationship executor.cores <-> .memory

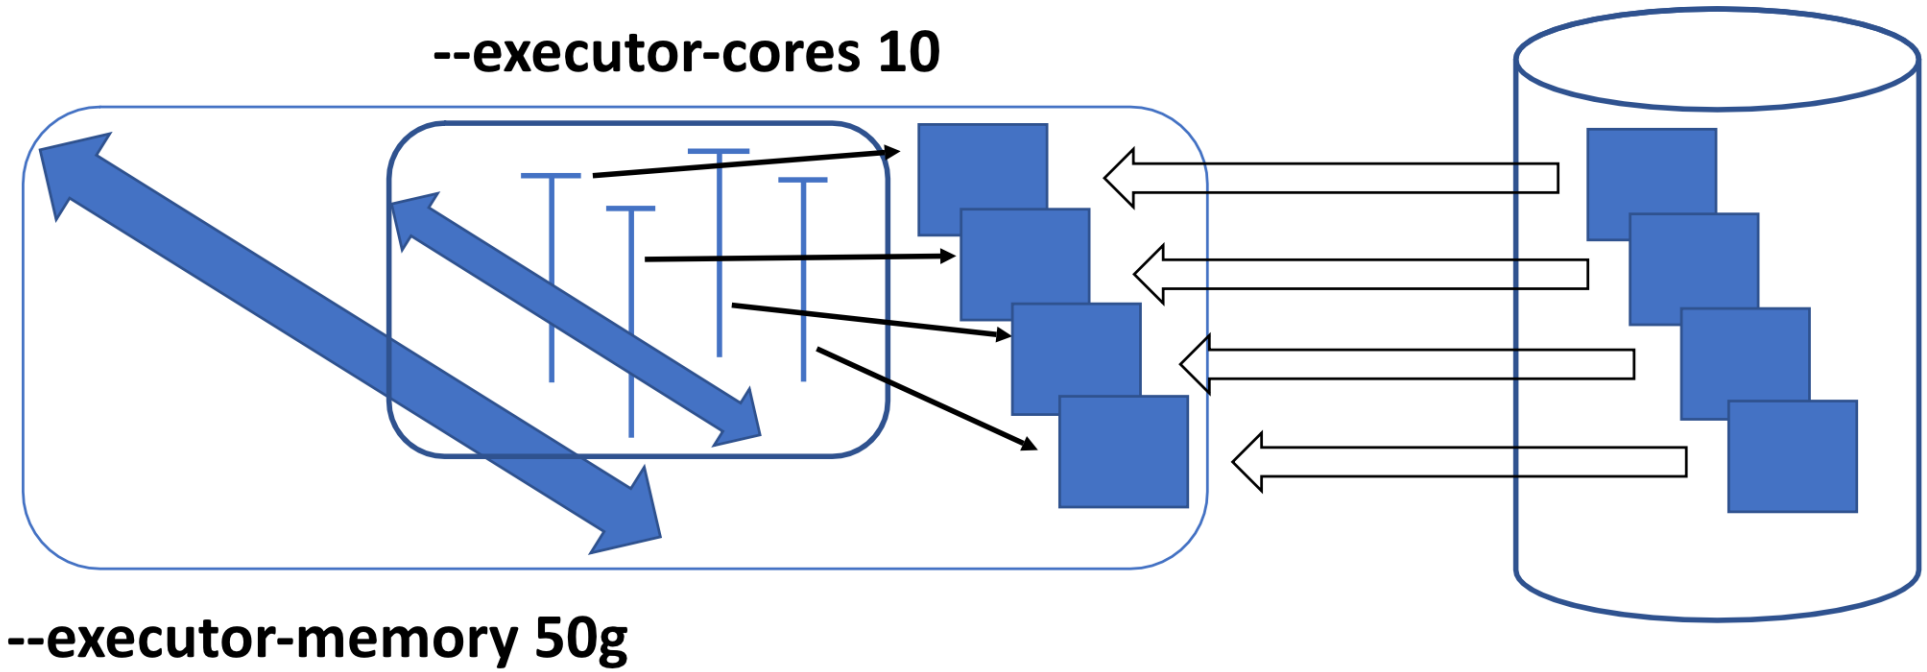

1 Active Core => 1 Task => read 1 File Block => 1 Dataset Partition

=> … maybe 1 x 256 MB (compressed) … 1 x GB ?

Summary main arguments

Driver

--driver-memory 1g

spark.driver.memory=1g

--num-executors 5

spark.executor.instances=5

Executor

--executor-cores N

spark.executor.cores=N

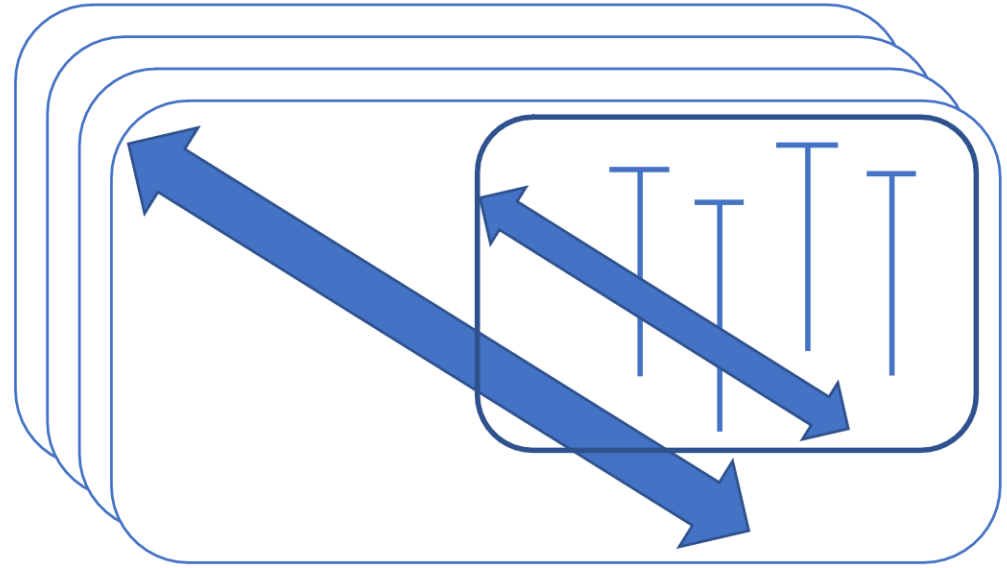

--executor-memory 50g

spark.executor.memory=50g

Outline

Reminder Spark-driver / spark-executors

Cluster Tuning Params

deployMode (local, spark://, yarn, k8s)

Other launchers (java embedded, livy, jupyter, pyspark ..)

```
$ spark-submit --master ${mode} ....
```

${mode}: one of

- **local[N]**
- **mesos (deprecated)**
- **spark://**
- **yarn**
- **k8s**

--master local[*]

```
# Run application locally on 8 cores
./bin/spark-submit \
  --class org.apache.spark.examples.SparkPi \
  --master local[8] \
  /path/to/examples.jar \
  100
```

Local mode: 1 JVM
client = driver = executors

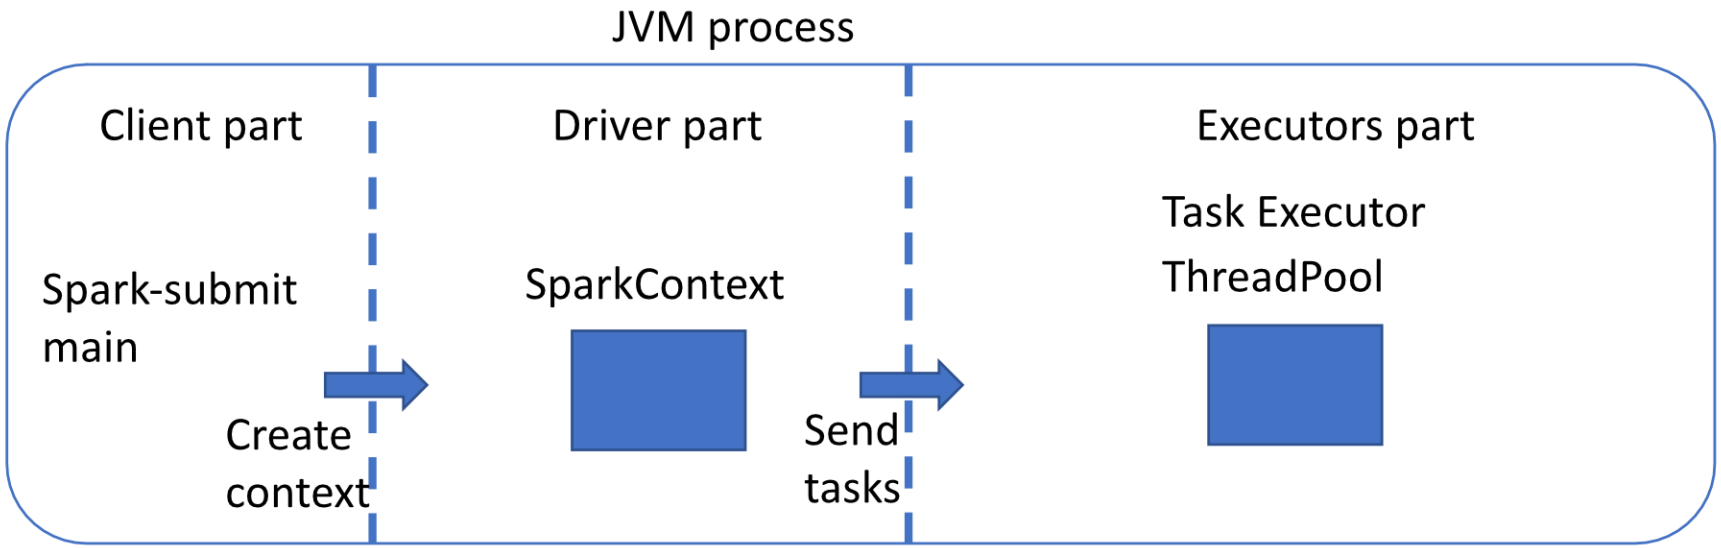

--master spark://<sparkServer>:<port>

```
# Run on a Spark standalone cluster in client deploy mode
./bin/spark-submit \
  --class org.apache.spark.examples.SparkPi \
  --master spark://207.184.161.138:7077 \
  --executor-memory 20G \
  --total-executor-cores 100 \
  /path/to/examples.jar \
  1000
```

Spark Standalone Cluster

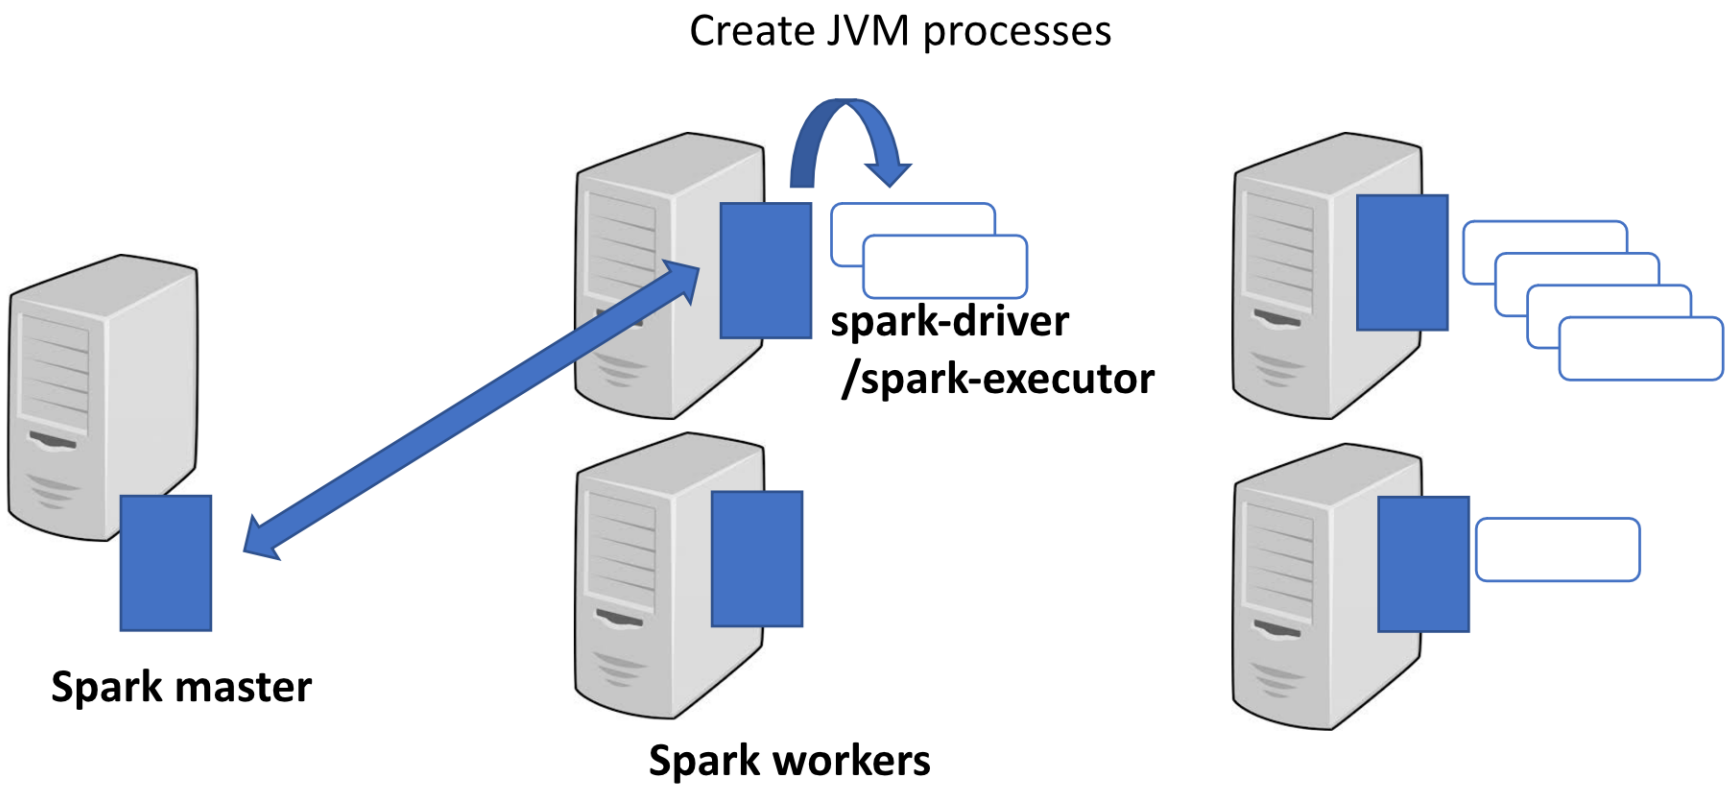

--master yarn--deploy-mode client

```
# Run on a YARN cluster in client deploy mode
export HADOOP_CONF_DIR=XXX
./bin/spark-submit \
  --class org.apache.spark.examples.SparkPi \
  --master yarn \
  --deploy-mode client \
  --executor-memory 20G \
  --num-executors 50 \
  /path/to/examples.jar \
  1000
```

Yarn Client Mode

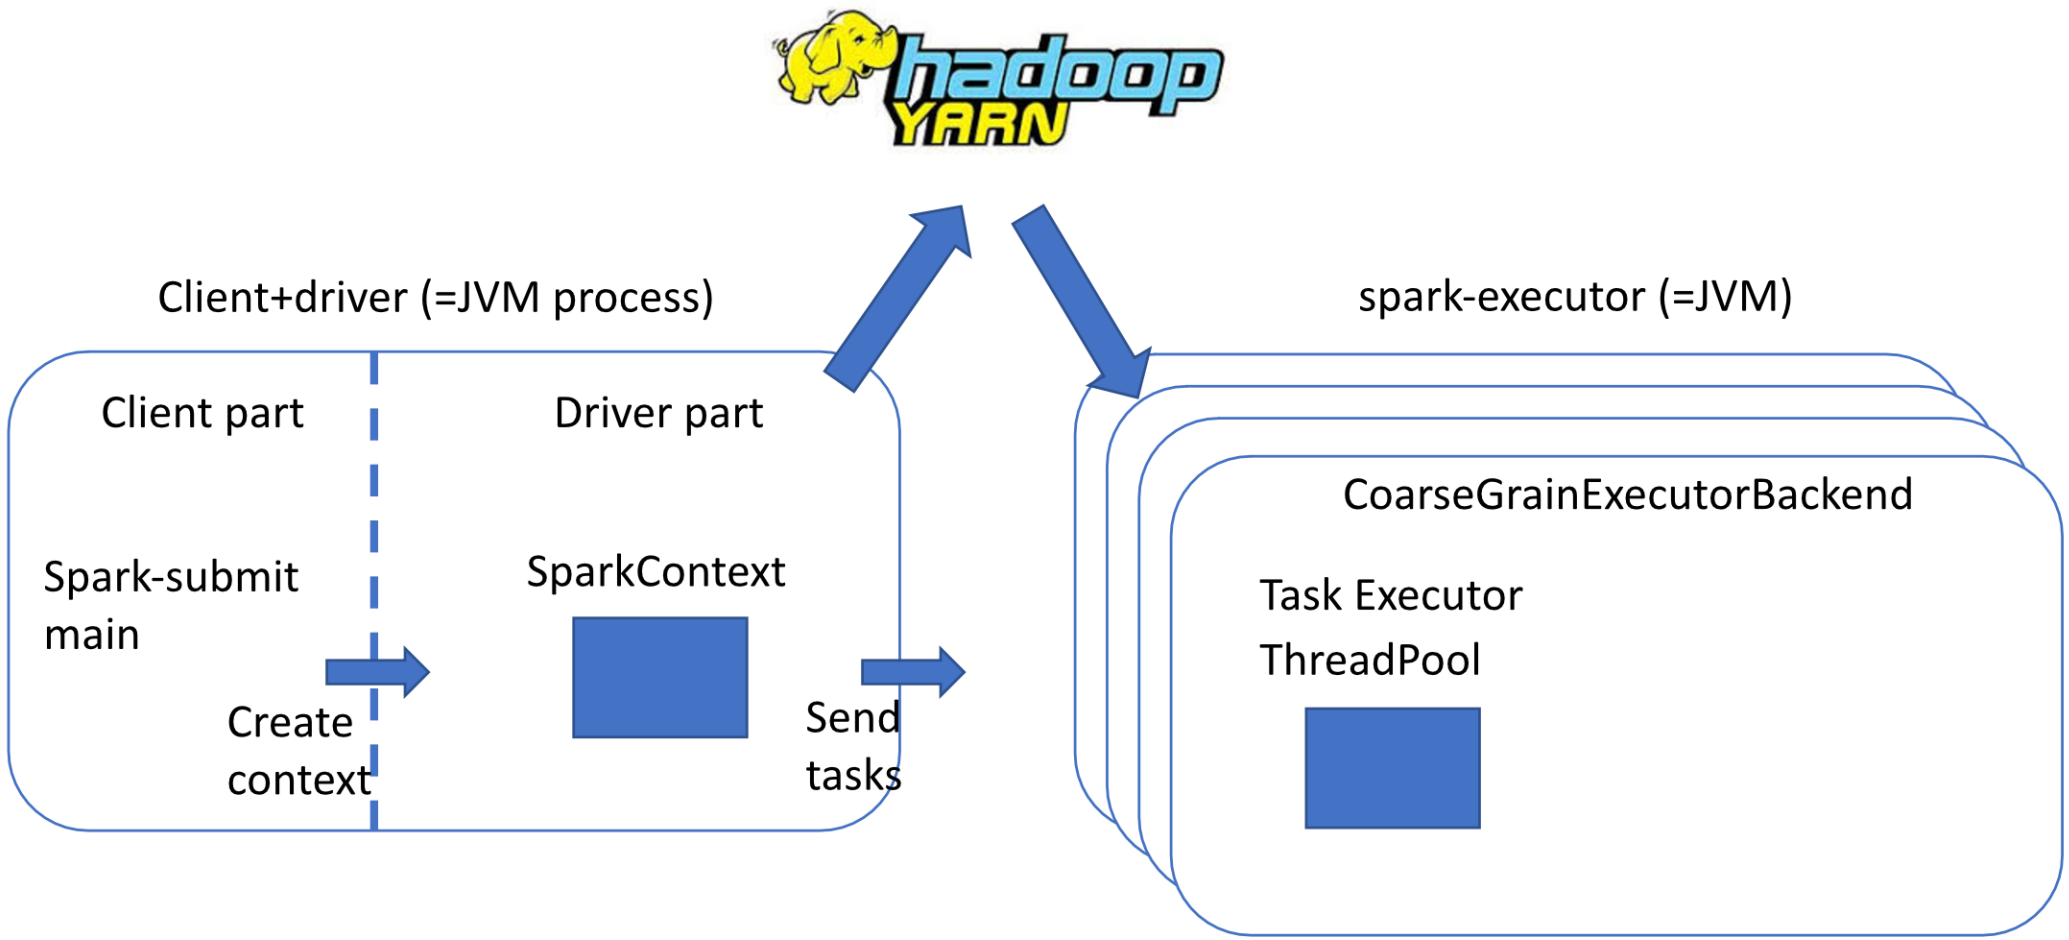

Pros/Cons

Easy to start: spark-submit.sh

Easy to see driver logs… as stdout !

Easy to kill   (Ctrl+C / kill -9)

Client = Driver   … within firewall

Need to choose 1 server of cluster + SSH on it
Or use Oozie

If driver crash.. END

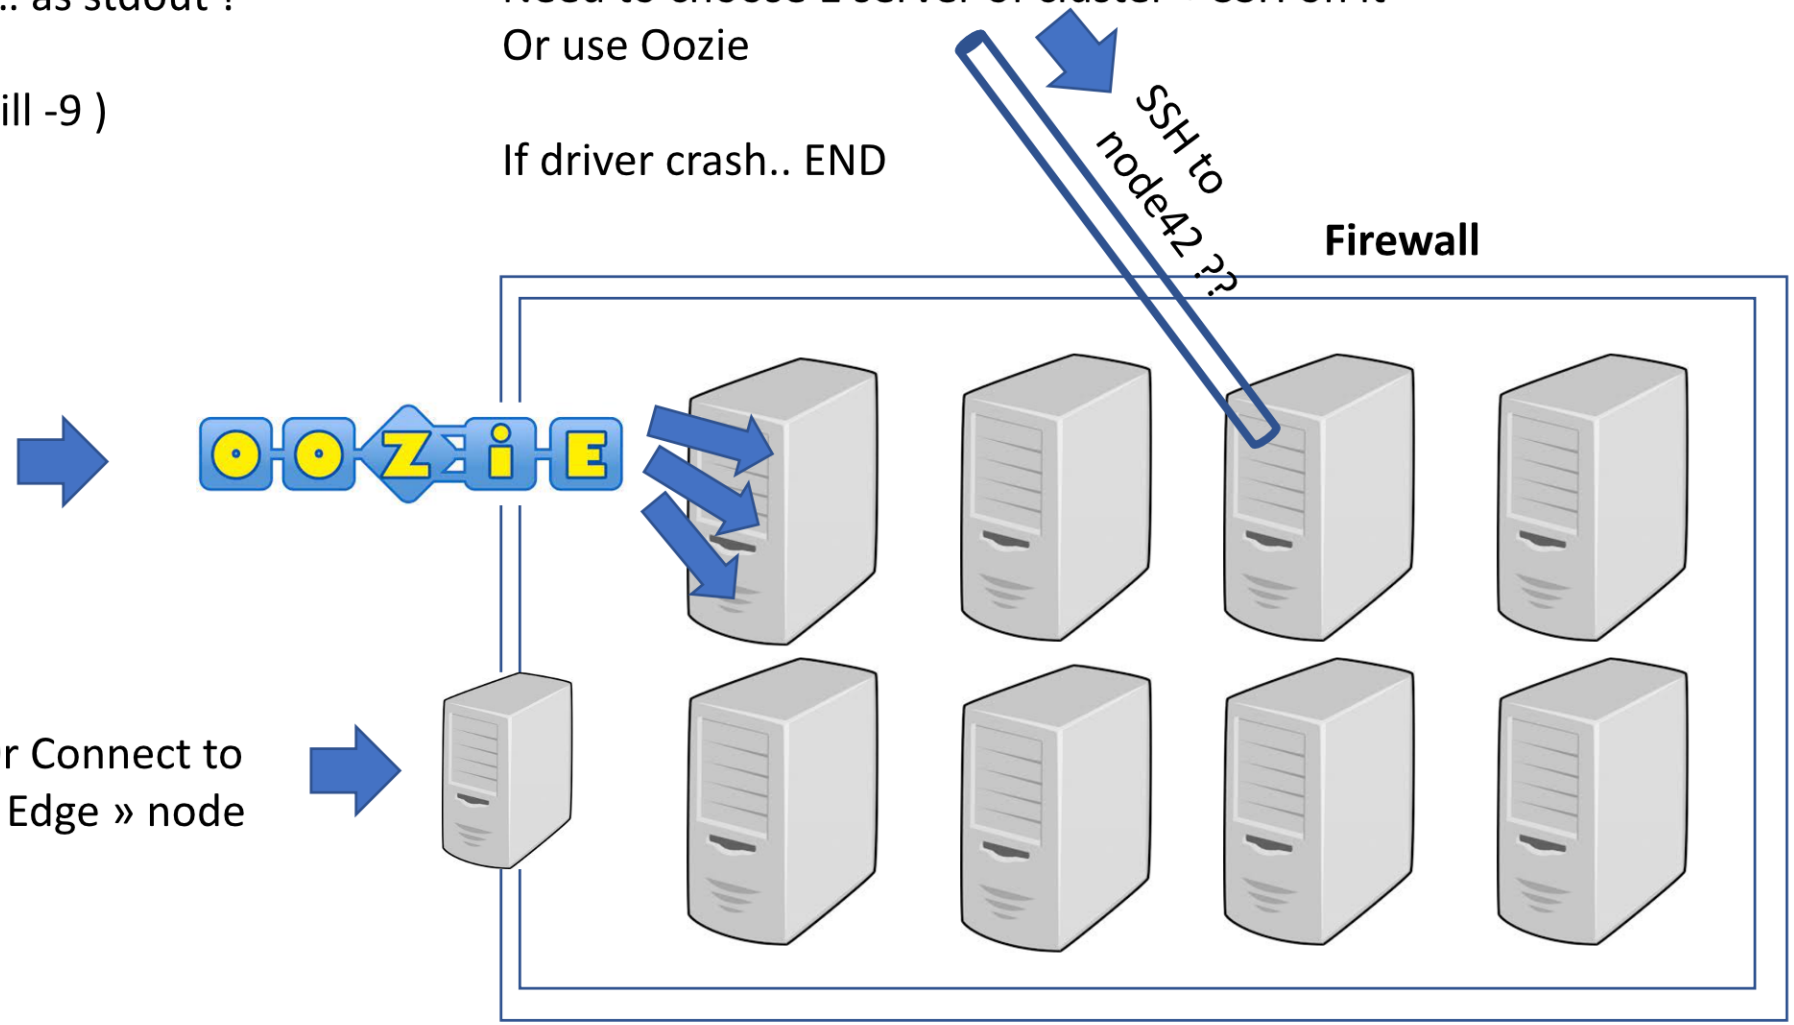

--master yarn--deploy-mode cluster

```
# Run on a YARN cluster in cluster deploy mode
export HADOOP_CONF_DIR=XXX
./bin/spark-submit \
  --class org.apache.spark.examples.SparkPi \
  --master yarn \
  --deploy-mode cluster \
  --executor-memory 20G \
  --num-executors 50 \
  /path/to/examples.jar \
  1000
```

Yarn Cluster Mode

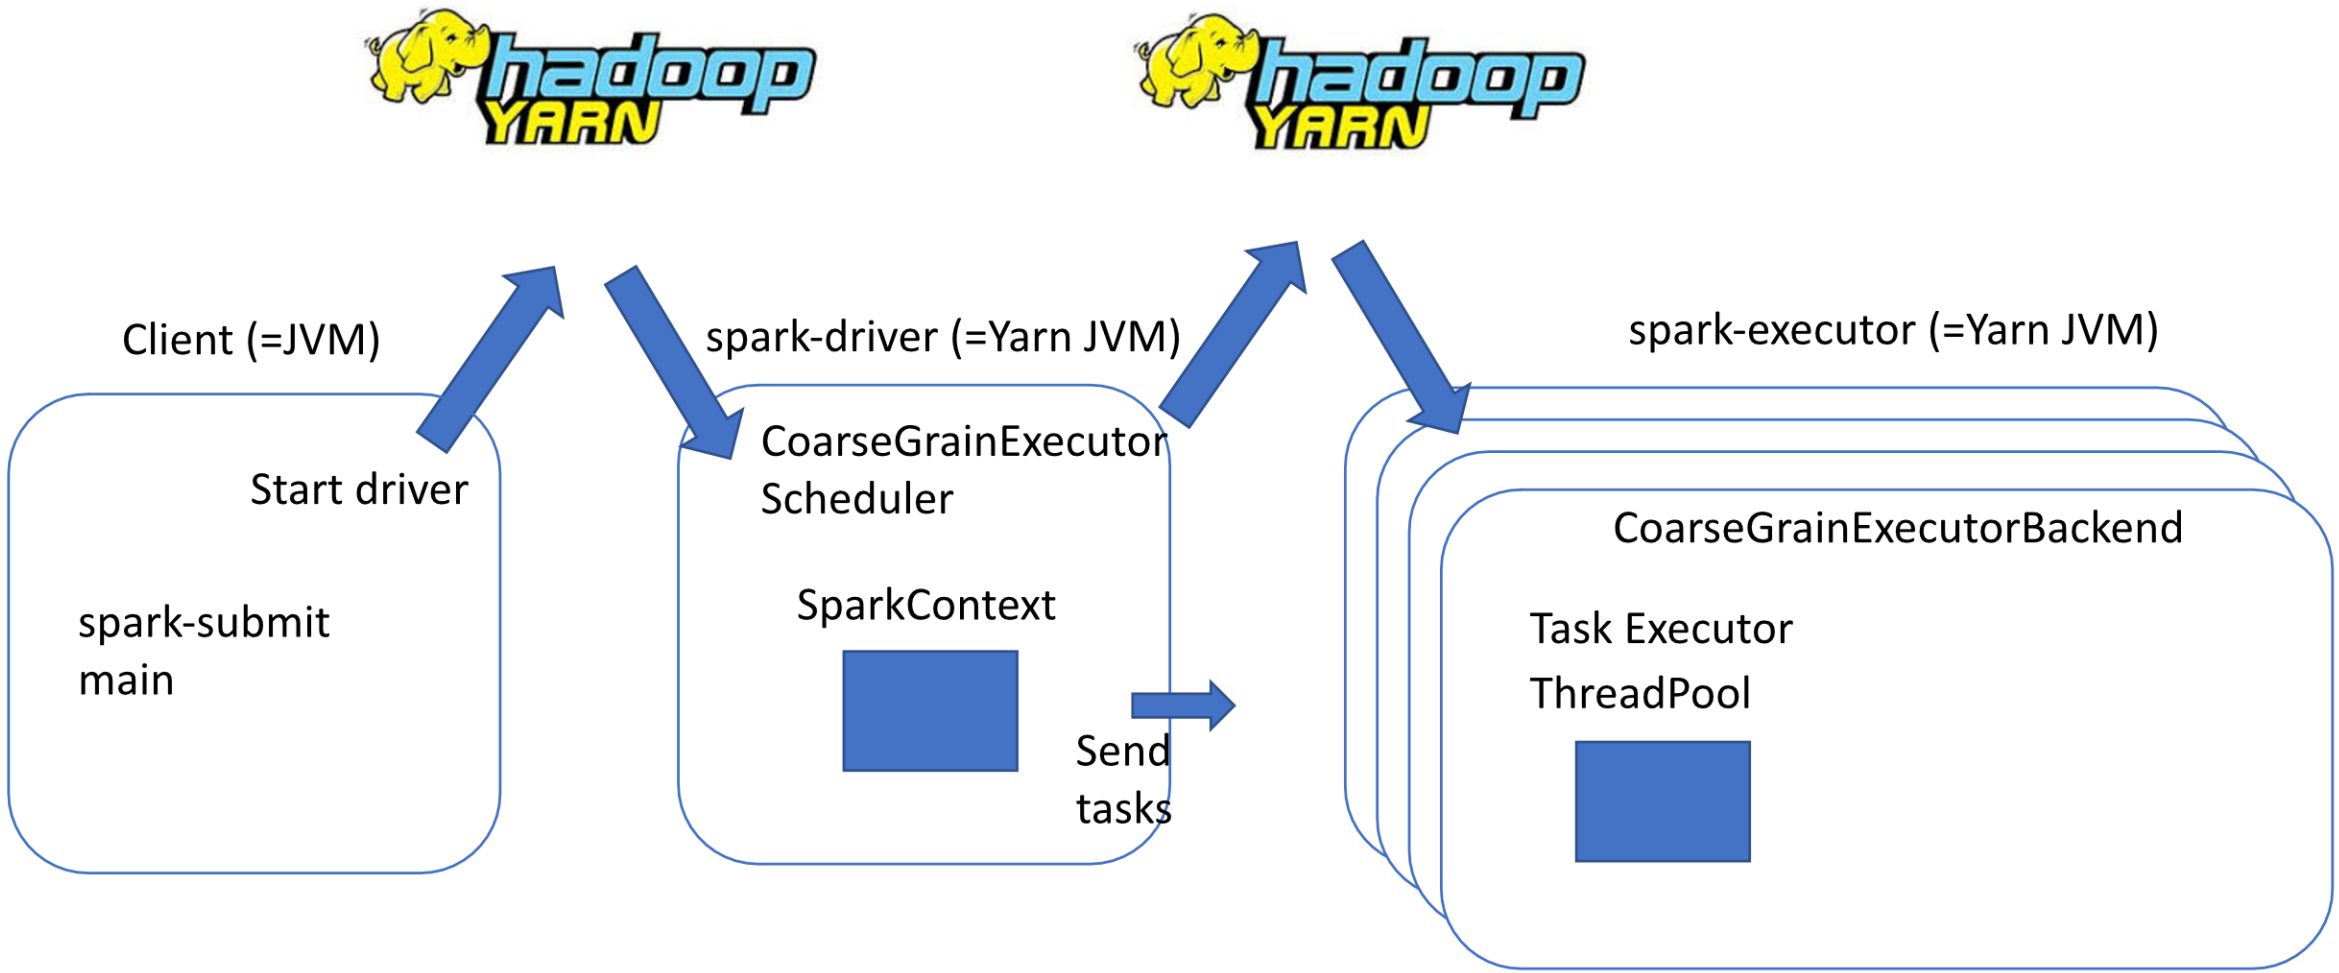

Pros-Cons Yarn « Cluster » vs « Client » mode

Client can run outside firewall / cluster

Client is lightweight
Many Clients can run on « Edge » Node

If Driver crashes… client will relaunch driver!

More difficult to see logs (Yarn Logs!!)

More difficult to kill : yarn app –kill ${applicationId}

Kill client => does not kill driver !!
(Zombie yarn driver ?)

--master k8s://<k8sApiServer>:443 (basic without K8s operator)

```
# Run on a Kubernetes cluster in cluster deploy mode
./bin/spark-submit \
  --class org.apache.spark.examples.SparkPi \
  --master k8s://xx.yy.zz.ww:443 \
  --deploy-mode cluster \
  --executor-memory 20G \
  --num-executors 50 \
  http://path/to/examples.jar \
  1000
```

Spark on Kubernetes … Docker Images/Jars

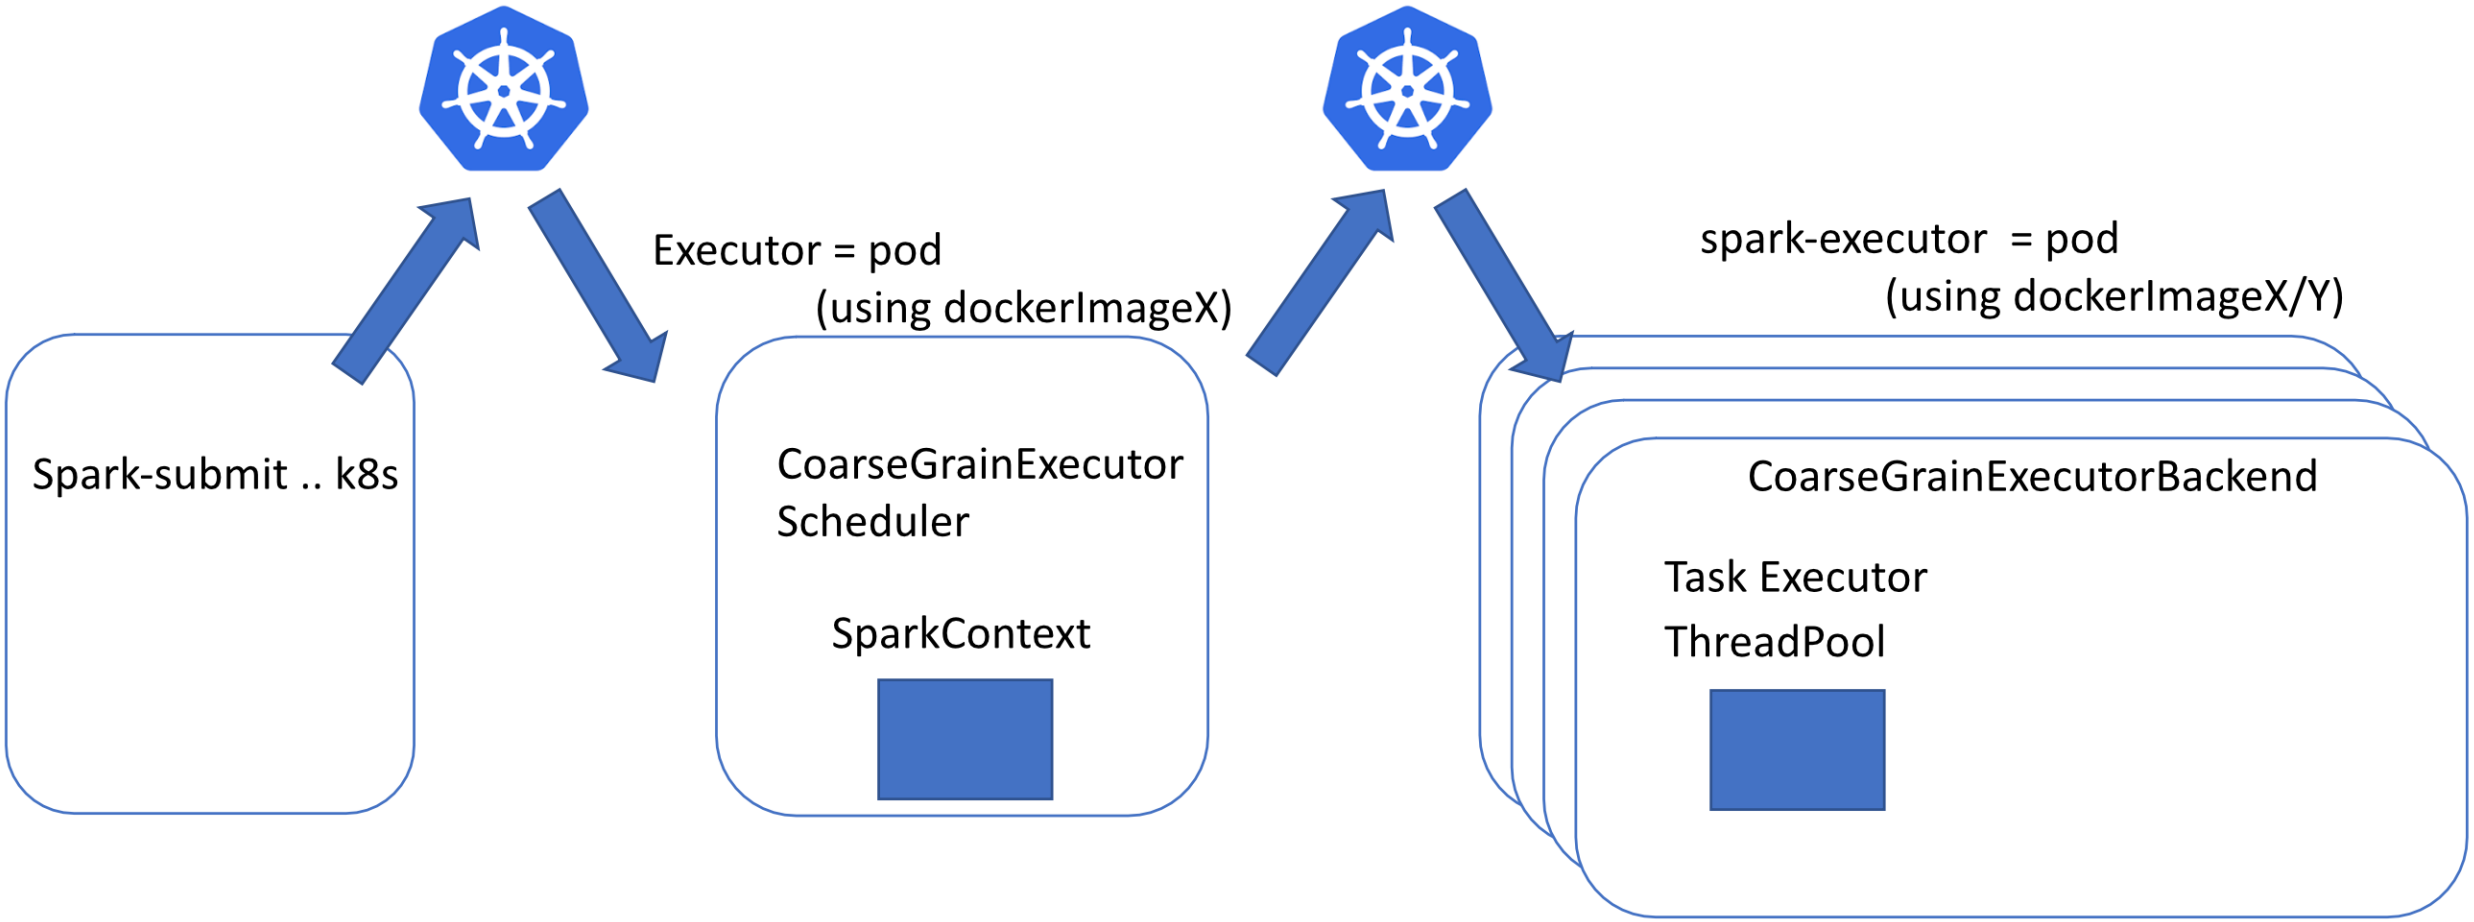

Spark on K8s with SparkOperator mode

```
$ kubectl apply -f examples/spark-pi.yaml
```

Spark on K8S, using K8s SparkOperator

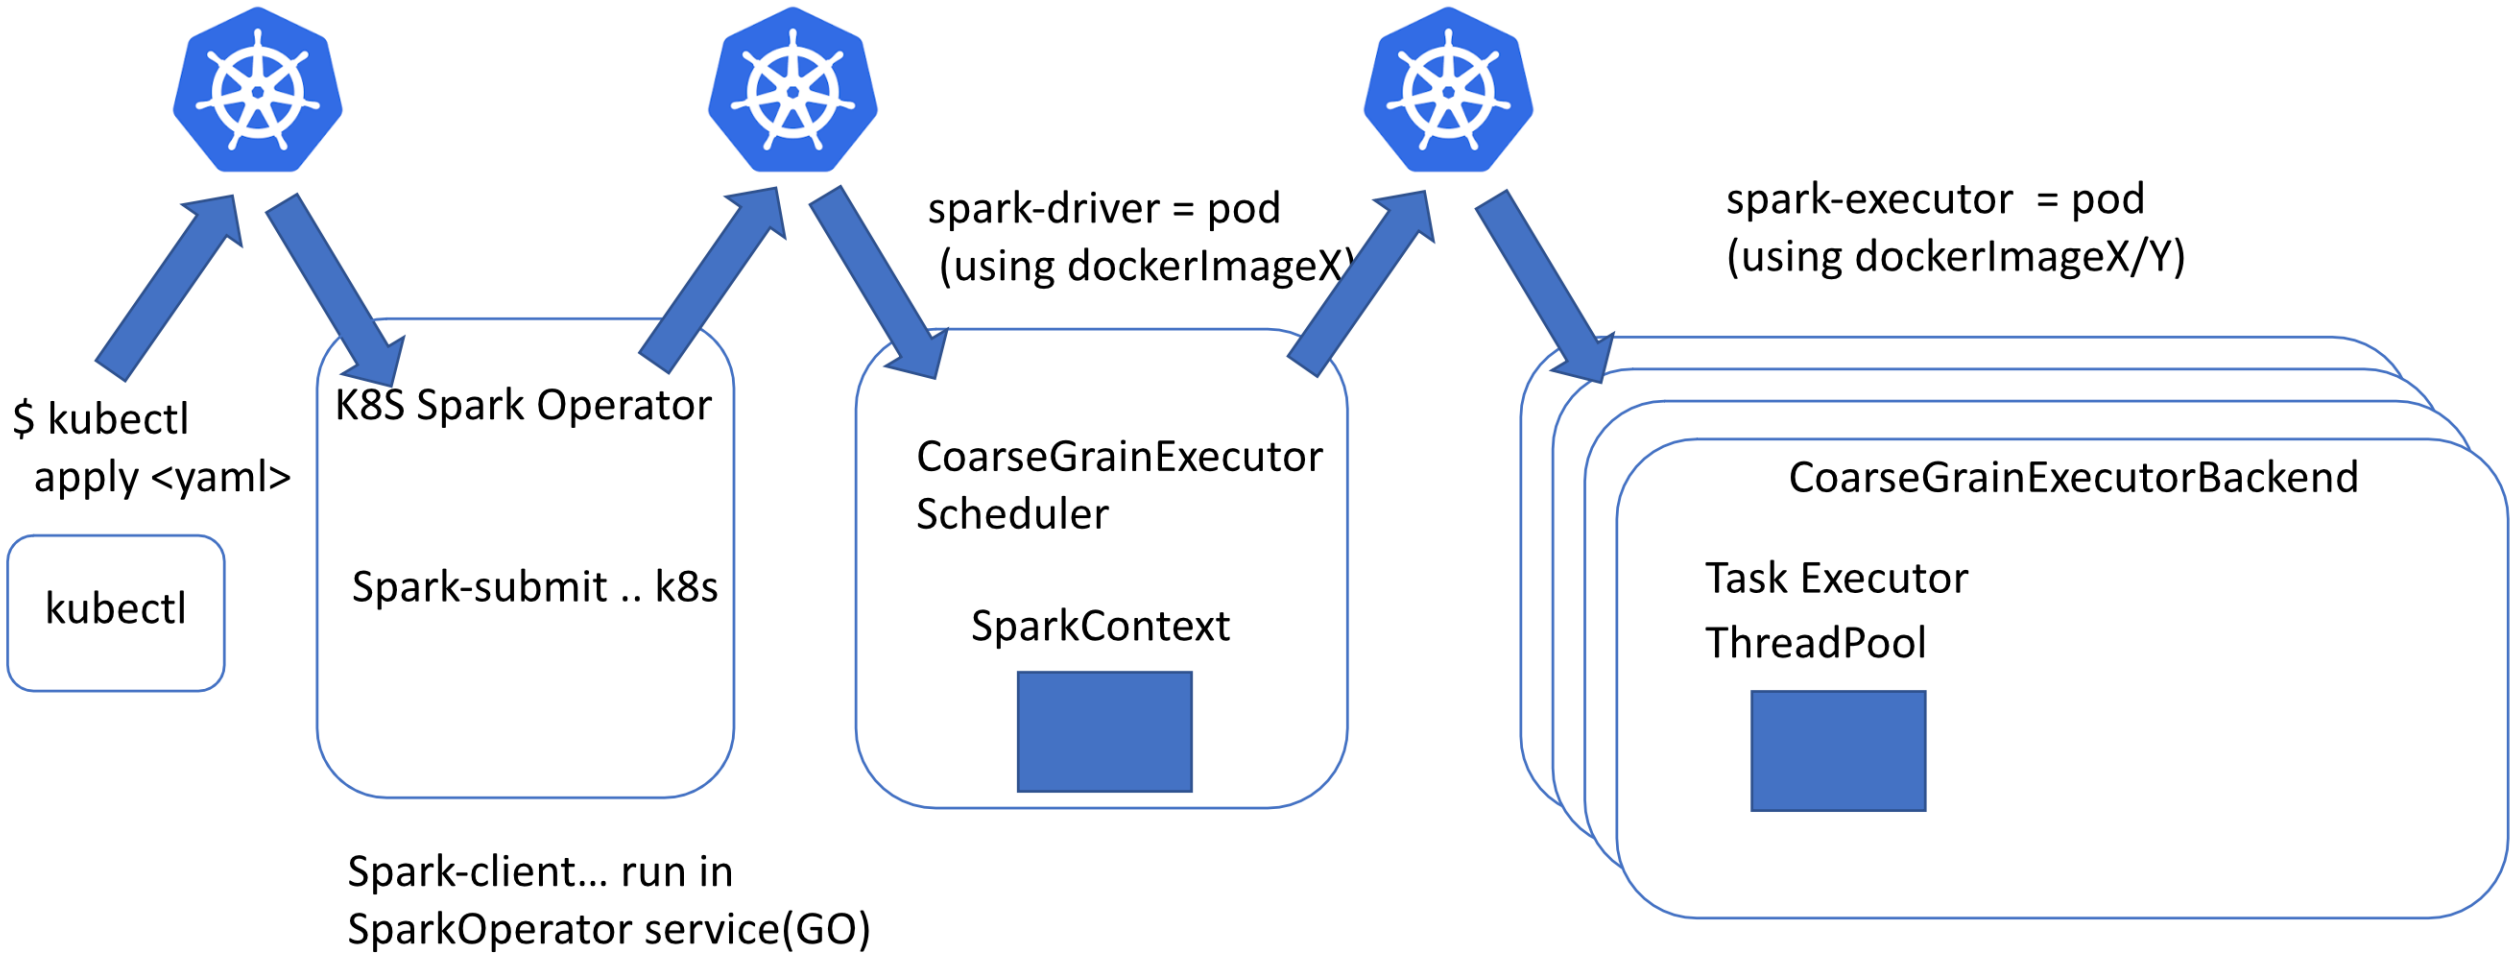

spark-on-k8s-operator ： K8s CRD from Google

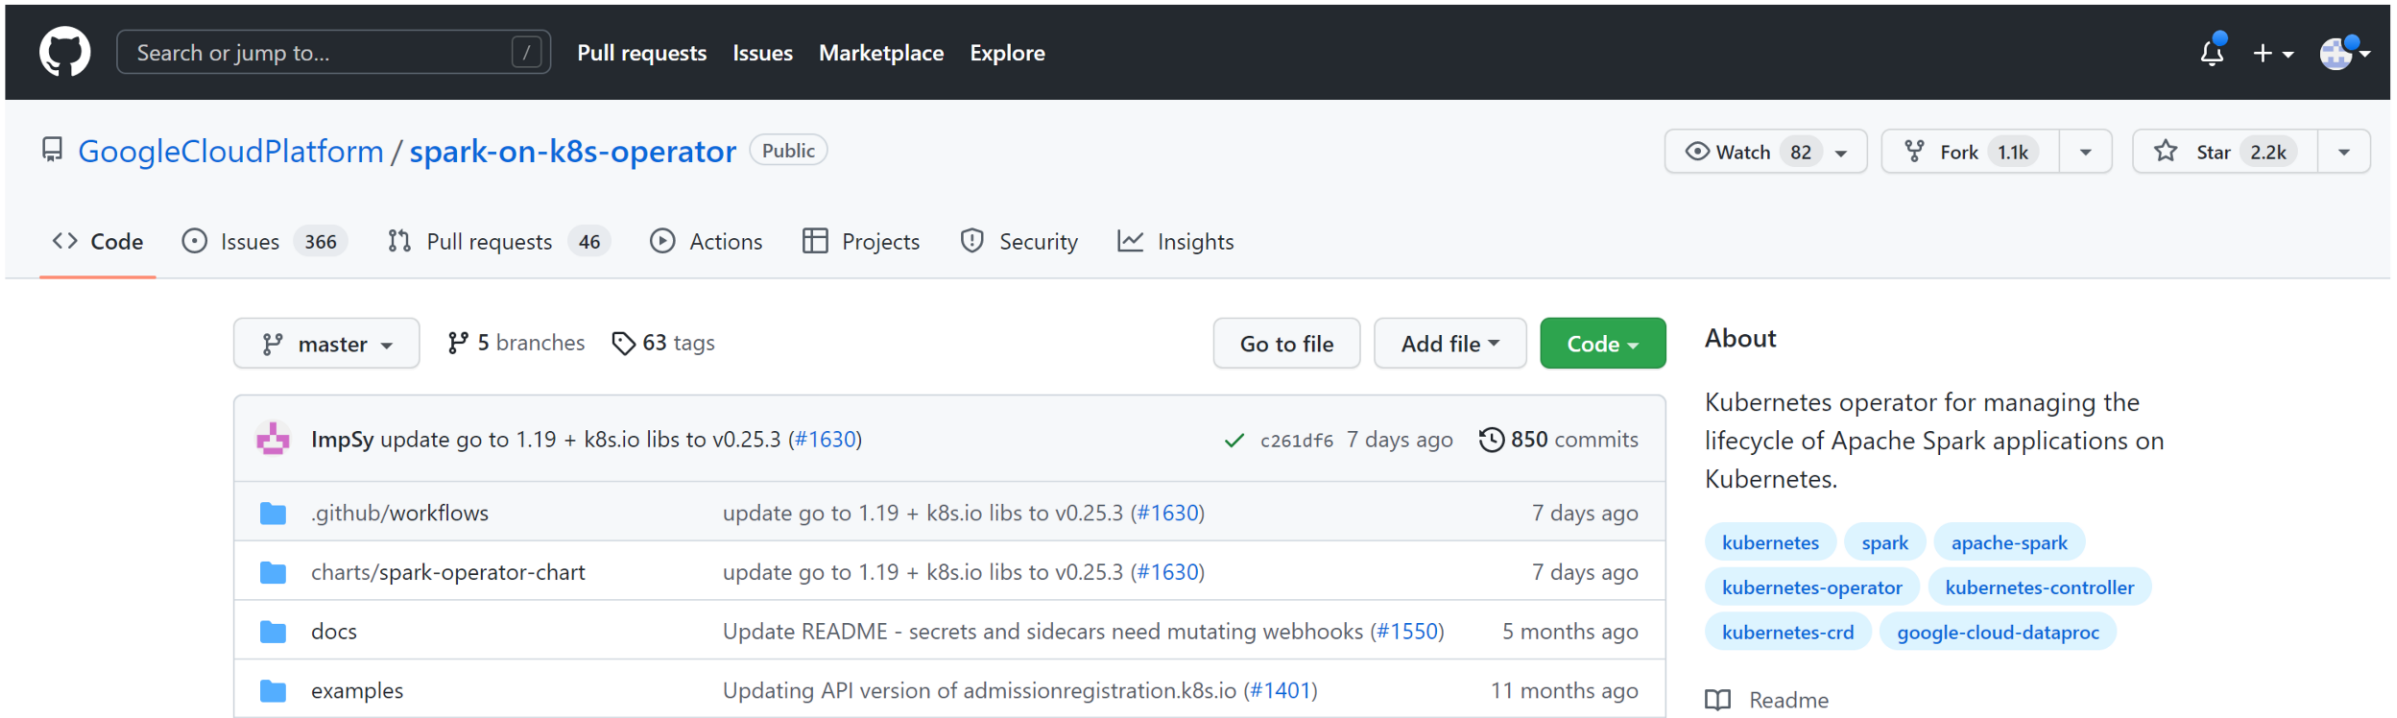

$ kubectl get sparkapplications ..-o=yaml

spec
- driver pod
- executor pod
- Spark args

status

```yaml
apiVersion: sparkoperator.k8s.io/v1beta2
kind: SparkApplication
metadata:
  ...
spec:
  deps: {}
  driver:
    coreLimit: 1200m
    cores: 1
    labels:
      version: 2.3.0
    memory: 512m
    serviceAccount: spark
  executor:
    cores: 1
    instances: 1
    labels:
      version: 2.3.0
    memory: 512m
  image: gcr.io/ynli-k8s/spark:v3.1.1
  mainApplicationFile: local:///opt/spark/examples/jars/spark-examples_2.12-3.1.1.jar
  mainClass: org.apache.spark.examples.SparkPi
  mode: cluster
  restartPolicy:
      type: OnFailure
      onFailureRetries: 3
      onFailureRetryInterval: 10
      onSubmissionFailureRetries: 5
      onSubmissionFailureRetryInterval: 20
  type: Scala
status:
  sparkApplicationId: spark-5f4ba921c85ff3f1cb04bef324f9154c9
  applicationState:
    state: COMPLETED
  completionTime: 2018-02-20T23:33:55Z
  driverInfo:
    podName: spark-pi-83ba921c85ff3f1cb04bef324f9154c9-driver
    webUIAddress: 35.192.234.248:31064
    webUIPort: 31064
    webUIServiceName: spark-pi-2402118027-ui-svc
    webUIIngressName: spark-pi-ui-ingress
    webUIIngressAddress: spark-pi.ingress.cluster.com
  executorState:
    spark-pi-83ba921c85ff3f1cb04bef324f9154c9-exec-1: COMPLETED
  LastSubmissionAttemptTime: 2018-02-20T23:32:27Z
```

Outline

Reminder Spark-driver / spark-executors

Cluster Tuning Params

deployMode (local, spark://, yarn, k8s)

Other launchers (java embedded, livy, jupyter, pyspark ..)

More ways to start spark

Plain old java main(String[] args) Spark « Hello World »

```java
import org.apache.spark.sql.Dataset;
import org.apache.spark.sql.Encoders;
import org.apache.spark.sql.SparkSession;

public class SimpleMain {

    public static void main(String[] args) {
        System.out.println("ensure HADOOP_HOME env var is set, (and Winutils on windows)");
        System.out.println("HADOOP_HOME: " + System.getenv("HADOOP_HOME"));

        System.out.println("start embedded SparkContext - SparkSession");
        SparkSession spark = SparkSession.builder()
                .master("local[*]")
                .appName("myapp")
                .getOrCreate();

        List<String> ls = Arrays.asList("Hello spark");
        Dataset<String> ds = spark.createDataset(ls, Encoders.STRING());
        ds.show();

        System.out.println("finish SparkContext");
        spark.close();
    }
}
```

Console × Problems Progress Error Log Debug Shell Search Call Hierarchy Search

<terminated> SimpleSparkMain [Java Application] C:\apps\jdk\jdk-8\bin\javaw.exe (19 nov. 2022, 20:20:06 – 20:20:14) [pid: 18140]

```
ensure HADOOP_HOME env var is set, (and Winutils on windows)
HADOOP_HOME: C:\apps\hadoop\spark-3.3.0-hadoop-3.3
start embedded SparkContext - SparkSession
+-----------+
|      value|
+-----------+
|Hello spark|
+-----------+

finish SparkContext
```

Spark maven pom.xml <dependency>

```
<dependency>
    <groupId>org.apache.spark</groupId>
    <artifactId>spark-core_${scala.version}</artifactId>
    <version>${spark.version}</version>
    <scope>provided</scope>
</dependency>
<dependency>
    <groupId>org.apache.spark</groupId>
    <artifactId>spark-sql_${scala.version}</artifactId>
    <version>${spark.version}</version>
    <scope>provided</scope>
</dependency>
```

Use <scope>provided</scope>
when deployed in HADOOP cluster

SparkContext / SparkSession ?

SparkContext = in module spark-core
singleton, instantiable only on Driver, multi-thread safe

SparkSession = in module spark-sql
also multi-thread safe (~ SparkContext + SQL UDFs, temporaryViews..)

When creating a SparkSession, it internally checks/creates SparkContext

It is possible to have multiple sessions
(for having different SQL UDF, temporaryView, .. per user session)

SparkConf – SparkContext- SparkSession

```
class SparkConf {
… key=value
+ hadoopConf
}
```

package org.apache.spark (module spark-core)

1

activeContext

0..1

```
class SparkContext {
…
}
```

new RDD + transforms

package org.apache.spark.sql (module spark-sql)

1

```
class SparkSession {
…
}
```

getOrCreate()

new Dataset<> + transforms

JavaSparkContext, JavaRDD, JavaFunction ?

Because Spark has a cleaner API

(less verbose, using implicits + operator overload),

some specific

helper classes have been added, for Java developers

some deprecated, ex: SQLContext

Oozie … http Api wrapper for Yarn + UI

$ curl http://<knox>/oozie/

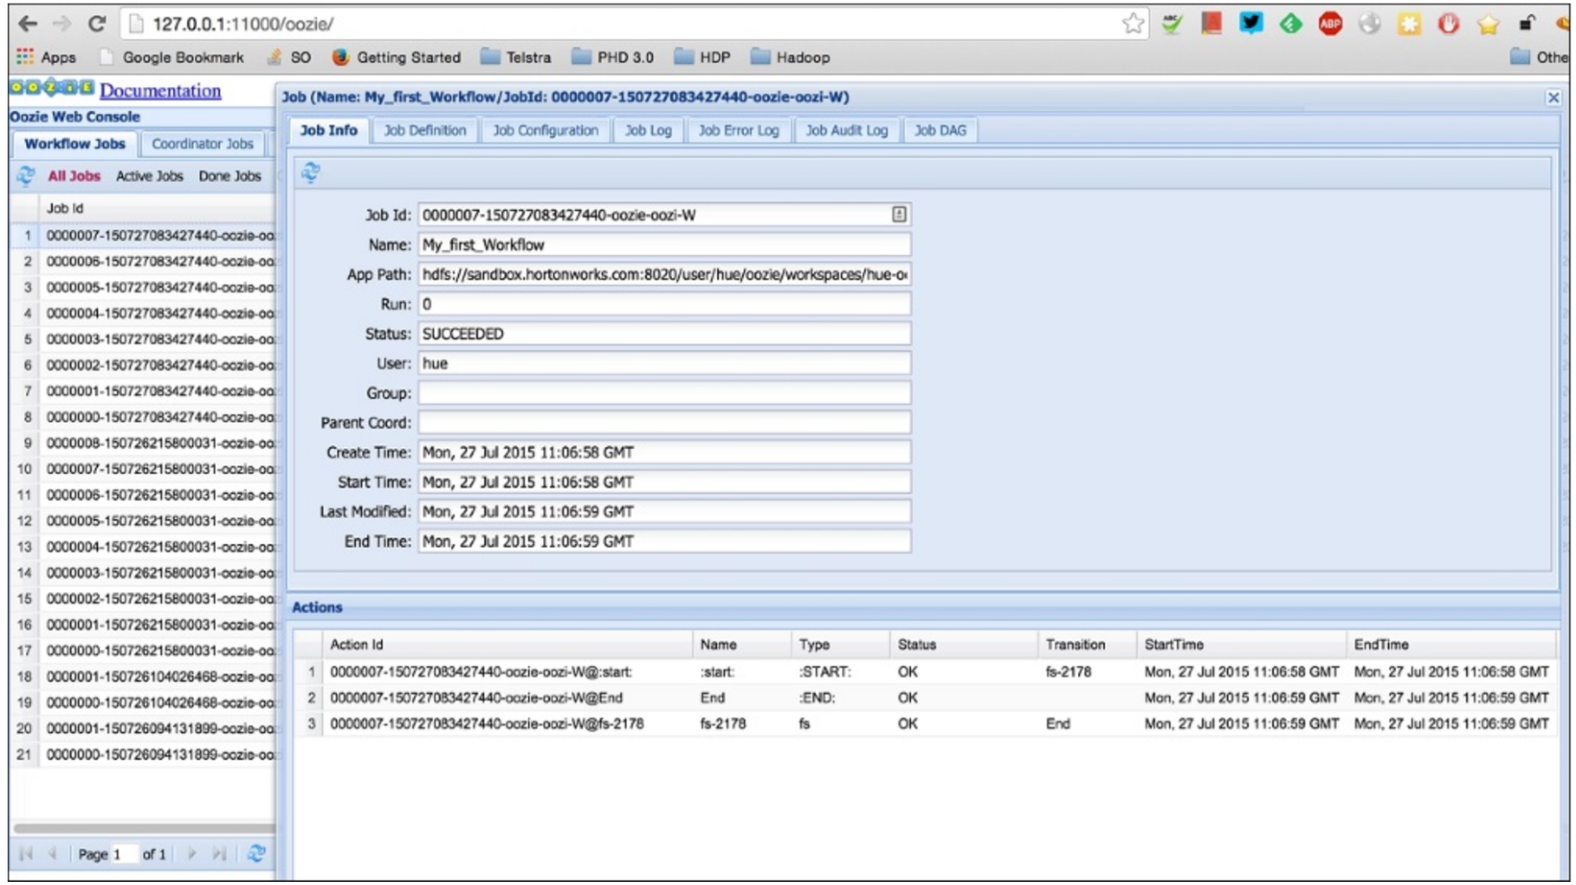

Oozie <shell> -> bash spark-submit.sh

$ **curl** -X POST /oozie/job -d @config.properties

config.properties

oozie.wf.application.path=hdfs://app/workflow.xml

workflow.xml

```
<workflow-app name="[WF-DEF-NAME]" xmlns="uri:oozie:workflow:1.0">
    ...
    <action name="[NODE-NAME]">
        <shell xmlns="uri:oozie:shell-action:1.0">
            <resource-manager>[RESOURCE-MANAGER]</resource-manager>
            <name-node>[NAME-NODE]</name-node>
            <exec>spark-submit.sh</exec>
            <argument>[ARG-VALUE]</argument>
            <env-var>[VAR1=VALUE1]</env-var>
                ...
            <env-var>[VARN=VALUEN]</env-var>
            <file>[FILE-PATH]</file>
            ...
        </shell>
        <ok to="[NODE-NAME]"/>
        <error to="[NODE-NAME]"/>
    </action>
    ...
</workflow-app>
```

Oozie Workflow.xml- <spark>

Oozie Spark Action Extension

- Spark Action
  - Spark Action Logging
  - Spark on YARN
  - PySpark with Spark Action
  - Using Symlink in <jar>
- Appendix, Spark XML-Schema
  - AE.A Appendix A, Spark XML-Schema
    - Spark Action Schema Version 1.0
    - Spark Action Schema Version 0.2
    - Spark Action Schema Version 0.1

Spark Action

The `spark` action runs a Spark job.

The workflow job will wait until the Spark job completes before continuing to the next action.

Syntax:

```
<workflow-app name="[WF-DEF-NAME]" xmlns="uri:oozie:workflow:1.0">
    ...
    <action name="[NODE-NAME]">
        <spark xmlns="uri:oozie:spark-action:1.0">
            <resource-manager>[RESOURCE-MANAGER]</resource-manager>
            <name-node>[NAME-NODE]</name-node>
            <prepare>
               <delete path="[PATH]"/>
               ...
               <mkdir path="[PATH]"/>
               ...
            </prepare>
            <job-xml>[SPARK SETTINGS FILE]</job-xml>
            <configuration>
                <property>
                    <name>[PROPERTY-NAME]</name>
                    <value>[PROPERTY-VALUE]</value>
                </property>
                ...
            </configuration>
            <master>[SPARK MASTER URL]</master>
            <mode>[SPARK MODE]</mode>
            <name>[SPARK JOB NAME]</name>
            <class>[SPARK MAIN CLASS]</class>
            <jar>[SPARK DEPENDENCIES JAR / PYTHON FILE]</jar>
            <spark-opts>[SPARK-OPTIONS]</spark-opts>
            <arg>[ARG-VALUE]</arg>
                ...
            <arg>[ARG-VALUE]</arg>
            ...
        </spark>
        <ok to="[NODE-NAME]"/>
        <error to="[NODE-NAME]"/>
    </action>
    ...
</workflow-app>
```

LIVY

https://livy.apache.org/

LIVY Apache Livy

Get Started | Documentation | Community | Apache

Submit Jobs from Anywhere

Livy enables programmatic, fault-tolerant, multi-tenant submission of Spark jobs from web/mobile apps (no Spark client needed). So, multiple users can interact with your Spark cluster concurrently and reliably.

Use Interactive Scala or Python

Livy speaks either Scala or Python, so clients can communicate with your Spark cluster via either language remotely. Also, batch job submissions can be done in Scala, Java, or Python.

No Code Changes Needed

Don't worry, no changes to existing programs are needed to use Livy. Just build Livy with Maven, deploy the configuration file to your Spark cluster, and you're off! Check out Get Started to get going.

What is Apache Livy?

Apache Livy is a service that enables easy interaction with a Spark cluster over a REST interface. It enables easy submission of Spark jobs or snippets of Spark code, synchronous or asynchronous result retrieval, as well as Spark Context management, all via a simple REST interface or an RPC client library. Apache Livy also simplifies the interaction between Spark and application servers, thus enabling the use of Spark for interactive web/mobile applications. Additional features include:

- Have long running Spark Contexts that can be used for multiple Spark jobs, by multiple clients
- Share cached RDDs or Dataframes across multiple jobs and clients
- Multiple Spark Contexts can be managed simultaneously, and the Spark Contexts run on the cluster (YARN/Mesos) instead of the Livy Server, for good fault tolerance and concurrency
- Jobs can be submitted as precompiled jars, snippets of code or via java/scala client API
- Ensure security via secure authenticated communication

To learn more, watch this tech session video from Spark Summit West 2016.

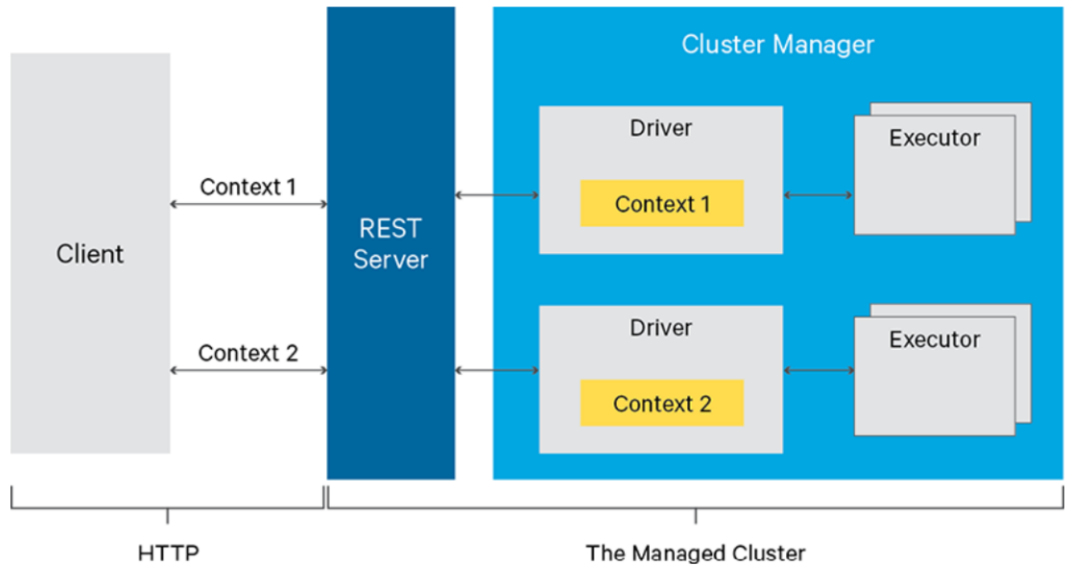

Livy … submit via http
Stateful batch / session

POST /batches
GET /batches
GET /batches/{batchId}
GET /batches/{batchId}/state
GET /batches/{batchId}/log
DELETE /batches/{batchId}

POST /sessions
GET /sessions
GET /sessions/{sessionId}
GET /sessions/{sessionId}/state
GET /sessions/{sessionId}/log
POST /sessions/{sessionId}/completion
DELETE /sessions/{sessionId}

POST /sessions/{sessionId}/statements
POST /sessions/{sessionId}/statements/{statementId}/cancel
GET /sessions/{sessionId}/statements
GET /sessions/{sessionId}/statements/{statementId}

Livy Batch Sample

```
spark-submit \
  --master yarn --deploy-mode cluster \
  --executor-memory 20G \
  --class org.apache.spark.examples.SparkPi \
  /path/to/examples.jar \
  1000
```

$ **curl -X POST http://<livyhost>:8998/batches -d @req.json**

req.json:

```
{
  "executorMemory": "20g",
   "className": "org.apache.spark.examples.SparkPi",
   "file": "/path/to/examples.jar",
   "args": [1000]
}
```

Livy (Interactive) Session- Statement

$ **response = $(curl -X POST http://<livyhost>:8998/sessions -d session1.json)**

session1.json:

```
{
  "kind": "spark",
  "executorMemory": "20g",
  "files": [ "/path/to/examples.jar"  ]
}
```

response:

id=$(.. **jq '.id'**)

```
{
 "id"  : 2 ,
 "appId" : " application_123_456",
 "owner" : "user", "kind": "spark",
 "log": [ « log line1 », « line2 » ],
 "state": "started",
 "appInfo": { … }
}
```

${id}

$ **statResp = $(curl -X POST http://<livyhost>:8998/sessions/ 2 /statements -d statement-req.json)**

statement-req.json

```
{
  "mode": "raw",
  "raw": "{
     \"code\": \" val ds = spark.sql(\"  SELECT * from table1 \") \"
  } "
}
```

Zeppelin

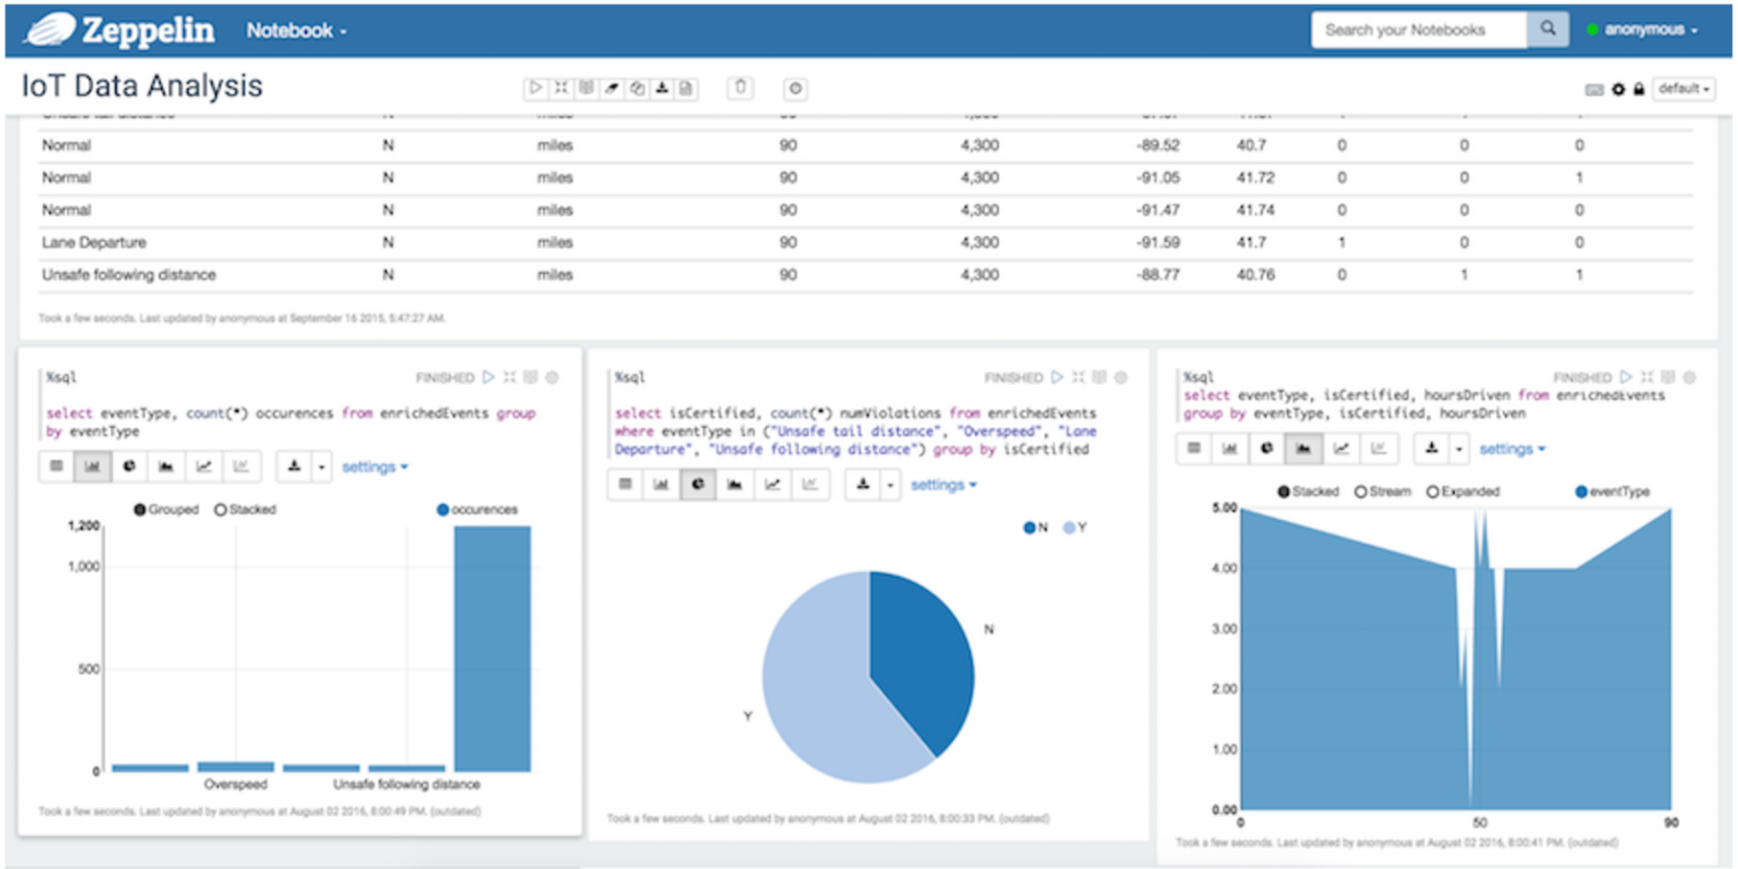
Zeppelin
Notebook
Search your Notebooks
anonymous
IoT Data Analysis
default
Normal N miles 90 4,300 -89.52 40.7 0 0 0
Normal N miles 90 4,300 -91.05 41.72 0 0 1
Normal N miles 90 4,300 -91.47 41.74 0 0 0
Lane Departure N miles 90 4,300 -91.59 41.7 1 0 0
Unsafe following distance N miles 90 4,300 -88.77 40.76 0 1 1
Took a few seconds. Last updated by anonymous at September 16 2015, 5:47:27 AM.
%sql
FINISHED
select eventType, count(*) occurences from enrichedEvents group by eventType
settings
Grouped
Stacked
occurences
1,200
1,000
500
0
Overspeed
Unsafe following distance
Took a few seconds. Last updated by anonymous at August 02 2016, 8:00:49 PM. (outdated)
%sql
FINISHED
select isCertified, count(*) numViolations from enrichedEvents where eventType in ("Unsafe tail distance", "Overspeed", "Lane Departure", "Unsafe following distance") group by isCertified
settings
N
Y
N
Y
Took a few seconds. Last updated by anonymous at August 02 2016, 8:00:33 PM. (outdated)
%sql
FINISHED
select eventType, isCertified, hoursDriven from enrichedEvents group by eventType, isCertified, hoursDriven
settings
Stacked
Stream
Expanded
eventType
5.00
4.00
3.00
2.00
1.00
0.00
0
50
90
Took a few seconds. Last updated by anonymous at August 02 2016, 8:00:41 PM. (outdated)

Zeppelin-> Livy Interactive Session

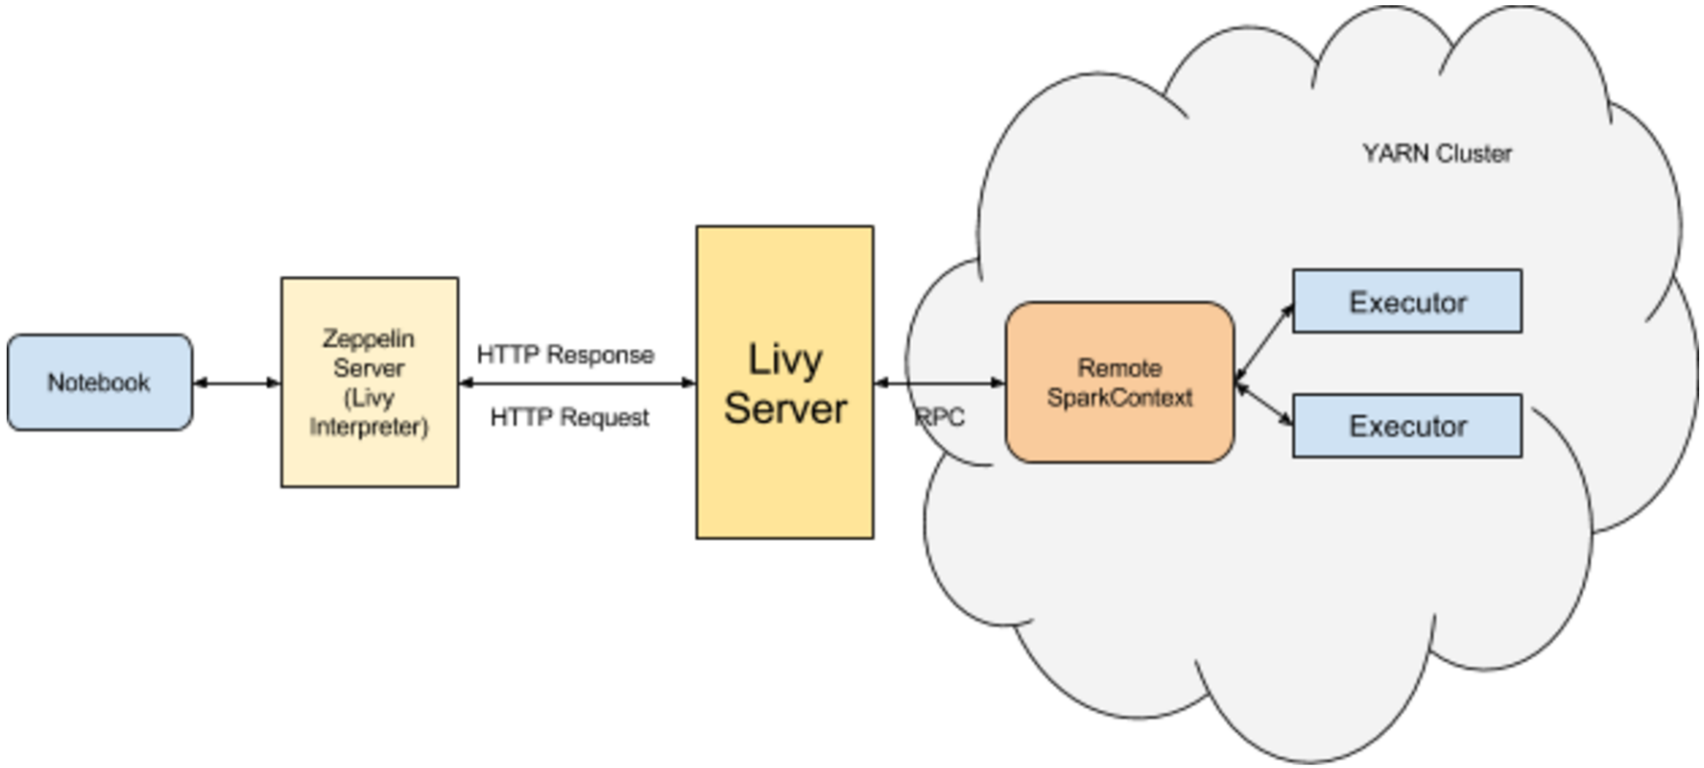

Config: livy.spark.*

Zeppelin 0.10.1 Quick Start Usage Setup Interpreter More

Configuration

We added some common configurations for spark, and you can set any configuration you want. You can find all Spark configurations in here. And instead of starting property with `spark.` it should be replaced with `livy.spark.`. Example: `spark.driver.memory` to `livy.spark.driver.memory`

Property	Default	Description
zeppelin.livy.url	http://localhost:8998	URL where livy server is running
zeppelin.livy.spark.sql.maxResult	1000	Max number of Spark SQL result to display.
zeppelin.livy.spark.sql.field.truncate	true	Whether to truncate field values longer than 20 characters or not
zeppelin.livy.session.create_timeout	120	Timeout in seconds for session creation
zeppelin.livy.displayAppInfo	true	Whether to display app info
zeppelin.livy.pull_status.interval.millis	1000	The interval for checking paragraph execution status
livy.spark.driver.cores		Driver cores. ex) 1, 2.
livy.spark.driver.memory		Driver memory. ex) 512m, 32g.
livy.spark.executor.instances		Executor instances. ex) 1, 4.
livy.spark.executor.cores		Num cores per executor. ex) 1, 4.
livy.spark.executor.memory		Executor memory per worker instance. ex) 512m, 32g.
livy.spark.dynamicAllocation.enabled		Use dynamic resource allocation. ex) True, False.
		Remove an executor which has cached data

Link Zeppelin-> SparkUI Application

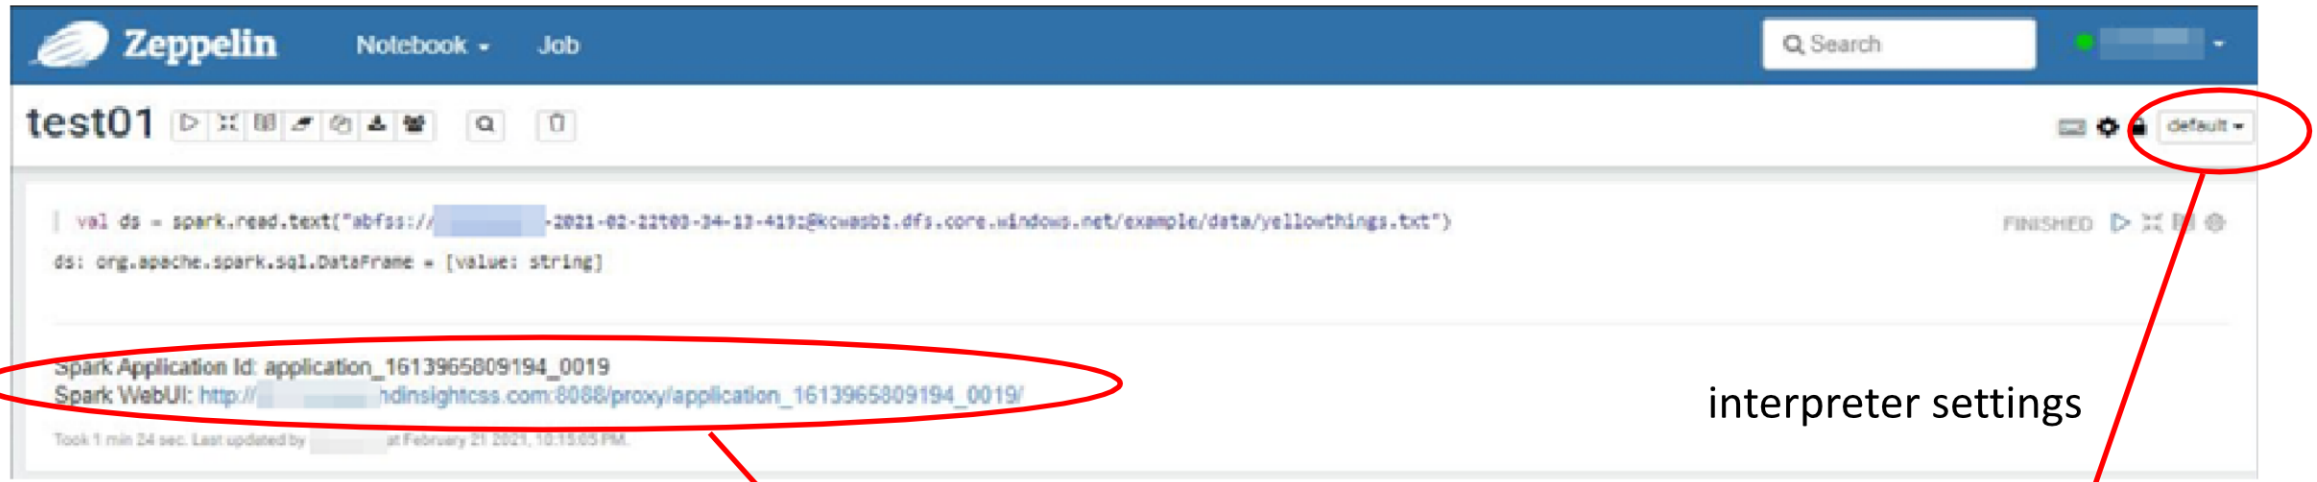

interpreter settings

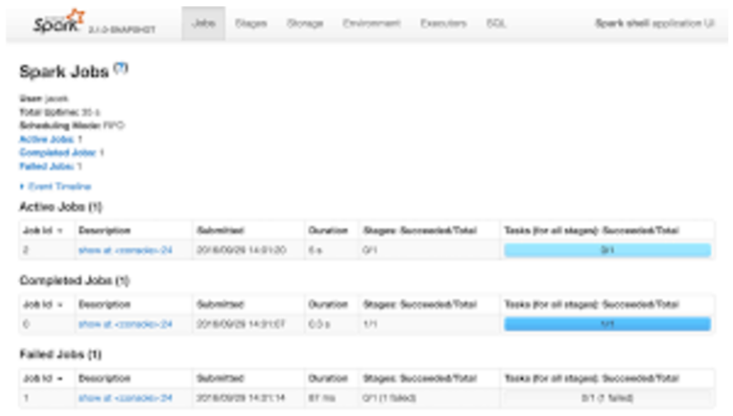

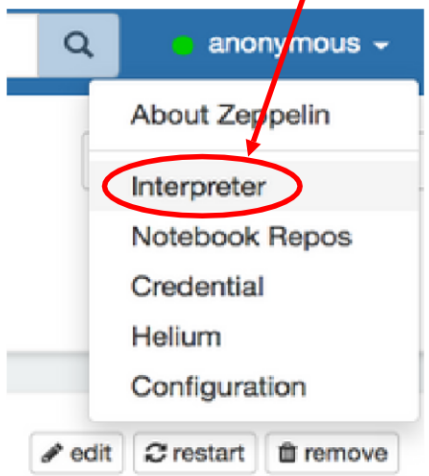

Jupyter Notebook

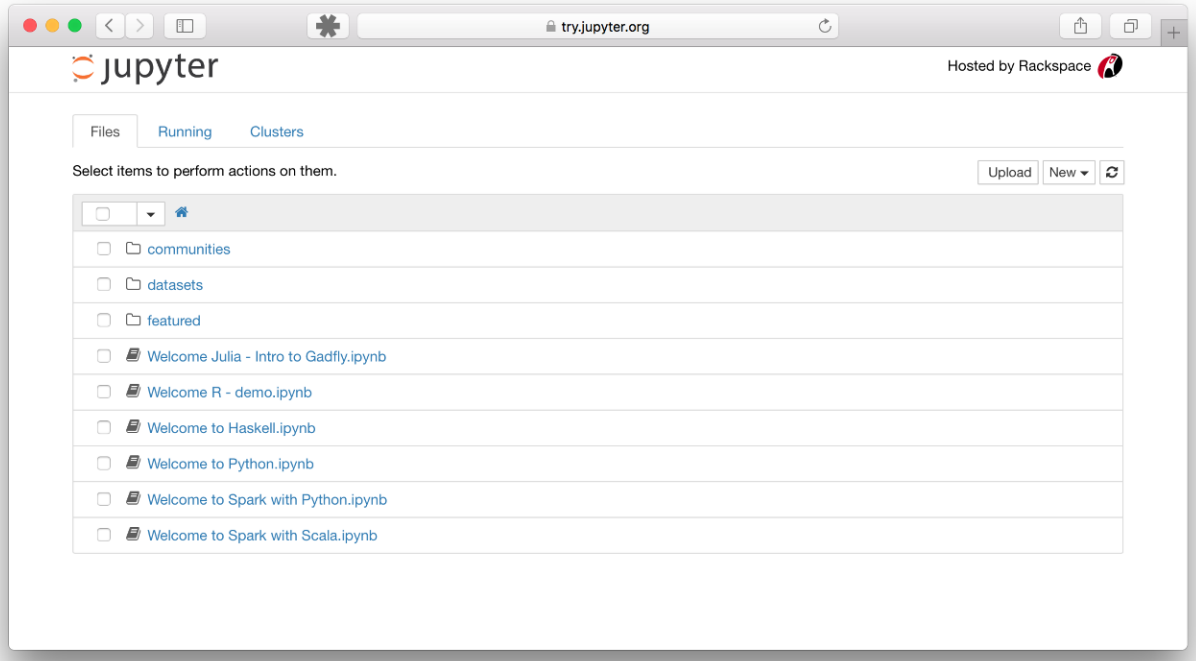

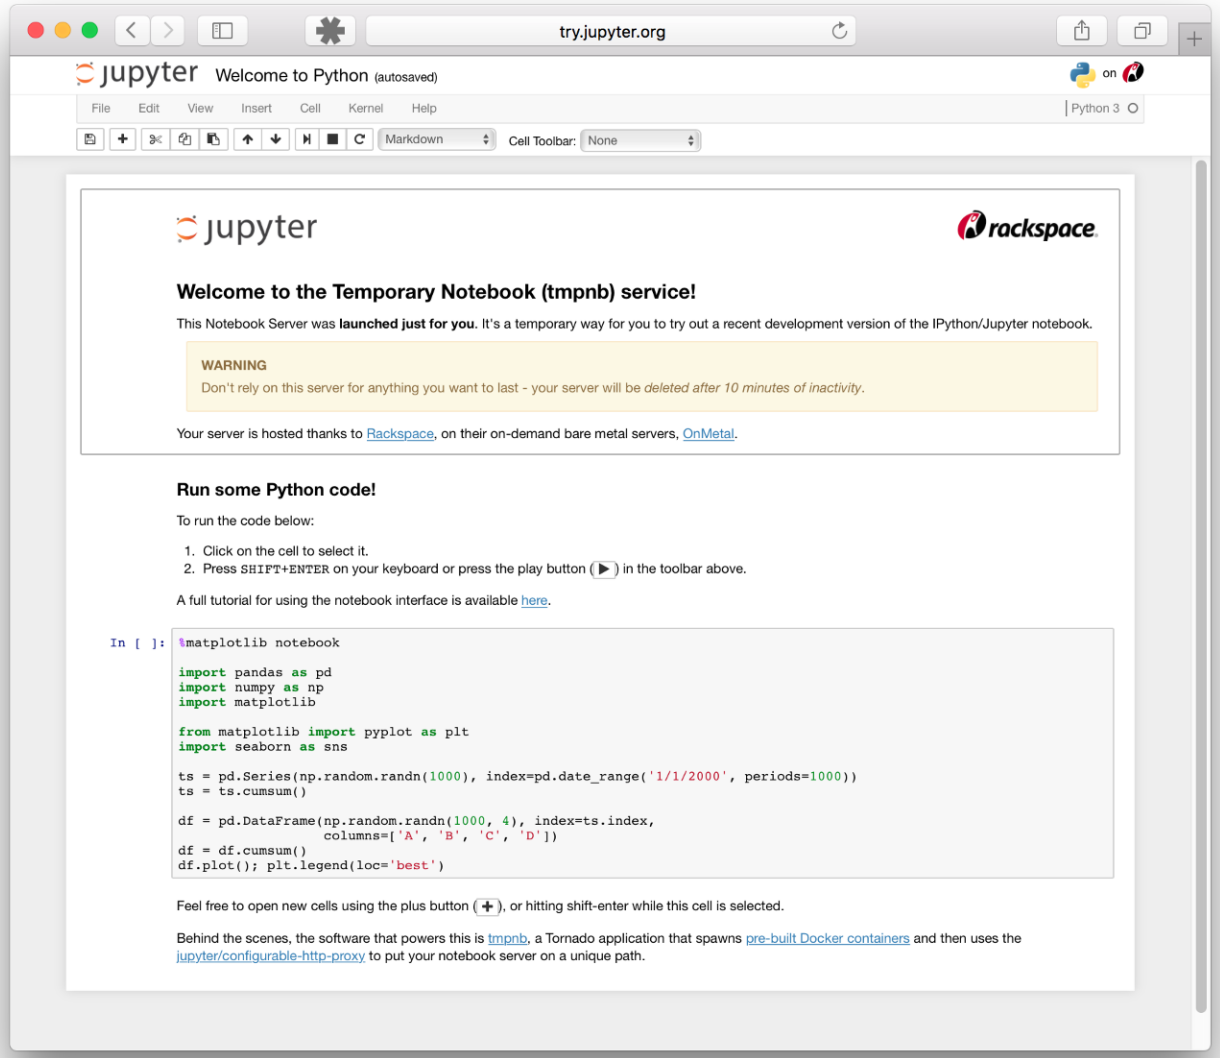

Jupyter – Kernel(s) architecture

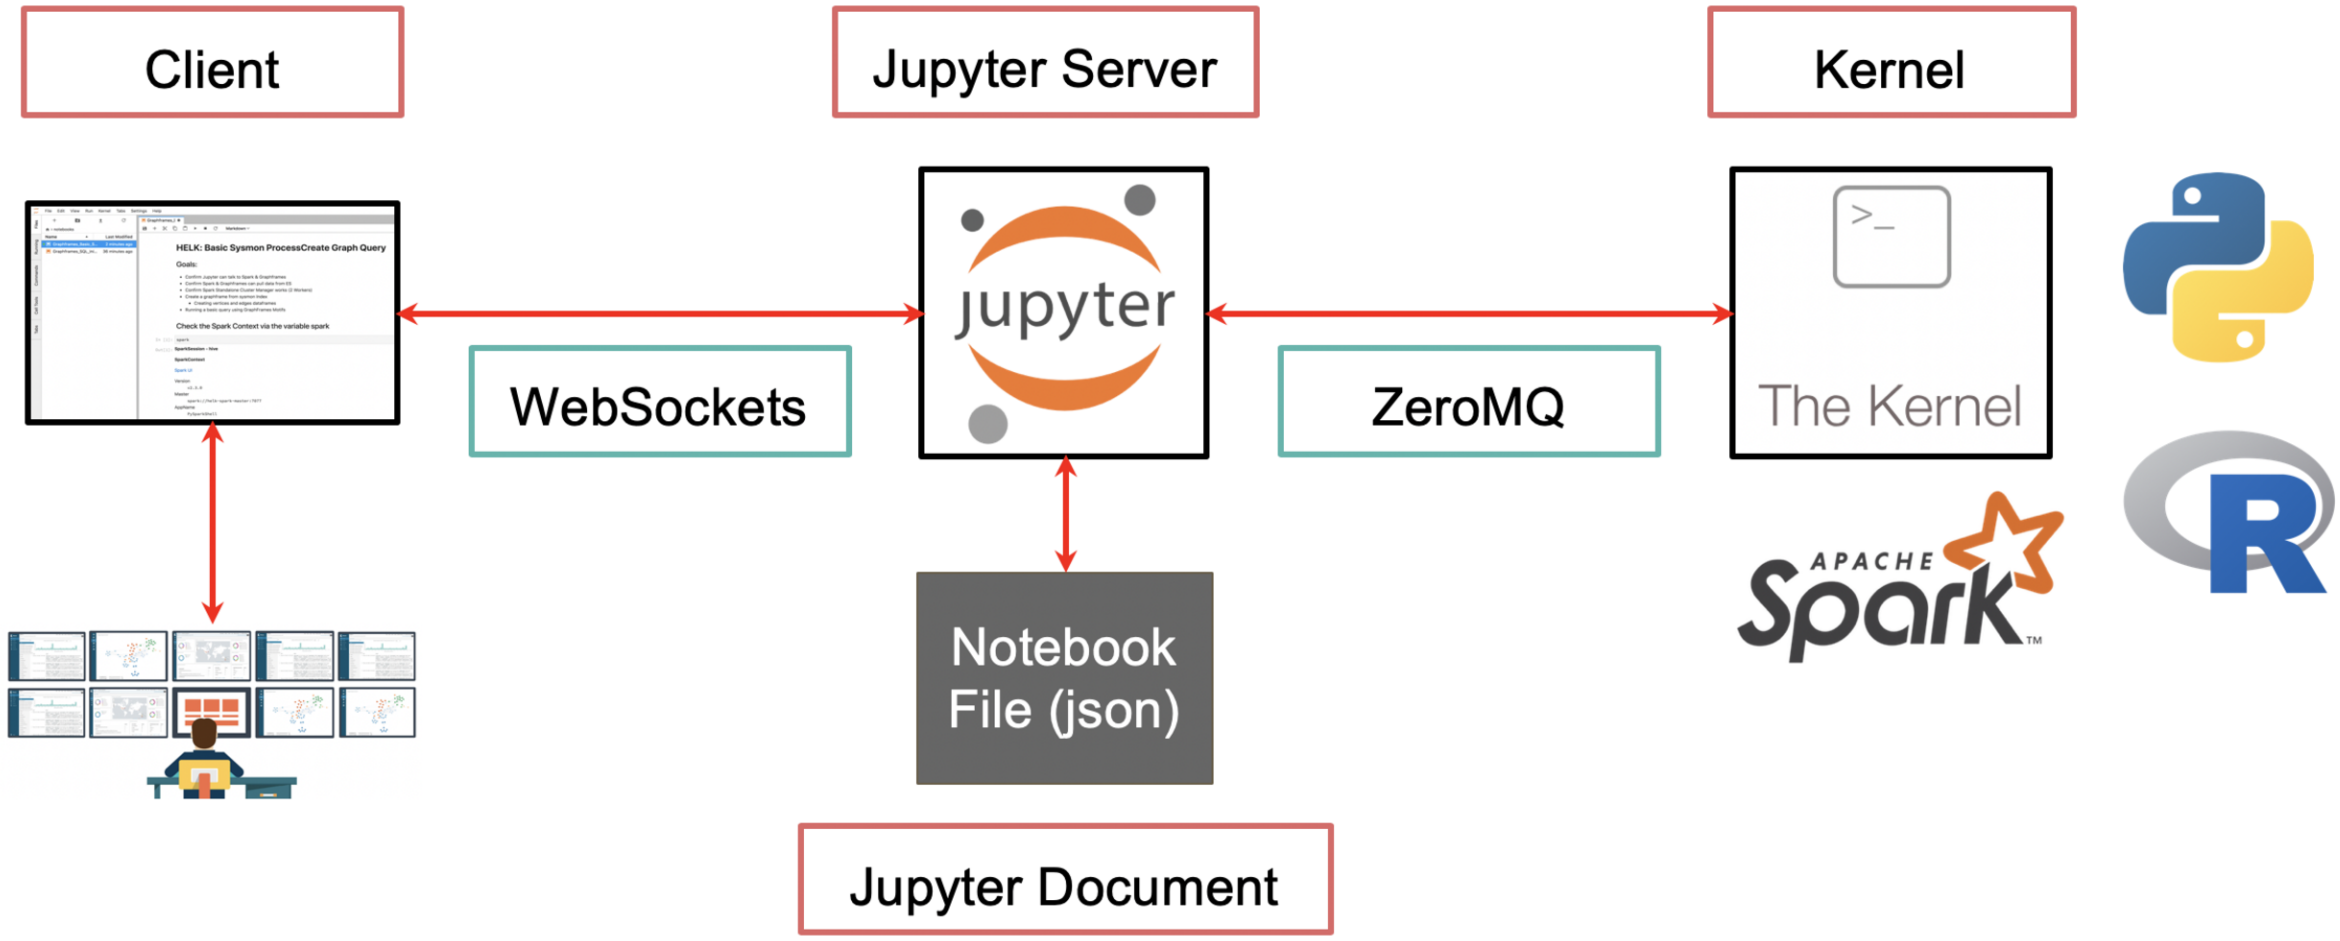

Jupyter + python kernel + pyspark

```
In [11]: from pyspark.context import SparkContext
         from pyspark.sql.session import SparkSession
         from pyspark.sql.types import *
         from pyspark.sql.functions import *

         sqlContext = SQLContext(sc)
         df  = sqlContext.sql("select 1 as A")
         df.show()
```

```
+---+
|  A|
+---+
|  1|
+---+
```

PySpark

More « Extensions »
backport API to other Languages (Python, R)

Reminder

class DataFrame:

… hundred methods …

def **method123** (self, x) -> Y
  self._jdf.method(x, self.sparkSession)

API

API

public class Dataset {

… hundred methods …

public Y **method123** (X x) {
  … internals sparkContext…
}

py4j

```
from py4j.java_gateway import JavaObject
from subprocess import Popen
from pyspark.context import SparkContext
..
pid = Popen(« spark-submit »  …)
Socket(.. )
```

Scala

SparkContext

Python + pyspark (GLUE package)
… = spark JVM process + socket

$ pyspark

SparkContext (spark-driver)

python.exe

Popen('java.exe', …)

socket

```
23  import socket
24  import platform
25  import tempfile
26  import time
27  from subprocess import Popen, PIPE
28
29  from py4j.java_gateway import java_import, JavaGateway, JavaObject, GatewayParameters
30  from py4j.clientserver import ClientServer, JavaParameters, PythonParameters
31  from pyspark.find_spark_home import _find_spark_home
32  from pyspark.serializers import read_int, write_with_length, UTF8Deserializer
```

Write python Object
to pickle format

read java Object
from pickle format

Py4J https://www.py4j.org/

Python UDF function

PySpark supports various UDFs and APIs to allow users to execute Python native functions. See also the latest Pandas UDFs and Pandas Function APIs. For instance, the example below allows users to directly use the APIs in a pandas Series within Python native function.

[21]:
```
import pandas as pd
from pyspark.sql.functions import pandas_udf

@pandas_udf('long')
def pandas_plus_one(series: pd.Series) -> pd.Series:
    # Simply plus one by using pandas Series.
    return series + 1

df.select(pandas_plus_one(df.a)).show()
```

```
+------------------+
|pandas_plus_one(a)|
+------------------+
|                 2|
|                 3|
|                 4|
+------------------+
```

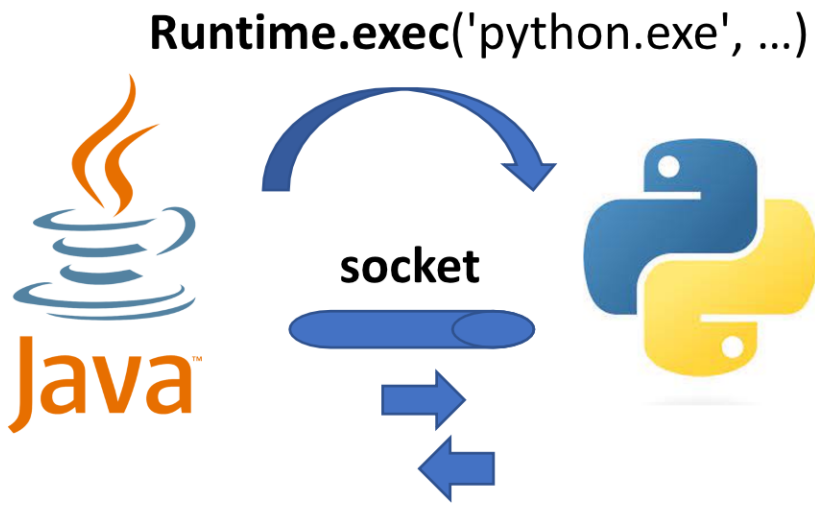

Pyspark – py4j – Spark Architecture

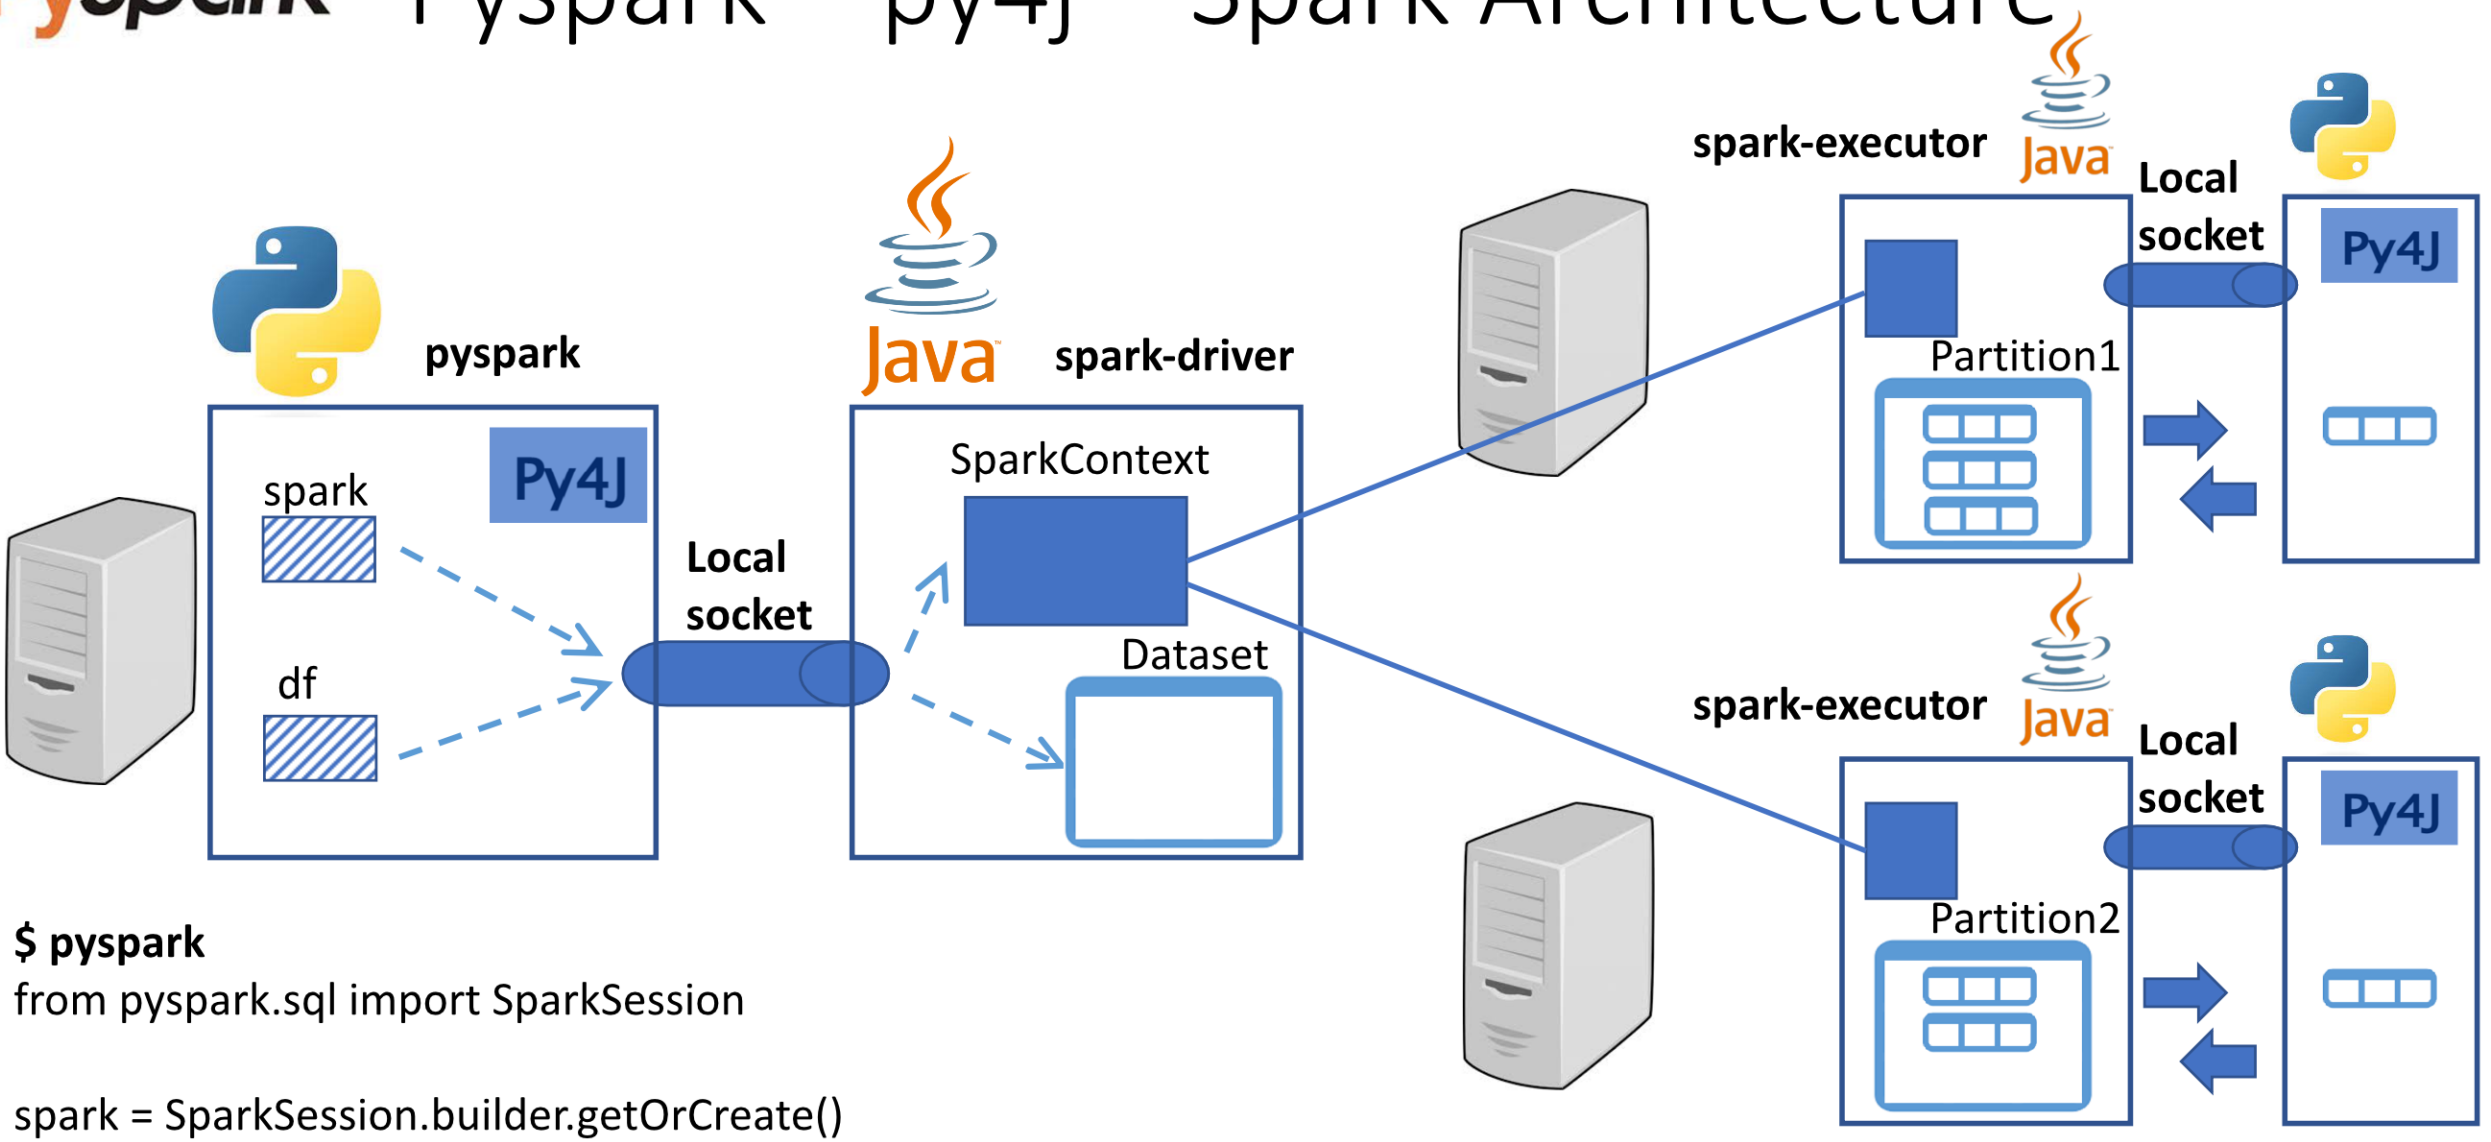

$ pyspark

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.getOrCreate()
df = spark.sql(« SELECT * FROM table1 »)
```

Python + Pickle = Poor Performance

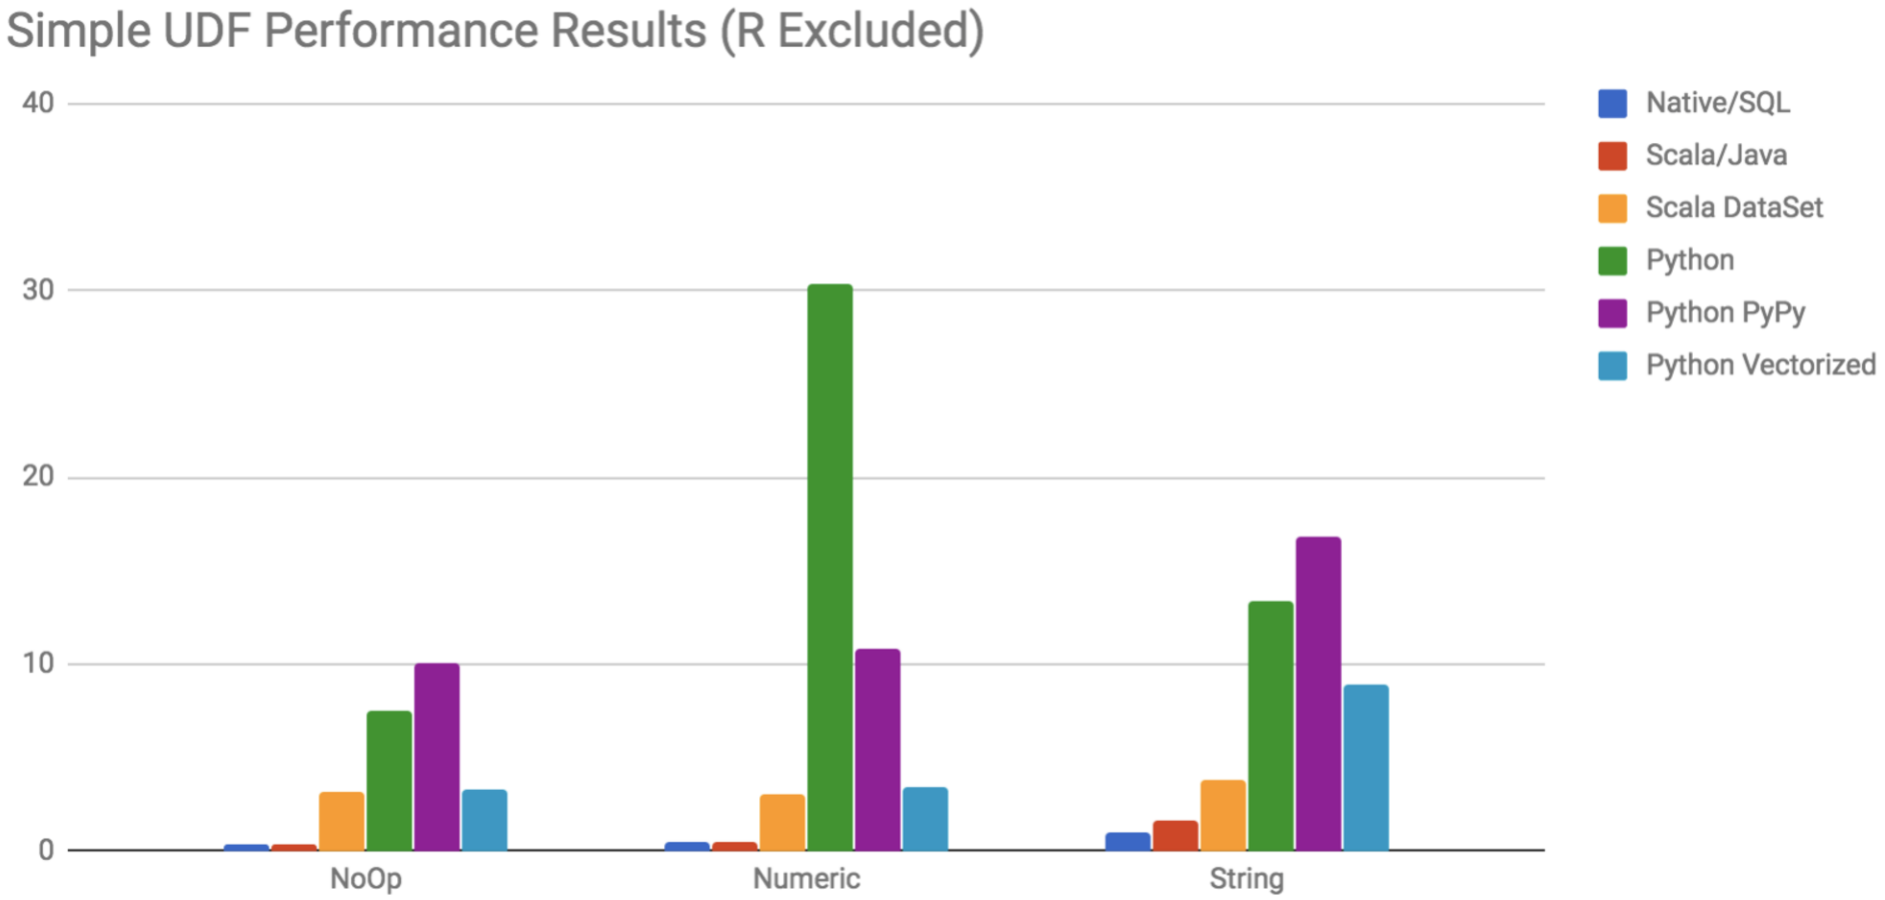

Jupyter + spylon (spark scala kernel)

```
pip install spylon-kernel
# or
conda install -c conda-forge spylon-kernel
```

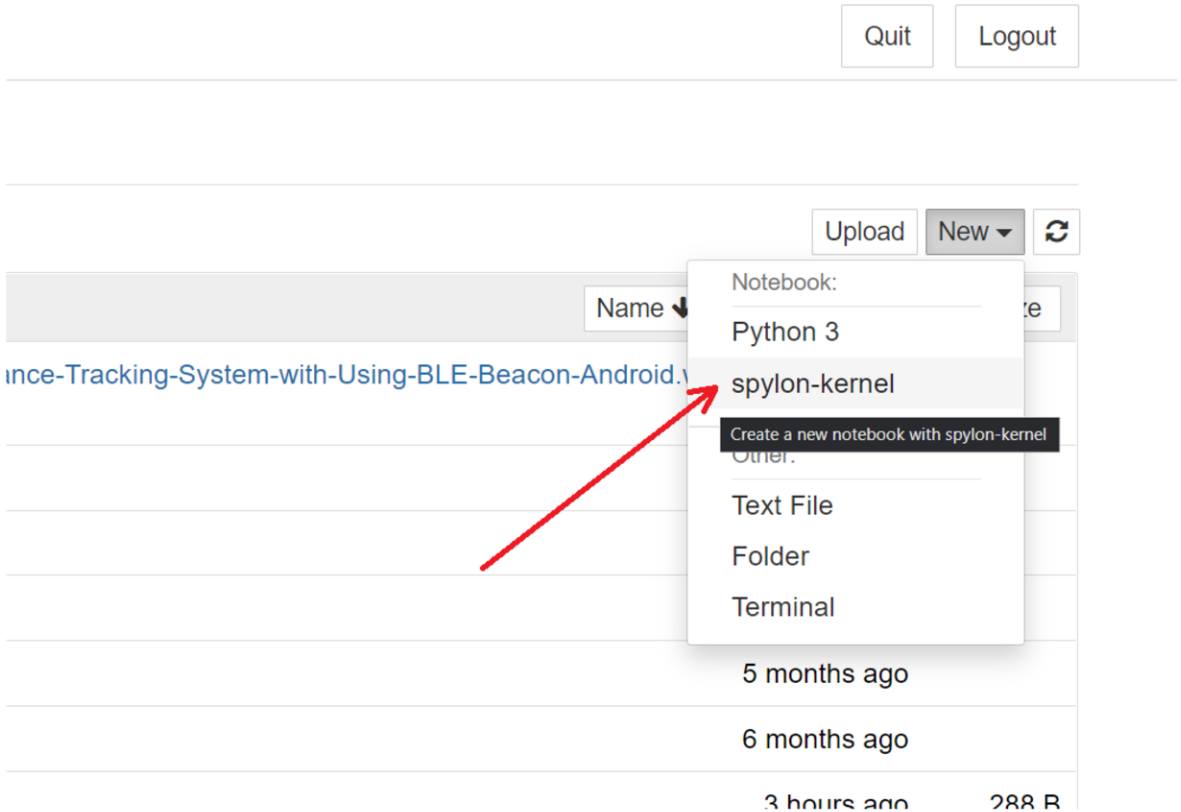

spylon : %%scala

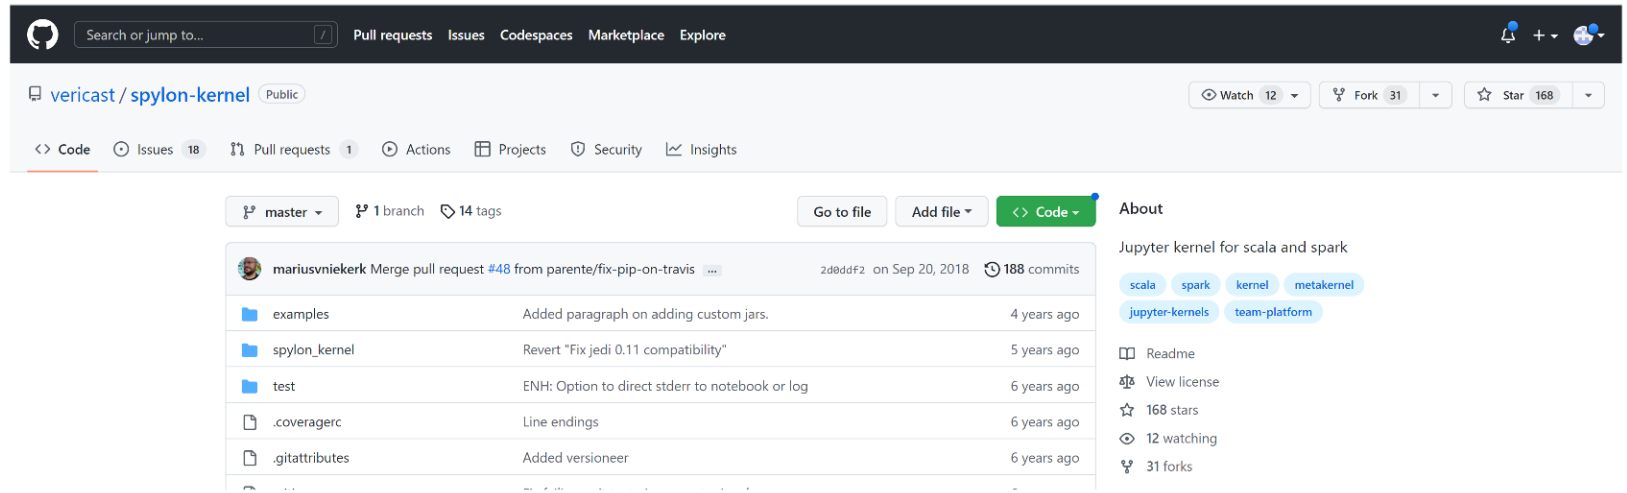

Using it as an IPython Magic

You can also use spylon-kernel as a magic in an IPython notebook. Do this when you want to mix a little bit of Scala into your primarily Python notebook.

```
from spylon_kernel import register_ipython_magics
register_ipython_magics()
```

```
%%scala
val x = 8
x
```

Using it as a Library

Finally, you can use spylon-kernel as a Python library. Do
Python script or shell.

```
from spylon_kernel import get_scala_interpreter

interp = get_scala_interpreter()

# Evaluate the result of a scala code block.
interp.interpret("""
    val x = 8
    x
""")

interp.last_result()
```

Jupyter + spylon architecture

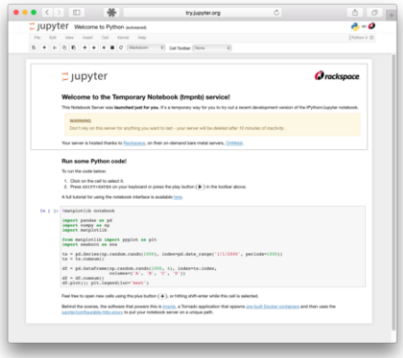

%%scala

Text

html

jupyter

%%scala

spylon-kernel

socket

spark-driver

SparkContext

Dataset

spark-executor

Partition1

spark-executor

Partition1

Spark Connect ... new since version 3.4

Spark Connect Overview

Building client-side Spark applications

In Apache Spark 3.4, Spark Connect introduced a decoupled client-server architecture that allows remote connectivity to Spark clusters using the DataFrame API and unresolved logical plans as the protocol. The separation between client and server allows Spark and its open ecosystem to be leveraged from everywhere. It can be embedded in modern data applications, in IDEs, Notebooks and programming languages.

To get started, see Quickstart: Spark Connect.

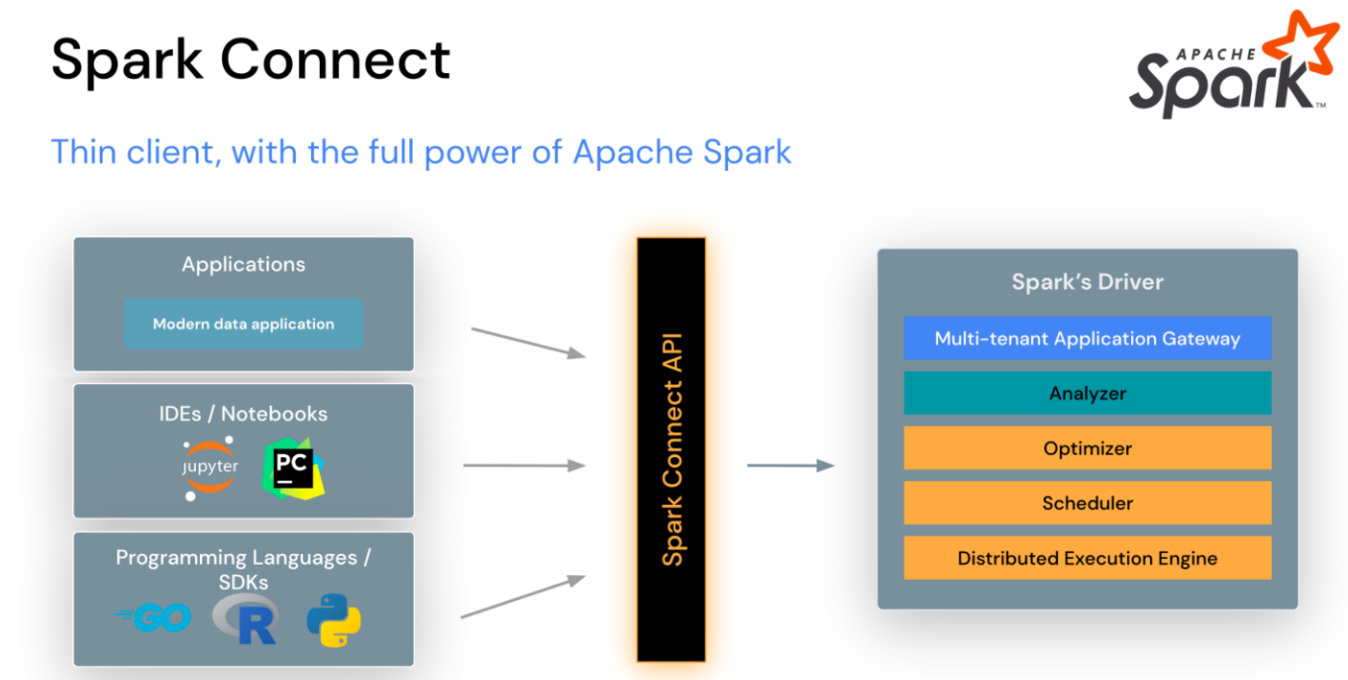

./sbin/start-connect-server.sh

client library.

Download and start Spark server with Spark Connect

In a terminal window, go to the `spark` folder in the location where you extracted Spark before and run the `start-connect-server.sh` script to start Spark server with Spark Connect, like in this example:

```
./sbin/start-connect-server.sh --packages org.apache.spark:spark-connect_2.12:3.5.4
```

Note that we include a Spark Connect package (`spark-connect_2.12:3.5.4`), when starting Spark server. This is required to use Spark Connect. Make sure to use the same version of the package as the Spark version you downloaded previously. In this example, Spark 3.5.4 with Scala 2.12.

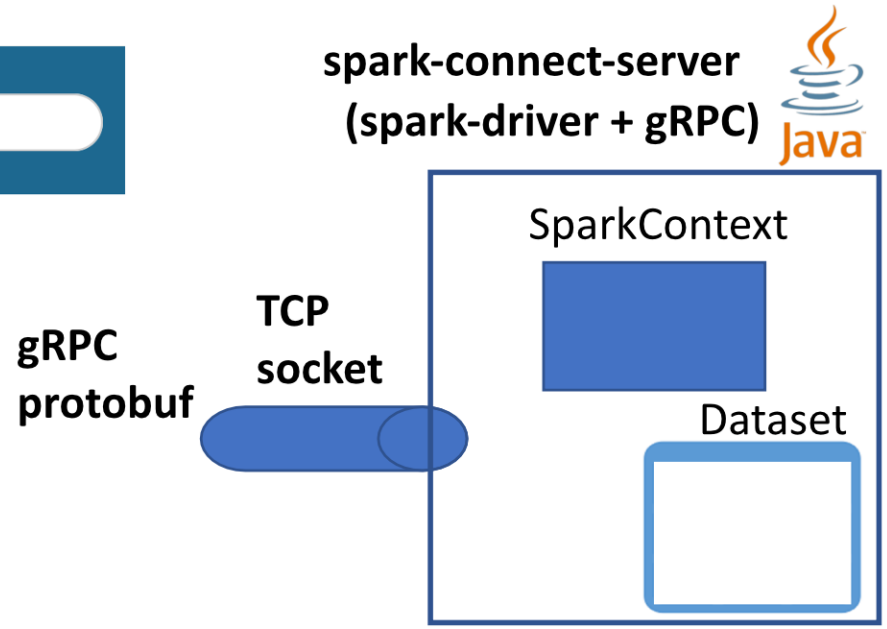

example: Spark Connect via python (no need JRE on client)

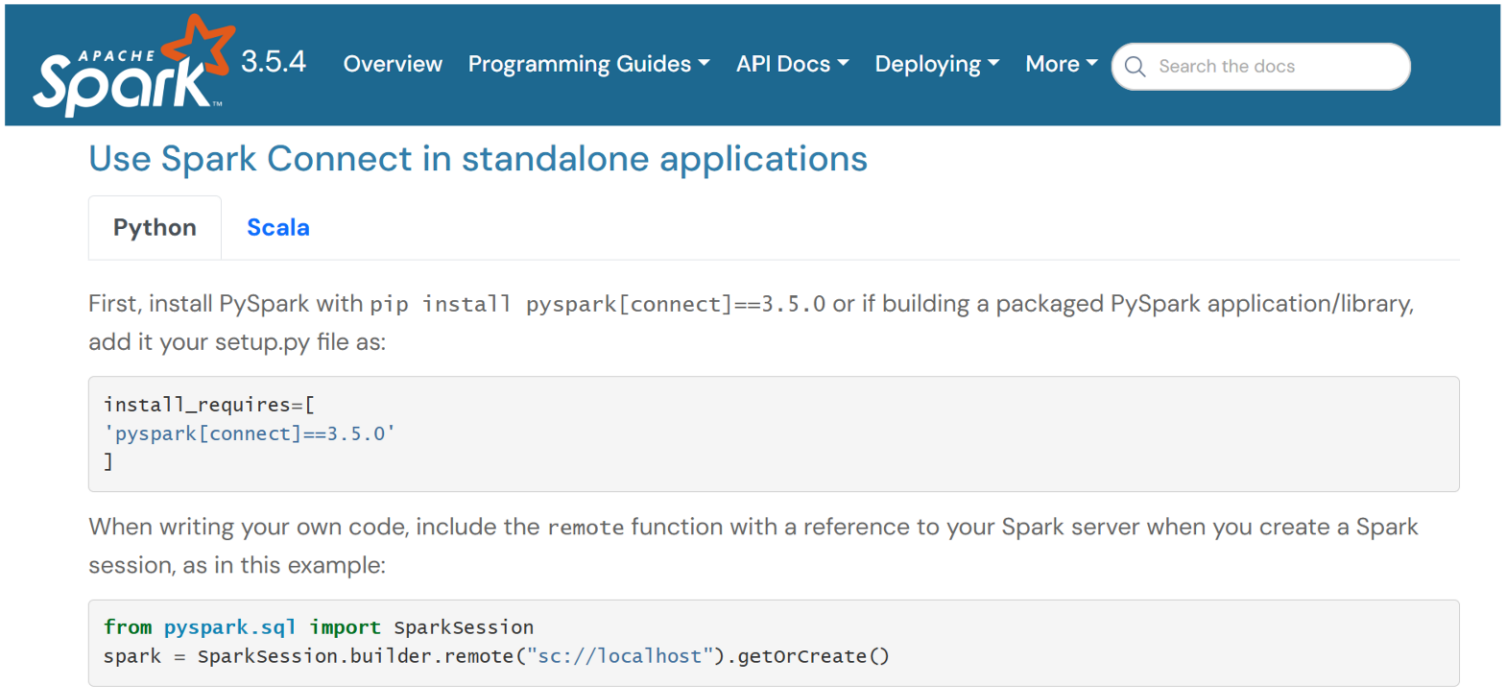

Enough for Today!

… but more to come

Questions are Welcome

Take Away

What did you learn?

MindMap

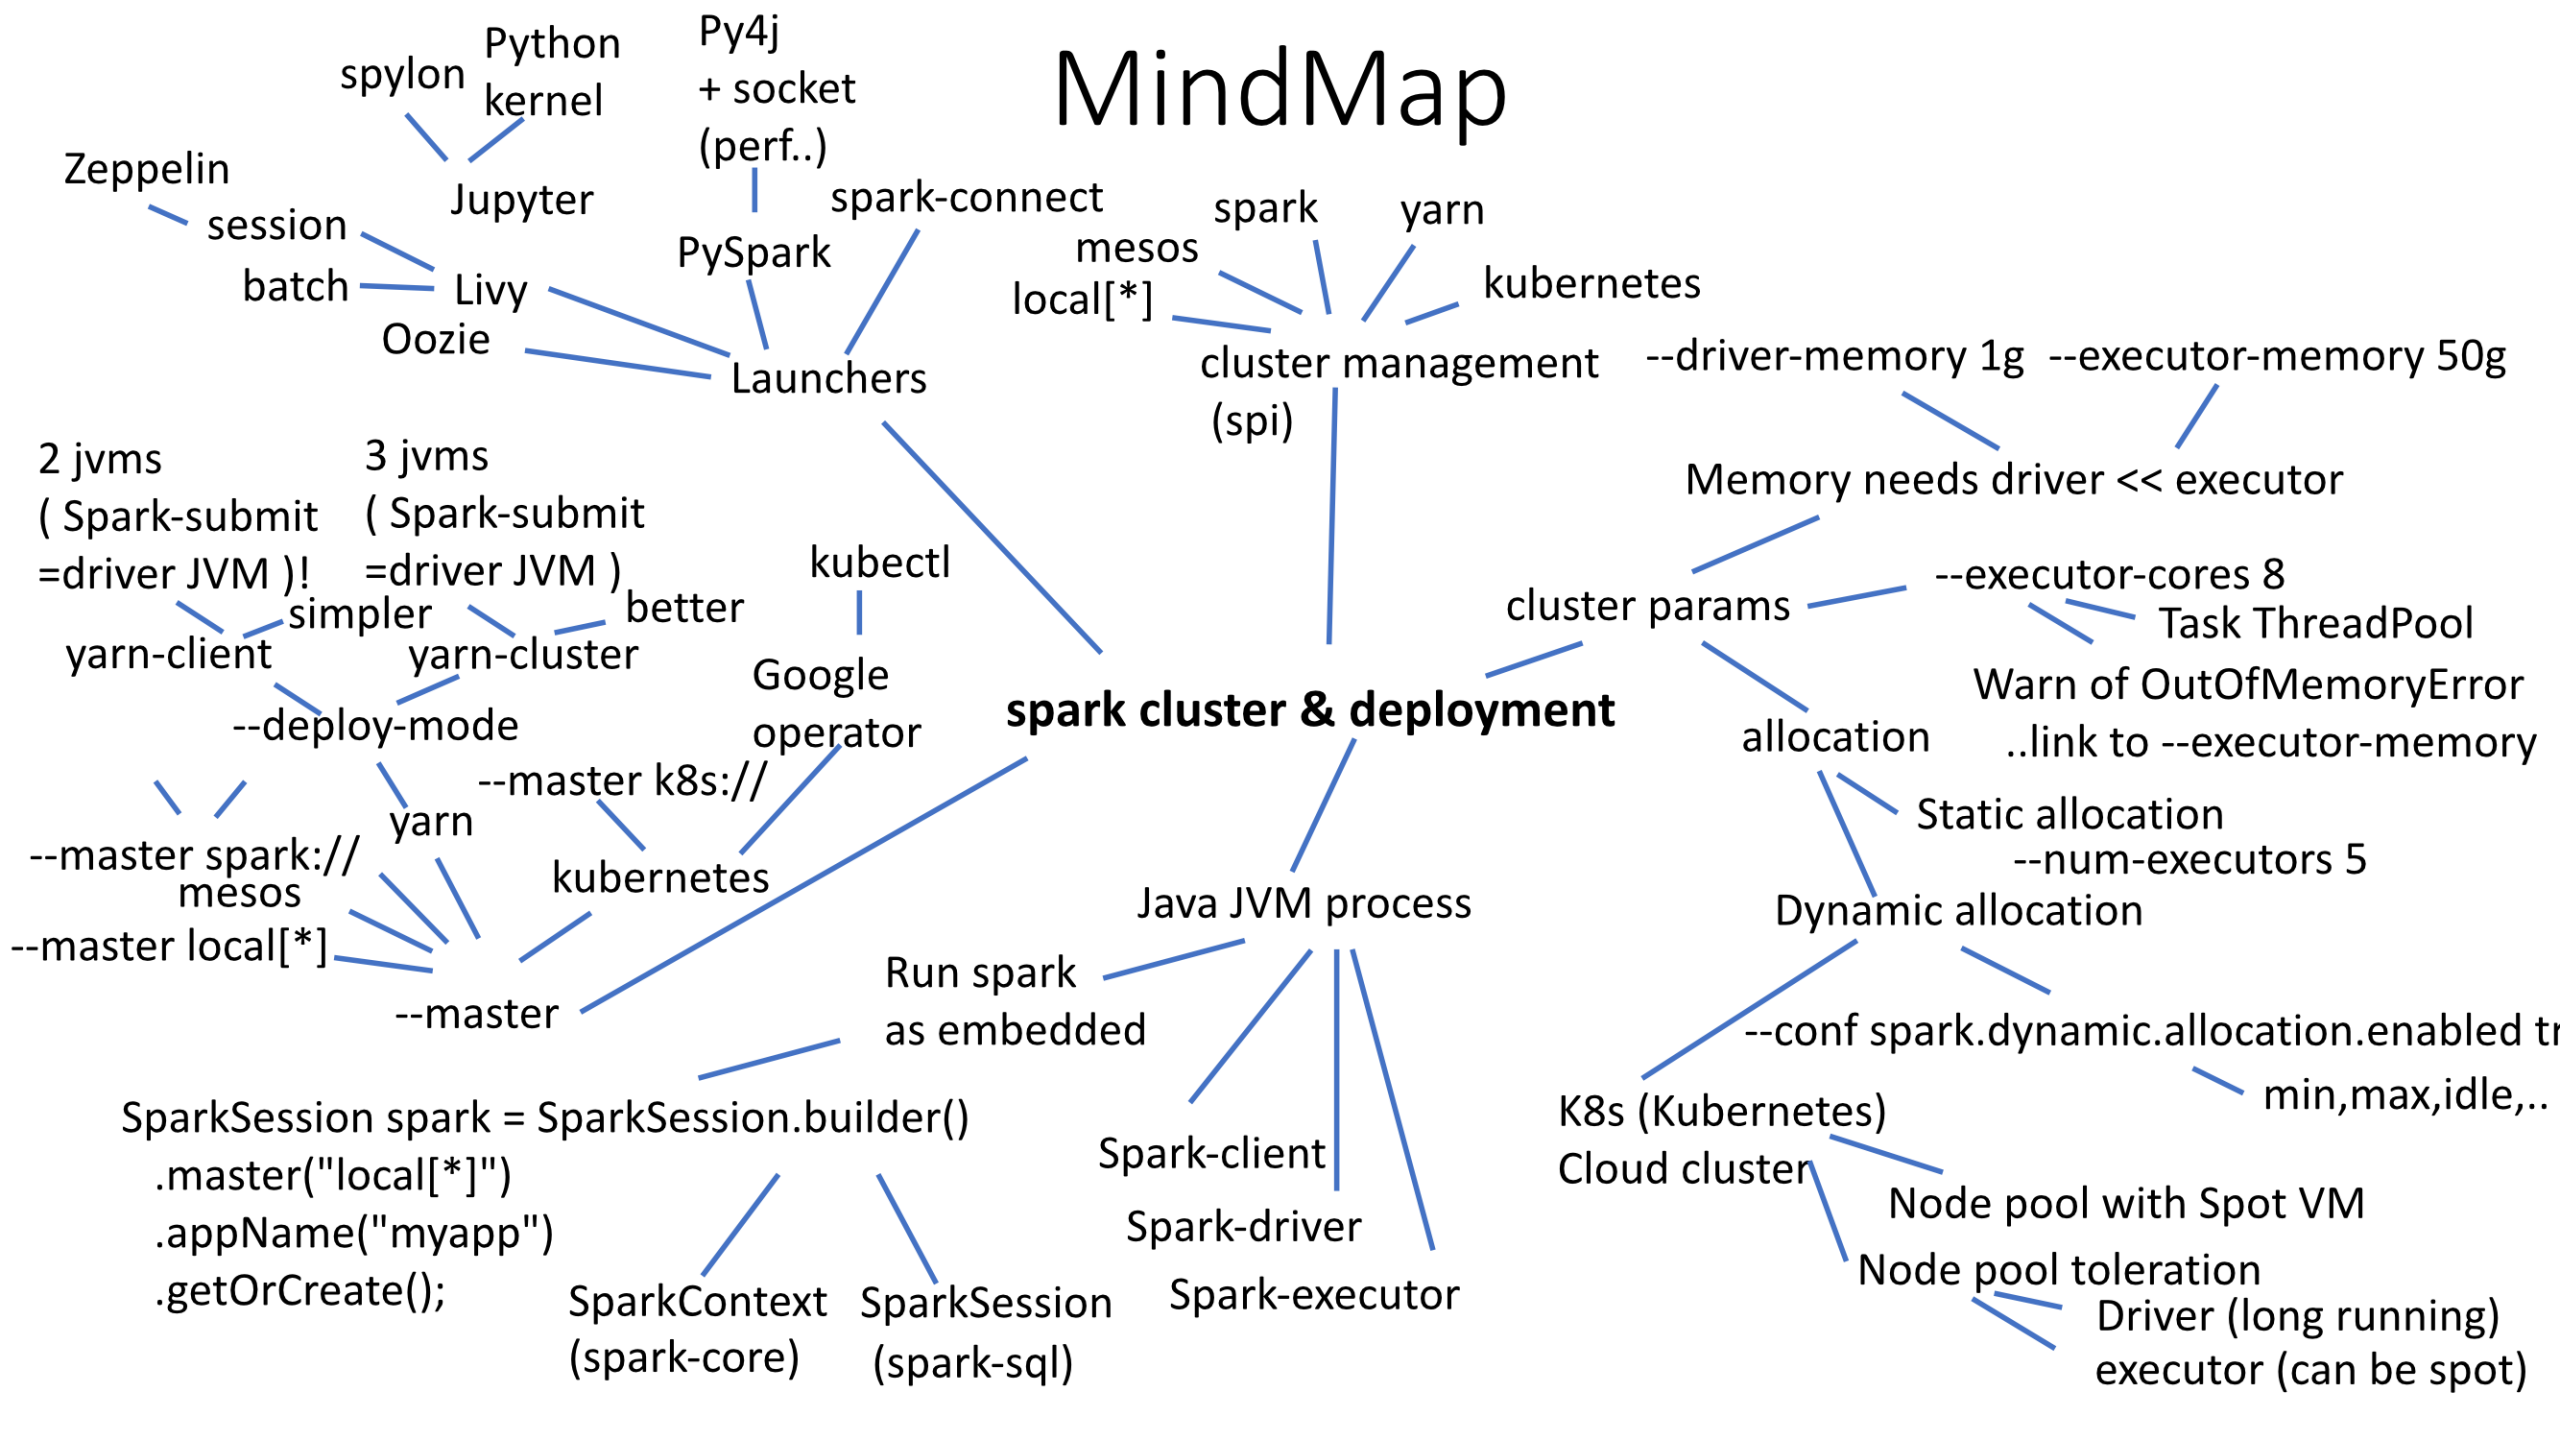
spylon
Python kernel
Py4j + socket (perf..)
Zeppelin
Jupyter
session
batch
Livy
PySpark
spark-connect
Oozie
Launchers
mesos
spark
yarn
local[*]
kubernetes
cluster management (spi)
--driver-memory 1g
--executor-memory 50g
Memory needs driver << executor
2 jvms (Spark-submit =driver JVM)!
3 jvms (Spark-submit =driver JVM)
simpler
better
yarn-client
yarn-cluster
kubectl
Google operator
--deploy-mode
spark cluster & deployment
cluster params
--executor-cores 8
Task ThreadPool
Warn of OutOfMemoryError
..link to --executor-memory
allocation
Static allocation
--num-executors 5
Dynamic allocation
--master k8s://
yarn
--master spark://
mesos
--master local[*]
kubernetes
--master
Java JVM process
Run spark as embedded
--conf spark.dynamic.allocation.enabled tr
min,max,idle,..
K8s (Kubernetes) Cloud cluster
Node pool with Spot VM
Node pool toleration
Driver (long running)
executor (can be spot)
SparkSession spark = SparkSession.builder()
.master("local[*]")
.appName("myapp")
.getOrCreate();
Spark-client
Spark-driver
Spark-executor
SparkContext (spark-core)
SparkSession (spark-sql)

Questions?

arnaud.nauwynck@gmail.com